

MAP 5611 Intro to Computational Finance HW1

Student: Sen Zhang

| | |
|-----------------|---|
| Problem 1 | 1 |
| Problem 2 | 5 |
| Problem 3 | 7 |
| Problem 4 | 7 |

Problem 1

I. Executive Summary

After comparing the computed machine epsilon, the largest and smallest floating point numbers with the IEEE standard, we checked that the computer is using the IEEE Standard. And from the basic and strong tests, it can be tested whether two floating point numbers are almost equal.

II. Statement of Problem

To check whether my computer uses the IEEE floating point numbers by writing a re-useable class defining a new floating point number type, FMfloat, that can be changed at compile time to either float or double, and can return the machine epsilon, the largest and smallest floating point numbers.

III. Description of The Mathematics

Today the computers use IEEE floating point numbers:

- For Single Precision: $FL(x) = (-1)^s 2^{E-127} 1.M$, where $0 < E < 255$

The digits of M (Mantissa) are 23, then the

single machine epsilon $\epsilon = 2^{-23} \approx 1.192093 * 10^{-7}$

single smallest number $= 2^{-126} * 1.00 \dots 0 \approx 1.175494 * 10^{-38}$

$$\text{single largest number} = 2^{127} * 1.11 \dots 1 = (2 - 2^{-23})2^{127} \approx 3.402823 * 10^{38}$$

- For Double Precision: $FL(x) = (-1)^s 2^{E-1023} 1.M$, where $0 < E < 2047$

The digits of M (Mantissa) are 52, then the

$$\text{double machine epsilon } \epsilon = 2^{-52} \approx 2.220446 * 10^{-16}$$

$$\text{double smallest number} = 2^{-1022} * 1.00 \dots 0 = 2^{-1022} \approx 2.225073 * 10^{-308}$$

$$\text{double largest number} = 2^{1023} * 1.11 \dots 1 = (2 - 2^{-52})2^{1023} \approx 1.797693 * 10^{308}$$

IV. Description of The Algorithm

To return the machine epsilon, the largest and smallest floating numbers, the functions in “float.h” is needed.

For the “bool almostequal(float, float)” and bool almostequal(double,double)”, the algorithm is as following:

```
bool falmostequal(float a, float b)
{
    const float epsilon = fepsilon();
    if (fabs(b-a) <= fabs(a)*epsilon && fabs(b-a) <= fabs(b)*epsilon)
        return 1;
    else
        return 0;
}

bool dalmostequal(double a, double b)
{
    const double epsilon = depsilon();
    if (fabs(b-a) <= fabs(a)*epsilon && fabs(b-a) <= fabs(b)*epsilon)
        return 1;
    else
        return 0;
}
```

V. Results

Table 1 Comparison of IEEE Standard and Computed Results for Single Precision

| Single Precision | IEEE Standard | Computed | Error |
|------------------|----------------------|-------------|-------|
| Machine Epsilon | $1.19209 * 10^{-7}$ | 1.19209e-7 | 0 |
| Smallest | $1.17549 * 10^{-38}$ | 1.17549e-38 | 0 |
| Largest | $3.40282 * 10^{38}$ | 3.40282e38 | 0 |

Table 2 Comparison of IEEE Standard and Computed Results for Double Precision

| Double Precision | IEEE Standard | Computed | Error |
|------------------|-----------------------|--------------|-------|
| Machine Epsilon | $2.22045 * 10^{-16}$ | 2.22045e-16 | 0 |
| Smallest | $2.22507 * 10^{-308}$ | 2.22507e-308 | 0 |
| Largest | $1.79769 * 10^{308}$ | 1.79769e308 | 0 |

Table 3 Example of the “bool almostequal()”

| Example for “almostequal” | |
|---------------------------|------------------|
| a | 2.11111111111112 |
| b | 2.11111111222222 |
| Result under strong test | Different |
| Result under basic test | Almost equal |

Graph 1 Results of running codes

```

C:\Windows\system32\cmd.exe
float epsilon: 1.19209e-007
float huge: 3.40282e+038
float tiny: 1.17549e-038
double epsilon: 2.22045e-016
double huge: 1.79769e+308
double tiny: 2.22507e-308
Please enter a number: 2.11111111111112
Please enter another number: 2.11111111222222
They are different under strong test.
Moreover,
They are almost equal under basic test.
Press any key to continue . . .

```

VI. Conclusions

As predicted, the computer is using the IEEE Standard after comparing it with the computed results for both single and double precisions. To return the machine epsilon, the largest and smallest floating numbers, the functions in “float.h” is used. For the “bool almostequal(float, float)” and bool almostequal(double, double)”, the basic test is under single precision, and the strong test is under double precision. When two numbers satisfied the formulas $|b - a| \leq \epsilon|a|$ and $|b - a| \leq \epsilon|b|$, it means that they are almost equal, if not, they are different.

VII. Program Listing

```
//FMfloat.h
class FMfloat
{
public:
    float fepsilon(float);
    double depsilon(double);
    float fhuge(float);
    double dhuge(double);
    float ftiny(float);
    double dtiny(double);
    bool falmostequal(float, float);
    bool dalmostequal(double, double);
};
```

```
//FMfloat.cpp
#include <iostream>
#include <cmath>
#include "float.h"
#include "FMfloat.h"
using namespace std;

float fepsilon()
{
    return FLT_EPSILON;
}

float fhuge()
{
    return FLT_MAX;
}

float ftiny()
{
    return FLT_MIN;
}

double depsilon()
{
    return DBL_EPSILON;
}

double dhuge()
{
    return DBL_MAX;
}

double dtiny()
{
    return DBL_MIN;
}
```

```

bool falmostequal(float a, float b)
{
    const float epsilon = fepsilon();
    if (fabs(b-a) <= fabs(a)*epsilon && fabs(b-a) <= fabs(b)*epsilon)
        return 1;
    else
        return 0;
}

bool dalmostequal(double a, double b)
{
    const double epsilon = depsilon();
    if (fabs(b-a) <= fabs(a)*epsilon && fabs(b-a) <= fabs(b)*epsilon)
        return 1;
    else
        return 0;
}

int main()
{
    cout << "float epsilon: " << fepsilon() << endl;
    cout << "float huge: " << fhuge() << endl;
    cout << "float tiny: " << ftiny() << endl;
    cout << "double epsilon: " << depsilon() << endl;
    cout << "double huge: " << dhuge() << endl;
    cout << "double tiny: " << dtiny() << endl;

    double a,b;
    cout<<"Please enter a number: ";
    cin>>a;
    cout<<"Please enter another number: ";
    cin>>b;

    if (dalmostequal(a,b))
        cout<<"They are almost equal under strong test.\n";
    else
        cout<<"They are different under strong test.\n";

    if (falmostequal(a,b))
        cout<< "Moreover, \n" <<"They are almost equal under basic test.\n";
    else
        cout<< "Moreover, \n" <<"They are different under basic test.\n";

    return 0;
}

```

Problem 2

The IEEE floating point number for Single Precision: $FL(x) = (-1)^s 2^{E-127} 1.M$, where $0 < E < 255$

The digits of M (Mantissa) are 23, then

$$\text{single machine epsilon } \epsilon = 2^{-23}$$

$$\text{single smallest number} = 2^{-126} * 1.00 \dots 0 = 2^{-126}$$

$$\text{single 2nd smallest number} = 2^{-126} * 1.00 \dots 1 = (1 + 2^{-23}) * 2^{-126}$$

$$\text{single largest number} = 2^{127} * 1.11 \dots 1 = (2 - 2^{-23})2^{127}$$

$$\text{single 2nd largest number} = 2^{127} * 1.11 \dots 0 = (2 - 2^{-23} - 2^{-23})2^{127}$$

$$1. \quad \text{The smallest spacing} = \text{the 2nd smallest number} - \text{the smallest number} = 2^{-23}2^{-126} = 1.40129 * 10^{-45}$$

$$2. \quad \text{The largest spacing} = \text{the largest number} - \text{the 2nd largest number} = 2^{-23}2^{127} = 2.02824 * 10^{31}$$

3. The smallest relative spacing is between

$$\left[\frac{\text{smallest spacing}}{\text{2nd smallest number}}, \frac{\text{smallest spacing}}{\text{smallest number}} \right] = \left[\frac{2^{-149}}{2^{-126}(1 + 2^{-23})}, \frac{2^{-149}}{2^{-126}} \right] = \left[\frac{2^{-23}}{1 + 2^{-23}}, 2^{-23} \right]$$

The largest relative spacing is between

$$\left[\frac{\text{largest spacing}}{\text{largest number}}, \frac{\text{largest spacing}}{\text{2nd largest number}} \right] = \left[\frac{2^{104}}{2^{127}(2 - 2^{-23})}, \frac{2^{104}}{2^{127}(2 - 2^{-22})} \right] = \left[\frac{2^{-23}}{2 - 2^{-23}}, \frac{2^{-23}}{2 - 2^{-22}} \right]$$

The smallest relative spacing and the largest relative spacing are different.

4. The largest and smallest (non-zero) absolute errors due to rounding in the IEEE number system are half of the largest and smallest spacing.

$$\text{largest absolute errors} = \frac{2.02824 * 10^{31}}{2} = 1.01412 * 10^{31}$$

$$\text{smallest absolute errors} = \frac{1.40129 * 10^{-45}}{2} = 7.00649 * 10^{-46}$$

5. $\epsilon * \text{the smallest number} = 2 * \text{smallest absolute errors}$

$\epsilon * \text{the largest number} = 2 * \text{largest absolute errors}$

Problem 3

Since the floating point numbers are approximations to the real numbers, then there exists some errors between them.

1. A false positive example in single precision:

Assume real numbers $a=0.12345673$, $b=0.12345671$

$$|a - b| = 0.00000002 \geq \epsilon * |a| = 1.47e - 8, \text{ and } |a - b| \geq \epsilon * |b|$$

However, the floating numbers $fl(a)$ and $fl(b)$

$$|fl(a) - fl(b)| = 0 \leq \epsilon * |fl(a)|, \text{ and } |fl(a) - fl(b)| \leq \epsilon * |fl(b)|$$

It's a contradiction.

2. A false negative example in single precision:

Assume real numbers $a=6.12345678$, $b=6.12345641$

$$|a - b| = 0.00000036 = 3.6e - 7 \leq \epsilon * |a| = 7.3e - 7, \text{ and } |a - b| \leq \epsilon * |b|$$

However, the floating numbers $fl(a)$ and $fl(b)$

$$|fl(a) - fl(b)| = 10e - 6 \geq \epsilon * |fl(a)|, \text{ and } |fl(a) - fl(b)| \geq \epsilon * |fl(b)|$$

It's also a contradiction.

Problem 4

No, floating point is not arithmetic distributive.

For an instance $\beta = 10, t = 2, L = -1, U = 2$:

$a=9.9, b=7.5, c=0.1$

$$(fl(a) + fl(b)) + fl(c) = 17 + 0.1 = 17,$$

$$fl(a) + (fl(b) + fl(c)) = 9.9 + 7.6 = 18$$

We can see that $(fl(a) + fl(b)) + fl(c) \neq fl(a) + (fl(b) + fl(c))$