

Programming with Shared Memory PART II

HPC Fall 2012

Prof. Robert van Engelen



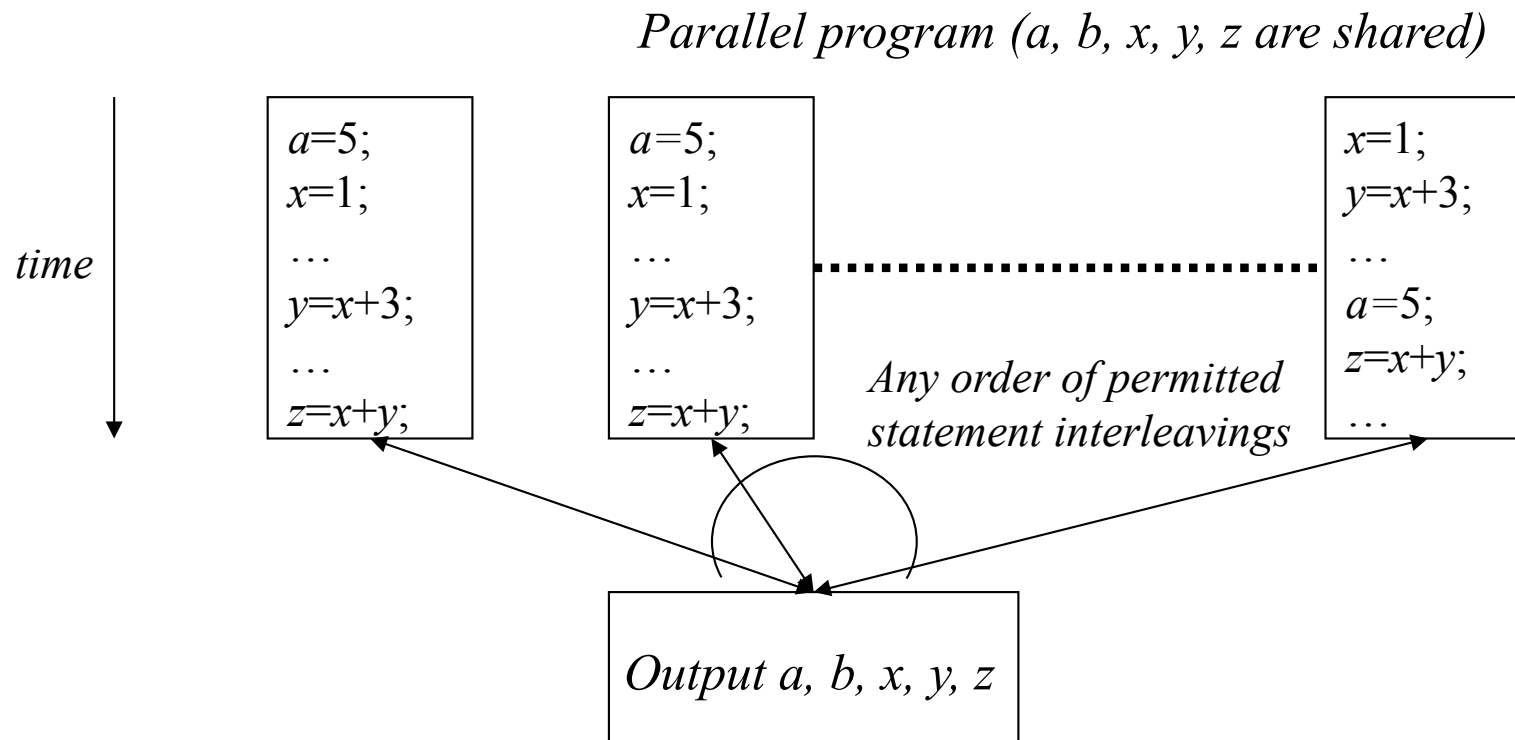


Overview


- Sequential consistency
- Parallel programming constructs
- Dependence analysis
- OpenMP
- Autoparallelization
- Further reading

Sequential Consistency

- *Sequential consistency*: the result of a parallel program is always the same as the sequential program, irrespective of the statement interleaving that is a result of parallel execution



Data Flow: Implicitly Parallel



Flow dependences determine the parallel execution schedule: each operation waits until operands are produced



Explicit Parallel Programming Constructs

- Declaring shared data, when private is implicit

`shared int x;`

A shared variable

`shared int *p;`


*Private pointer to
a shared value*

- Declaring private data, when private is explicit

`private int x;`

`private int *p;`

Would this make any sense?




Explicit Parallel Programming Constructs

- The **par** construct

```
par {  
  S1;  
  S2;  
  ...  
  Sn;  
}
```

*Statements in the body are
executed concurrently*






Explicit Parallel Programming Constructs

- The `forall` construct (also called `parfor`)

```
forall (i=0; i<n; i++) {  
    S1;  
    S2;  
    ...  
    Sm;  
}
```

Statements in the body are executed in serial order by n threads $i=0..n-1$ in parallel





Explicit Parallel: Many Choices, Which is Safe?

```
par {  
  a=5;  
  x=1;  
}  
...  
y=x+3;  
...  
z=x+y;
```

```
x=1;  
...  
par {  
  a=5;  
  y=x+3;  
}  
...  
z=x+y;
```

```
x=1;  
...  
par {  
  a=5;  
  y=x+3;  
  z=x+y;  
}  
...
```

Think about data flow: each operation requires completion of operands first

Data dependences preserved ➡ *sequential consistency guaranteed*



Bernstein's Conditions

- I_i is the set of memory locations read by process P_i
- O_j is the set of memory locations altered by process P_j
- Processes P_1 and P_2 can be executed concurrently if all of the following conditions are met

$$I_1 \cap O_2 = \emptyset$$

$$I_2 \cap O_1 = \emptyset$$

$$O_1 \cap O_2 = \emptyset$$



Bernstein's Conditions Verified by Dependence Analysis

- *Dependence analysis* performed by a compiler determines that Bernstein's conditions are not violated when optimizing and/or parallelizing a program

independent

$P_1: \mathbf{A} = \mathbf{x} + \mathbf{y};$
 $P_2: \mathbf{B} = \mathbf{x} + \mathbf{z};$

$$I_1 \cap O_2 = \emptyset$$

$$I_2 \cap O_1 = \emptyset$$

$$O_1 \cap O_2 = \emptyset$$

RAW

$P_1: \mathbf{A} = \mathbf{x} + \mathbf{y};$
 $P_2: \mathbf{B} = \mathbf{x} + \mathbf{A};$

$$I_1 \cap O_2 = \emptyset$$

$$I_2 \cap O_1 = \{\mathbf{A}\}$$

$$O_1 \cap O_2 = \emptyset$$

WAR

$P_1: \mathbf{A} = \mathbf{x} + \mathbf{B};$
 $P_2: \mathbf{B} = \mathbf{x} + \mathbf{z};$

$$I_1 \cap O_2 = \{\mathbf{B}\}$$

$$I_2 \cap O_1 = \emptyset$$

$$O_1 \cap O_2 = \emptyset$$

WAW

$P_1: \mathbf{A} = \mathbf{x} + \mathbf{y};$
 $P_2: \mathbf{A} = \mathbf{x} + \mathbf{z};$

$$I_1 \cap O_2 = \emptyset$$

$$I_2 \cap O_1 = \emptyset$$

$$O_1 \cap O_2 = \{\mathbf{A}\}$$



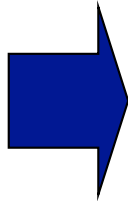
Bernstein's Conditions Verified by Dependence Analysis

- *Dependence analysis* performed by a compiler determines that Bernstein's conditions are not violated when optimizing and/or parallelizing a program

independent

$P_1: \mathbf{A} = \mathbf{x} + \mathbf{y};$
 $P_2: \mathbf{B} = \mathbf{x} + \mathbf{z};$

$I_1 \cap O_2 = \emptyset$
 $I_2 \cap O_1 = \emptyset$
 $O_1 \cap O_2 = \emptyset$



`par {`
 $P_1: \mathbf{A} = \mathbf{x} + \mathbf{y};$
 $P_2: \mathbf{B} = \mathbf{x} + \mathbf{z};$
`}`

Recall:
instruction scheduling for
instruction-level parallelism
(ILP)



Bernstein's Conditions in Loops

- A parallel loop is valid when any ordering of its parallel body yields the same result

```
forall (I=4;I<7;I++)  
S1:  A[I] = A[I-3]+B[I];
```

S1(4):	A[4]	=	A[1]+B[4];
S1(5):	A[5]	=	A[2]+B[5];
S1(6):	A[6]	=	A[3]+B[6];

S1(4):	A[4]	=	A[1]+B[4];
S1(6):	A[6]	=	A[3]+B[6];
S1(5):	A[5]	=	A[2]+B[5];

S1(5):	A[5]	=	A[2]+B[5];
S1(4):	A[4]	=	A[1]+B[4];
S1(6):	A[6]	=	A[3]+B[6];

S1(6):	A[6]	=	A[3]+B[6];
S1(5):	A[5]	=	A[2]+B[5];
S1(4):	A[4]	=	A[1]+B[4];

S1(5):	A[5]	=	A[2]+B[5];
S1(6):	A[6]	=	A[3]+B[6];
S1(4):	A[4]	=	A[1]+B[4];

S1(6):	A[6]	=	A[3]+B[6];
S1(4):	A[4]	=	A[1]+B[4];
S1(5):	A[5]	=	A[2]+B[5];



OpenMP

- OpenMP is a portable implementation of common parallel constructs for shared memory machines
- OpenMP in C

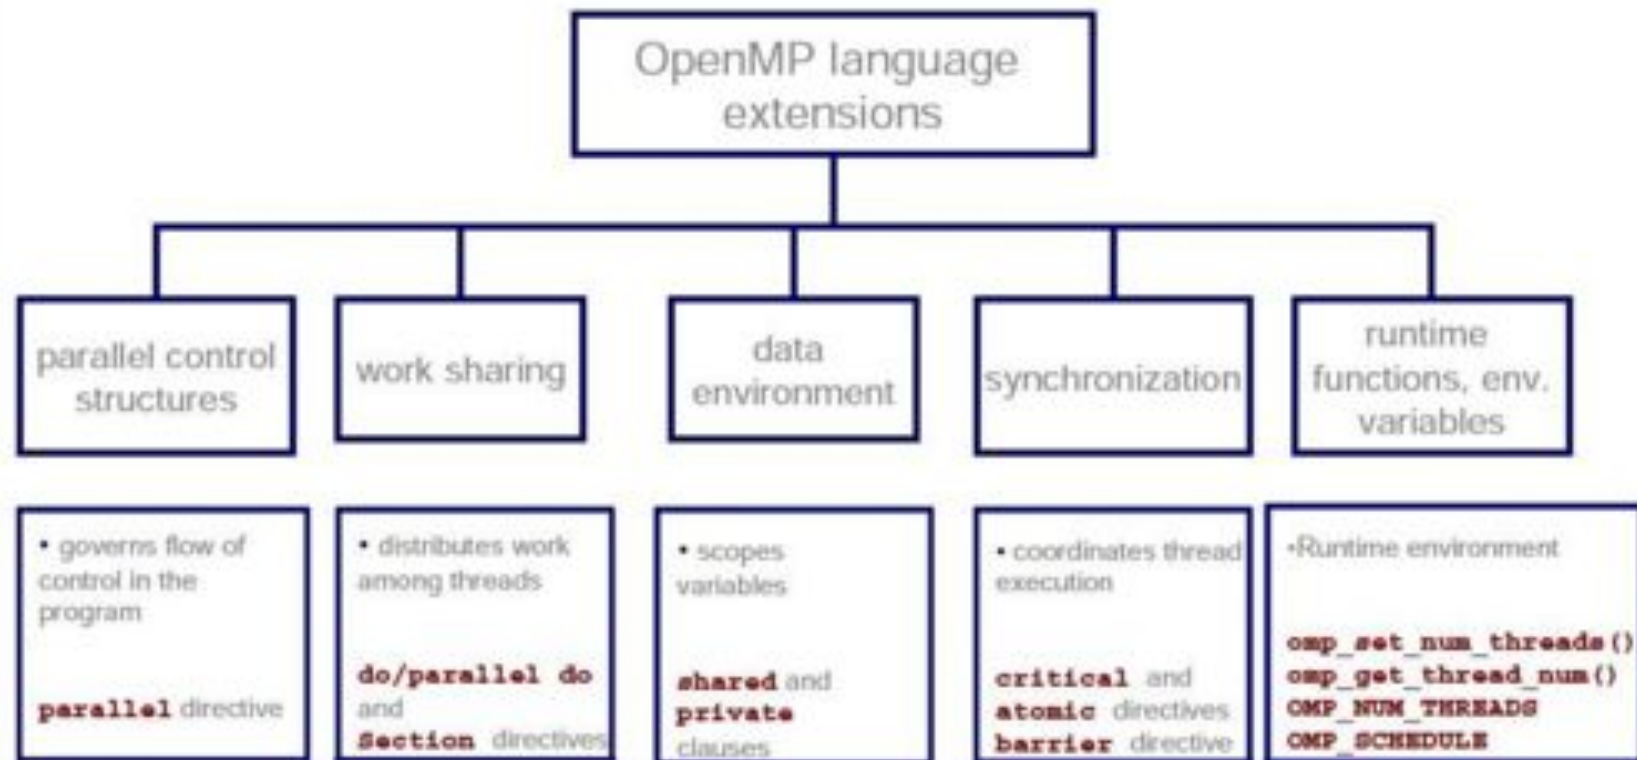
```
#pragma omp directive_name  
    statement_block
```

- OpenMP in Fortran

```
!$OMP directive_name  
    statement_block  
!$OMP end directive name
```

OpenMP

OpenMP Constructs




OpenMP Parallel

- The `parallel` construct in OpenMP is not the same as `par`

```
#pragma omp parallel
{
    S1;
    S2;
    ...
    Sm;
}
```

A team of threads all execute the body statements and joins when done




OpenMP Parallel

- The `parallel` construct

```
#pragma omp parallel default(none) shared(vars)
{
    S1;
    S2;
    ...
    Sm;
}
```

This specifies that variables should not be assumed to be shared by default





OpenMP Parallel

- The `parallel` construct

```
#pragma omp parallel private(n, i)
{
    n = omp_get_num_threads();
    i = omp_get_thread_num();
    ...
}
```

Use **private** to declare private data

omp_get_num_threads()

returns the number of threads that are currently being used

omp_get_thread_num()

returns the thread id (0 to n-1)

OpenMP Parallel with Reduction

- The `parallel` construct with reduction clause

operation: +, *, -, &, ^, |, &&, ||

```
#pragma omp parallel reduction(+:var)
```

```
{
```

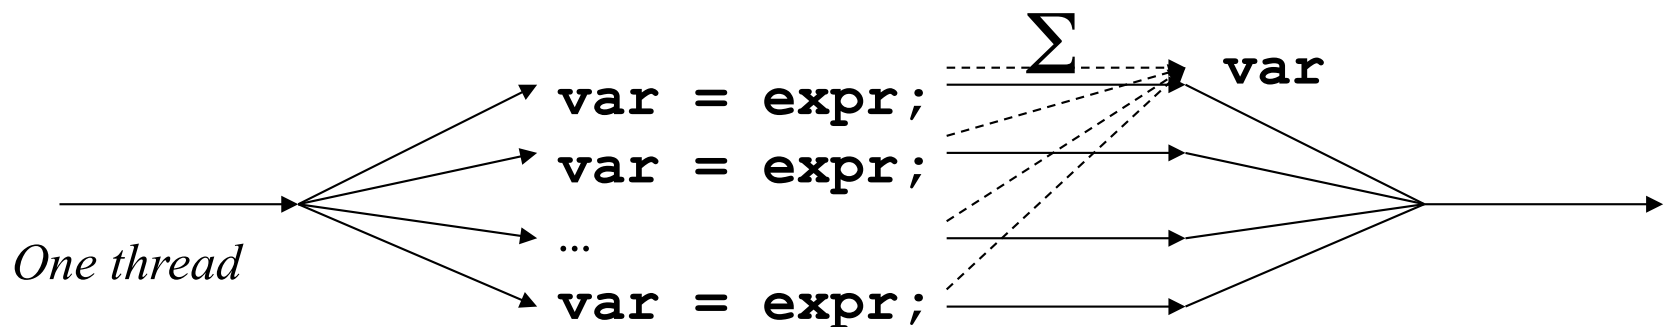
```
    var = expr;
```

```
    ...
```

```
}
```

```
... = var;
```

Performs a global reduction operation over privatized variable(s) and assigns final value to master's private variable(s) or to the shared variable(s) when shared




OpenMP Parallel

- The `parallel` construct

```
#pragma omp parallel num_threads(n)
{
    S1;
    S2;
    ...
    Sm;
}
```

↑
Number of threads we want

*Alternatively, use `omp_set_num_threads()`
or set environment variable `OMP_NUM_THREADS`*




OpenMP Parallel Sections

- The `sections` construct is for *work-sharing*, where a current team of threads is used to execute statements concurrently

```
#pragma omp parallel
...
#pragma omp sections
{
    #pragma omp section
        S1;
    #pragma omp section
        S2;
    ...
    #pragma omp section
        Sm;
}
```

Statements in the sections are executed concurrently





OpenMP Parallel Sections

- The `sections` construct is for *work-sharing*, where a current team of threads is used to execute statements concurrently

```
#pragma omp parallel
```

```
...
```

```
#pragma omp sections nowait
```

```
{
```

```
    #pragma omp section
```

```
        S1;
```

```
    #pragma omp section
```

```
        S2;
```

```
    ...
```


```
    #pragma omp section
```

```
        Sm;
```

```
}
```

*Use **nowait** to remove the implicit barrier*

*n threads
executing m sections*




OpenMP Parallel Sections

- The `sections` construct is for *work-sharing*, where a current team of threads is used to execute statements concurrently

```
#pragma omp parallel sections
{
    #pragma omp section
    S1;
    #pragma omp section
    S2;
    ...
    #pragma omp section
    Sm;
}
```

Use **parallel sections** to
combine **parallel** with **sections**



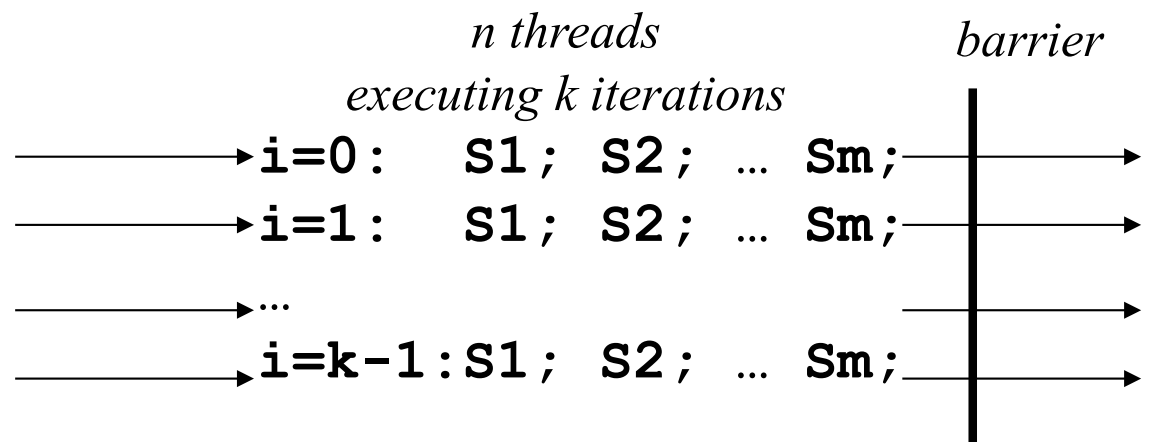
OpenMP For/Do

- The **for** construct (**do** in Fortran) is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
...
#pragma omp for
for (i=0; i<k; i++)
{
    S1;
    S2;
    ...
    Sm;
}
```

Loop iterations are executed concurrently by n threads

*Use **nowait** to remove the implicit barrier*





OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
```

```
...
```

```
#pragma omp for schedule(dynamic)
```

```
for (i=0; i<k; i++)
```

```
{
```

```
    S1;
```


```
    S2;
```

```
    ...
```

```
    Sm;
```

```
}
```

When $k > n$, threads execute randomly chosen loop iterations until all iterations are completed



OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
```

```
...
```

```
#pragma omp for schedule(static)
```

```
for (i=0; i<4; i++)
```

```
{
```

```
    S1;
```


```
    S2;
```

```
    ...
```

```
    Sm;
```

```
}
```

When $k > n$, threads are assigned to $\lceil k/n \rceil$ chunks of the iteration space



OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
```

```
...
```

```
#pragma omp for schedule(static, 2)
```

```
for (i=0; i<8; i++)
```

```
{
```

```
  S1;
```

```
  S2;
```

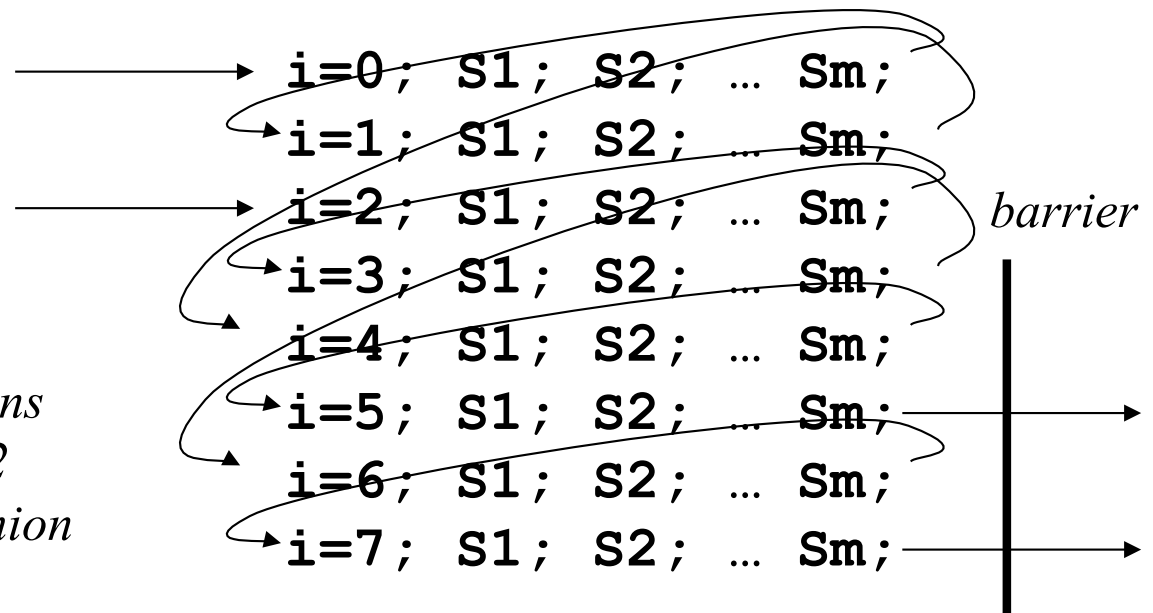
```
  ...
```

```
  Sm;
```

```
}
```

*2 threads
executing 8 iterations
using chunk size 2
in a round-robin fashion*

Chunk size





OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
```

```
...
```

```
#pragma omp for schedule(guided, 4)
```

```
for (i=0; i<k; i++)
```

```
{
```

```
    S1;
```

```
    S2;
```

```
    ...
```

```
    Sm;
```

```
}
```



Chunk size



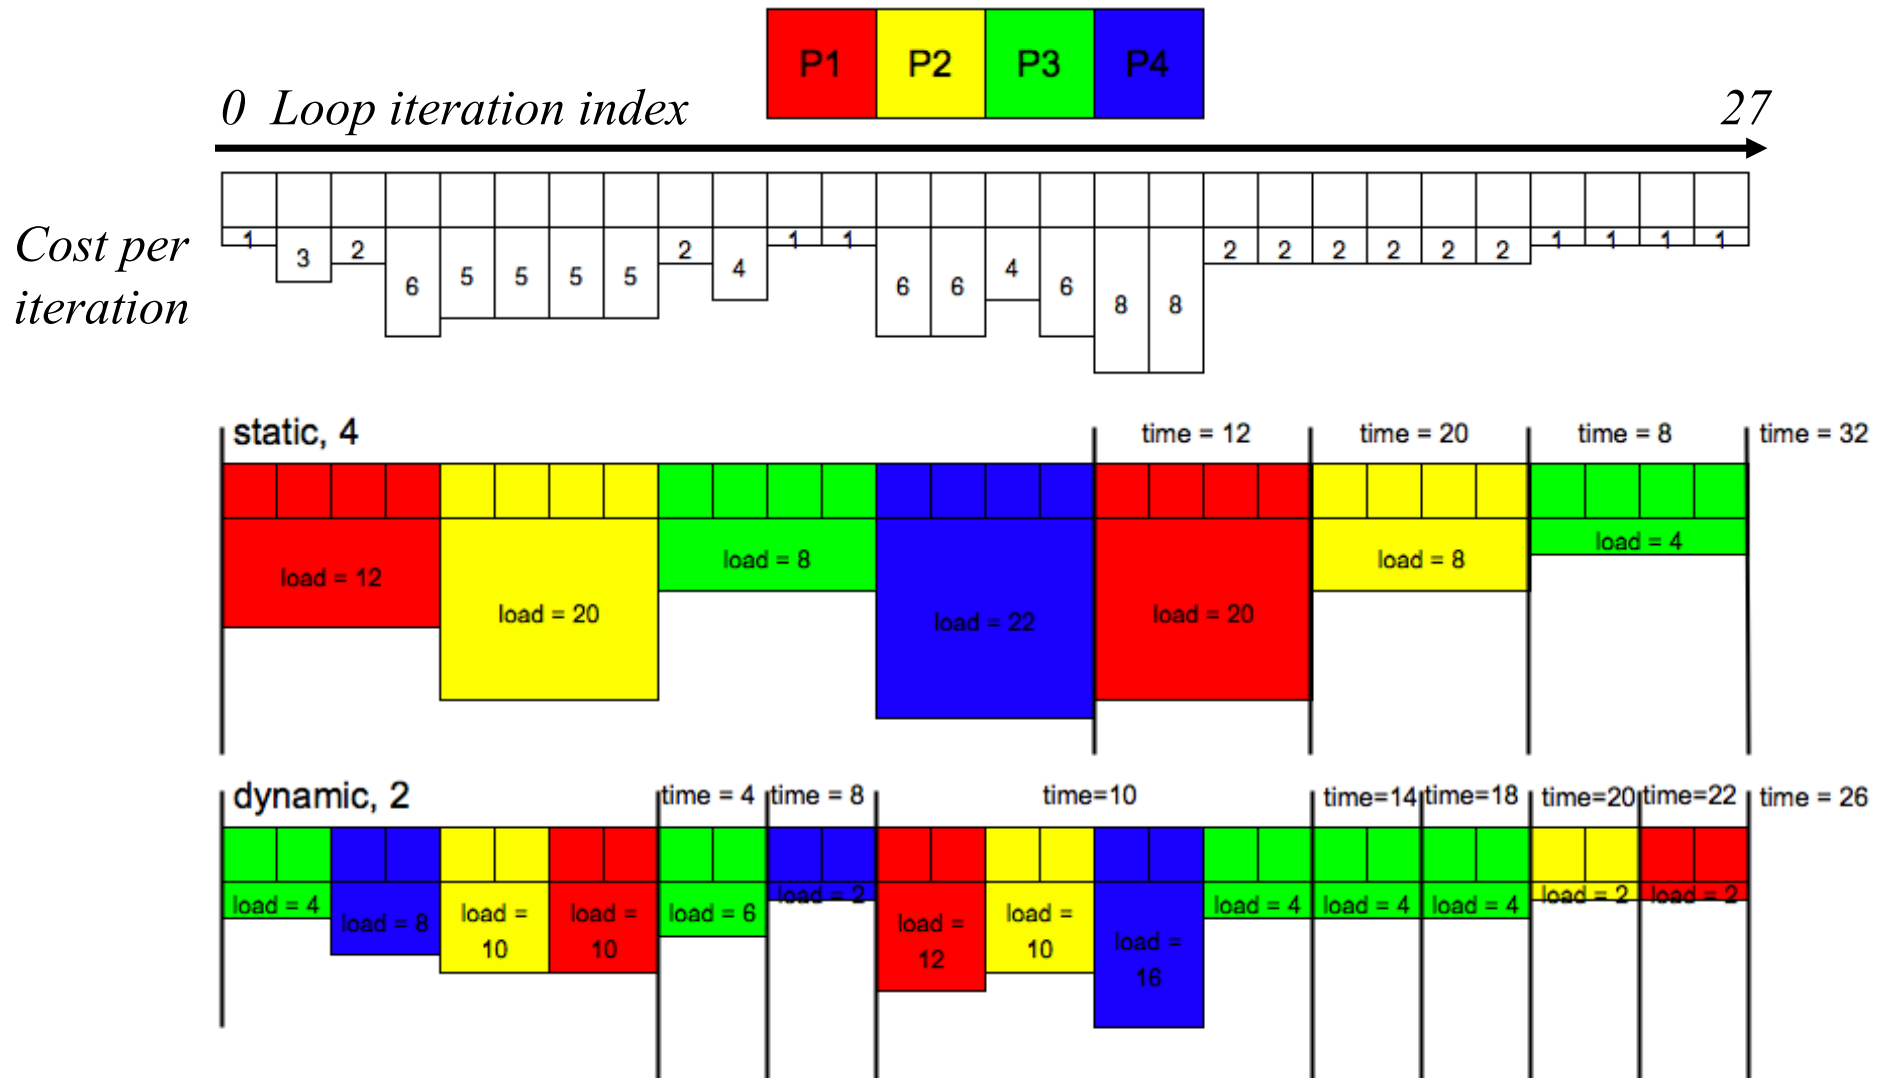
*Exponentially decreasing chunk size,
for example: 4, 2, 1*



OpenMP For/Do Scheduling Comparison



OpenMP For/Do Scheduling with Load Imbalances





OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
```

```
...
```

```
#pragma omp for schedule(run_time)
```

```
for (i=0; i<k; i++)
```

```
{
```

```
    S1;
```

```
    S2;
```

```
    ...
```

```
    Sm;
```

```
}
```

Controlled by environment variable **OMP_SCHEDULE:**

```
setenv OMP_SCHEDULE "dynamic"
```

```
setenv OMP_SCHEDULE "static"
```

```
setenv OMP_SCHEDULE "static,2"
```

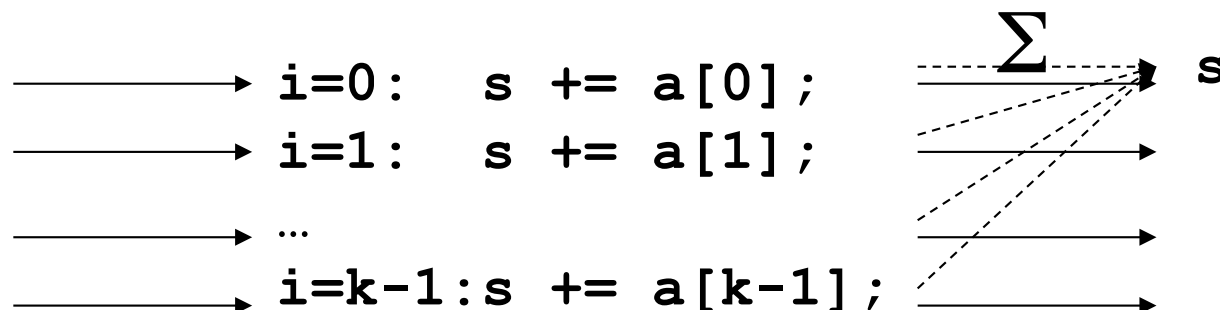
OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel
...
#pragma omp for reduction(+:s)
for (i=0; i<k; i++)
{
    s += a[i];
}
```

operation: +, *, -, &, ^, |, &&, ||

Performs a global reduction operation over privatized variables and assigns final value to master's private variable(s) or to the shared variable(s) when shared



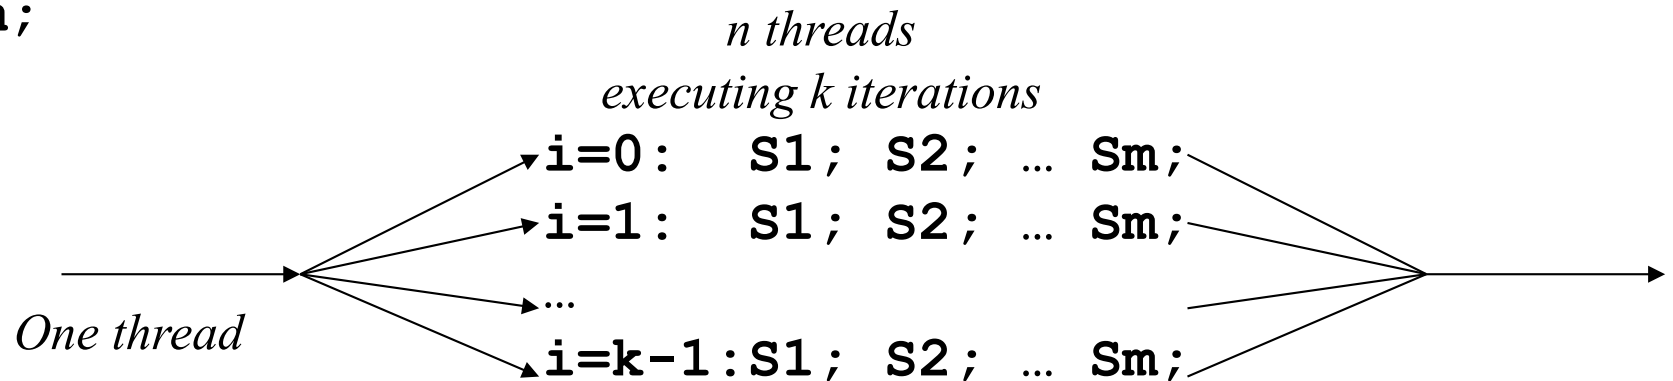


OpenMP For/Do

- The **for** construct is for *work-sharing*, where a current team of threads is used to execute a loop concurrently

```
#pragma omp parallel for
for (i=0; i<k; i++)
{
    S1;
    S2;
    ...
    Sm;
}
```

Use **parallel for** to combine
parallel with **for**





OpenMP Firstprivate and Lastprivate

- The `parallel` construct with `firstprivate` and/or `lastprivate` clause

```
x = ...;
#pragma omp parallel firstprivate(x) lastprivate(y)
{ x = x + ...;
  #pragma omp for
  for (i=0; i<k; i++)
  { ...
    y = i;
  }
}
... = y;
```

*Use **firstprivate** to declare private variables that are initialized with the main thread's value of the variables*

*Likewise, use **lastprivate** to declare private variables whose values are copied back out to main thread's variables by the thread that executes the last iteration of a parallel for loop, or the thread that executes the last parallel section*



OpenMP Single

- The **single** construct selects one thread of the current team of threads to execute the body

```
#pragma omp parallel
```

```
...
```

```
#pragma omp single
```

```
{
```

```
  S1;
```

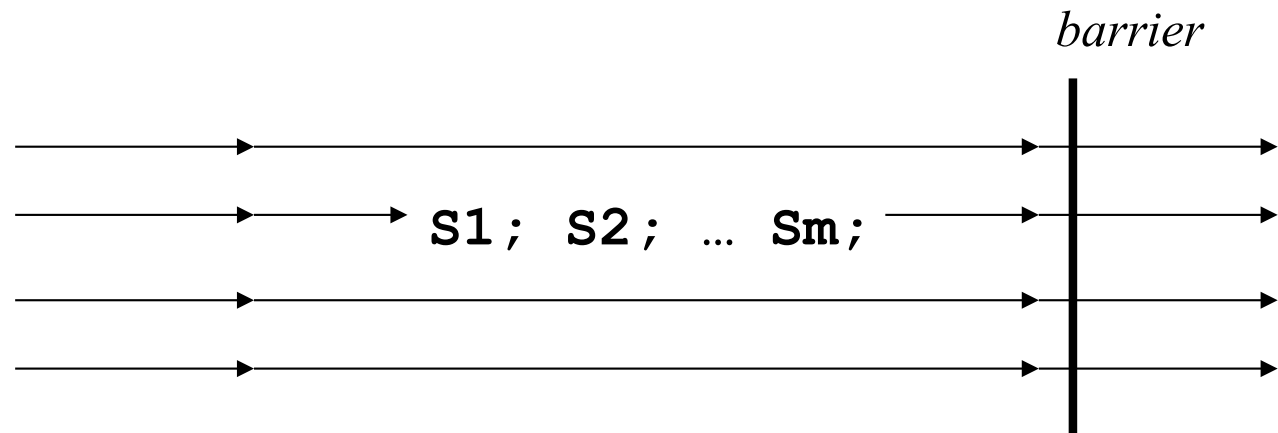
```
  S2;
```

```
  ...
```

```
  Sm;
```

```
}
```

One thread executes the body





OpenMP Master

- The **master** construct selects the master thread of the current team of threads to execute the body

```
#pragma omp parallel
```

```
...
```

```
#pragma omp master
```

```
{
```

```
    S1;
```


```
    S2;
```

```
    ...
```

```
    Sm;
```

```
}
```

The “master” thread executes the body, no barrier is inserted



OpenMP Critical

- The `critical` construct defines a critical section

```
#pragma omp parallel
```

```
...
```

```
#pragma omp critical name
```

```
{
```

```
    S1;
```


```
    S2;
```

```
    ...
```

```
    Sm;
```

```
}
```

Mutual exclusion is enforced on the body using a named lock



OpenMP Critical

- The `critical` construct defines a critical section

```
#pragma omp parallel
```

```
...
```


```
#pragma omp critical qlock  
{ enqueue(job) ;  
}
```

One thread is here

```
...
```

```
#pragma omp critical qlock  
{ dequeue(job) ;  
}
```

Another thread is here



OpenMP Critical

- The `critical` construct defines a critical section

```
#pragma omp parallel
```

```
...
```

```
#pragma omp critical
```

```
{
```

```
  S1;
```


```
  S2;
```

```
  ...
```

```
  Sm;
```

```
}
```

Mutual exclusion is enforced on the body using an anonymous lock






OpenMP Barrier

- The **barrier** construct synchronizes the current team of threads

```
#pragma omp parallel
```

```
...
```

```
#pragma omp barrier
```






OpenMP Atomic

- The **atomic** construct executes an expression atomically (expressions are restricted to simple updates)

```
#pragma omp parallel
```

```
...
```

```
#pragma omp atomic  
expression;
```





OpenMP Atomic

- The **atomic** construct executes an expression atomically (expressions are restricted to simple updates)

```
#pragma omp parallel
```

```
...
```

```
#pragma omp atomic
```

```
n = n+1;
```


```
...
```

```
#pragma omp atomic
```

```
n = n-1;
```

One thread is here

Another thread is here





OpenMP Flush


- The **flush** construct flushes shared variables from local storage (registers, cache) to shared memory
- OpenMP adopts a *relaxed consistency model* of memory

```
#pragma omp parallel
```

```
...
```

```
#pragma omp flush(variables)
```

b = 3, but there is no
guarantee that **a** will be 3






OpenMP Relaxed Consistency Memory Model

- *Relaxed consistency* means that memory updates made by one CPU may not be immediately visible to another CPU
 - Data can be in registers
 - Data can be in cache
(cache coherence protocol is slow or non-existent)
- Therefore, the updated value of a shared variable that was set by a thread may not be available to another
- An OpenMP *flush* is automatically performed at
 - Entry and exit of `parallel` and `critical`
 - Exit of `for`
 - Exit of `sections`
 - Exit of `single`
 - Barriers

OpenMP Thread Scheduling

- Controlled by environment variable **OMP_DYNAMIC**
- When set to **FALSE**
 - Same number of threads used for every parallel region
- When set to **TRUE**
 - The number of threads is adjusted for each parallel region
 - `omp_get_num_threads()` returns actual number of threads
 - `omp_get_max_threads()` returns **OMP_NUM_THREADS**





OpenMP Threadprivate

- The **threadprivate** construct declares variables in a global scope private to a thread across multiple parallel regions
 - Must use when variables should stay private, even outside of the current scope, e.g. across function calls

```
int counter;           Global counter
#pragma omp threadprivate(counter)

#pragma omp parallel
{ counter = 0;         Each thread has a local copy of counter
  ...
}
...
#pragma omp parallel
{ counter++;           Each thread has a local copy of counter
  ...
}
```



OpenMP Locks

- Mutex locks, with additional “nestable” versions of locks

```
omp_lock_t lck;  
  
omp_init_lock(&lck);  
omp_set_lock(&lck);  
...  
... critical section ...  
...  
omp_unset_lock(&lck);  
omp_destroy_lock(&lck);
```

omp_lock_t
the lock type

omp_init_lock()
initialization

omp_set_lock()
blocks until lock is acquired

omp_unset_lock()
releases the lock

omp_destroy_lock()
deallocates the lock



Compiler Options for OpenMP

- GOMP project for GCC 4.2 (C and Fortran)
- Use `#include <omp.h>`
 - Note: the `_OPENMP` define is set when compiling with OpenMP
- Intel compiler:
`icc -openmp ...`
`ifort -openmp`
- Sun compiler:
`suncc -xopenmp ...`
`f95 -xopenmp`



Autoparallelization

- Compiler applies dependence analysis and parallelizes a loop (or entire loop nest) automatically when possible
 - Typically task-parallelizes *outer loops* (more parallel work), possibly after loop interchange, fusion, etc.
 - Similar to adding `#pragma parallel for` to loop(s), with appropriate `private` and `shared` clauses
- Intel compiler:
`icc -parallel ...`
`ifort -parallel ...`
- Sun compiler:
`suncc -xautopar ...`
`f95 -xautopar ...`



Further Reading

- [PP2] pages 248-271
- Optional:
OpenMP tutorial at Lawrence Livermore
<http://www.llnl.gov/computing/tutorials/openMP/>