

Parallel Programming Languages

HPC Fall 2012

Prof. Robert van Engelen






Overview

- Partitioned Global Address Space (PGAS)
- A selection of PGAS parallel programming languages
 - CAF
 - UPC
- Further reading



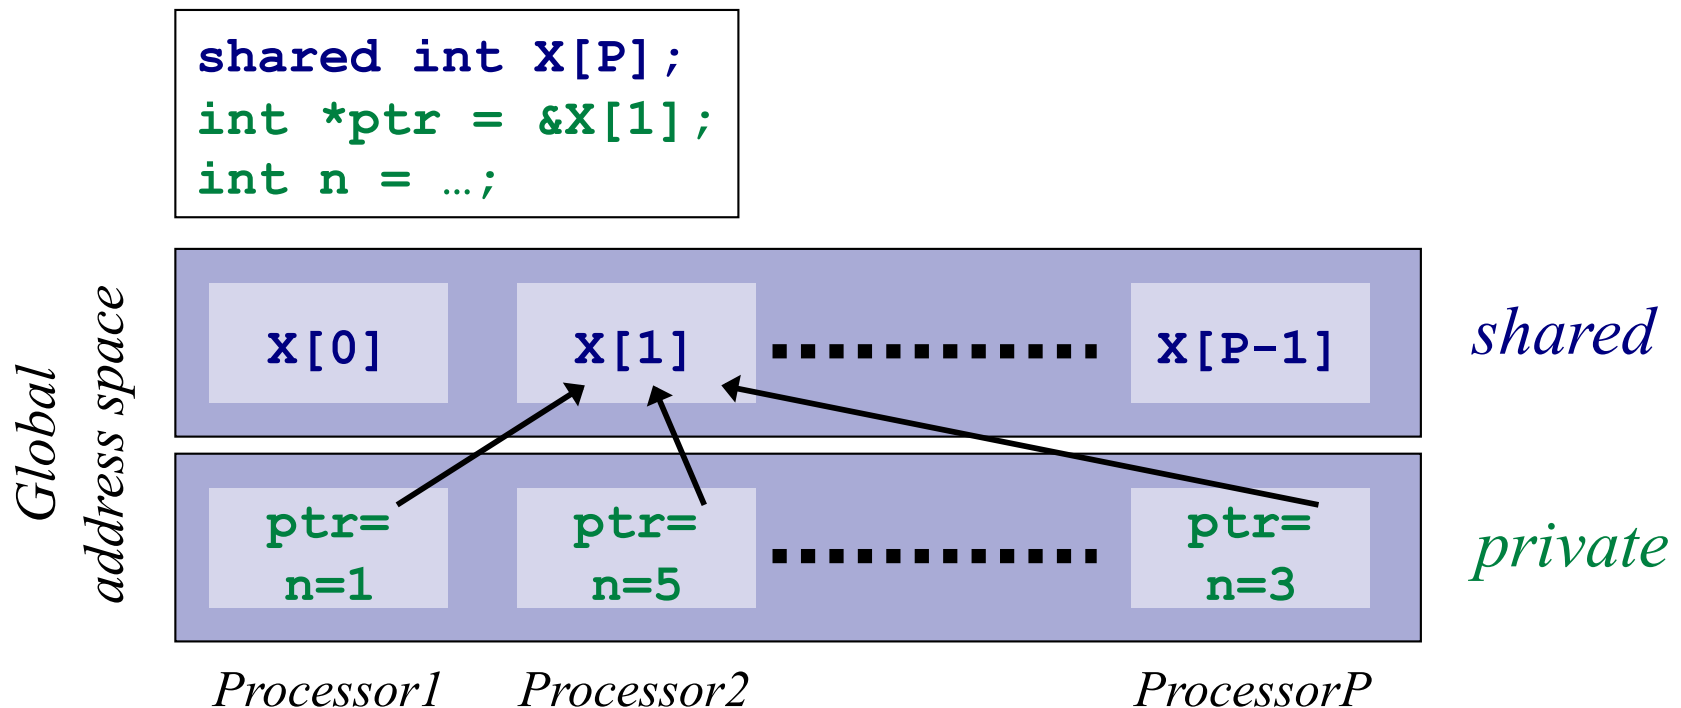
Global Address Space (GAS)

- Global address space languages take advantage of
 - Ease of programmability of shared memory parallel
 - SPMD parallelism
 - Allow local-global distinction of data, because data layout matters for performance
- Partitioned global address space is logically shared, physically distributed
 - Shared arrays are distributed over processor memories
 - Implicit communication for remote data access



Partitioned Global Address Space (PGAS)

- Global address space with two-level model that supports locality management
 - Local memory (private variables)
 - Remote memory (shared variables)





Partitioned Global Address Space (PGAS) Model

- Global address space with two-level memory model that supports locality management
 - Local memory (private variables)
 - Remote memory (shared variables)
- Programmer controls critical decisions
 - Data partitioning (by data placement in PGAS memory)
 - Communication (implicitly, via remote PGAS memory access)
- Suitable for mapping to a range of parallel architectures
 - Shared memory, message passing, and hybrid
- Languages: CAF (Fortran), UPC (C), X10 (Java), Titanium (Java)



PGAS Model vs Implementation

- PGAS is an abstract model
- Implementations differ with respect to details:
 - Address space partitioned by processors
 - Physically: at the memory address level (= DSM, e.g. Cray T3D/E)
 - Logically: at the variable level, where each variable can be arbitrarily placed in local memory on remote processor
 - Local caching of remote memory?
 - Coherence protocol
 - Communication
 - One-sided, e.g. DMA, is usually faster
 - Two-sided, e.g. MPI send/recv
 - Bulk memory copy operations or individual copies




Co-Array Fortran (CAF)

- Explicitly-parallel extension of Fortran 90/95
 - Commercial compiler from Cray/SGI
 - Open source compiler from Rice University
- Partitioned global address space SPMD with two-level model that supports locality management
 - Local memory (private variables)
 - Remote memory (shared variables)
- As usual, programmer controls critical decisions
 - Data partitioning
 - Communication

CAF: Co-Arrays

- A co-array is a partitioned array with an *image* dimension

```
REAL, DIMENSION(N) [*] :: X,Y  
X(:) = Y(:) [Q]
```





CAF: Array Syntax and Implicit Remote Memory Operations


```
REAL, DIMENSION(N)           :: X ! array
REAL, DIMENSION(N) [*]       :: Y ! co-array
REAL, DIMENSION(N,P) [*]     :: Z ! co-array
```

```
X          = Y[PE]           ! get X(1..N) from Y(1..N)[PE]
Y[PE]      = X               ! put X(1..N) into Y(1..N)[PE]
Y[:]       = X               ! broadcast X(1..N) to Y(1..N)
Y[LIST]    = X               ! broadcast X(1..N) over subset
                           ! of PEs in array LIST
Z(:)       = Y[:]            ! all-gather, collect Y(1..N)
                           ! over PEs in Z(1..N,1..P)
S = MINVAL(Y[:])             ! min (reduce) Y(1..N) over PEs
Z[:]       = S               ! S scalar, promoted to array
                           ! of shape (1:N,1:P)
```

CAF: Synchronization

```
COMMON/XCTILB4/ B(N,4) [*]  
SAVE /XCTILB4/
```

```
ME = THIS_IMAGE()  
IF (ME > 1 .AND. ME < NUM_IMAGES()) THEN  
  CALL SYNC_ALL( WAIT=(/ME-1,ME+1/) ) ← Wait for  
  B(:,1) = B(:,3) [ME-1]                processors on the  
  B(:,4) = B(:,2) [ME+1]                left and right  
  CALL SYNC_ALL( WAIT=(/ME-1,ME+1/) )  
ENDIF
```



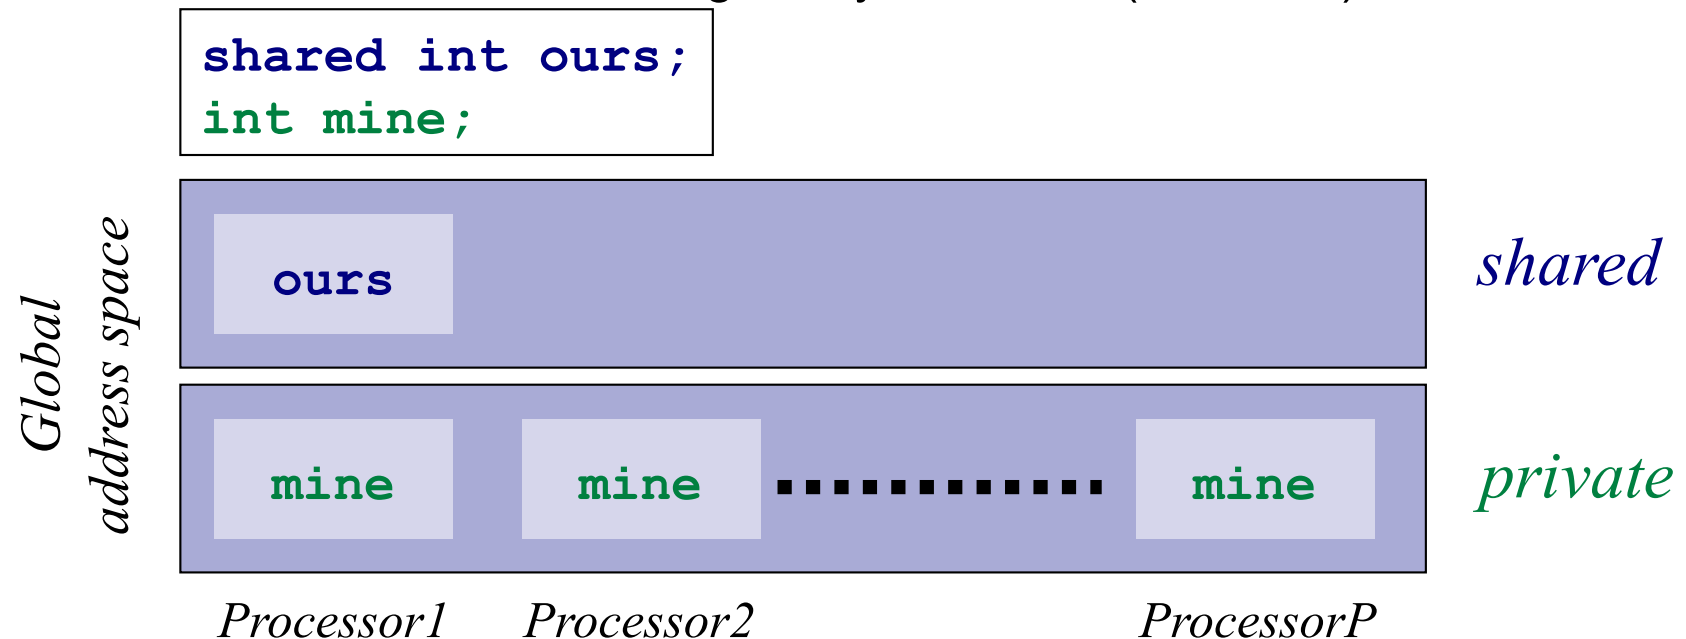



Unified Parallel C (UPC)

- UPC is an explicit extension of ANSI C
 - Commercial compilers from Cray/SGI, HP
 - Open source compiler from LBNL/UCB/MTU/UF and GCC-UPC project
- Follows the C language philosophy
 - Programmers are clever and careful and may need to work close to the hardware level
 - to get performance,
 - but allows you to get into trouble, just like programming low level C!
 - Concise and efficient syntax
- UPC is a PGAS language
 - Global address space with private and shared variables
 - Private/shared pointers to private/shared variables
 - Array data distributions (block/cyclic)
 - Forall worksharing loops
 - Barriers and locks
 - Bulk copy operations between shared and private memory

UPC: Shared Variables

- Private by default
 - C variables and objects are allocated in private memory space for each thread
- Shared variables are explicitly declared and allocated once (by thread 0)
 - Shared variables must be “globally” declared (i.e. static)







UPC: Simple Example Monte Carlo pi Calculation

```
int hit()
{
    int const rand_max = 0xFFFFFFFF;
    double x = ((double) rand()) / RAND_MAX;
    double y = ((double) rand()) / RAND_MAX;
    return ((x*x + y*y) <= 1.0);
}
```

Randomly throw darts at (x,y) positions in a unit circle, if $x^2 + y^2 \leq 1$, then point is inside circle

*Compute ratio of points inside/total, then $\pi = 4 * \text{ratio}$*






UPC: Simple Example Monte Carlo pi Calculation

```
#include <upc.h>
shared int hits = 0;
main()
{ int i;
  int my_trials, trials = ...;
  my_trials = (trials + THREADS - 1) / THREADS;
  srand(MYTHREAD*17);
  for (i=0; i < my_trials; i++)
    hits += hit();
  if (MYTHREAD == 0)
    printf("pi estimated to %g\n",
          4*(double)hits/(double)trials);
}
```

Divide the work

Score hits

What can go wrong?



UPC: Simple Example Monte Carlo pi Calculation

```
shared int hits = 0;
main()
{ int i, my_trials, trials = ...;
  upc_lock_t *hit_lock = upc_all_lock_alloc();
  my_trials = (trials + THREADS - 1)/THREADS;
  srand(MYTHREAD*17);


  for (i=0; i < my_trials; i++)
  { upc_lock(hit_lock);
    hits += hit();
    upc_unlock(hit_lock);
  }

  upc_barrier;
  if (MYTHREAD == 0)
    printf("pi estimated to %g\n",
          4*(double)hits/(double)trials);
  upc_lock_free(hit_lock);
}
```

Score hits

Synchronize

*Anything
wrong
here...?*



UPC: Simple Example Monte Carlo pi Calculation

```
shared int hits[THREADS] = { 0 };
main()
{ int i, my_trials, trials = ...;
  my_trials = (trials + THREADS - 1)/THREADS;
  srand(MYTHREAD*17);

  for (i=0; i < my_trials; i++)
    hits[MYTHREAD] += hit();

  upc_barrier;
  if (MYTHREAD == 0)
  { for (i=1; i < THREADS; i++)
    hits[0] += hits[i];

    tot_trials = THREADS*my_trials;
    printf("pi estimated to %g\n",
           4*(double)hits[0]/(double)tot_trials);
  }
}
```

Score hits

Sync

Sum hits

Corrected

UPC: Forall Work Sharing


```
shared int v1[N], v2[N], sum[N];  
  
int i;  
upc_forall(i=0; i<N; i++; i)  
    sum[i] = v1[i] + v2[i];
```

← *Default distribution: cyclic*

↙ *Affinity: here it forces
owner-computes rule*

```
for(i=0; i<N; i++)  
    if (MYTHREAD == i%THREADS)  
        sum[i] = v1[i] + v2[i];
```

Assume THREADS=4



Elements with affinity to processor 0 are shown in red

UPC: Pointers


Where does the pointer reside?

		Local	Shared
<i>Where does the referenced value reside?</i>	Private	PP (p1)	PS (p3)
	Shared	SP (p2)	SS (p4)

```
int *p1;           /* private pointer to local memory */
shared int *p2;    /* private pointer to shared space */
int *shared p3;    /* shared pointer to local memory */
shared int *shared p4; /* shared pointer to
                        shared space */
```

Shared pointer to private local memory is not recommended

UPC: Pointers



```
int *p1;           /* private pointer to local memory */
shared int *p2;    /* private pointer to shared space */
int *shared p3;    /* shared pointer to local memory */
shared int *shared p4; /* shared pointer to
                        shared space */
```




UPC: Pointer Example

```
shared int v1[N], v2[N], sum[N];  
  
int i;  
shared int *p1, *p2;  
  
p1 = v1;  
p2 = v2;  
upc_forall(i=0; i<N; i++, p1++, p2++; i)  
    sum[i] = *p1 + *p2;
```



UPC: Pointers

- In UPC pointers to shared objects have three fields:
 - thread number
 - local address of block (for blocked data distributions)
 - phase (specifies position in the block)





UPC: Shared Variable Layout

- Non-array shared variables have affinity with thread 0
- Array layouts are cyclic or blocked:

```
shared double x[n];      /* cyclic */  
shared [b] double y[n];  /* blocked */
```

where **b** is the block size

- For blocked layouts, element *i* has affinity with thread:

$$(i/b) \% \text{ THREADS}$$

therefore use *i/b* in forall (owner-computes):

```
upc_forall(i=0; i<N; i++; i/b) y[i] = ...
```



UPC: Consistency Model

- The consistency model of shared memory accesses are controlled by qualifiers

- ☐ Strict: will always appear in order
- ☐ Relaxed: may appear out of order to other threads

```
strict: {  
    x = y;  
    z = y+1;  
}
```

- Use strict on variables that are used as synchronization

Thread1

```
flag = 0;  
data = ...;  
flag = 1;
```

Thread2

```
while (flag)  
    ... = data;
```

- Select the default consistency model with:

- ☐ `#include <upc_strict.h>`
- ☐ `#include <upc_relaxed.h>`



UPC: Fence

- UPC provides a fence construct

- Syntax

- ```
upc_fence;
equivalent to a null strict reference
strict { }
```

- Ensures that all shared references issued before the `upc_fence` are complete





# Further Reading

- CAF: [www.co-array.org](http://www.co-array.org)
- UPC: [upc.gwu.edu](http://upc.gwu.edu)