

Programming with Message Passing PART II: MPI

HPC Fall 2012

Prof. Robert van Engelen





Overview

- MPI overview
- MPI process creation
- MPI point-to-point communication
- MPI collective communications
- MPI groups and communicators
- MPI virtual topologies
- MPE and Jumpshot
- Further reading



MPI

- Message Passing Interface (MPI) is an *API and protocol standard with portable implementations*
 - MPICH, LAM-MPI, OpenMPI, ...
- Hardware platforms:
 - Distributed Memory
 - Shared Memory
 - Particularly SMP / NUMA architectures
 - Hybrid
 - SMP clusters, workstation clusters, and heterogeneous networks
- Parallelism is explicit
 - Programmer is responsible for correctly identifying parallelism and implementing parallel algorithms using MPI constructs
- SPMD model with static process creation
 - MPI-2 allows dynamic process creation: **MPI_Comm_spawn()**



MPI Process Creation

- Static process creation
 - Start a N processes running the same program *prog1*
`mpirun prog1 -np N`
but this does not specify where the prog1 copies run
 - Use batch processing tools, such as SGE, to run MPI programs on a cluster

- MPI-2 supports dynamic process creation:
`MPI_Comm_spawn()`



MPI SPMD Computational Model

```
main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);

    doWork();

    MPI_Finalize();
}
```

All processes execute this work

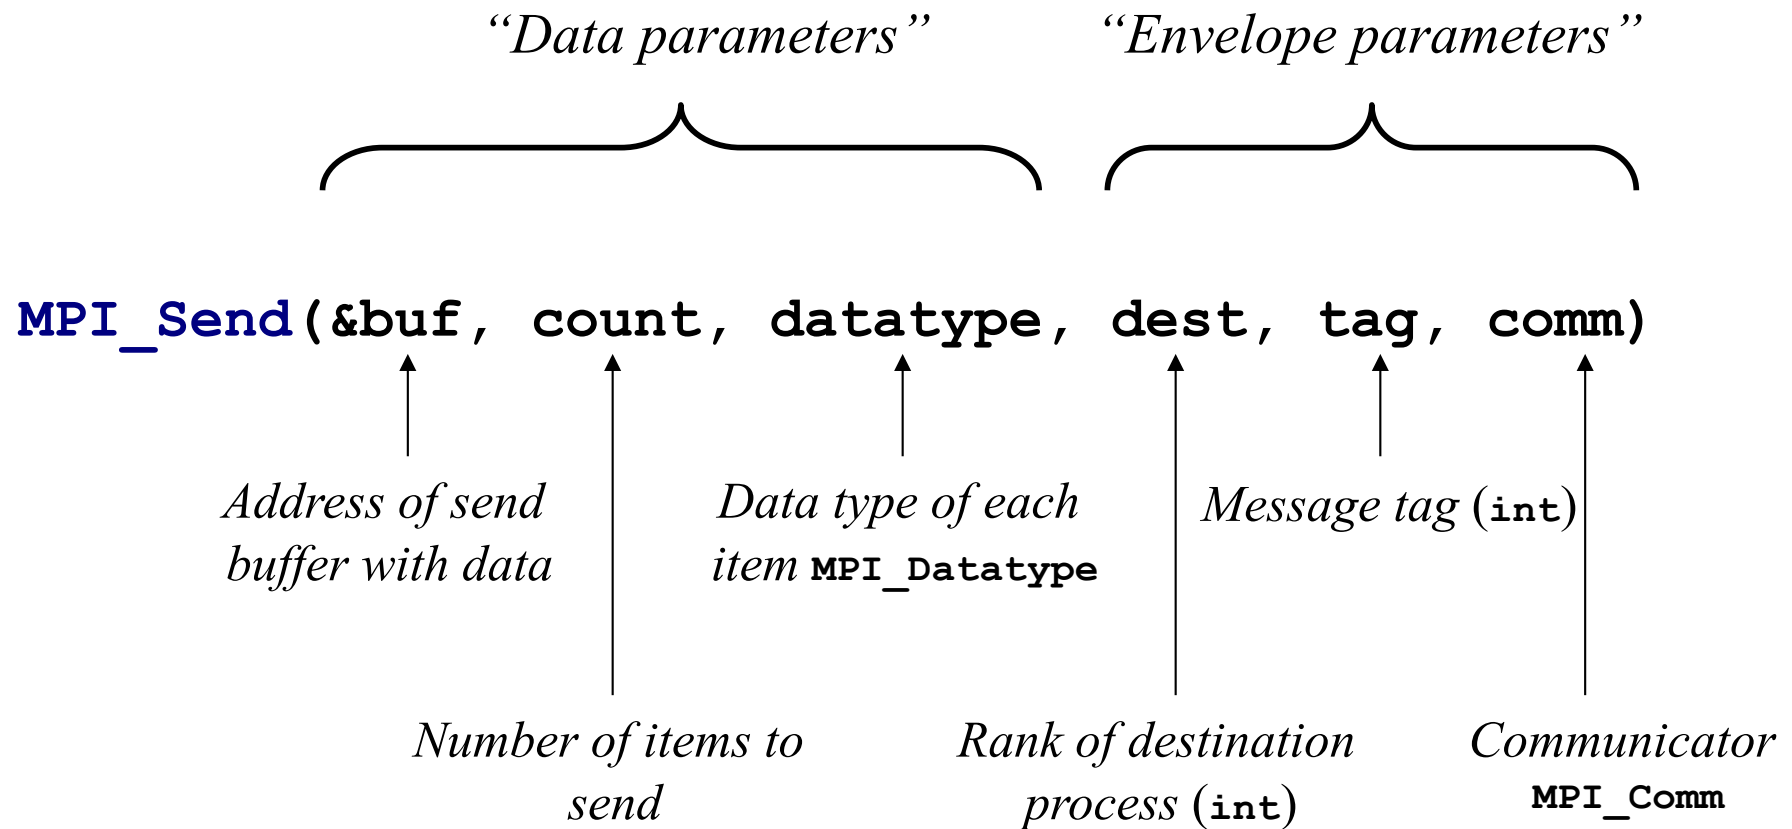
*All processes belong to the
MPI_COMM_WORLD communicator,
and each process has a rank from 0
to P-1 in the communicator*

```
doWork()
{
    int myrank;
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    if (myrank == 0)
        printf("I'm the master process");
}
```

*Use rank numbers
to differentiate work*

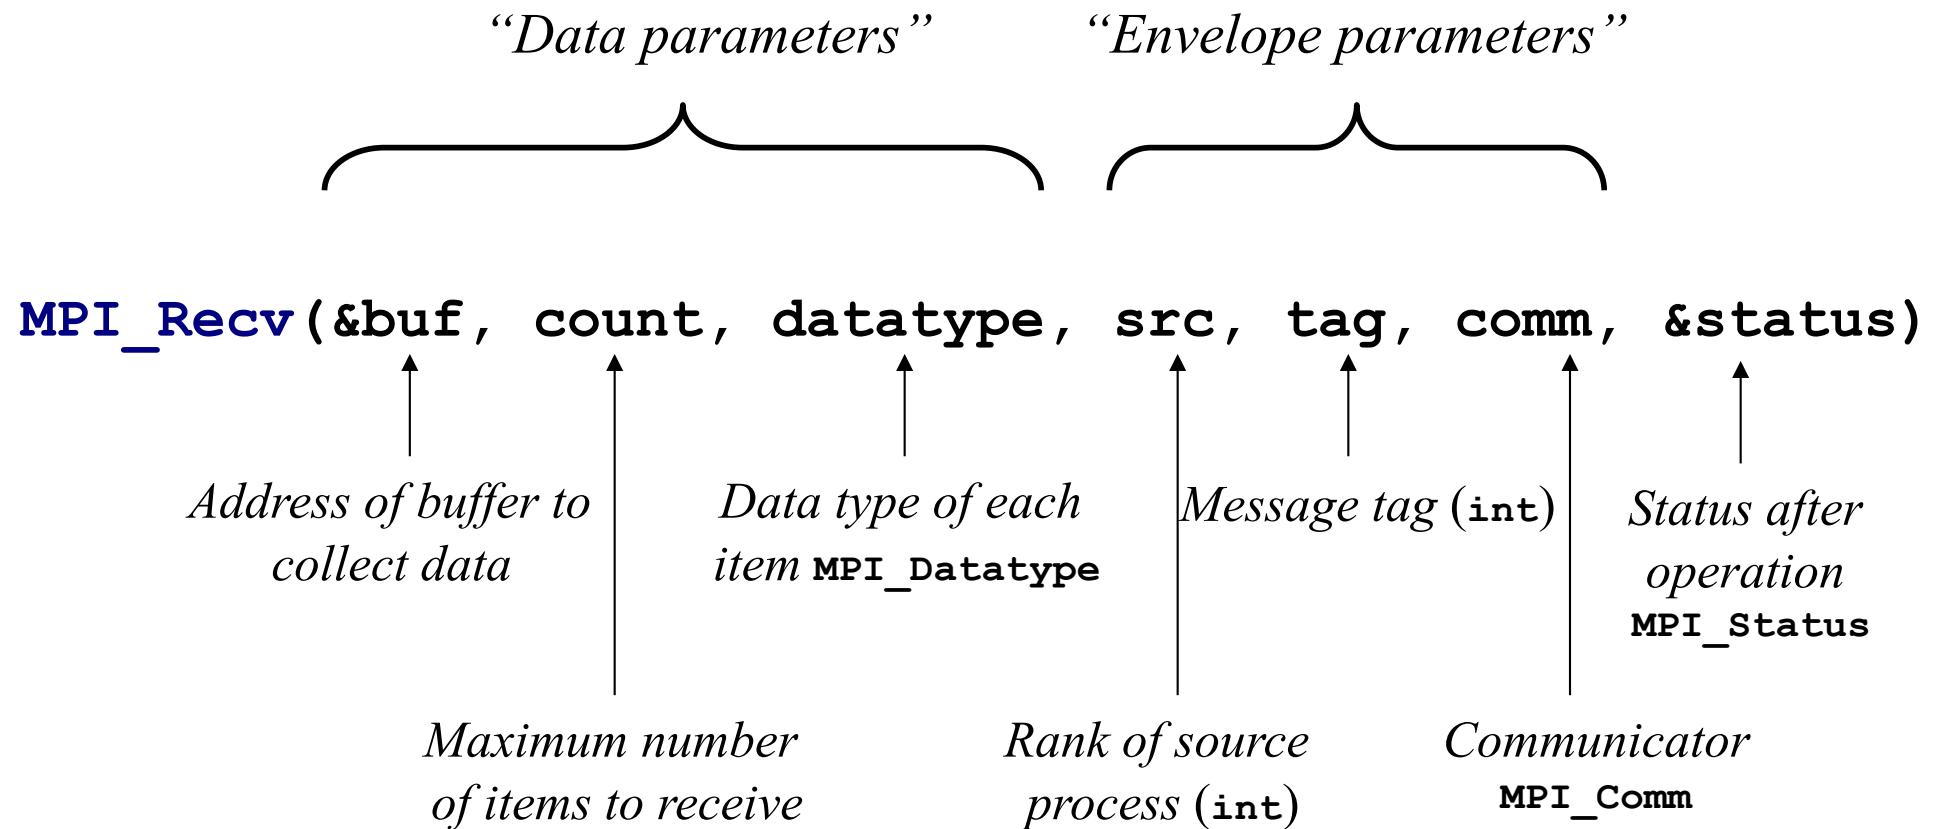


MPI Point-to-point Send Format





MPI Point-to-point Recv Format





Example: Send-Recv Between two Processes

```
main(int argc, char *argv[])
{
    int myrank;
    int value = 123;
    MPI_Status status;

    MPI_Init(&argc, &argv);

    MPI_Comm_Rank(MPI_COMM_WORLD, &myrank);

    if (myrank == 0)
        MPI_Send(&value, 1, MPI_INT, 1, MPI_ANY_TAG, MPI_COMM_WORLD);
    else if (myrank == 1)
        MPI_Recv(&value, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);

    MPI_Finalize();
}
```

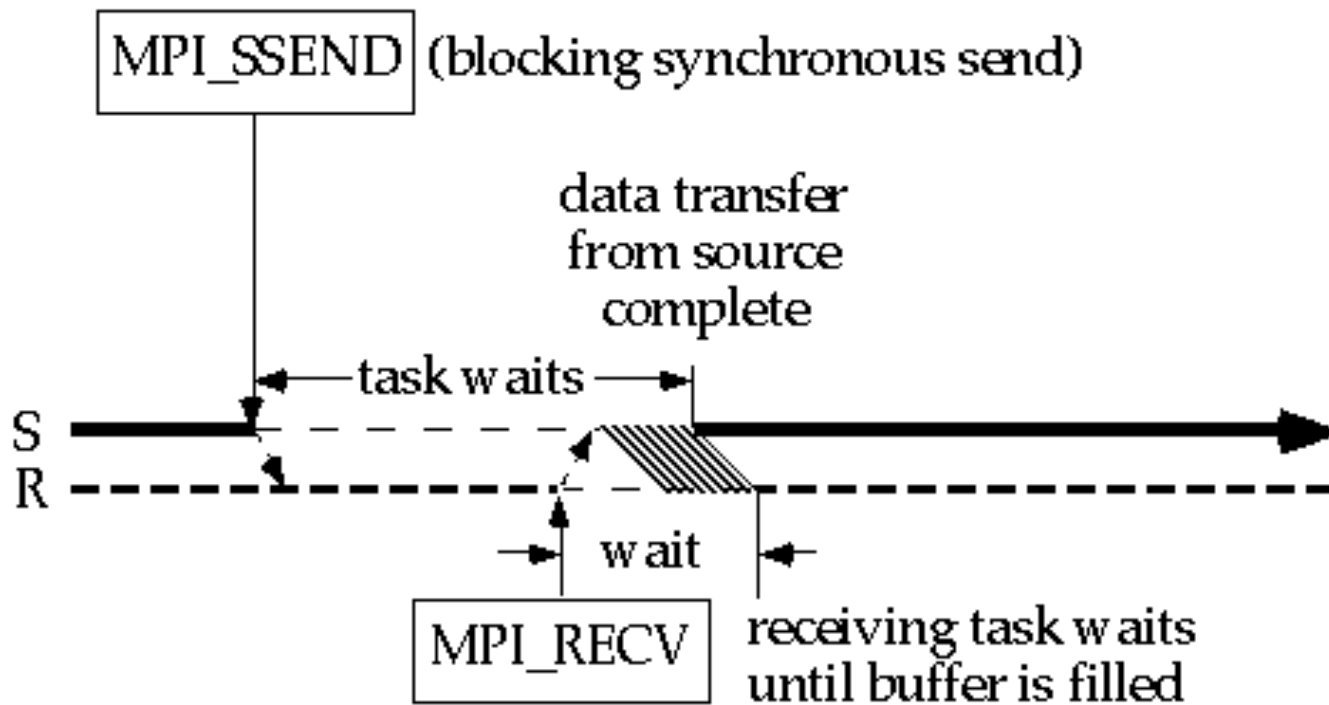



MPI Point-to-point Communication Modes

Communication mode	Blocking routines	Nonblocking routines
<i>Synchronous</i>	<code>MPI_Ssend</code>	<code>MPI_Issend</code>
<i>Ready</i>	<code>MPI_Rsend</code>	<code>MPI_Irsend</code>
<i>Buffered</i>	<code>MPI_Bsend</code>	<code>MPI_Ibsend</code>
<i>Standard</i>	<code>MPI_Send</code>	<code>MPI_Isend</code>
	<code>MPI_Recv</code>	<code>MPI_Irecv</code>
	<code>MPI_Sendrecv</code>	
	<code>MPI_Sendrecv_replace</code>	

Roughly speaking: $\text{MPI_Xsend} = \text{MPI_Ixsend} + \text{MPI_Wait}$
 $\text{MPI_Recv} = \text{MPI_Irecv} + \text{MPI_Wait}$

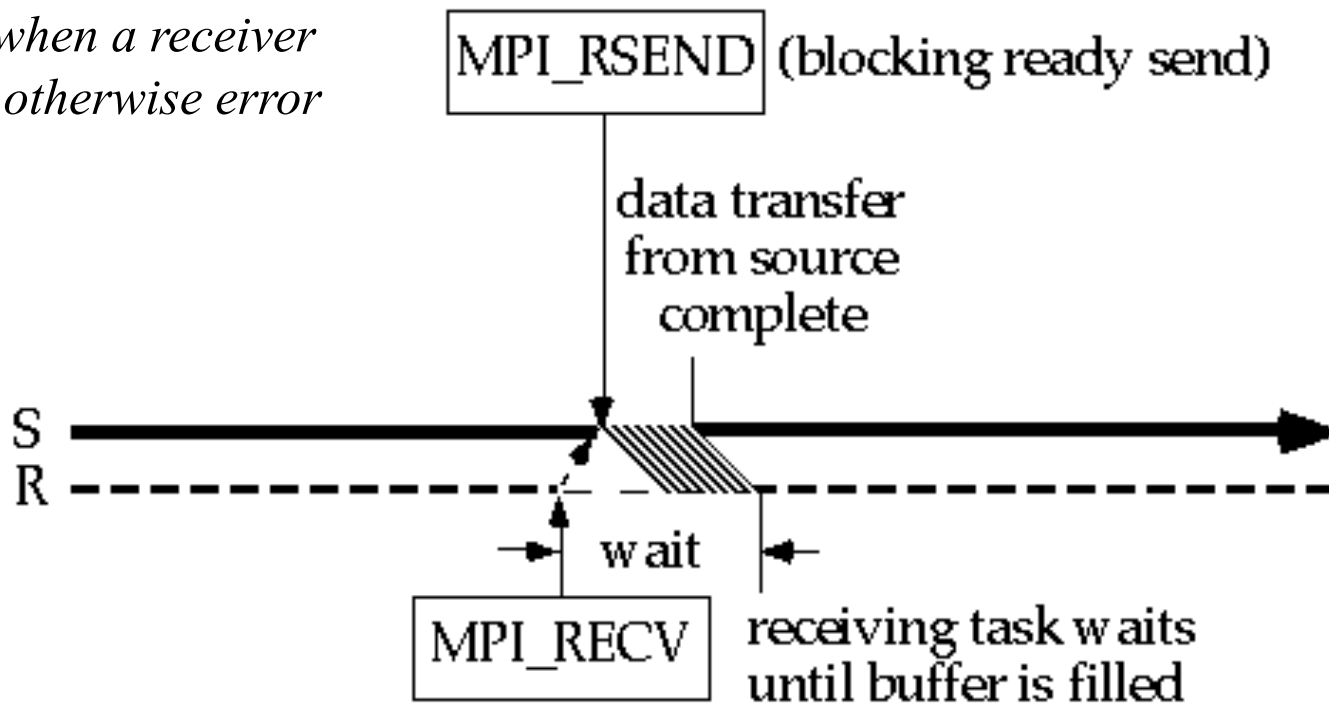
MPI Synchronous Send



```
MPI_Ssend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```

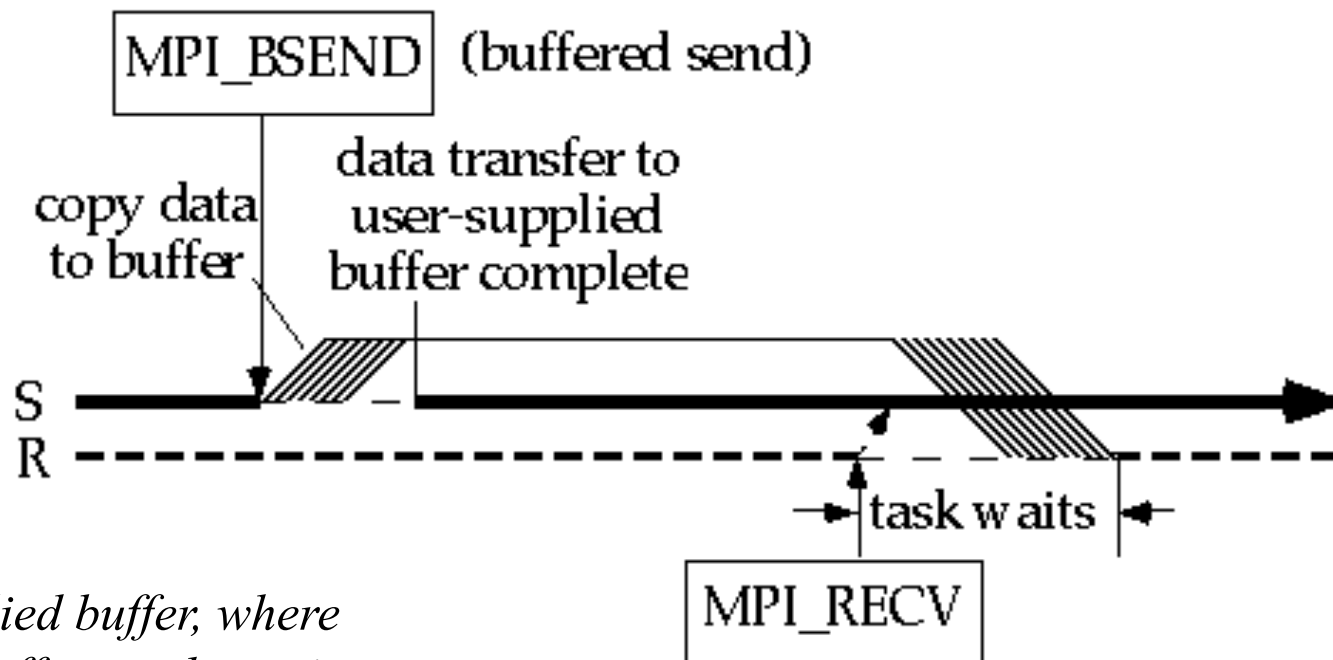
MPI Blocking Ready Send

Only sends when a receiver was posted, otherwise error



```
MPI_Rsend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```

MPI Buffered Send



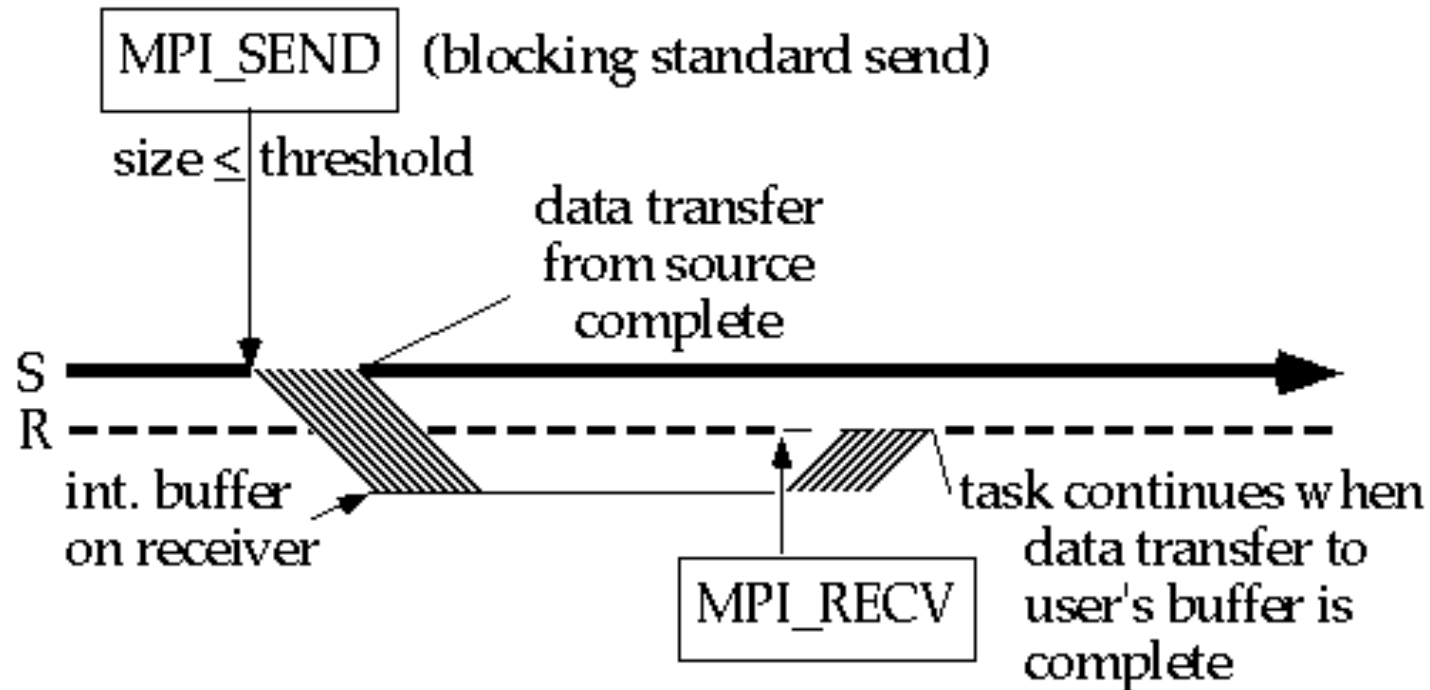
*User-supplied buffer, where
only one buffer can be active*

*Note: buffer size must be the maximum data size + **MPI_BSEND_OVERHEAD** + 7*

```
MPI_Buffer_attach(void *buf, int size)
```

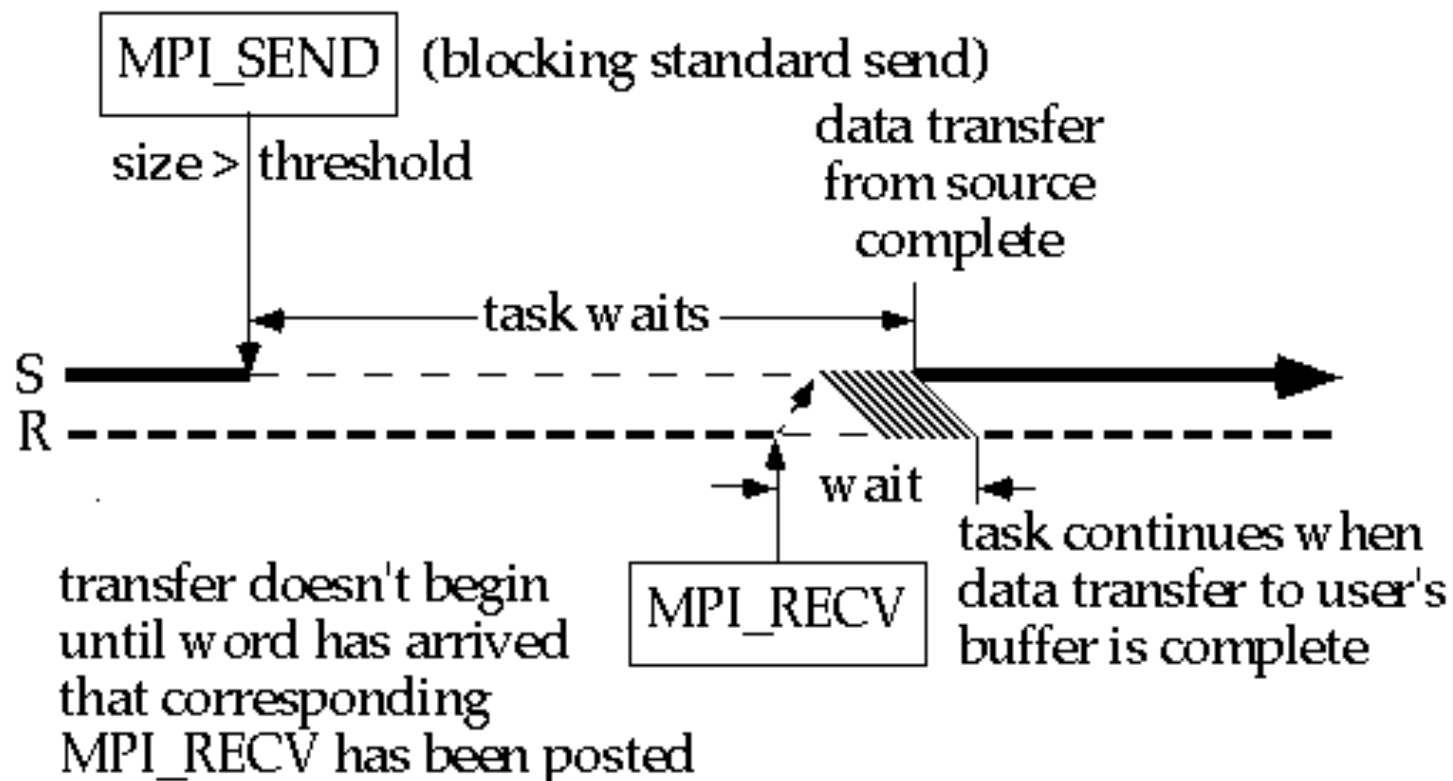
```
MPI_Bsend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```

MPI Blocking Standard Send Small Message Size



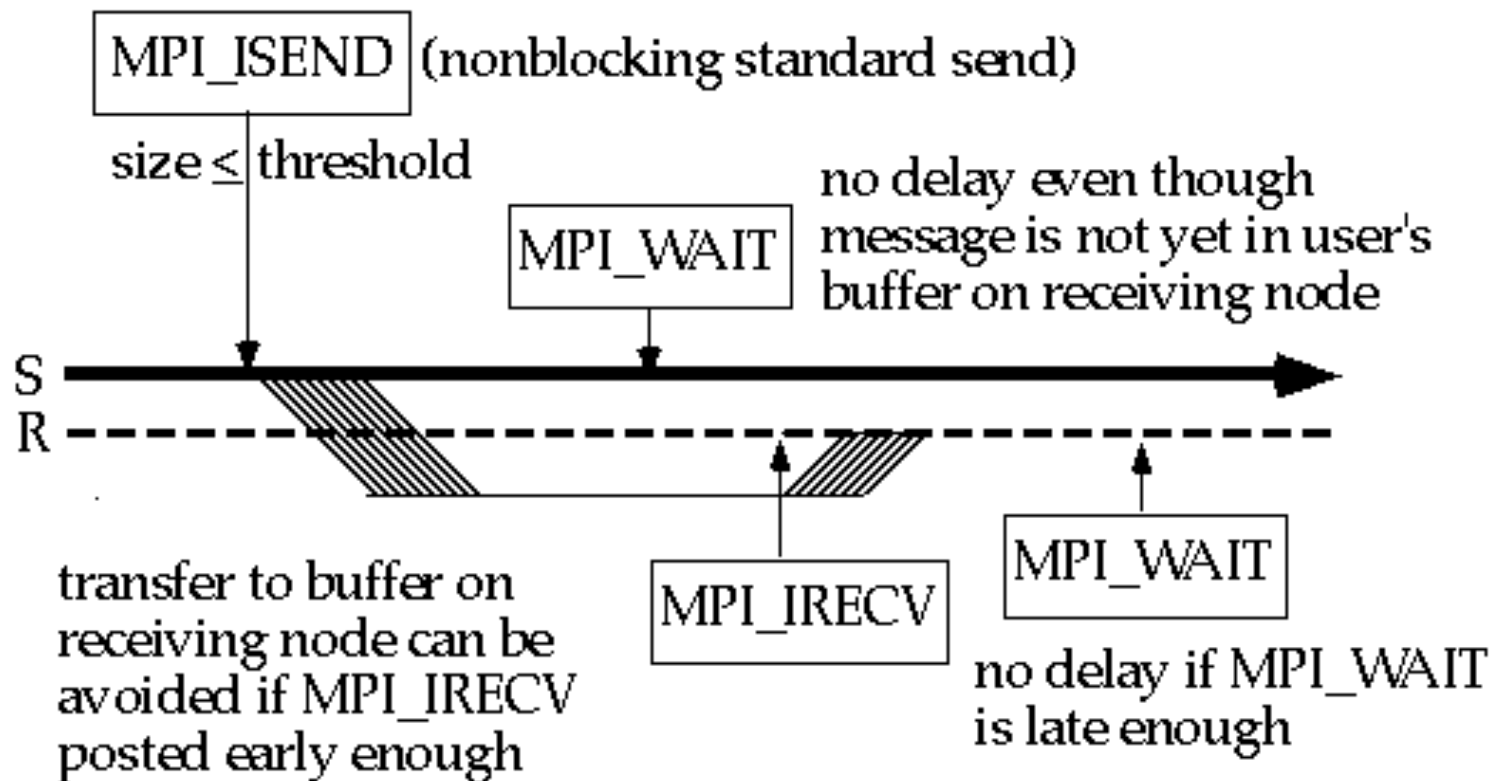
The threshold value (also called the “eager limit”) differs between systems

MPI Blocking Standard Send Large Message Size



The threshold value (also called the “eager limit”) differs between systems

MPI Nonblocking Standard Send Small Message Size

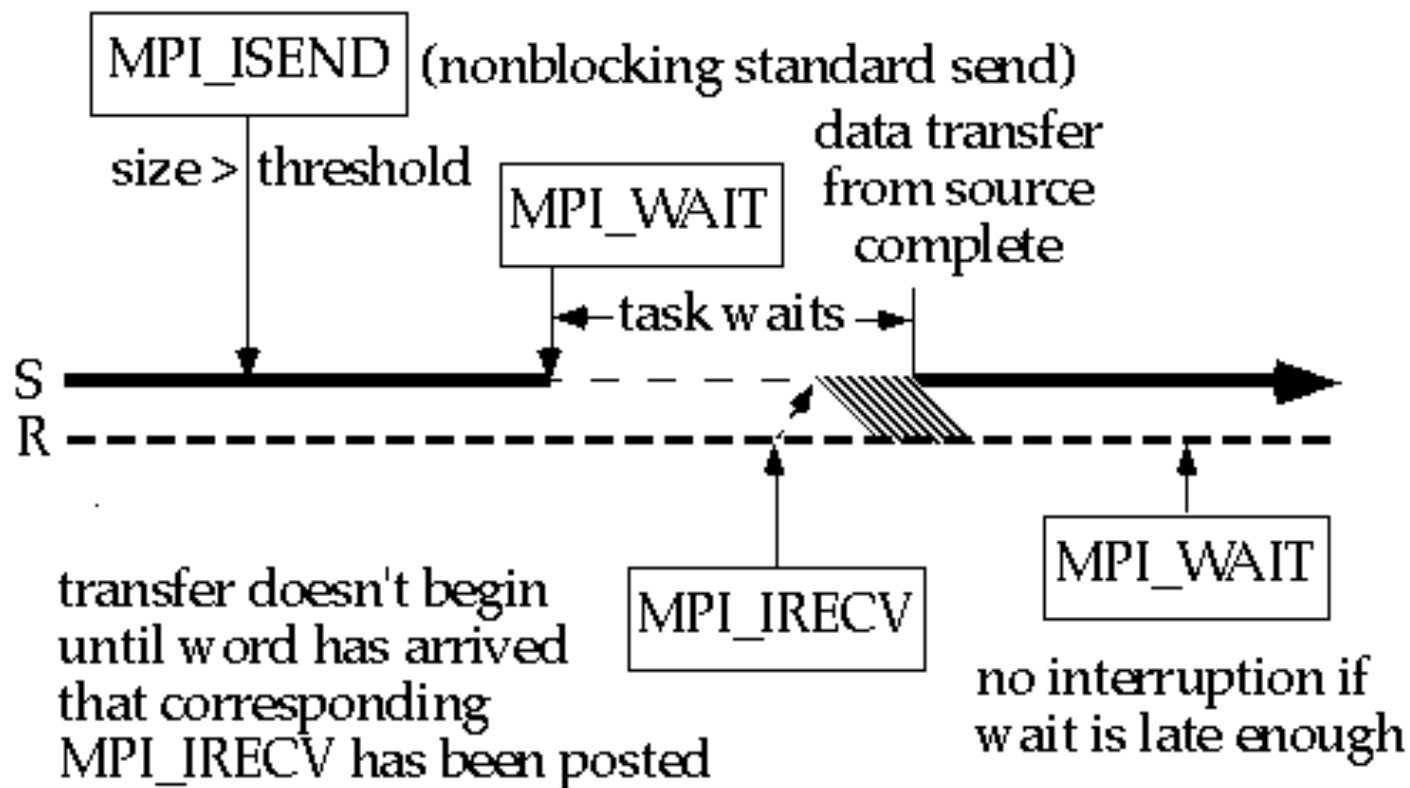


```
MPI_Isend(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Request *request)
```

```
MPI_Wait(MPI_Request *request, MPI_Status *status)
```

MPI Nonblocking Standard Send Large Message Size





MPI Communication Modes

Communication mode	Advantages	Disadvantages
<i>Synchronous</i>	<i>Safest, and therefore most portable SEND/RECV order not critical Amount of buffer space irrelevant</i>	<i>Can incur substantial synchronization overhead</i>
<i>Ready</i>	<i>Lowest total overhead SEND/RECV handshake not required</i>	<i>RECV must precede SEND</i>
<i>Buffered</i>	<i>Decouples SEND from RECV No sync overhead on SEND Order of SEND/RECV irrelevant Programmer can control size of buffer space</i>	<i>Additional system overhead incurred by copy to buffer</i>
<i>Standard</i>	<i>Good for many cases</i>	<i>May not be suitable for your program</i>



Example 1: Nonblocking Send and Blocking Recv

```
main(int argc, char *argv[])
{
    int myrank;
    int value = 123;
    MPI_Status status;
    MPI_Request req;

    MPI_Init(&argc, &argv);

    MPI_Comm_Rank(MPI_COMM_WORLD, &myrank);

    if (myrank == 0)
    { MPI_Isend(&value, 1, MPI_INT, 1, MPI_ANY_TAG, MPI_COMM_WORLD, req);
      doWork(); /* do not modify value */
      MPI_Wait(&req, &status);
    }
    else if (myrank == 1)
        MPI_Recv(&value, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);

    MPI_Finalize();
}
```



Example 2:

Nonblocking Send/Recv

```
main(int argc, char *argv[])
{
    int myrank;
    int value = 123;
    MPI_Status status;
    MPI_Request req;

    MPI_Init(&argc, &argv);
    MPI_Comm_Rank(MPI_COMM_WORLD, &myrank);
    if (myrank == 0)
    { MPI_Isend(&value, 1, MPI_INT, 1, MPI_ANY_TAG, MPI_COMM_WORLD, &req);
      doWork(); /* do not modify value */
      MPI_Wait(&req, &status);
    }
    else if (myrank == 1)
    { MPI_Irecv(&value, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &req);
      doWork(); /* do not read or modify value */
      MPI_Wait(&req, &status);
      /* using value is OK now */
    }
    MPI_Finalize();
}
```



MPI Wait

MPI_Wait(MPI_Request *request, MPI_Status *status)

Waits until pending send/rcv request is completed, sets status

MPI_Waitall(int count, MPI_Request *array_of_requests,
MPI_Status *array_of_statuses)

Wait for all pending send/rcv requests to be completed, sets array of status values

MPI_Waitany(int count, MPI_Request *array_of_requests, int *index,
MPI_Status *status)

Wait until any of the pending send/rcv requests is completed, sets index and status

MPI_Waitsome(int incount, MPI_Request *array_of_requests,
int *outcount, int* array_of_indices,
MPI_Status *array_of_statuses)

Wait until some of the pending send/rcv requests is completed, sets arrays of index and status



MPI Test

MPI_Test(MPI_Request *request, int *flag, MPI_Status *status)

Check if pending send/rcv request is completed (flag=1), or pending (flag=0), sets status

MPI_Testall(int count, MPI_Request *array_of_requests, int *flag, MPI_Status *array_of_statuses)

Check if all pending send/rcv requests are completed (flag=1), sets array of status values

MPI_Testany(int count, MPI_Request *array_of_requests, int *index, int *flag, MPI_Status *status)

Check if any of the pending send/rcv requests is completed (flag=1), sets index and status

MPI_Testsome(int incount, MPI_Request *array_of_requests, int *outcount, int* array_of_indices, MPI_Status *array_of_statuses)

Checks if some of the pending send/rcv requests is completed, sets arrays of index and status



MPI Combined Sendrecv

```
MPI_Sendrecv(void *sendbuf, int sendcount, MPI_Datatype sendtype,  
              int dest, int sendtag,  
              void *recvbuf, int recvcount,  
              MPI_Datatype recvtype, int source, int recvtag,  
              MPI_Comm comm, MPI_Status *status)
```

Combined send and receive operation (blocking), with disjoint send and receive buffers

```
MPI_Sendrecv_replace(void *buf, int count, MPI_Datatype datatype,  
                      int dest, int sendtag, int source, int recvtag,  
                      MPI_Comm comm, MPI_Status *status)
```

Combined send and receive operation (blocking), with shared send and receive buffers



Example 3: Combined Sendrecv

```
main(int argc, char *argv[])
{
    int myrank, hisrank;
    float value;
    MPI_Status status;

    MPI_Init(&argc, &argv);

    MPI_Comm_Rank(MPI_COMM_WORLD, &myrank);
    hisrank = 1-myrank;
    if (myrank == 0)
        value = 3.14;
    else
        value = 1.41;

    MPI_Sendrecv_replace(&value, 1, MPI_FLOAT, hisrank, MPI_ANY_TAG,
                        hisrank, MPI_ANY_TAG, MPI_COMM_WORLD, &status);

    printf("Process %d got value %g\n", myrank, value);

    MPI_Finalize();
}
```



MPI Basic Datatypes

<code>MPI_CHAR</code>	<i>signed char</i>
<code>MPI_SHORT</code>	<i>signed short int</i>
<code>MPI_INT</code>	<i>signed int</i>
<code>MPI_LONG</code>	<i>signed long int</i>
<code>MPI_UNSIGNED_CHAR</code>	<i>unsigned char</i>
<code>MPI_UNSIGNED_SHORT</code>	<i>unsigned short int</i>
<code>MPI_UNSIGNED</code>	<i>unsigned int</i>
<code>MPI_UNSIGNED_LONG</code>	<i>unsigned long int</i>
<code>MPI_FLOAT</code>	<i>float</i>
<code>MPI_DOUBLE</code>	<i>double</i>
<code>MPI_LONG_DOUBLE</code>	<i>long double</i>
<code>MPI_BYTE</code>	
<code>MPI_PACKED</code>	



MPI Collective Communications

- Coordinated communication within a group of processes identified by an MPI communicator
- Collective communication routines block until locally complete
- Amount of data sent must exactly match amount of data specified by receiver
- No message tags are needed



MPI Collective Communications Performance Considerations

- Communications are hidden from user
 - Communication patterns depend on implementation and platform on which MPI runs
 - In some cases, the root process originates or receives all data
 - Performance depends on implementation of MPI
- Communications may, or may not, be synchronized (implementation dependent)
 - Not always the best choice to use collective communication
 - There may be forced synchronization, which can be avoided with nonblocking point-to-point communications



MPI Collective Communications Classes

Three classes:

1. Synchronization
 - ☐ Barrier synchronization
2. Data movement
 - ☐ Broadcast
 - ☐ Scatter
 - ☐ Gather
 - ☐ All-to-all
3. Global computation
 - ☐ Reduce
 - ☐ Scan



MPI Barrier

```
MPI_Barrier(MPI_Comm comm)
```

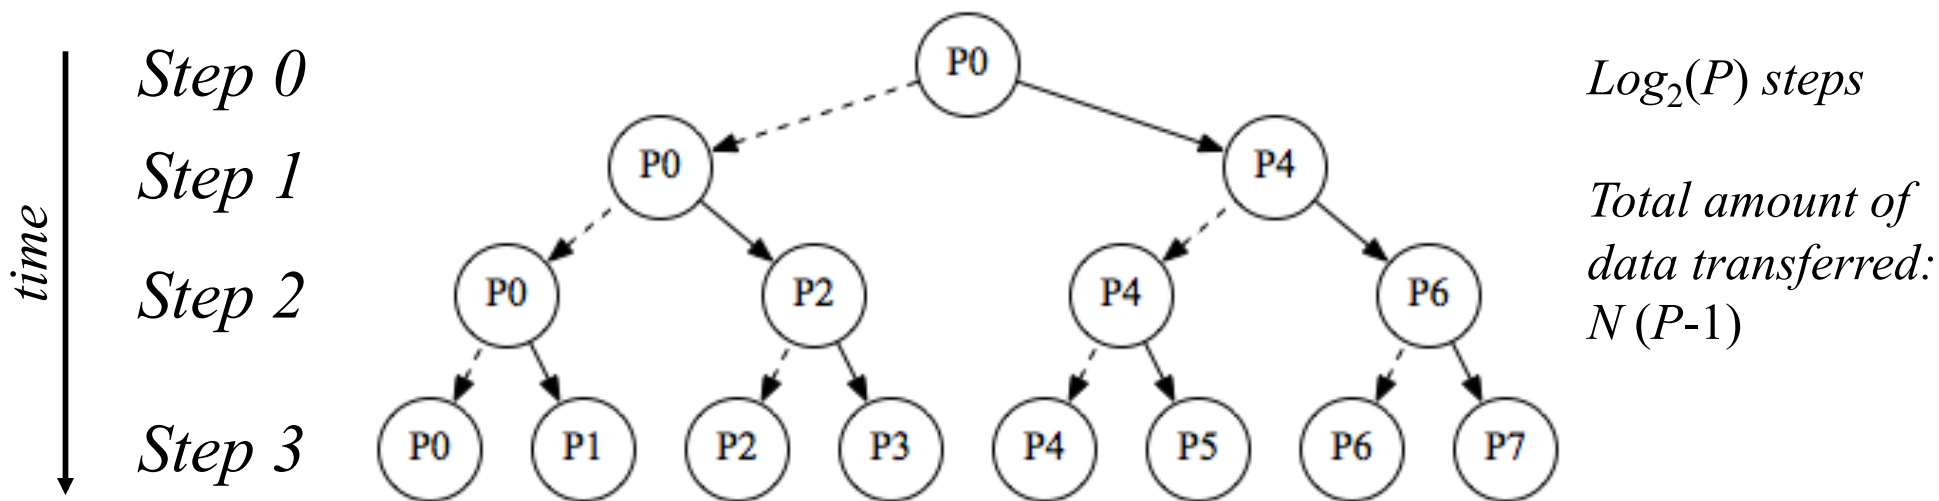
A node invoking the barrier routine will be blocked until all the nodes within the group (communicator) have invoked it

MPI Broadcast


```
MPI_Broadcast(void* buffer,  
              int count,  
              MPI_Datatype datatype,  
              int root,  
              MPI_Comm comm)
```

Simple broadcast implementation: root sends data to all processes, which is efficient on a bus-based parallel machine

More efficient on a network: broadcast as a tree operation




MPI Broadcast Example



```
float A[N][N], Ap[N/P][N], b[N], c[N], cp[N/P];  
...  
root = 0;  
MPI_Broadcast(b, N, MPI_Float, root, MPI_COMM_WORLD);  
...
```

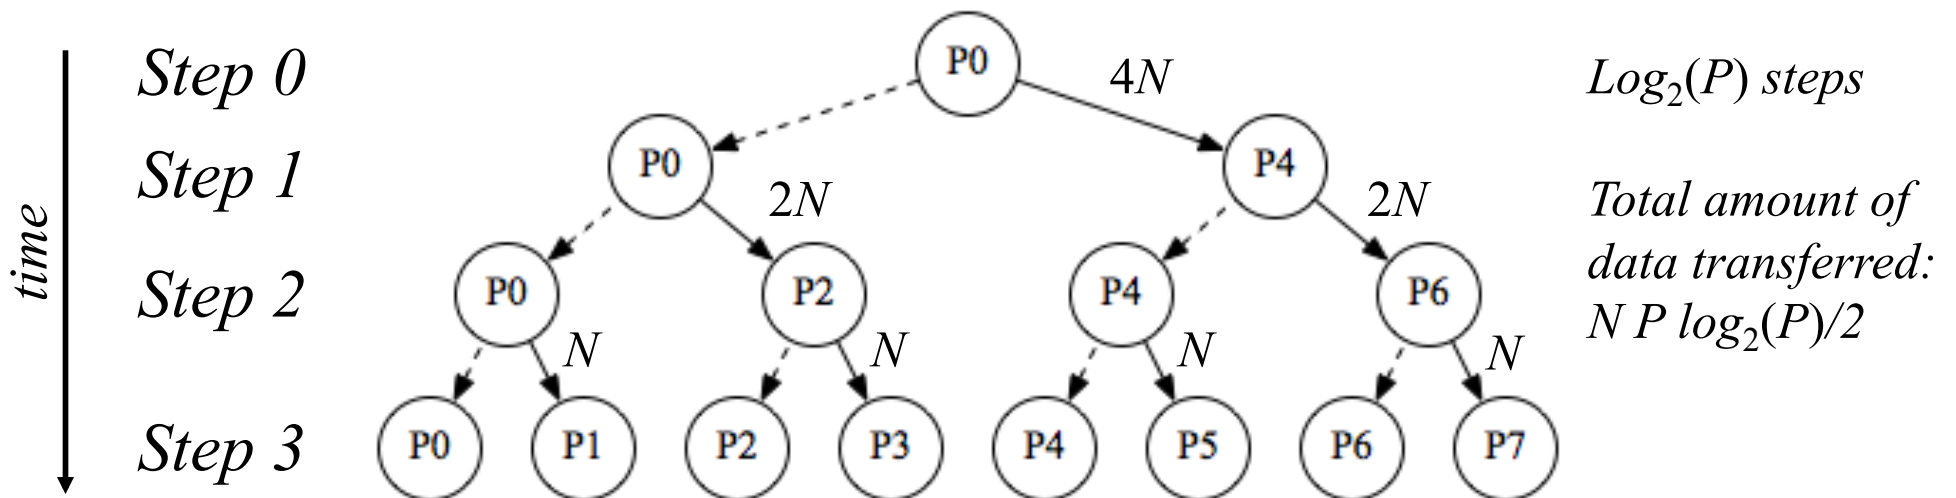
MPI Scatter

```
MPI_Scatter(void *sbuf,  
            int scount,  
            MPI_Datatype stype,  
            void *rbuf,  
            int rcount,  
            MPI_Datatype rtype,  
            int root,  
            MPI_Comm comm)
```




MPI Scatter

```
MPI_Scatter(void *sbuf,  
            int  scount,  
            MPI_Datatype stype,  
            void *rbuf,  
            int  rcount,  
            MPI_Datatype rtype,  
            int  root,  
            MPI_Comm comm)
```



MPI Scatter Example



```
float A[N][N], Ap[N/P][N], b[N], c[N], cp[N/P];
```

```
...
```


```
root = 0;
```

```
...
```


```
MPI_Scatter(A, N/P*N, MPI_Float, Ap, N/P*N, MPI_Float, root,  
            MPI_COMM_WORLD);
```

MPI Gather

```
MPI_Gather(void *sbuf,  
          int  scount,  
          MPI_Datatype stype,  
          void *rbuf,  
          int  rcount,  
          MPI_Datatype rtype,  
          int  root,  
          MPI_Comm comm)
```



MPI Gather Example



```
float A[N][N], Ap[N/P][N], b[N], c[N], cp[N/P];
```

```
...
```

```
for (i = 1; i < N/P; i++)
```

```
{
```

```
    cp[i] = 0;
```

```
    for (k = 0; k < N; k++)
```

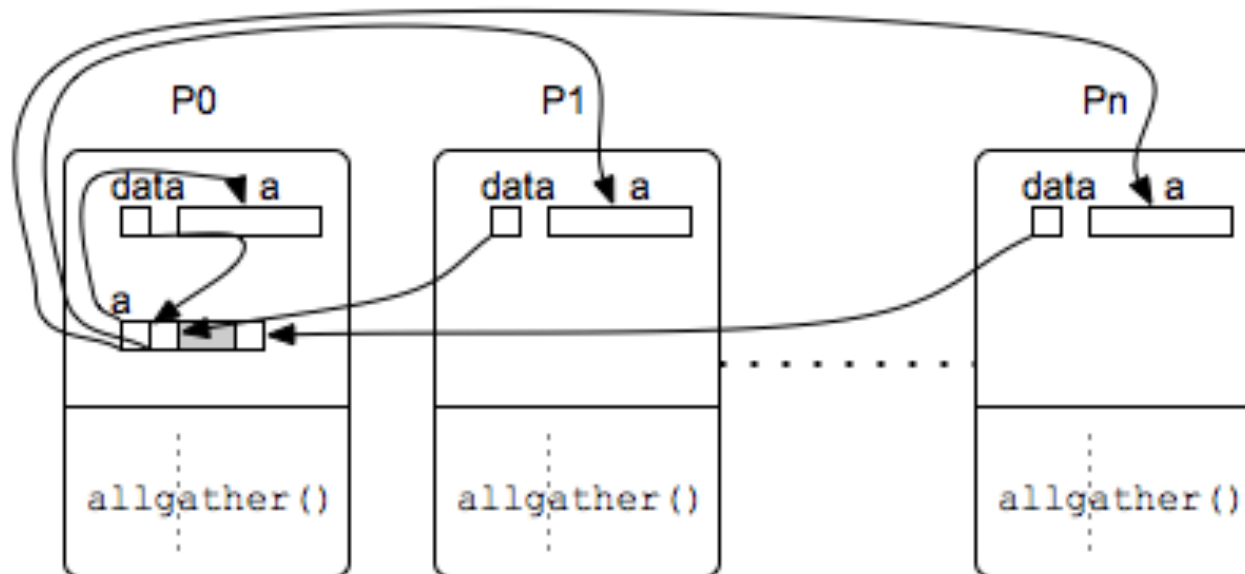
```
        cp[i] = cp[i] + Ap[i][k] * b[k];
```

```
}
```

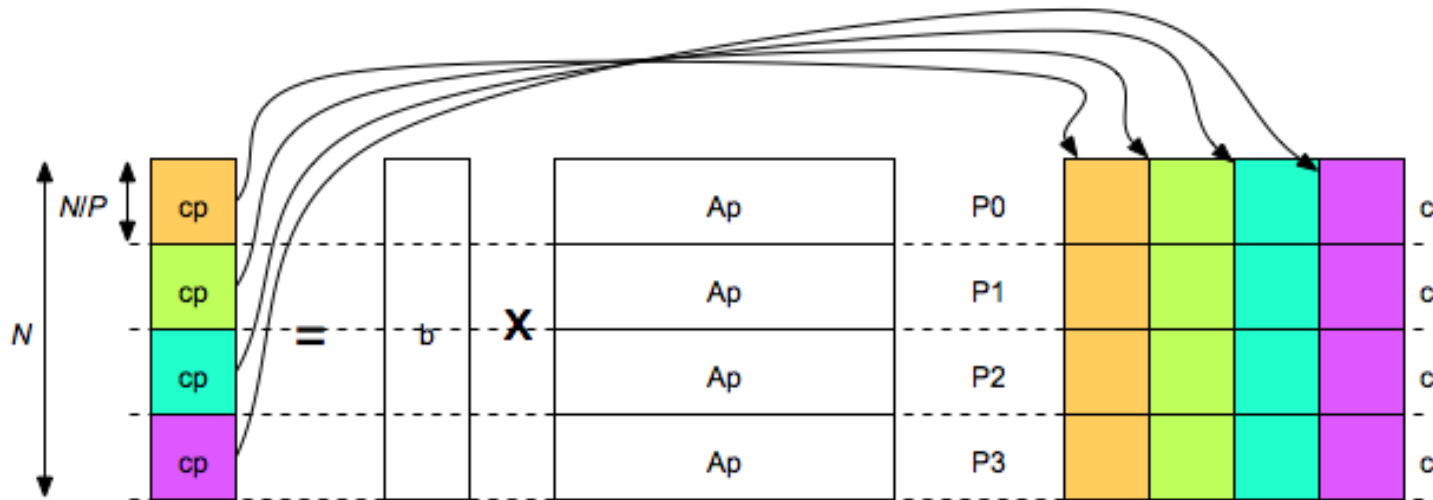
```
MPI_Gather(cp, N/P, MPI_Float, c, N/P, MPI_Float, root,  
          MPI_COMM_WORLD);
```

MPI AllGather

```
MPI_AllGather(void *sbuf,  
             int scount,  
             MPI_Datatype stype,  
             void *rbuf,  
             int rcount,  
             MPI_Datatype rtype,  
             MPI_Comm comm)
```



MPI AllGather Example



```
float A[N][N], Ap[N/P][N], b[N], c[N], cp[N/P];
```

```
...
```

```
for (i = 1; i < N/P; i++)
```

```
{
```

```
    cp[i] = 0;
```

```
    for (k = 0; k < N; k++)
```

```
        cp[i] = cp[i] + Ap[i][k] * b[k];
```

```
}
```

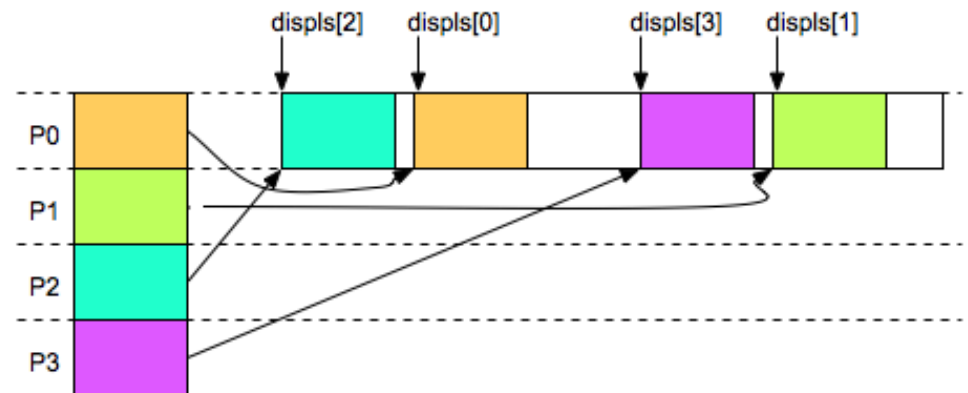
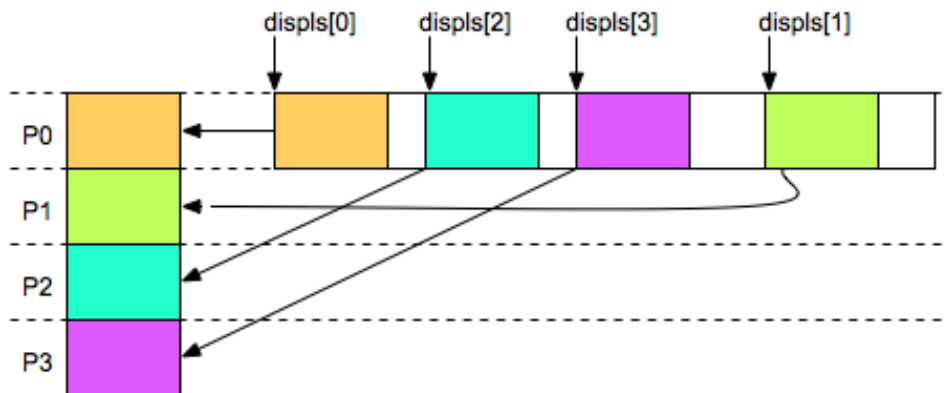
```
MPI_AllGather(cp, N/P, MPI_Float, c, N/P, MPI_Float,  
             MPI_COMM_WORLD);
```

MPI Scatterv and (All)Gatherv

```
MPI_Scatterv(void *sbuf,  
            int *scounts,  
            int *displs,  
            MPI_Datatype stype,  
            void *rbuf,  
            int rcount,  
            MPI_Datatype rtype,  
            int root,  
            MPI_Comm comm)
```

```
MPI_Gatherv(void *sbuf,  
            int scount,  
            MPI_Datatype stype,  
            void *rbuf,  
            int *rcounts,  
            int *displs,  
            MPI_Datatype rtype,  
            int root,  
            MPI_Comm comm)
```

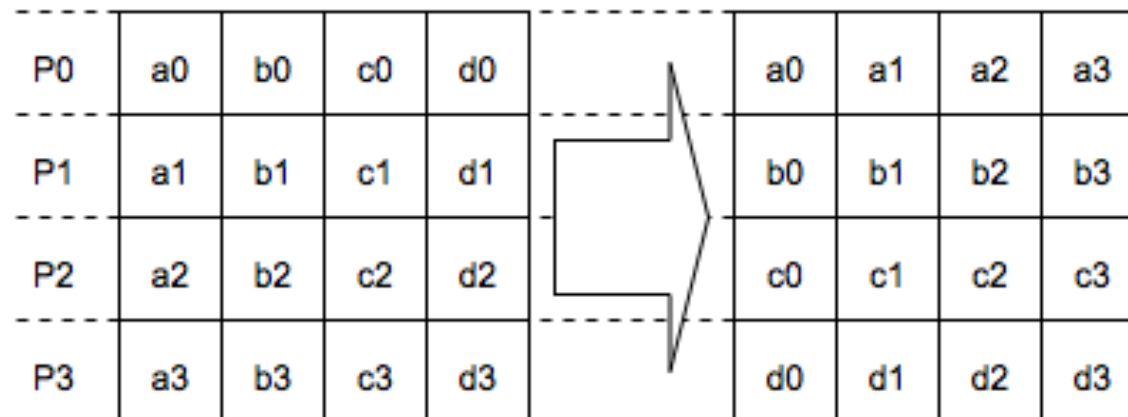
Adds displacements to scatter/gather operations and counts may vary per process



MPI All to All


```
MPI_Alltoall(void *sbuf,  
             int scount,  
             MPI_Datatype stype,  
             void *rbuf,  
             int rcount,  
             MPI_Datatype rtype,  
             MPI_Comm comm)
```

Global transpose: the j th block from processor i is received by processor j and stored in i th block



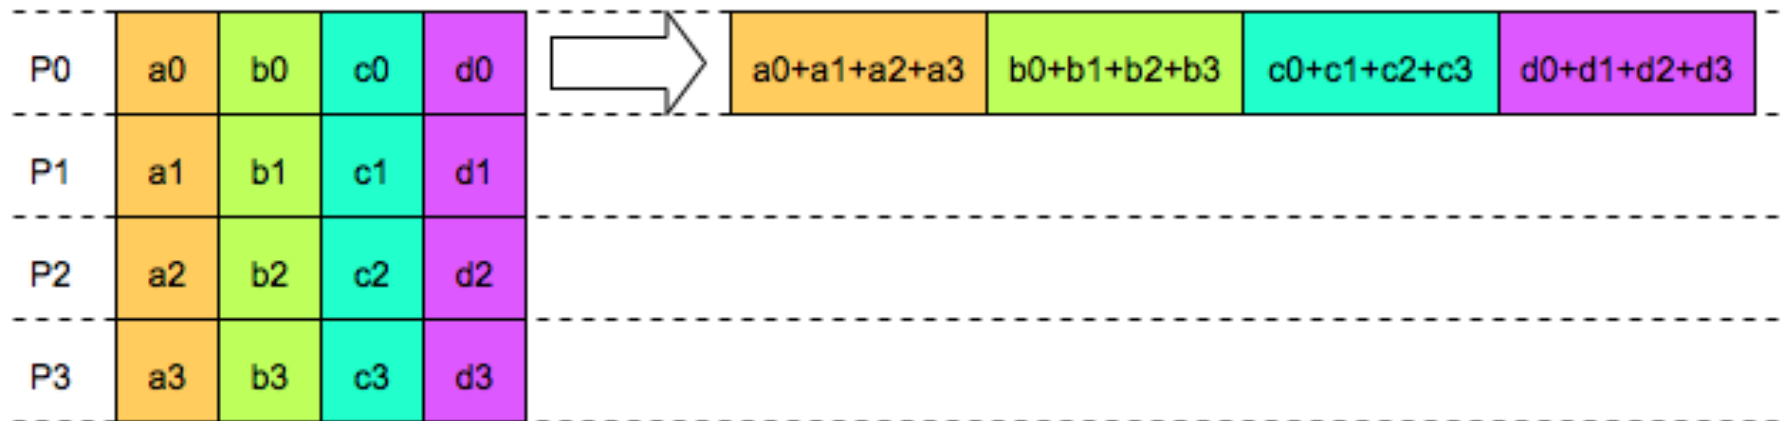
MPI Reduce

```
MPI_Reduce(void *sbuf,  
           void *rbuf,  
           int count,  
           MPI_Datatype stype,  
           MPI_Op op,  
           int root,  
           MPI_Comm comm)
```





MPI Reduce Example




```
float abcd[4], sum[4];
```

```
...
```


```
MPI_Reduce(abcd, sum, 4, MPI_Float, root, MPI_SUM,  
           MPI_COMM_WORLD);
```

MPI AllReduce

```
MPI_AllReduce(void *sbuf,  
              void *rbuf,  
              int count,  
              MPI_Datatype stype,  
              MPI_Op op,  
              MPI_Comm comm)
```



MPI AllReduce Example



```
float abcd[4], sum[4];
```

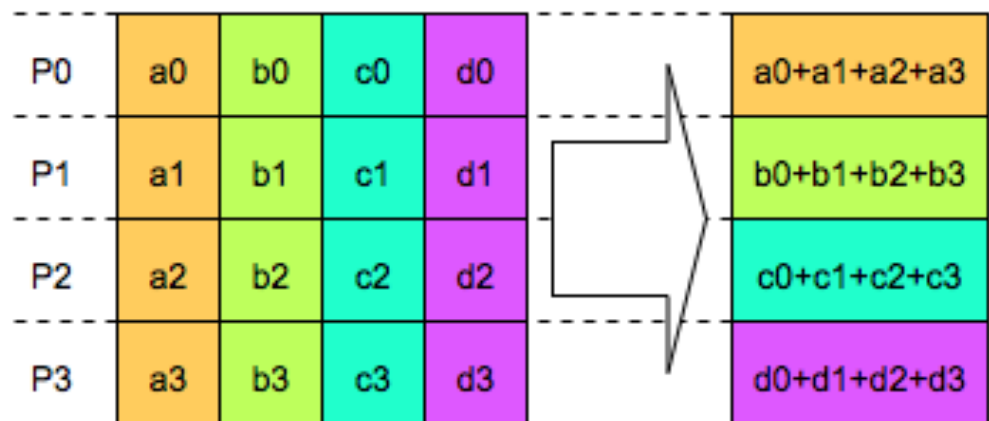
```
...
```

```
MPI_AllReduce(abcd, sum, 4, MPI_Float, MPI_SUM,  
               MPI_COMM_WORLD);
```

MPI Reduce_scatter

```
MPI_Reduce_scatter(void *sbuf,  
                  void *rbuf,  
                  int *rcounts,  
                  MPI_Datatype stype,  
                  MPI_Op op,  
                  MPI_Comm comm)
```

Same as Reduce followed by Scatter

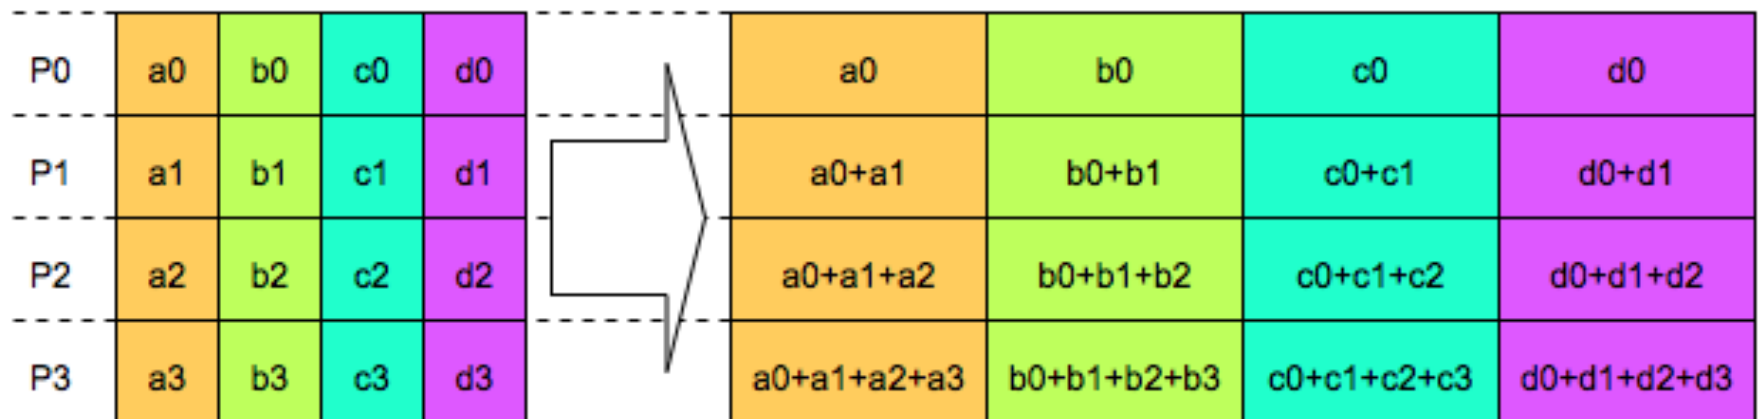


Note: rcounts = number of elements received, which is >1 when $N > P$



MPI Scan

```
MPI_Scan(void *sbuf,  
         void *rbuf,  
         int count,  
         MPI_Datatype stype,  
         MPI_Op op,  
         MPI_Comm comm)
```





MPI Reduce and Scan Reduction Operators

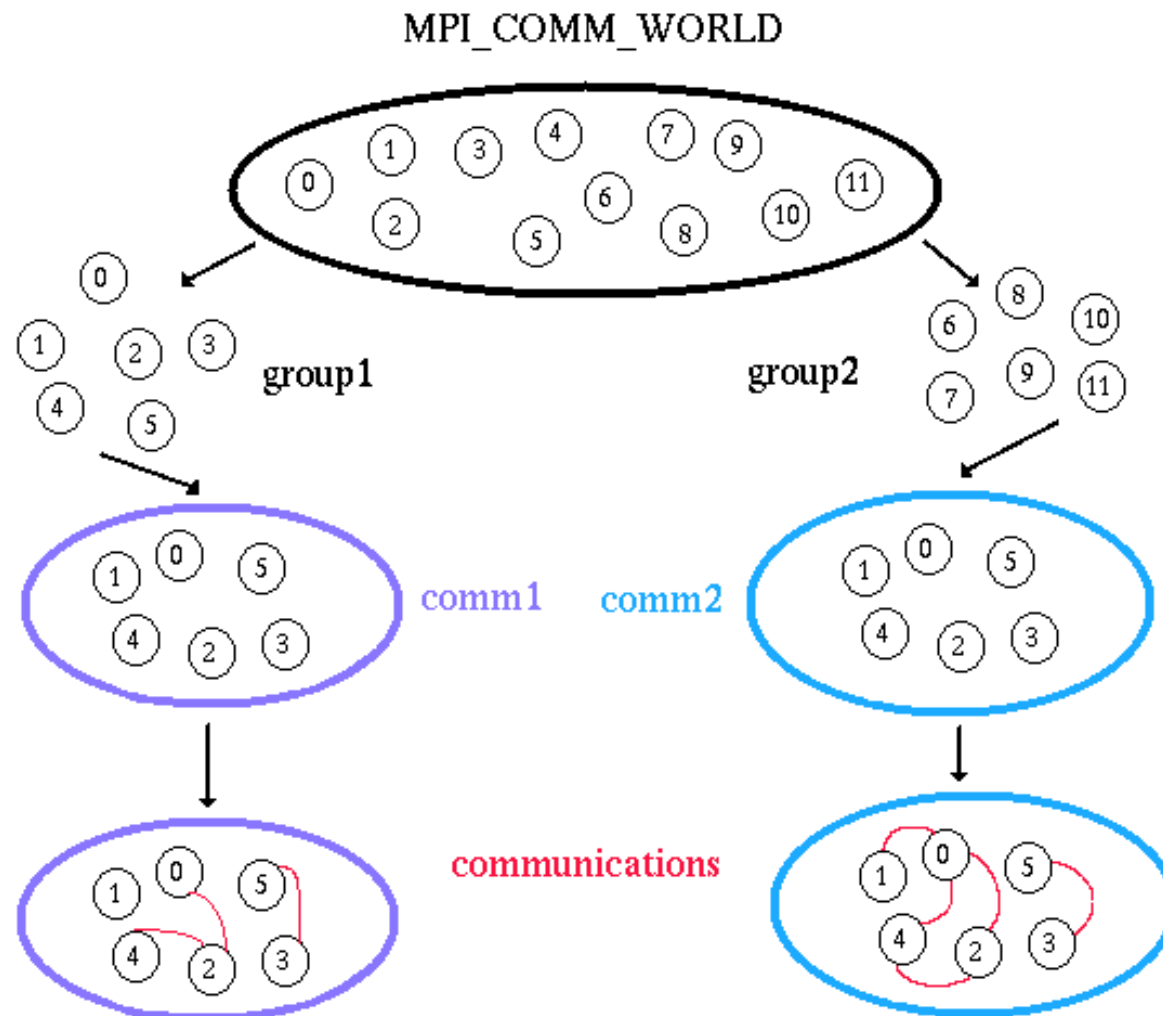
MPI_OP	Operation	C	Fortran
MPI_MAX	maximum	integer, float	integer, real, complex
MPI_MIN	minimum	integer, float	integer, real, complex
MPI_SUM	sum	integer, float	integer, real, complex
MPI_PROD	product	integer, float	integer, real, complex
MPI_LAND	logical and	integer	logical
MPI_BAND	bit-wise and	integer, MPI_BYTE	integer, MPI_BYTE
MPI_LOR	logical or	integer	logical
MPI_BOR	bit-wise or	integer, MPI_BYTE	integer, MPI_BYTE
MPI_LXOR	logical xor	integer	logical
MPI_BXOR	bit-wise xor	integer, MPI_BYTE	integer, MPI_BYTE
MPI_MAXLOC	max val and loc	float, double, long double	real, complex, double precision
MPI_MINLOC	min val and loc	float, double, long double	real, complex, double precision



MPI Groups and Communicators

- A group is an ordered set of processes
 - Each process in a group is associated with a unique integer rank between 0 and $P-1$, with P the number of processes in the group
- A communicator encompasses a group of processes that may communicate with each other
 - Communicators can be created for specific groups
 - Processes may be in more than one group/communicator
- Groups/communicators are dynamic and can be setup and removed at any time
- From the programmer's perspective, a group and a communicator are the same

MPI Groups and Communicators





MPI Group Operations

MPI_Comm_group

returns the group associated with a communicator

MPI_Group_union

creates a group by combining two groups

MPI_Group_intersection

creates a group from the intersection of two groups

MPI_Group_difference

creates a group from the difference between two groups

MPI_Group_incl

creates a group from listed members of an existing group

MPI_Group_excl

creates a group excluding listed members of an existing group

MPI_Group_range_incl

creates a group according to first rank, stride, last rank

MPI_Group_range_excl

creates a group by deleting according to first rank, stride, last rank

MPI_Group_free

marks a group for deallocation



MPI Communicator Operations

MPI_Comm_size

returns number of processes in communicator's group

MPI_Comm_rank

returns rank of calling process in communicator's group

MPI_Comm_compare

compares two communicators

MPI_Comm_dup

duplicates a communicator

MPI_Comm_create

creates a new communicator for a group

MPI_Comm_split

splits a communicator into multiple, non-overlapping communicators

MPI_Comm_free

marks a communicator for deallocation



MPI Virtual Topologies

- A *virtual topology* can be associated with a communicator
- Two types of topologies supported by MPI
 - Cartesian (grid)
 - Graph
- Increases efficiency of communications
 - Some MPI implementations may use the physical characteristics of a given parallel machine
 - Introduces locality: low communication overhead with nodes that are “near” (few message hops), while distant nodes impose communication penalties (many hops)



MPE and Jumpshot

MPE_Log_get_state_eventIDs(int *startID, int *finalID)

Get a pair of event numbers to describe a state (see next)

MPE_Describe_state(int startID, int finalID, const char *name
const char *color)

Describe the state with a name and color (e.g. “red”, “blue”, “green”)

MPE_Log_get_solo_eventID(int *eventID)

Get an event number to describe an event (see next)

MPE_Describe_event(int eventID, const char *name, const char *color)

Describe the event with a name and color (e.g. “red”, “blue”, “green”)

MPE_Log_event(int event, int data, const char *bytebuf)

Record event in the log, data is unused and bytebuf carries informational data are should be NULL.

Use two calls, one with the startID and the other with the finalID to log the time of a state.



MPE and Jumpshot

```
int event1, state1s, state1f;
int myrank;

MPI_Comm_rank(MPI_COMM_WORLD, &myrank);

MPE_Log_get_solo_eventIDs(&event1);
MPE_Log_get_state_eventIDs(&state1s, &state1f);

if (myrank == 0)
{ MPE_Describe_event(event1, "Start", "green");
  MPE_Describe_state(state1s, state1f, "Computing", "red");
}

MPE_Log_event(event1, 0, NULL);
MPI_Bcast(...);
...
MPE_Log_event(state1s, 0, NULL);
doComputation();
MPE_Log_event(state1f, 0, NULL);
MPI_Barrier();
```



MPI+MPE Compilation, Linking, and Run

- Compilation with MPE requires `mpe` and `lmpe` libs:

- ☐ `mpicc -o myprog myprog.c -lmpe -llmpe`
- ☐ Note: without `-llmpe` you must init and finalize the logs

- Run directly...

- ☐ `mpirun -np 4 ./myprog`

- ... or in batch mode (e.g. with SGE)


- ☐ `qsub runmyprog.sh`

- This generates the log file (typically in home dir)

- ☐ `myprog.clog2`

- Display log file with

- ☐ `jumpshot`

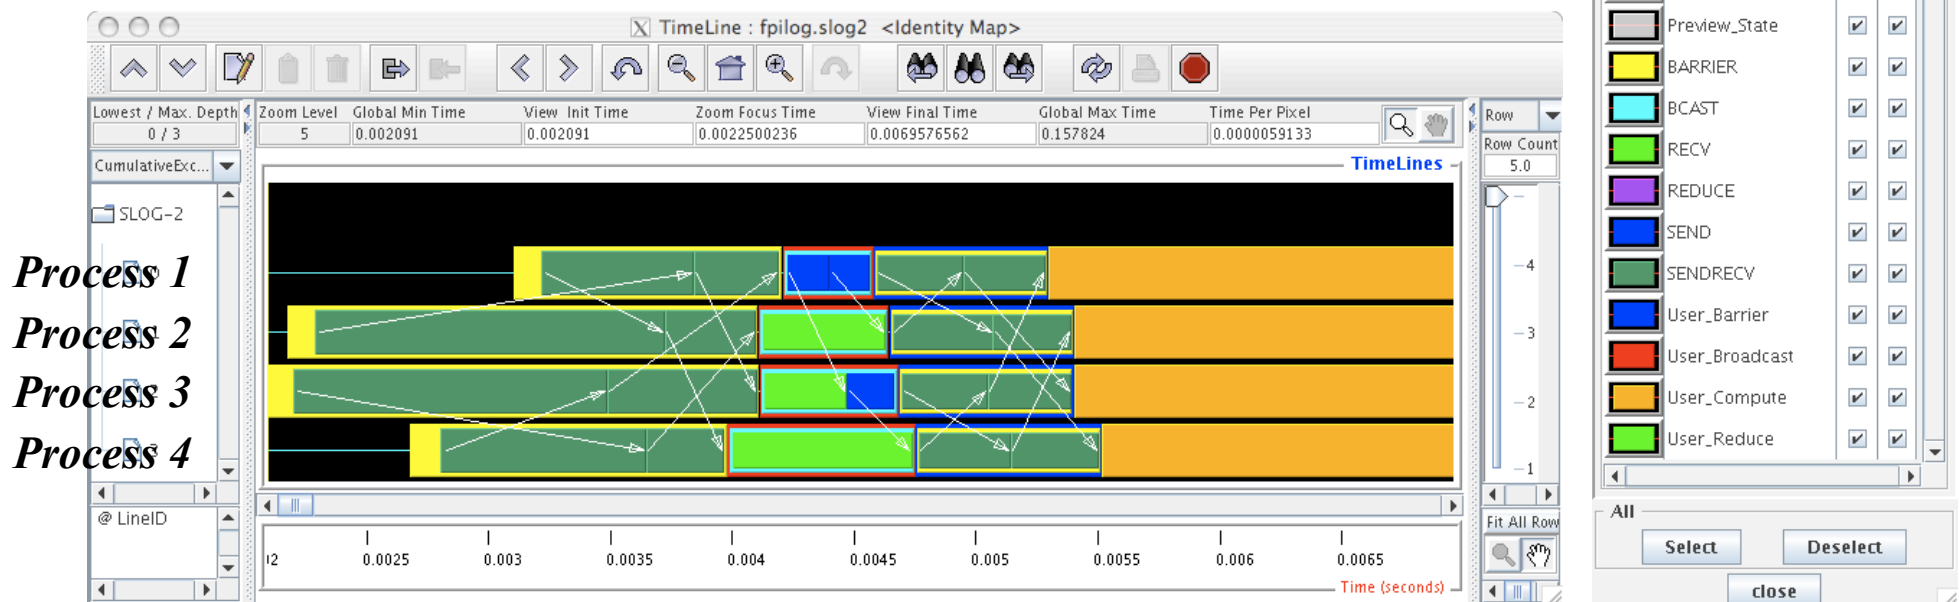


```
#!/bin/bash
# All environment variables active within the qsub
# utility to be exported to the context of the job:
#$ -v
# use openmpi with 4 nodes
#$ -pe openmpi_4 4
# use current directory
#$ -cwd
mpirun -np $NSLOTS ./myprog
```

Jumpshot

- Jumpshot generates a space-time diagram from a log file
 - Supports clog2 (through conversion) and newer slog2 formats
 - Shows user-defined states (start-end) and single point events
 - www.cs.fsu.edu/~engelen/courses/HPC/cpilog.c
 - www.cs.fsu.edu/~engelen/courses/HPC/fpilog.f

Space-time diagram





Further Reading

- [PP2] pages 52-61
- Optional:
Lawrence Livermore National Laboratories MPI Tutorial
<http://www.llnl.gov/computing/tutorials/mpi/>