

# Programming with Shared Memory PART I

**HPC Fall 2012**

*Prof. Robert van Engelen*







# Overview

- Shared memory machines
- Programming strategies for shared memory machines
- Allocating shared data for IPC
- Processes and threads
- MT-safety issues
- Coordinated access to shared data
  - Locks
  - Semaphores
  - Condition variables
  - Barriers
- Further reading


# Shared Memory Machines



*Shared memory UMA machine with a single bus*




*Shared memory NUMA machine with memory banks*



**DSM**

- *Single address space*
- *Shared memory*
  - *Single bus (UMA)*
  - *Interconnect with memory banks (NUMA)*
  - *Cross-bar switch*
- *Distributed shared memory (DSM)*
  - *Logically shared, physically distributed*

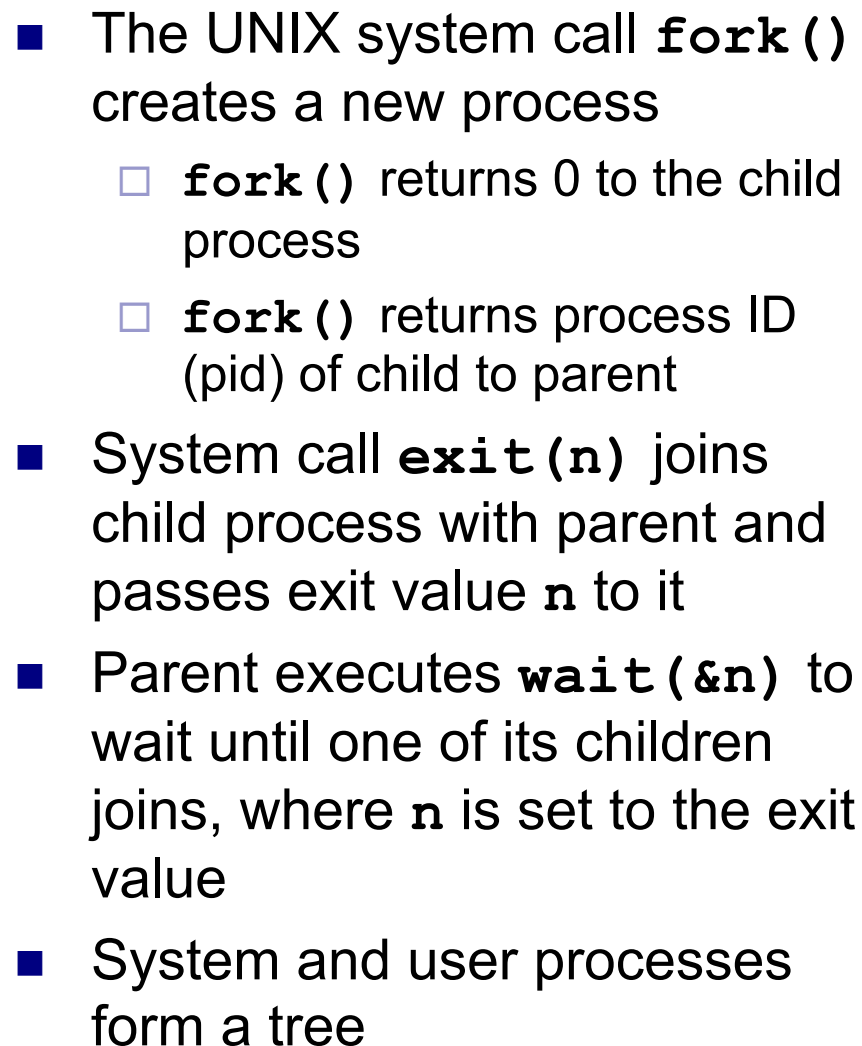


*Shared memory multiprocessor with cross-bar switch*



# Programming Strategies for Shared Memory Machines

- Use a *specialized programming language* for parallel computing
  - For example: HPF, UPC
- Use *compiler directives* to supplement a sequential program with parallel directives
  - For example: OpenMP
- Use *libraries*
  - For example: ScaLapack (though ScaLapack is primarily designed for distributed memory)
- Use *heavyweight processes and a shared memory API*
- Use *threads*
- Use a *parallelizing compiler* to transform (part of) a sequential program into a parallel program





# Fork-Join

*Process 1*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*SPMD program*



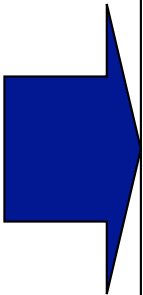
# Fork-Join

*Process 1*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*SPMD program*

*Process 2*



```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*Copy of program, data, and file descriptors  
(operations by the processes on open files  
will be independent)*



# Fork-Join

*Process 1*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*SPMD program*

*Process 2*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*Copy of program and data*





# Fork-Join

*Process 1*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*SPMD program*

*Process 2*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*Copy of program and data*



# Fork-Join

*Process 1*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*SPMD program*

*Process 2*

```
...  
...  
pid = fork();  
if (pid == 0)  
{ ... // code for child  
  exit(0);  
} else  
{ ... // parent code continues  
  wait(&n); // join  
}  
... // parent code continues  
...
```

*Terminated*



# Creating Shared Data for IPC

---

## **shmget ()**

returns the shared memory identifier for a given key (key is for naming and locking)

## **shmat ()**

attaches the segment identified by a shared memory identifier and returns the address of the memory segment

## **shmctl ()**

deletes the segment with **IPC\_RMID** argument

---

## **mmap ()**

returns the address of a mapped object described by the file id returned by **open ()**

## **munmap ()**

deletes the mapping for a given address

---

- *Interprocess communication (IPC) via shared data*
- Processes do not automatically share data
- Use files to share data
  - Slow, but portable
- Unix system V **shmget ()**
  - Allocates shared pages between two or more processes
- BSD Unix **mmap ()**
  - Uses file-memory mapping to create shared data in memory
  - Based on the principle that files are shared between processes



# shmget vs mmap

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

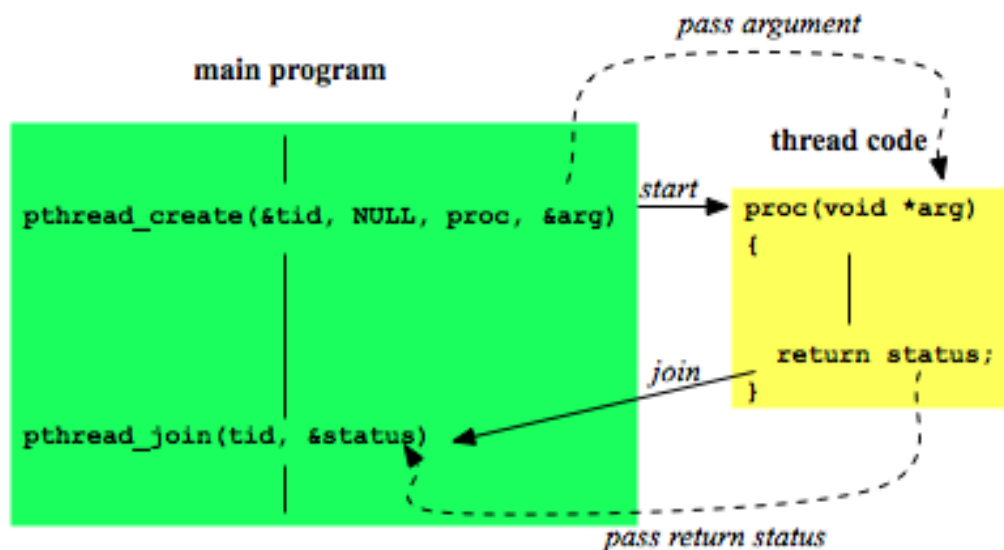
size_t len; // size of data we want
void *buf; // to point to shared data
int shmidx;
key_t key = 9876; // or IPC_PRIVATE
shmidx = shmget(key,
                len,
                IPC_CREAT|0666);
if (shmidx == -1) ... // error
buf = shmat(shmidx, NULL, 0);
if (buf == (void*)-1) ... // error
...
fork(); // parent and child use buf
...
wait(&n);
shmctl(shmidx, IPC_RMID, NULL);
```

```
#include <sys/types.h>
#include <sys/mman.h>

size_t len; // size of data we want
void *buf; // to point to shared data
buf = mmap(NULL,
            len,
            PROT_READ|PROT_WRITE,
            MAP_SHARED|MAP_ANON,
            -1, // fd=-1 is unnamed
            0);
if (buf == MAP_FAILED) ... // error
...
fork(); // parent and child use buf
...
wait(&n);
munmap(buf, len);
...
```

Tip: use **ipcs** command to display  
IPC shared memory status of a system

# Threads



## *Thread creation and join*

- *Threads of control* operate in the same memory space, sharing code and data
  - Data is implicitly shared
  - Consider data on a thread's stack private
- Many OS-specific thread APIs
  - Windows threads, Linux threads, Java threads, ...
- POSIX-compliant Pthreads:

---

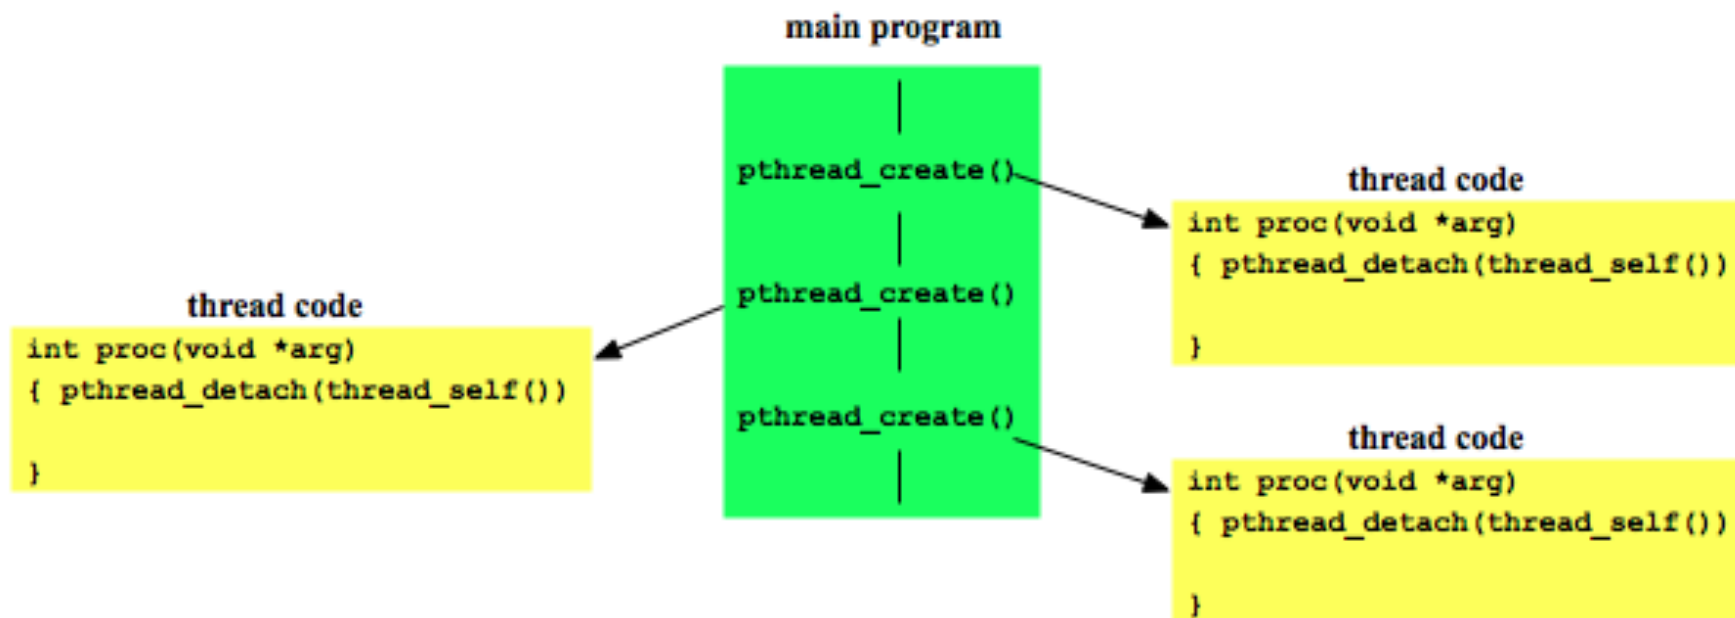
**pthread\_create()**  
start a new thread

**pthread\_join()**  
wait for child thread to join


**pthread\_exit()**  
stop thread

# Detached Threads


- Detached threads do not join
- Use `pthread_detach(thread_id)`
- Detached threads are more efficient
- Make sure that all detached threads terminate before program terminates



# Process vs Threads



*Process*

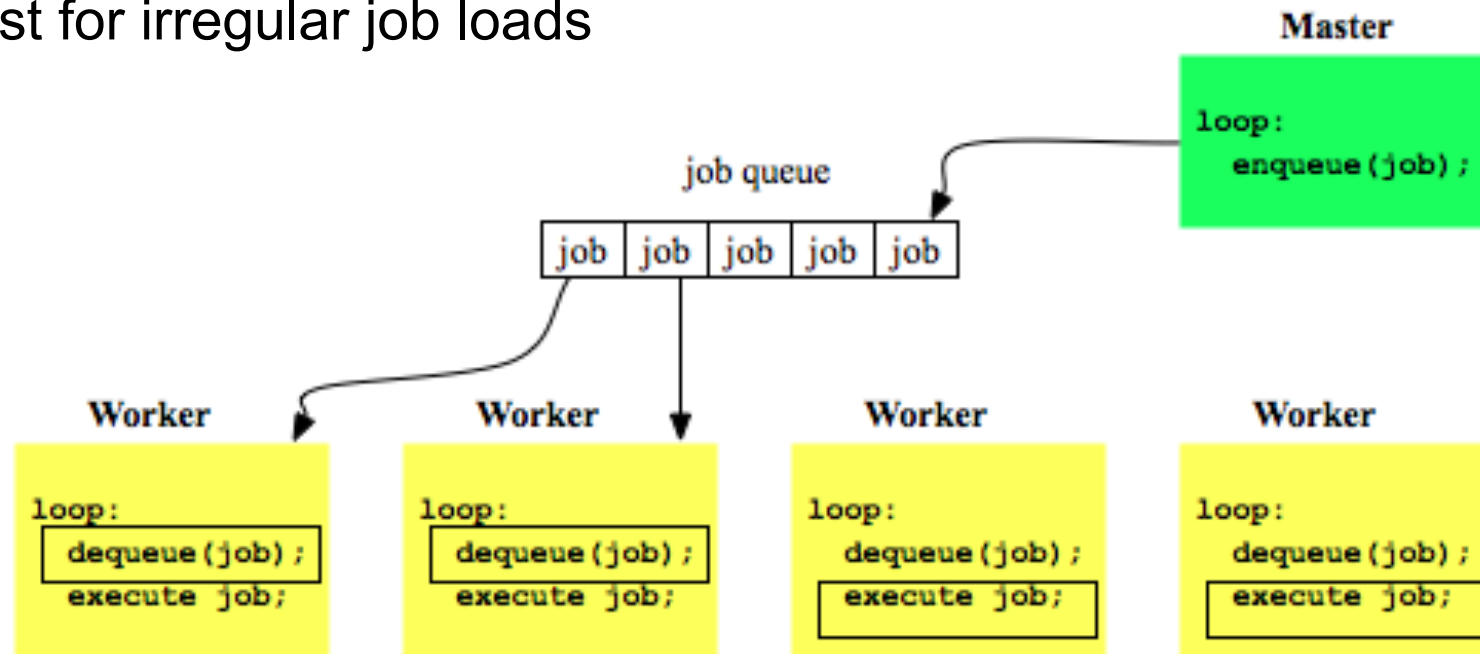


*Process with two threads*

*What happens when we fork a process that executes multiple threads?  
Does fork duplicate only the calling thread or all threads?*

# Thread Pools

- *Thread pooling* (or *process pooling*) is an efficient mechanism
- One *master thread* dispatches jobs to worker threads
- *Worker threads* in pool never terminate and keep accepting new jobs when old job done
- Jobs are communicated to workers via a *shared job queue*
- Best for irregular job loads







# MT-Safety

```
time_t clk = clock();
char *txt = ctime(&clk);
printf("Current time: %s\n", txt);
```

*Use of a non-MT-safe routine*

```
time_t clk = clock();
char txt[32];
ctime_r(&clk, txt);
printf("Current time: %s\n", txt);
```

*Use of the reentrant version of ctime*

```
static int counter = 0;
int count_events()
{ return counter++;
}
```

*Is this routine MT-safe?*

*What can go wrong?*

- Routines must be *multi-thread safe* (MT-safe) when invoked by more than one thread
- Non-MT-safe routines must be placed in a critical section, e.g. using a mutex lock (see later)
- Many C libraries are not MT-safe
  - Use *libroutine\_r()* versions that are “reentrant”
  - When building your own MT-safe library, use `#define _REENTRANT`
- Always make your routines MT-safe for reuse in a threaded application
- Use locks when necessary (see next slides)

# Coordinated Access to Shared Data (Such as Job Queues)


- Reading and writing shared data by more than one thread or process requires coordination with *locking*
- Cannot update shared variables simultaneously by more than one thread

```
static int counter = 0;
int count_events()
{ return counter++;
}
```



```
reg1 = M[counter] = 3
reg2 = reg1 + 1 = 4
M[counter] = reg2 = 4
return reg1 = 3
```


*Thread 1*



```
reg1 = M[counter] = 3
reg2 = reg1 + 1 = 4
M[counter] = reg2 = 4
return reg1 = 3
```

*Thread 2*

```
static int counter = 0;
int count_events()
{ pthread_mutex_lock(&lock);
  counter++;
  pthread_mutex_unlock(&lock);
  return counter-1;
}
```



```
acquire lock
reg1 = M[counter] = 3
reg2 = reg1 + 1 = 4
M[counter] = reg2 = 4
release lock
```

*Thread 1*



```
acquire lock
...
... wait
...
...
reg1 = M[counter] = 4
reg2 = reg1 + 1 = 5
M[counter] = reg2 = 5
release lock
```

*Thread 2*



# Spinlocks

- *Spin locks* use *busy waiting* until a condition is met
- Naïve implementations are almost always incorrect

```
// initially lock = 0
```

```
while (lock == 1)
    ; // do nothing
lock = 1;
... critical section ...
lock = 0;
```

} *Acquire lock*

} *Release lock*

*Two or more threads want to enter the critical section,  
what can go wrong?*



# Spinlocks

- *Spin locks use busy waiting* until a condition is met
- Naïve implementations are almost always incorrect

*Thread 1*

```
while (lock == 1)
    ; // do nothing
lock = 1;
... critical section ...
lock = 0;
```

*Thread 2*

```
while (lock == 1)
    ...
    ...
lock = 1;
... critical section ...
lock = 0;
```

*This ordering works*

# Spinlocks

- *Spin locks use busy waiting* until a condition is met
- Naïve implementations are almost always incorrect

*Thread 1*

```
while (lock == 1)
    ; // do nothing
lock = 1;
... critical section ...
lock = 0;
```

*Not set in time  
before read*

*Thread 2*

```
while (lock == 1)
    ...
lock = 1;
... critical section ...
lock = 0;
```

*This statement  
interleaving leads  
to failure*

*Both threads end up executing the  
critical section!*

# Spinlocks

- Spin locks use *busy waiting* until a condition is met
- Naïve implementations are almost always incorrect

*Thread 1*

```
while (lock == 1)
    ; // do nothing
lock = 1;
    ... critical section ...
lock = 0;
```

***Compiler optimizes  
the code!***

*Useless  
assignment  
removed*

*Assignment  
can be  
moved by  
compiler*

*Thread 2*

```
while (lock == 1)
    ...
lock = 1;
    ... critical section ...
lock = 0;
```

*Atomic operations such as atomic “test-and-set” instructions must be used  
(these instructions are not reordered or  
removed by compiler)*



# Spinlocks

- Advantage of spinlocks is that the kernel is not involved
- Better performance when acquisition waiting time is short
- Dangerous to use in a uniprocessor system, because of *priority inversion*
- No guarantee of *fairness* and a thread may wait indefinitely in the worst case, leading to *starvation*

```
void spinlock_lock(spinlock *s)
{ while (TestAndSet(s))
    while (*s == 1)
        ;
}
```

```
void spinlock_unlock(spinlock *s)
{ *s = 0;
}
```

*Correct and efficient spinlock operations using atomic **TestAndSet** assuming hardware supports cache coherence protocol*

Note: **TestAndSet(int \*n)** sets **n** to 1 and returns old value of **n**



# Semaphores

- A semaphore is an integer-valued *counter*
- The counter is *incremented* and *decremented* by two operations *signal* (or *post*) and *wait*, respectively
  - Traditionally called *V* and *P* (Dutch “verhogen” and “probeer te verlagen”)
- When the counter  $\leq 0$  the *wait* operation blocks and waits until the counter  $> 0$

```
sem_post(sem_t *s)
{  (*s)++;
}
```

*Note: actual implementations of POSIX semaphores use atomic operations and a queue of waiting processes to ensure fairness*

```
sem_wait(sem_t *s)
{ while (*s <= 0)
    ; // do nothing
  (*s)--;
}
```



# Semaphores

- A two-valued (= binary) semaphore provides a mechanism to *implement mutual exclusion* (mutex)
- POSIX semaphores are *named* and have *permissions*, allowing use across a set processes

	<i>Unique name</i>	<i>Permissions</i>	<i>Initial value</i>
<code>#include "semaphore.h"</code>	↓	↓	↓
<code>sem_t *mutex = sem_open("lock371", O_CREAT, 0600, 1);</code>			
<code>...</code>			
<code>sem_wait(mutex);</code>			
<code>...</code>			
<code>... critical section ...</code>			
<code>...</code>			
<code>sem_post(mutex);</code>			
<code>...</code>			
<code>sem_close(mutex);</code>			

*// sem\_trywait() to poll state*

Tip: use **ipcs** command to display IPC semaphore status of a system



# Pthread Mutex Locks

```
pthread_mutex_t mylock;  
  
pthread_mutex_init(&mylock, NULL);  
...  
pthread_mutex_lock(&mylock);  
...  
... critical section ...  
...  
pthread_mutex_unlock(&mylock);  
...  
pthread_mutex_destroy(&mylock);
```

- POSIX mutex locks for thread synchronization
  - Threads share user space, processes do not
- Pthreads is available for Unix/  
Linux and Windows ports
  - `pthread_mutex_init()`  
initialize lock
  - `pthread_mutex_lock()`  
lock
  - `pthread_mutex_unlock()`  
unlock
  - `pthread_mutex_trylock()`  
check if lock can be acquired



# Using Mutex Locks

- Locks are used to synchronize shared data access from any part of a program, not just the same routine executed by multiple threads
- Multiple locks should be used, each for a set of shared data items that is disjoint from another set of shared data items (no single lock for everything)

```
pthread_mutex_lock(&array_A_lck);  
... A[i] = A[i] + 1 ...  
pthread_mutex_unlock(&array_A_lck);
```

```
pthread_mutex_lock(&array_A_lck);  
... A[i] = A[i] + B[i]  
pthread_mutex_unlock(&array_A_lck);
```

*Lock operations on array A*

```
pthread_mutex_lock(&queue_lck);  
... add element to shared queue ...  
pthread_mutex_unlock(&queue_lck);  
  
pthread_mutex_lock(&queue_lck);  
... remove element from shared queue ...  
pthread_mutex_unlock(&queue_lck);
```

*Lock operations on a queue*

*What if threads may or may not update some of the same elements of an array, should we use a lock for **every** array element?*



# Condition Variables


- *Condition variables* are associated with mutex locks
- Provide signal and wait operations *within* critical sections

*Process 1*

```
lock(mutex)
if (cannot continue)
    wait(mutex, event)
...
unlock(mutex)
```

*Process 2*

```
lock(mutex)
...
signal(mutex, event)
...
unlock(mutex)
```



*Can't use semaphore wait and signal here:  
what can go wrong when using semaphores?*



# Condition Variables

**signal** releases one waiting thread (if any)

*Process 1*

```
lock(mutex)
if (cannot continue)
    wait(mutex, event)
...
unlock(mutex)
```

*Process 2*

```
lock(mutex)
...
signal(mutex, event)
...
unlock(mutex)
```

**wait** blocks until a signal is received  
When blocked, it releases the mutex lock,  
and reacquires the lock when wait is over



# Producer-Consumer Example

- Producer adds items to a shared container, when **not full**
- Consumer picks an item from a shared container, when **not empty**

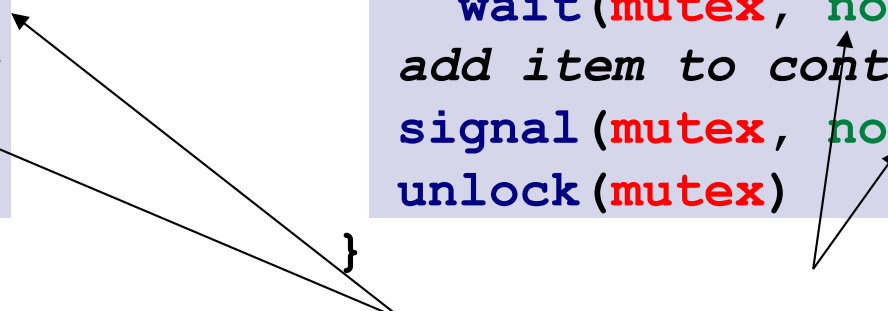
*A consumer*

```
while (true)
{ lock(mutex)
  if (container is empty)
    wait(mutex, notempty)
  get item from container
  signal(mutex, notfull)
  unlock(mutex)
}
```

*A producer*

```
while (true)
{ lock(mutex)
  if (container is full)
    wait(mutex, notfull)
  add item to container
  signal(mutex, notempty)
  unlock(mutex)
}
```

*Condition variables  
associated with mutex*





# Semaphores versus Condition Variables

## ■ Semaphores:

- Semaphores must have matching signal-wait pairs, that is, the semaphore counter must stay balanced
- One too many waits: one waiting thread is indefinitely blocked
- One too many signals: two threads may enter critical section that is guarded by semaphore locks

## ■ Condition variables:

- A signal can be executed at any time
- When there is no wait, signal does nothing
- If there are multiple threads waiting, signal will release one

## ■ Both provide:

- Fairness: waiting threads will be released with equal probability
- Absence of starvation: no thread will wait indefinitely



# Pthreads Condition Variables

- Pthreads supports condition variables
- A condition variable is always used in combination with a lock, based on the principle of “*monitors*”

*Declarations*

```
pthread_mutex_t mutex;  
pthread_cond_t notempty, notfull;
```

*Initialization*

```
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&notempty, NULL);  
pthread_cond_init(&notfull, NULL);
```

*A consumer*

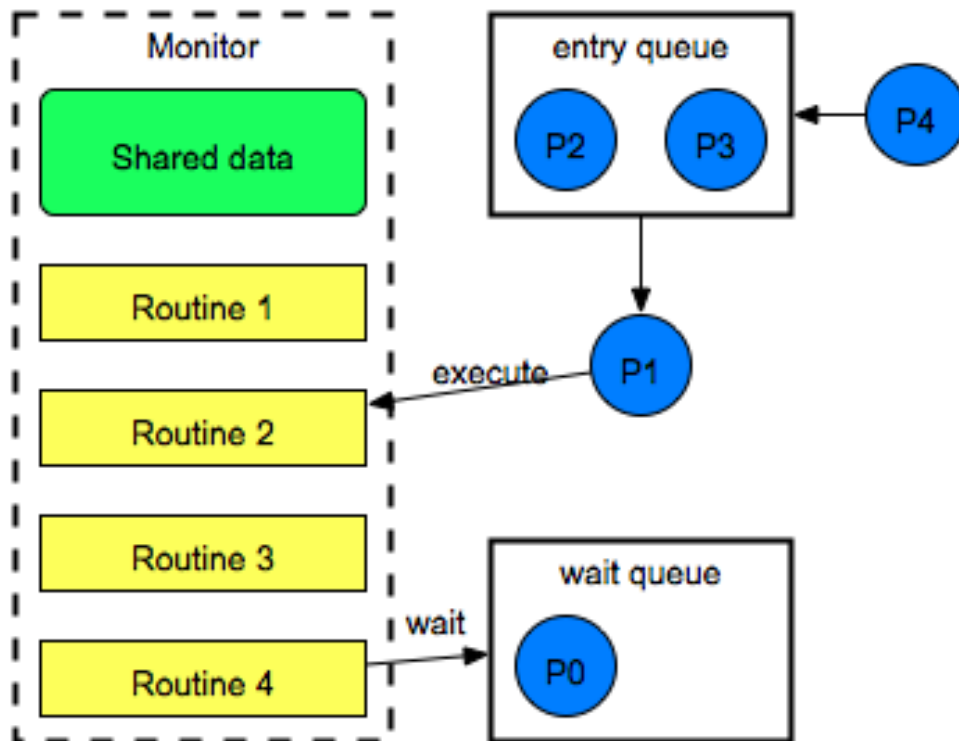
```
while (1)  
{ pthread_mutex_lock(&mutex);  
  if (container is empty)  
    pthread_cond_wait(&mutex, &notempty);  
  get item from container  
  pthread_cond_signal(&mutex, &notfull);  
  pthread_mutex_unlock(&mutex);  
}
```

*A producer*

```
while (1)  
{ pthread_mutex_lock(&mutex);  
  if (container is full)  
    pthread_cond_wait(&mutex, &notfull);  
  add item to container  
  pthread_cond_signal(&mutex, &notempty);  
  pthread_mutex_unlock(&mutex);  
}
```



# Monitor with Condition Variables



*Only P1 executes a routine, P0 waits on a signal, and P2, P3 are in the entry queue to execute next when P1 is done (or moved to the wait queue)*

- A *monitor* is a concept
- A monitor combines a set of shared variables and a set of routines that operate on the variables
- Only one process may be active in a monitor at a time
  - All routines are synchronized by implicit locks (like an entry queue)
  - Shared variables are safely modified under mutex
- Condition variables are used for signal and wait within the monitor routines
  - Like a wait queue



# Barriers

- A *barrier* synchronization statement in a program blocks processes until all processes have arrived at the barrier
- Frequently used in data parallel programming (implicit or explicit) and an essential part of BSP

*Each process produces part of shared data  $X$*

**barrier**

*Processes use shared data  $X$*



# Two-Phase Barrier with Semaphores for $P$ Processes

```
sem_t *mutex = sem_open("mutex-492", O_CREAT, 0600, 1);
sem_t *turnstile1 = sem_open("ts1-492", O_CREAT, 0600, 0);
sem_t *turnstile2 = sem_open("ts2-492", O_CREAT, 0600, 1);
int count = 0;
```

```
...
sem_wait(mutex);
    if (++count == P)
    { sem_wait(turnstile2);
      sem_signal(turnstile1);
    }
sem_signal(mutex);
sem_wait(turnstile1);
sem_signal(turnstile1);
sem_wait(mutex);
    if (--count == 0)
    { sem_wait(turnstile1);
      sem_signal(turnstile2);
    }
sem_signal(mutex);
sem_wait(turnstile2);
sem_signal(turnstile2);
```

*Rendezvous*

*Critical point*

*Barrier sequence*



# Pthread Barriers

```
pthread_barrier_t barrier;  
  
pthread_barrier_init(  
    barrier,  
    NULL,    // attributes  
    count); // number of threads  
...  
pthread_barrier_wait(barrier);  
...
```

- Barrier using POSIX pthreads (advanced realtime threads)
- Specify number of threads involved in barrier syncs in initialization

---

```
pthread_barrier_init()  
initialize barrier with thread count
```

```
pthread_barrier_wait()  
barrier synchronization
```

---



# Further Reading

- [PP2] pages 230-247
- [HPC] pages 191-218
- Optional:
  - [HPC] pages 219-240
  - Pthread manuals (many online)
  - “*The Little Book of Semaphores*” by Allen Downey  
<http://www.greenteapress.com/semaphores/downey05semaphores.pdf>