

High-Performance Libraries and Tools

HPC Fall 2012

Prof. Robert van Engelen





Overview

- Dense matrix
 - BLAS (serial)
 - ATLAS (serial/threaded)
 - LAPACK (serial)
 - Vendor-tuned LAPACK (shared memory parallel)
 - ScaLAPACK/PLAPACK (distributed memory parallel)
 - FLAME (an algorithm derivation framework)
- Sparse matrix
 - PETSc
- Further reading



BLAS

- The *Basic Linear Algebra Subprograms* (BLAS) consist of a set of lower-level linear algebra operations
- Level 1: vector-vector
 - $O(n)$ operations on $O(n)$ data
 - Bandwidth to memory is a limiting factor
- Level 2: matrix-vector
 - $O(n^2)$ operations on $O(n^2)$ data
 - Vectors kept in cache
- Level 3: matrix-matrix
 - $O(n^3)$ operations on $O(n^2)$ data
 - Blocked matrices kept in cache
- Netlib's BLAS is a reference implementation

Examples

$$y \leftarrow \alpha x + y$$

$$y \leftarrow \alpha Ax + \beta y$$

$$Tx = y \quad (\text{Triangular } T)$$

$$C \leftarrow \alpha AB + \beta C$$

$$B \leftarrow \alpha T^{-1} B \quad (\text{Triangular } T)$$



GotoBlas and Vendor-Tuned BLAS

- Implemented by Kazushige Goto
- Optimized for cache and Translation Lookaside Buffer (TLB)
- Restrictive open-source license
- Licensed to vendors for vendor-tuned BLAS libraries

- Vendor-tuned BLAS
 - Accelerate framework (Apple)
 - MLK (Intel)
 - ACML (AMD)
 - ESSL (IBM)
 - MLIB (HP)
 - Sun performance library



ATLAS

- The Automatically Tuned Linear Algebra Software (ATLAS) is a self-tuned BLAS version
- Installation tests numerical kernels and (other parts of) the code to determine which parameters are best for a particular machine, e.g. blocking, loop unrolling, ...
- Faster than the reference implementation
- Freely available



DGEMM




Image source: Robert van de Geijn (TACC)

11/26/12

HPC Fall 2012

DGEMM




Image source: Robert van de Geijn (TACC)

11/26/12

HPC Fall 2012



DGEMM




Image source: Robert van de Geijn (TACC)

11/26/12

HPC Fall 2012



LAPACK

- Linear Algebra PACKage (LAPACK) written in Fortran
- Built on BLAS
- Standard API (Application Programming Interface)
 - Data type: real and complex, single and double precision
 - Matrix shapes: general dense, diagonal, bidiagonal, tridiagonal, banded, trapeziodal, Hessenberg
 - Matrix properties: general, orthogonal, positive definite, Hermitian, symmetric
 - Linear least squares, eigenvalue problems, singular value decomposition, matrix factorizations (LU, QR, Cholesky, Schur)
- Reference implementation from Netlib
- Vendor-tuned versions available
 - Some for shared memory parallel



ScaLAPACK/PLAPACK

- ScaLAPACK/PLAPACK are versions of LAPACK for distributed memory MIMD parallel machines
 - Subset of LAPACK routines
- ScaLAPACK is built on BLAS and MPI
- ScaLAPACK reference implementation from Netlib
- PLAPACK is a project at UT Austin (TACC)



FLAME

- Formal Linear Algebra Methods Environment (FLAME)
- LAPACK code is hard to write/read/maintain/alter
- “Transform the development of dense linear algebra libraries from an art reserved for experts to a science that can be understood by novice and expert alike”
 - Notation for expressing algorithms
 - A methodology for systematic derivation of algorithms using loop invariants
 - Application Program Interfaces (APIs) for representing the algorithms in code
 - Tools for mechanical derivation, implementation and analysis of algorithms and implementations

FLAME-Derived Blocked LU

Algorithm: $[A] := \text{LU_BLK_VAR5}(A)$

Partition $A \rightarrow \left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right)$

where A_{TL} is 0×0

while $m(A_{TL}) < m(A)$ **do**

Determine block size b

Repartition

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} A_{00} & A_{01} & A_{02} \\ \hline A_{10} & A_{11} & A_{12} \\ \hline A_{20} & A_{21} & A_{22} \end{array} \right)$$

where A_{11} is $b \times b$

$$A_{11} = \text{LU}(A_{11})$$

$$A_{12} = \text{TRILU}(A_{11})^{-1} A_{12}$$

$$A_{21} = A_{21} \text{TRIU}(A_{11})^{-1}$$

$$A_{22} = A_{22} - A_{21} A_{12}$$

Continue with

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} A_{00} & A_{01} & A_{02} \\ \hline A_{10} & A_{11} & A_{12} \\ \hline A_{20} & A_{21} & A_{22} \end{array} \right)$$

endwhile

```

FLA_Part_2x2( A,      &ATL, &ATR,
              &ABL, &ABR,      0, 0, FLA_TL );

while (FLA_Obj_length( ATL ) < FLA_Obj_length( A )){

    b = min( FLA_Obj_length( ABR ), nb_alg );

    FLA_Repart_2x2_to_3x3
        ( ATL, /**/ ATR,      &A00, /**/ &A01, &A02,
          /* ***** */      /* ***** */
          ABL, /**/ ABR,      &A10, /**/ &A11, &A12,
          b, b, FLA_BR );
    /*-----*/
    LU_unb_var5( A11 );

    FLA_Trsm( FLA_LEFT, FLA_LOWER_TRIANGULAR,
              FLA_NO_TRANSPOSE, FLA_UNIT_DIAG,
              FLA_ONE, A11, A12 );

    FLA_Trsm( FLA_RIGHT, FLA_UPPER_TRIANGULAR,
              FLA_NO_TRANSPOSE, FLA_NONUNIT_DIAG,
              FLA_ONE, A11, A21 );

    FLA_Gemm( FLA_NO_TRANSPOSE, FLA_NO_TRANSPOSE,
              FLA_MINUS_ONE, A21, A12, FLA_ONE, A22 );
    /*-----*/
    FLA_Cont_with_3x3_to_2x2
        ( &ATL, /**/ &ATR,      A00, A01, /**/ A02,
          A10, A11, /**/ A12,
          /* ***** */      /* ***** */
          &ABL, /**/ &ABR,      A20, A21, /**/ A22,
          FLA_TL );
}
    
```

AutoFLAME

Mathematica 5.0 - [LAPACK.nb * -STUDENT VERSION-]

File Edit Cell Format Input Kernel Find Window Help

LAPACK.nb * -STUDENT VERSION-

1

Operation: $[L] = \text{TrinvLVar1}(L)$

Partition
 $L \rightarrow \left(\begin{array}{c|c} \hat{L}_{mn} & 0 \\ \hline \hat{L}_m & \hat{L}_{ms} \end{array} \right)$
 where \hat{L}_{mn} is empty

loop invariant:
 $\left(\begin{array}{c|c} \hat{L}_{mn} & 0 \\ \hline \hat{L}_m & \hat{L}_{ms} \end{array} \right) = \left(\begin{array}{c|c} \hat{L}_{mn}^{-1} & 0 \\ \hline \hat{L}_m & \hat{L}_{ms} \end{array} \right)$

While $\hat{L}_{mn} \neq L$

Repartition
 $\left(\begin{array}{c|c} \hat{L}_{mn} & 0 \\ \hline \hat{L}_m & \hat{L}_{ms} \end{array} \right) \rightarrow \left(\begin{array}{c|c} \hat{L}_{11} & 0 \\ \hline \hat{L}_{21} & \hat{L}_{22} \end{array} \right)$

loop invariant before the updates:
 $\left(\begin{array}{c|c} \hat{L}_{11} & 0 \\ \hline \hat{L}_{21} & \hat{L}_{22} \end{array} \right) = \left(\begin{array}{c|c} \hat{L}_{11}^{-1} & 0 \\ \hline \hat{L}_{21} & \hat{L}_{22} \end{array} \right)$

$$\begin{aligned} L_{11} &:= \hat{L}_{11}^{-1} \\ L_{10} &:= -\hat{L}_{11} \cdot \hat{L}_{10} \cdot \hat{L}_{00} \end{aligned}$$

Continue with
 $\left(\begin{array}{c|c} \hat{L}_{11} & 0 \\ \hline \hat{L}_{21} & \hat{L}_{22} \end{array} \right) \rightarrow \left(\begin{array}{c|c} \hat{L}_{mn} & 0 \\ \hline \hat{L}_m & \hat{L}_{ms} \end{array} \right)$

loop invariant after the updates:
 $\left(\begin{array}{c|c} \hat{L}_{11} & 0 \\ \hline \hat{L}_{21} & \hat{L}_{22} \end{array} \right) = \left(\begin{array}{c|c} \hat{L}_{11}^{-1} & 0 \\ \hline -\hat{L}_{21}^{-1} \cdot \hat{L}_{11} \cdot \hat{L}_{11}^{-1} & \hat{L}_{22}^{-1} \end{array} \right)$

end while

75%

LU w/ Pivoting on 8 Cores

4 x AMD 2.4GHz dual-core Opteron 880




Image source: Robert van de Geijn (TACC)

11/26/12

HPC Fall 2012

14

QR Factorization on 8 Cores

4 x AMD 2.4GHz dual-core Opteron 880

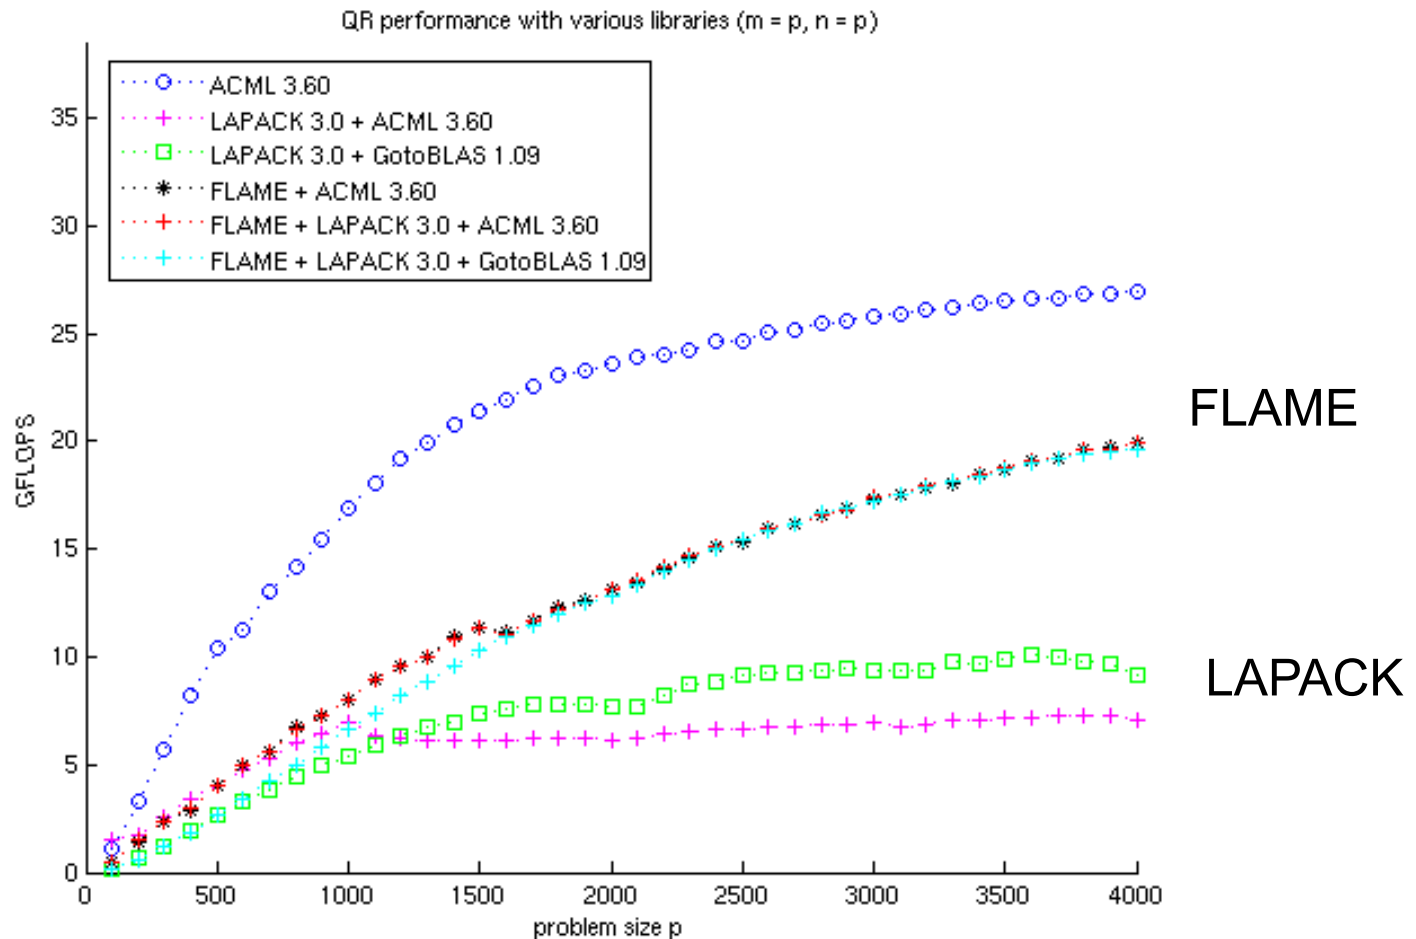


Image source: Robert van de Geijn (TACC)



Cholesky on 8 Cores

4 x AMD 2.4GHz dual-core Opteron 880

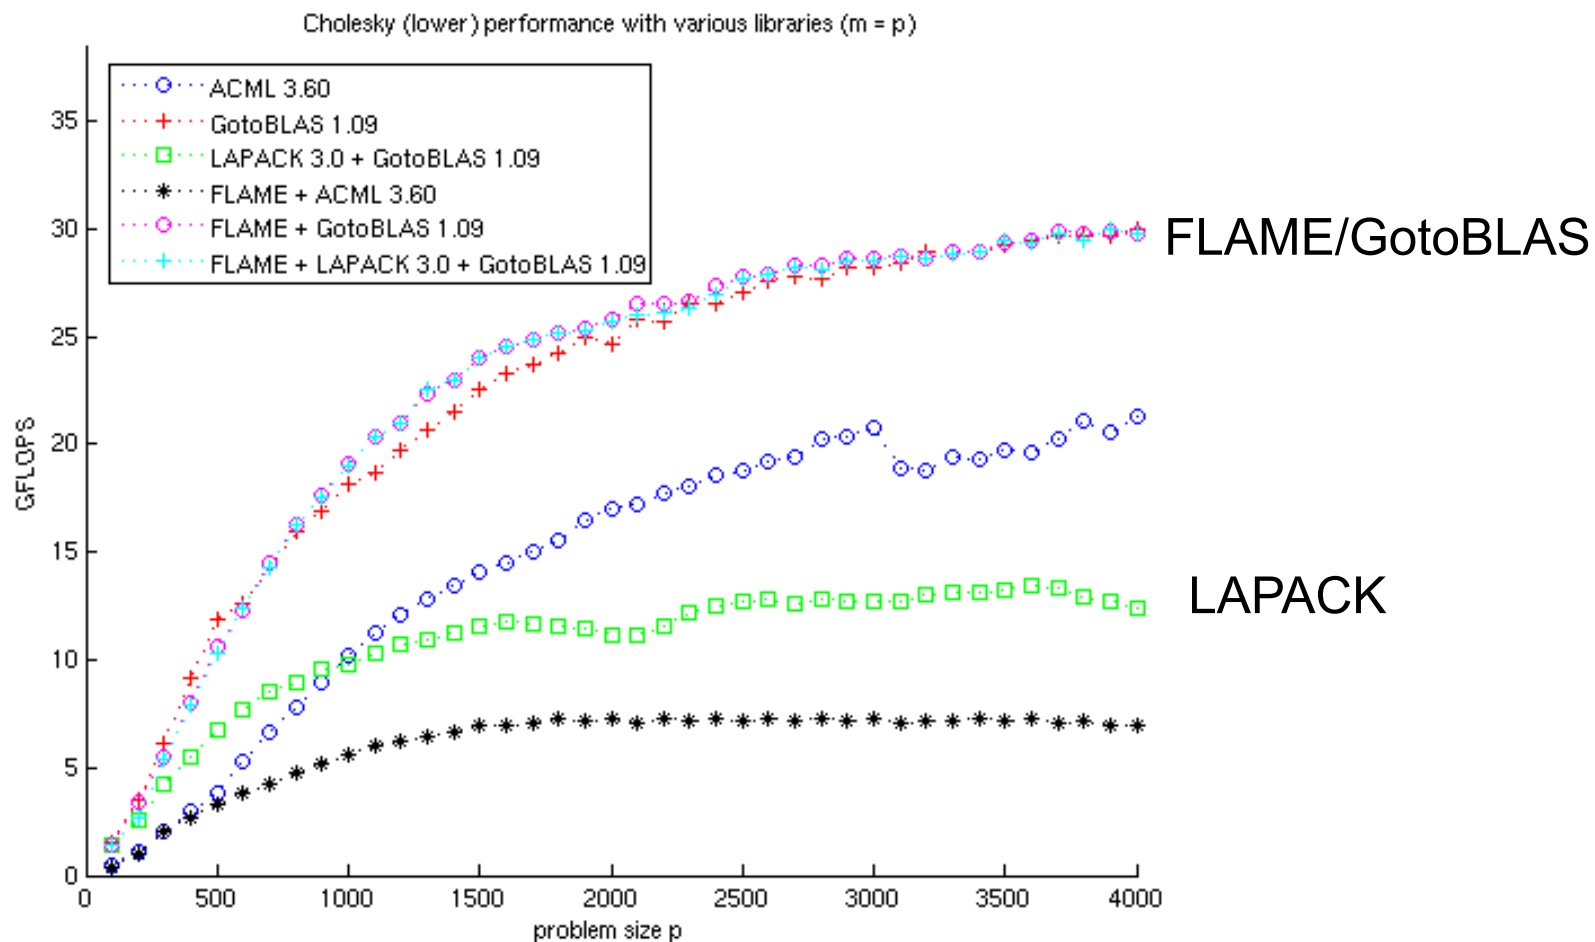


Image source: Robert van de Geijn (TACC)



PETSc


- Portable, Extensible Toolkit for Scientific Computation (PETSc) for distributed memory MIMD parallel machines
 - Vector/matrix formats and array operations (serial and parallel)
 - Linear and nonlinear solvers
 - Limited ODE integrators
 - Limited grid/data management (serial and parallel)
- Built on BLAS, LAPACK, and MPI
- Basically a solver library for general sparse matrices
 - User writes main() program
 - User orchestrates computation via object creations
 - User controls the basic flow of the PETSc program
 - PETSc propagates errors from underlying libs

PETSc Numerical Components

Nonlinear Solvers (SNES)				Time Steppers (TS)			
Newton-based Methods		Other		Euler	Backward Euler	Pseudo Time Stepping	Other
Line Search	Trust Region						
Krylov Subspace Methods (KSP)							
GMRES	CG	CGS	Bi-CG-STAB	TFQMR	Richardson	Chebyshev	Other
Preconditioners (PC)							
Additive Schwartz	Block Jacobi	Jacobi	ILU	ICC	LU (Sequential only)	Others	
Matrices (Mat)							
Compressed Sparse Row (AIJ)	Blocked Compressed Sparse Row (BAIJ)		Block Diagonal (BDIAG)	Dense	Matrix-free	Other	
Distributed Arrays(DA)				Index Sets (IS)			
Vectors (Vec)				Indices	Block Indices	Stride	Other

Image source: PE'20

Image source: PETSc project



PETSc Linear Solver Example

$Ax = b$

```
KSP ksp; /* linear solver context */
Mat    A; /* matrix */
Vec x, b; /* solution, RHS vectors */
int n;    /* problem dimension */

MatCreate(PETSC_COMM_WORLD, PETSC_DECIDE, PETSC_DECIDE, n, n, &A);
MatSetFromOptions(A);
/* (user-defined code to assemble matrix A not shown) */
VecCreate(PETSC_COMM_WORLD, &x);
VecSetSizes(x, PETSC_DECIDE, n);
VecSetFromOptions(x);
VecDuplicate(x, &b);
/* (user-defined code to assemble RHS vector b not shown) */
KSPCreate(PETSC_COMM_WORLD, &ksp);
KSPSetOperators(ksp, A, A, DIFFERENT_NONZERO_PATTERN);
KSPSetFromOptions(ksp);
KSPSolve(ksp, b, x);
KSPDestroy(ksp);
```

PETSc Flow of Control for PDEs

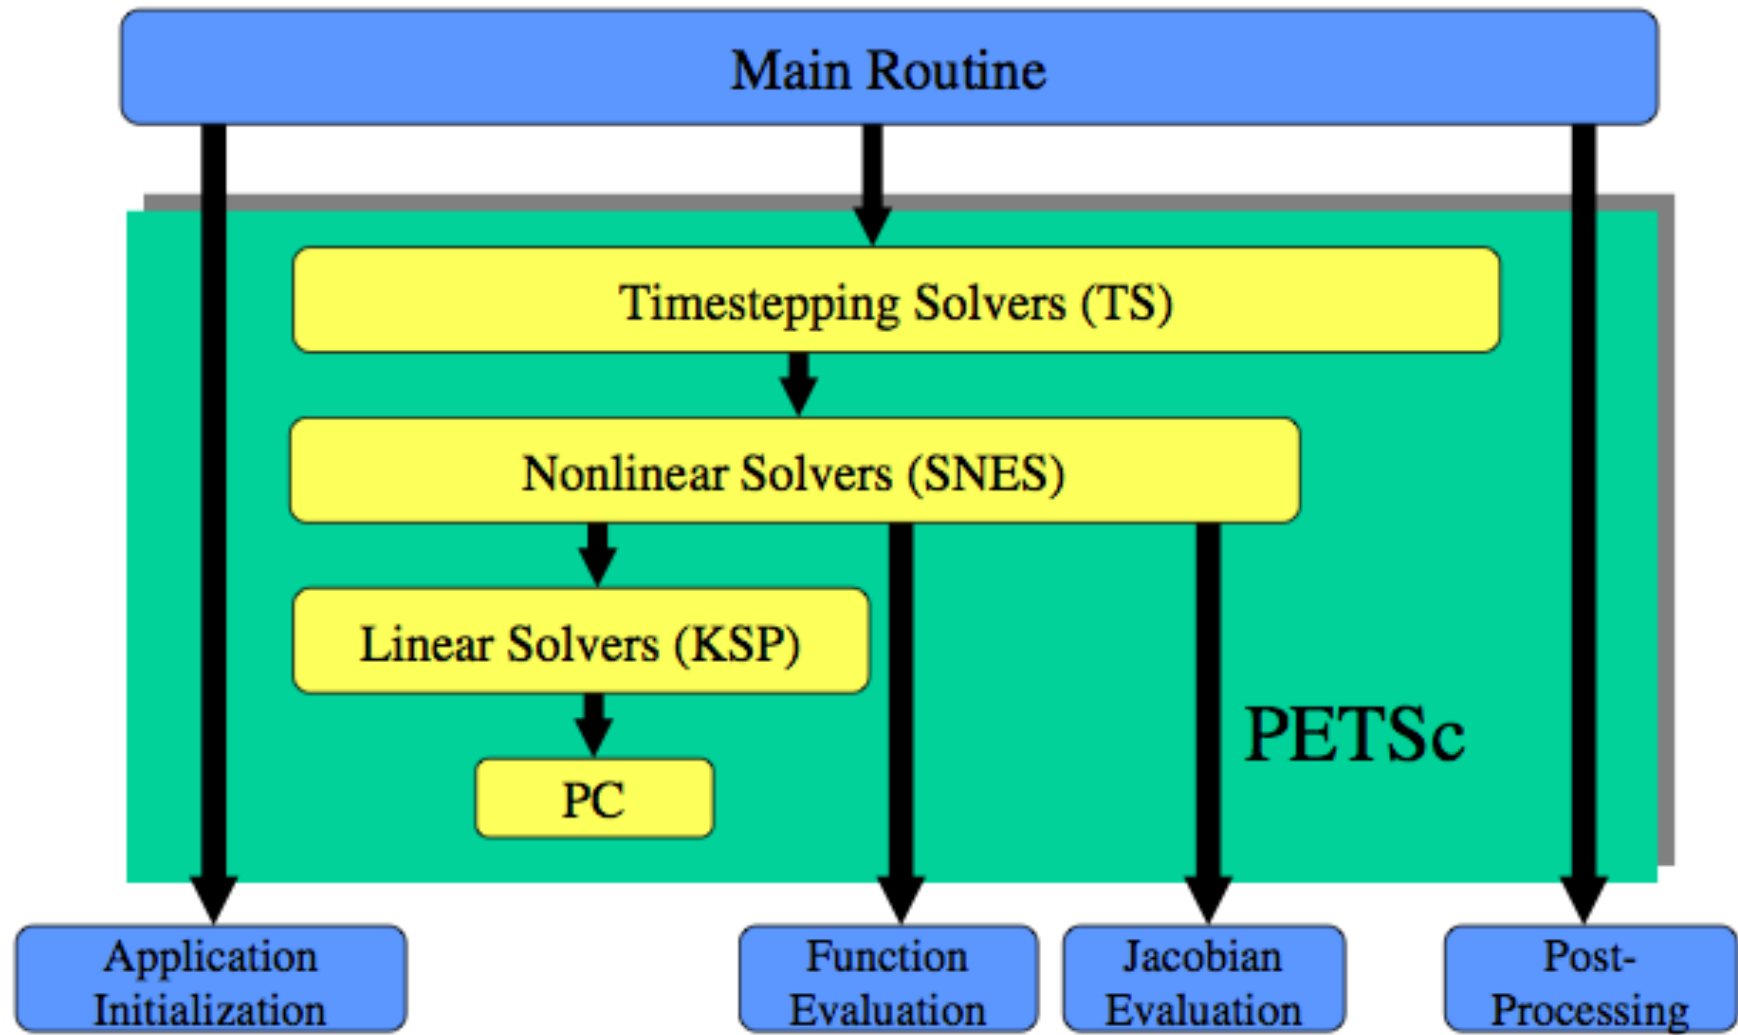


Image source: PETSc project

11/26/12

◆ User code

◆ PETSc code

HPC Fall 2012



PETSc Nonlinear Solver Interface: SNES

- For problems arising from PDEs
- Uses Newton-based methods
 - (Approximately) solve $F'(u_k) = -F(u_k)$
 - Update $u_{k+1} = u_k + \Delta u_k$
- Support the general solution to $F(u) = 0$
- User provides:
 - Code to evaluate $F(u)$
 - Code to evaluate Jacobian of $F(u)$
 - Or use (built-in) first-order sparse finite difference approximation
 - Or use automatic differentiation, e.g. ADIFOR and ADIC



PETSc Nonlinear Solver Example

```
SNES snes; /* nonlinear solver context */
Mat      J; /* Jacobian matrix */
Vec x, f; /* solution, RHS vectors */
int n, its; /* problem dimension, number of iterations */
ApptCtx uc; /* user-defined application context */

MatCreate(PETSC_COMM_WORLD, n, n, &J);
VecCreate(PETSC_COMM_WORLD, n, &x);
VecDuplicate(x, &f);

SNESCreate(PETSC_COMM_WORLD, SNES_NONLINEAR_EQUATIONS, &snes);
SNESSetFunction(snes, f, EvaluateFunction, uc);
SNESSetJacobian(snes, J, EvaluateJacobian, uc);
SNESSetFromOptions(snes);

SNESSolve(snes, x, &its);

SNESDestroy(snes);
```

PETSc Meshes

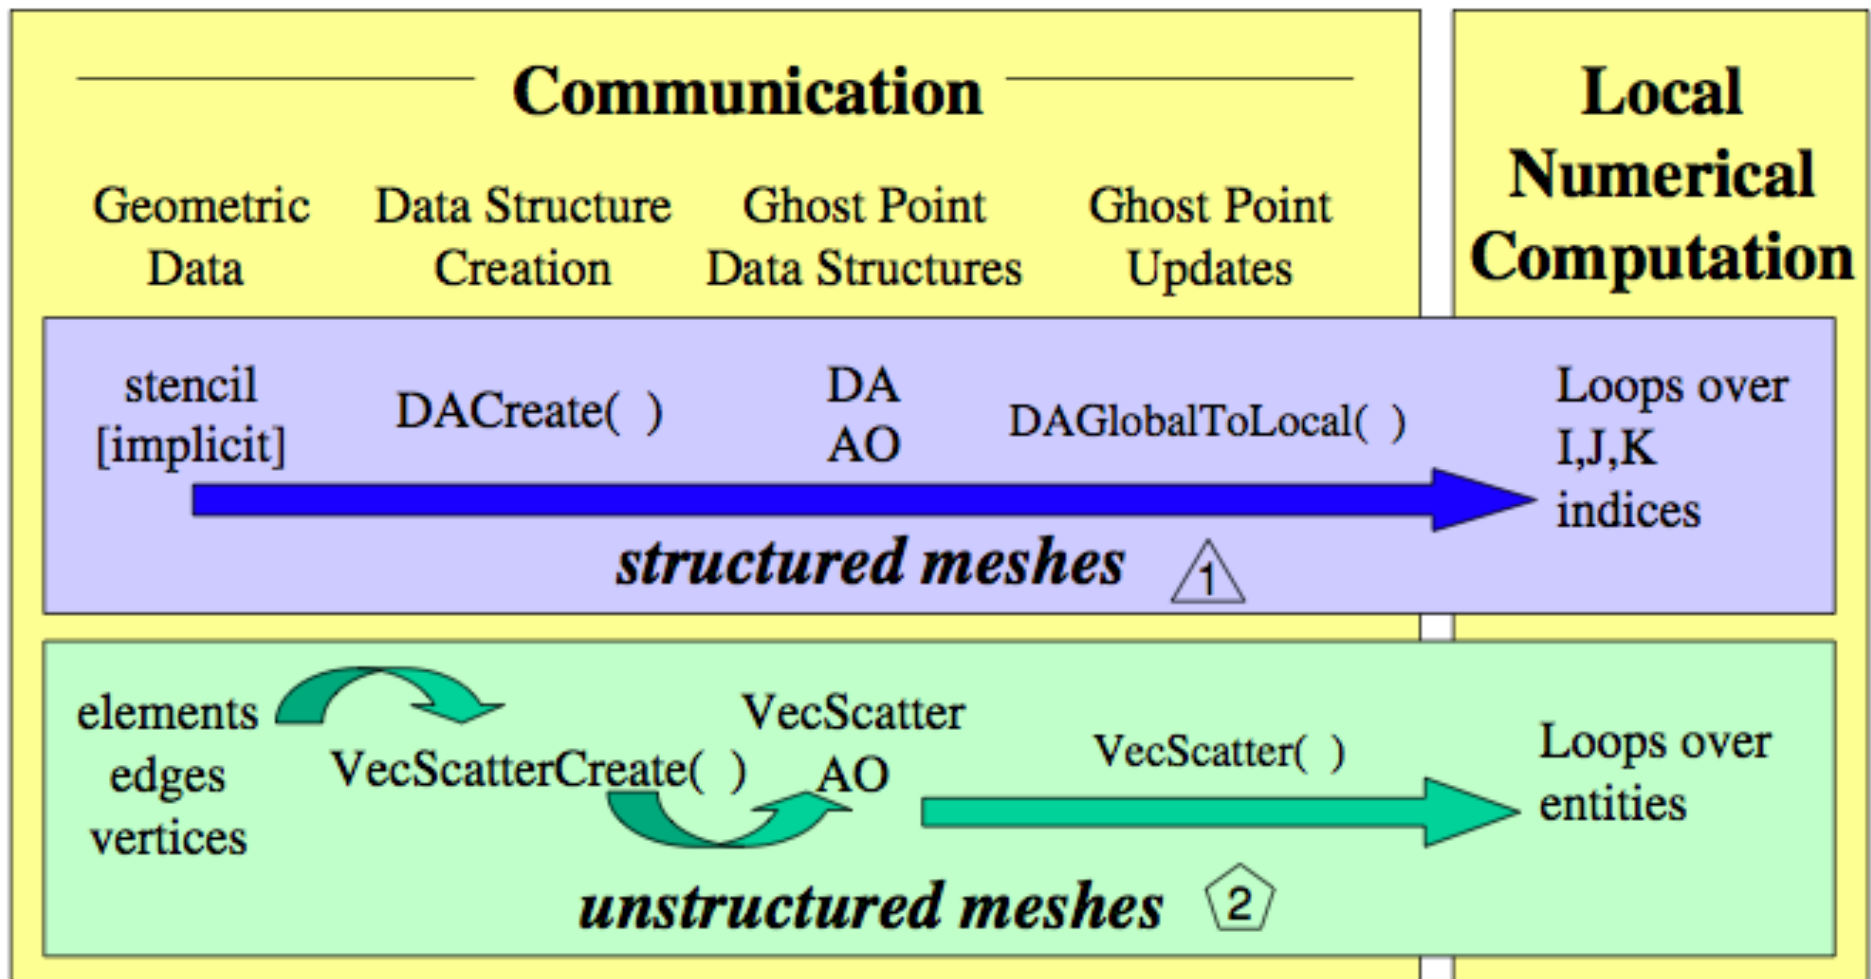

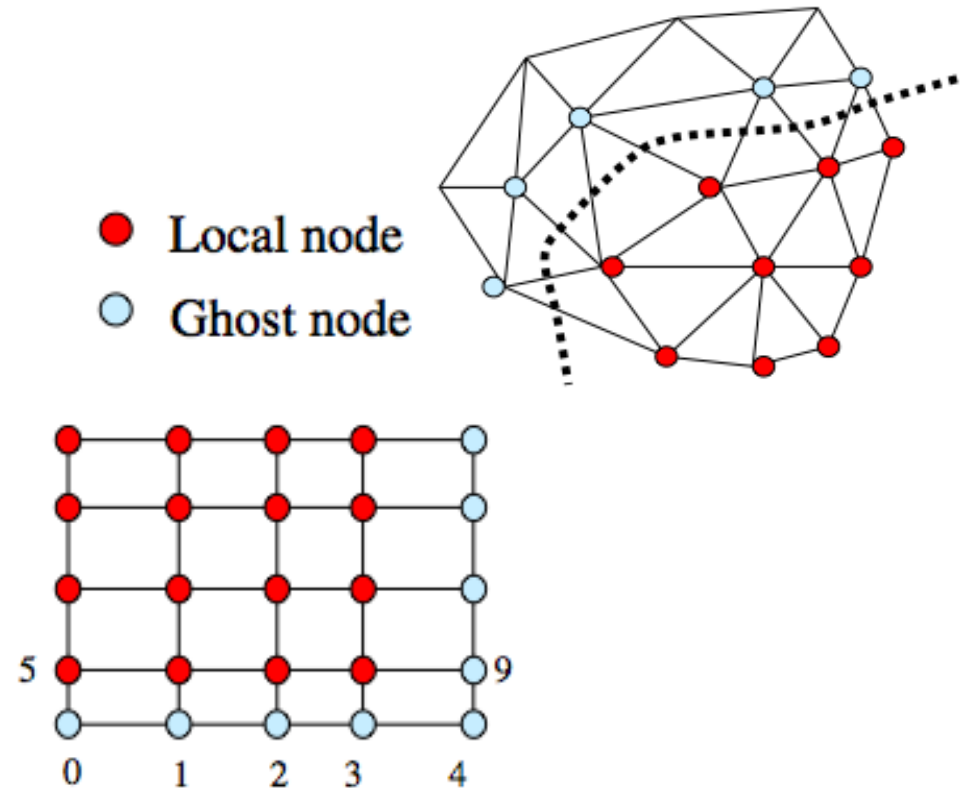


Image source: PETSc project

PETSc Global vs Local Meshes



Global: each process stores a unique **local set of vertices** (and each vertex is owned by exactly one process)



Local: each process stores a unique **local set of vertices as well as ghost nodes** from neighboring processes

Image source: PETSc project



PETSc Distributed Arrays

- Form a DA:
 - `DACreate1d(..., DA*)`
 - `DACreate2d(..., DA*)`
 - `DACreate3d(..., DA*)`
- Create the corresponding PETSc vectors
 - `DACreateGlobalVector(DA, Vec*)`
 - `DACreateLocalVector(DA, Vec*)`
- Update ghost points (scatter global vector into local parts, including ghost points)
 - `DAGlobalToLocalBegin(DA, ...)`
 - `DAGlobalToLocalEnd(DA, ...)`



Further Reading

- [SRC] pages 621-647
- Netlib organization: www.netlib.org
- FLAME project: www.cs.utexas.edu/users/flame
- PETSc project: www.mcs.anl.gov/petsc
- Linear algebra Wiki: www.linearalgebrawiki.org