



Software Profiling with TAU

Prasad Maddumage
mhemantha@fsu.edu

Research Computing Center
Florida State University



Software Profiling



- Dynamic program analysis using various measures related to code execution
 - CPU/memory utilization, frequency of function calls, I/O, MPI library usage, hardware counters, etc.
- Profilers *instrument* source or binary to obtain such measures during runtime
 - Instrumenting is inserting probes and replacing or wrapping function calls (eg: MPI calls, I/O) with modified calls of a source code
- Analyzing the results will help programmers/scientists to improve code performance



Profile vs Trace

- Profile: statistical summary of all metrics measured
 - Shows how much total time/resources each call utilized
- Trace: timeline of runtime events took place
 - Shows when each event happened and where

Metric: TIME
Value: Exclusive
Units: seconds






Why use TAU?

- Tuning and Analysis Utilities (20+ year project)
 - Actively developed by Univ. of Oregon, ANL, LANL, Julich
- Comprehensive performance profiling and tracing
 - Integrated, scalable, flexible, portable
 - Targets all parallel programming/execution paradigms
- Integrated performance toolkit
 - Instrumentation, measurement, analysis, visualization
 - Performance data management and data mining
 - Open source
- Easy to integrate in application frameworks
- Well documented

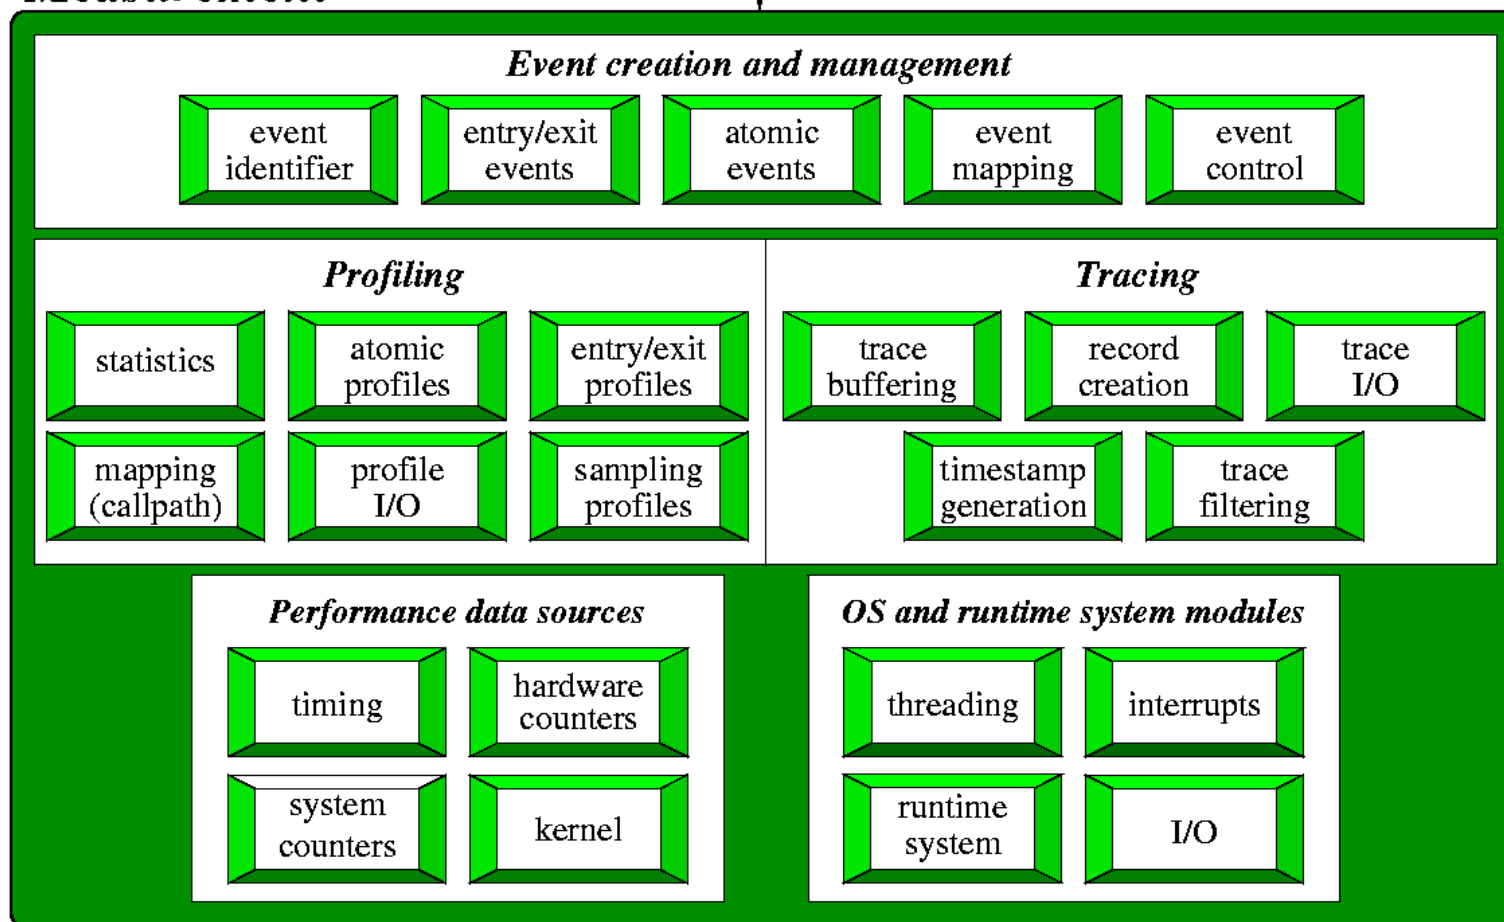


How does TAU work?

Instrumentation



Measurement





How does TAU work?

- Instrumentation
 - Source code instrumentation using pre-processors and compiler scripts
 - Wrapping external libraries (I/O, MPI, Memory, CUDA, OpenCL, pthread)
 - Rewriting the binary executable
- Measurement
 - Direct: interval events, Indirect: collect samples to profile statement execution
 - Per-thread storage of performance data
 - Throttling and runtime control of low-level events



How does TAU work?

- Analysis
 - TAU creates one profile file per node in a single location
 - Profile file names look like,
profile.0.0.0, profile.1.0.0, ...
 - 2D and 3D visualization of profile data using pprof and **paraprof**
 - Trace conversion & display in external visualizers such as **Jumpshot**



TAU Event Types

- Interval: start-stop events (eg: function call)
- Atomic: trigger at a single point with data (eg: memory allocation)
 - Measures total, samples, min/max/mean/std. deviation statistics
- Context: atomic events with executing context
 - Measures total, samples, min/max/mean/std. deviation statistics



TAU event types

profile.0.0.0

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	1:18.355	1:18.561	1	1818	78561006	.TAU application
0.3	202	202	1814	0	112	read()
0.0	3	3	2	0	1607	open()
0.0	0.004	0.004	2	0	2	lseek()

Interval events

USER EVENTS: profile.0.0.0

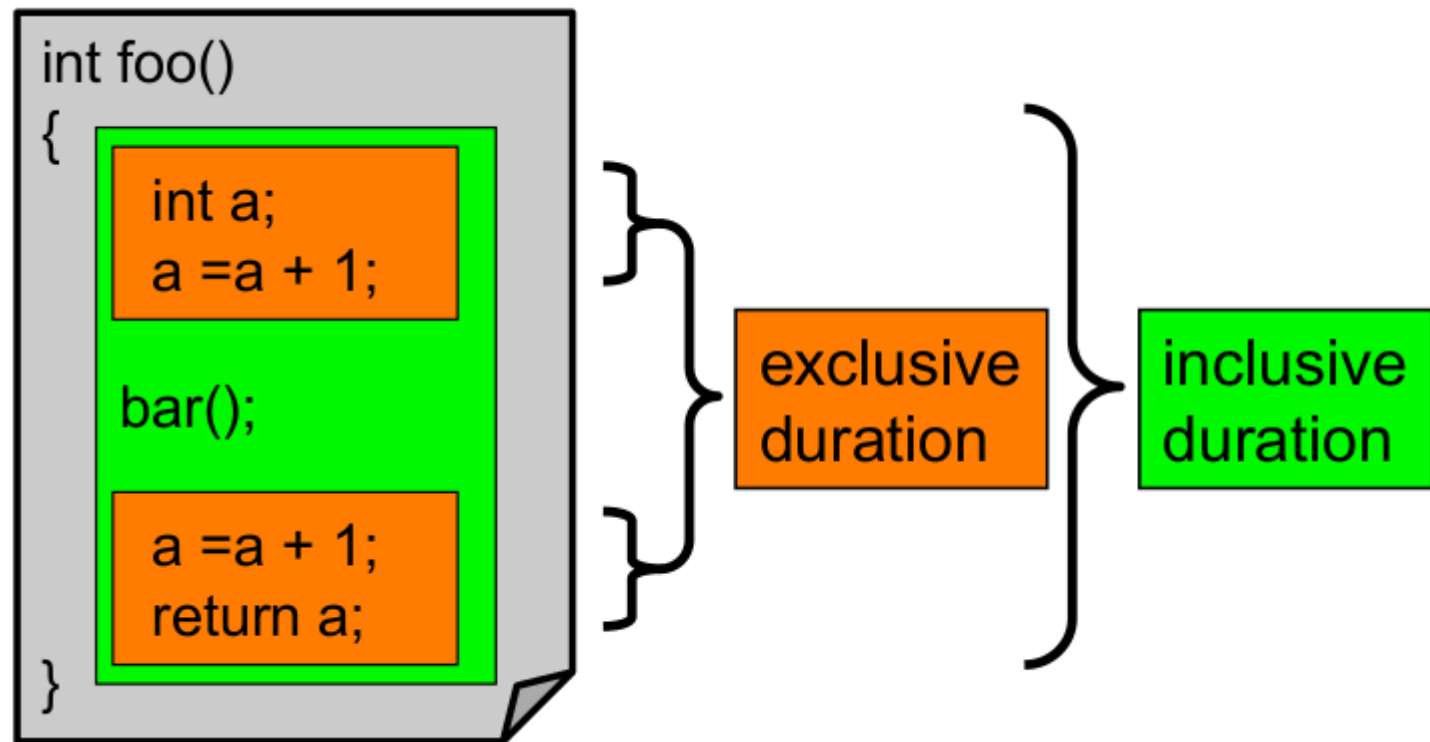
NumSamples	MaxValue	MinValue	MeanValue	Std. Dev.	Event Name
1812	8192	2174	8186	179.6	Bytes Read
1812	8192	2174	8186	179.6	Bytes Read : .TAU application => read()
906	8192	2174	8185	199.8	Bytes Read <file=data1.dat> : .TAU application => read()
906	8192	3467	8187	156.9	Bytes Read <file=data2.dat> : .TAU application => read()
1812	1170	0.113	913.9	124.7	Read Bandwidth (MB/s)

Atomic event

Context events



Exclusive vs Inclusive time





TAU at RCC

- Currently available on HPC for all six compilers
 - GNU, Intel, PGI - OpenMPI and MVAPICH2
- To use serial version

```
module load tau-serial
```

- To use parallel version

```
module load tau
```

- Documentation: <https://rcc.fsu.edu/software/tau>



TAU Instrumentation

- Library interposition (dynamic instrumentation)
 - No need to recompile your code
 - `mpirun -np 4 tau_exec <options> <your binary>`
 - Can profile MPI (default), memory use, I/O, ...
 - Cannot track user functions

Still buggy!

```
$ gfortran -o gauss gauss.f90
$ module load tau-serial
$ tau_exec -T serial -io ./gauss
$ pprof -s
```

Reading Profile files in profile.0.0.0.*

FUNCTION SUMMARY (total):

%Time	Exclusive Msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	1:16.607	1:16.627	1	1818	76627147	.TAU application
0.0	19	19	1814	0	11	read()
0.0	0.026	0.026	2	0	13	open()
0.0	0.001	0.001	2	0	0	lseek()



TAU Instrumentation

- Scripted Compilation
 - Use `tau_f90.sh`, `tau_cc.sh`, and `tau_cxx.sh` to instrument and compile Fortran, C, and C++ programs
 - Compiler Based Instrumentation
 - Use the compiler itself for instrumenting
 - Provides more detailed profiles than dynamic approach
 - Cannot profile user functions
 - Needs recompilation of the code

```
$ tau_cc.sh -tau_options=-optCompInst samplecprogram.c
```



TAU Instrumentation

- Source based instrumentation
 - Uses PDT (Program Database Toolkit) to fully instrument the source code
 - Able to generate complete profiles by measuring low level events (loops, hardware counters, etc.)
 - Needs recompilation of the code (Simply switch **CC** or **FC** with **tau_cc.sh** or **tau_f90.sh**)

```
$ module load gnu-openmpi
$ module load tau
$ tau_f90.sh -o mat_mul_par mat_mul_par.f90
$ msub mat_mul_par.sh
$ pprof
```



pprof

Reading Profile files in profile.*

NODE 0;CONTEXT 0;THREAD 0:

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	0.139	55,725	1	1	55725516	.TAU application
100.0	26,947	55,725	1	7	55725377	MAT_MUL_PAR
49.5	27,590	27,590	1	0	27590687	MPI_Gather()
1.0	541	541	1	0	541913	MPI_Init()
0.9	488	488	1	0	488823	MPI_Bcast()
0.2	94	94	1	0	94155	MPI_Scatter()
0.1	62	62	1	0	62281	MPI_Finalize()
0.0	0.001	0.001	1	0	1	MPI_Comm_size()
0.0	0	0	1	0	0	MPI_Comm_rank()


USER EVENTS Profile :NODE 0, CONTEXT 0, THREAD 0

NumSamples	MaxValue	MinValue	MeanValue	Std. Dev.	Event Name
1	6.4E+07	6.4E+07	6.4E+07	0	Message size for broadcast
1	4E+06	4E+06	4E+06	0	Message size for gather
1	4E+06	4E+06	4E+06	0	Message size for scatter




paraprof

Metric: TIME
Value: Exclusive



Poor
load balancing!

Metric: TIME
Value: Exclusive
Units: seconds



How much time
spent on
each operation

Only MPI calls
are profiled



Source based Instrumentation

- There is more...
 - The TAU module picks a “Makefile” for you, depending on the compiler you are using
 - It is stored in the variable **TAU_MAKEFILE**
 - Eg: Default for gnu-openmpi is **Makefile.tau-papi-mpi-pdt**
 - Makefiles can be changed by user depending on the purpose

```
Makefile.tau-communicators-papi-mpi-pdt
Makefile.tau-headroom-papi-mpi-pdt
Makefile.tau-memory-papi-mpi-pdt
Makefile.tau-papi-mpi-pdt
Makefile.tau-papi-mpi-pdt-trace
Makefile.tau-phase-papi-mpi-pdt
```



Using Different Makefiles

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	0.012	1:06.231	1	1	66231389	.TAU application
100.0	2,021	1:06.231	1	1	66231377	GAUSS
96.9	19,828	1:04.209	1	6000	64209853	GAUSSJ
45.8	30,348	30,348	4000	0	7587	OUTERPROD
21.2	14,023	14,023	1000	0	14023	OUTERAND
0.0	9	9	1000	0	10	SWAP

USER EVENTS: profile.-1.0.0

NumSamples	MaxValue	MinValue	MeanValue	Std. Dev.	Event Name
1	0	0	0	0	.TAU application - Heap Memory Used (KB)
1	0	0	0	0	GAUSS - Heap Memory Used (KB)
1	0	0	0	0	GAUSSJ - Heap Memory Used (KB)
1000	0	0	0	0	OUTERAND - Heap Memory Used (KB)
4000	0	0	0	0	OUTERPROD - Heap Memory Used (KB)
1000	0	0	0	0	SWAP - Heap Memory Used (KB)



Selective Instrumentation

- Not all functions need to be profiled in large applications

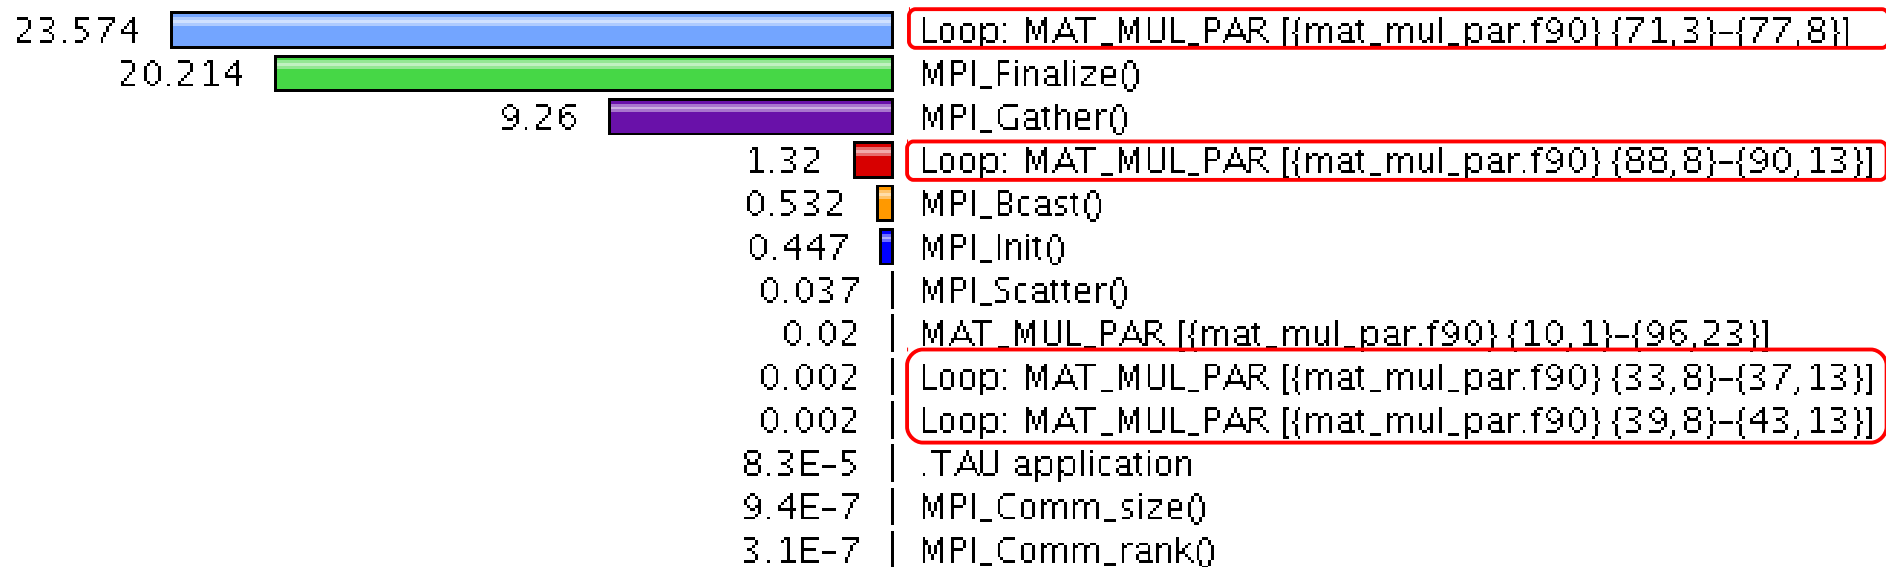
```
export TAU_OPTIONS="-optTauSelectFile=select.tau"  
  
cat select.tau  
BEGIN_INSTRUMENT_SECTION  
loops file="mat_mul_par.f90" routine="#"  
END_INSTRUMENT_SECTION
```

Only need to profile outer loops of the given file



Selective Instrumentation

Metric: TIME
Value: Exclusive
Units: seconds





Selective Instrumentation

```
BEGIN_EXCLUDE_LIST
void quicksort(int *, int, int)
# The next line excludes all functions beginning with "sort_"
# and having arguments "int *"
void sort_#(int *)
void interchange(int *, int *)
END_EXCLUDE_LIST
#Exclude these files from profiling
BEGIN_FILE_EXCLUDE_LIST
*.so
END_FILE_EXCLUDE_LIST
BEGIN_INSTRUMENT_SECTION
# instrument all the outer loops in this routine
loops file="loop_test.cpp" routine="multiply"
# tracks memory allocations/deallocations as well as
# potential leaks
memory file="foo.f90" routine="INIT"
# tracks the size of read, write and print statements in
# this routine
io file="foo.f90" routine="RINB"
```



Using Optional TAU Compiler Options

- By setting **TAU_OPTIONS** variable or directly using TAU compiler options while compiling will change its behavior
 - **-optTrackIO** will profile I/O operations
 - **-optHeaderInst** will enable instrumentation of headers
 - For a full list, use **tau_compiler.sh -help** command




Runtime Environment Variables

Environment Variable	Default	Description
TAU_PROFILE	1	Set 0 to stop profiling (eg: for tracing)
PROFILEDIR	.	Location for profile files
TAU_TRACE	0	Set 1 for tracing
TAU_TRACK_HEAP TAU_TRACK_HEADROOM	0	Set 1 to track heap memory or headroom available
TAU_CALLPATH	0	Set 1 to start callpath profiling
TAU_COMM_MATRIX	0	Set 1 to generate communication matrix data
TAU_COMPENSATE	0	Set 1 to compensate instrumentation overhead
TAU_THROTTLE	1	Skip instrumenting functions called frequently
TRACEDIR	.	Location for tracing data




Real World Examples






Call Path Graph






Communication Matrix






How and what each node is doing?





Flat profile of a real world case

Metric: TIME
Value: Exclusive
Units: seconds



MPI_Send and MPI_Wait seems to be the culprit for slowdown
But why?
How do we find?



Hardware Counters

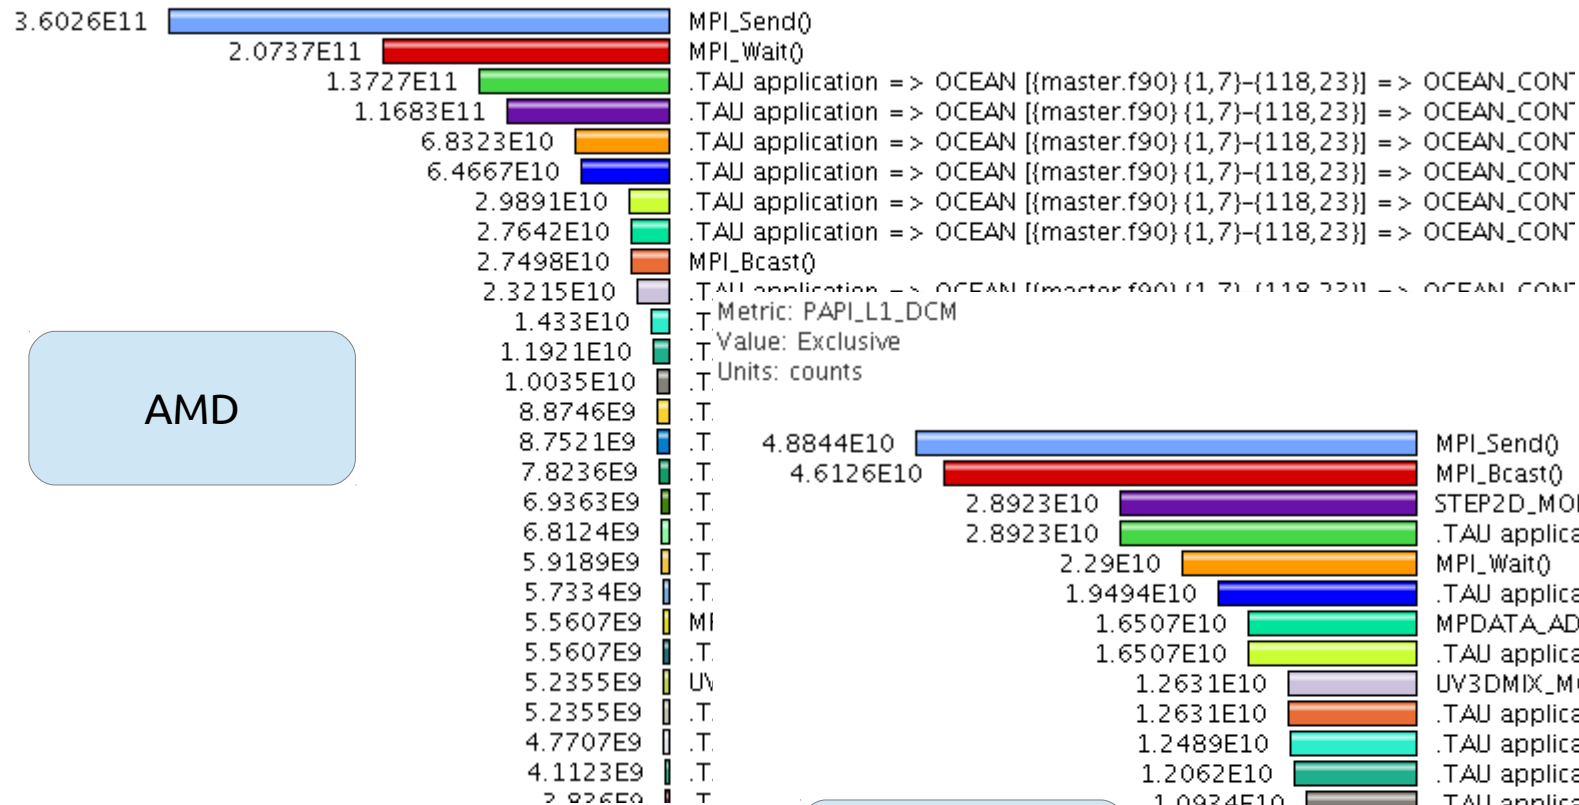
- TAU allows integration with other tools such as PAPI (Performance API)
- PAPI is installed on every HPC node and can be used to instrument a code using hardware counters as the metric
- **papi_avail** command will give you a complete list of available hardware counters on a specific node

```
export COUNTER1=GET_TIME_OF_DAY #To measure runtime
export COUNTER2=PAPI_L1_DCM #To find level 1 cache miss
export COUNTER3=PAPI_L2_DCM #To find level 2 cache miss
export COUNTER4=PAPI_FLOPS #To measure FLOPS
```




Hardware Counters

Metric: PAPI_L1_DCM
Value: Exclusive
Units: counts



AMD



Intel



Measuring FLOPS

MULTI__PAPI_FLOPS/profile.1.0.0

%Time	Exclusive	Inclusive	#Call	#Subrs	Count/Call	Name
	Count	total counts				
100.0	7.366E+06	5.984E+07	1	2	59837472	MATMUL_CACHE
43.8	7739	2.624E+07	1	1	26235629	CACHE_MISS
43.8	8066	2.624E+07	1	1	26235546	CACHE_NO_MISS
43.8	2.623E+07	2.623E+07	1	0	26227890	Loop: CACHE_MISS
43.8	2.623E+07	2.623E+07	1	0	26227480	Loop: CACHE_NO_MISS

MULTI__GET_TIME_OF_DAY/profile.-1.0.0

%Time	Exclusive	Inclusive	#Call	#Subrs	Inclusive	Name
	msec	total msec			usec/call	
100.0	2	79	1	2	79073	MATMUL_CACHE
88.6	0.046	70	1	1	70078	CACHE_MISS
88.6	70	70	1	0	70032	Loop: CACHE_MISS
7.9	0.02	6	1	1	6230	CACHE_NO_MISS
7.9	6	6	1	0	6210	Loop: CACHE_NO_MISS

CACHE_MISS 375 MFLOPS

CACHE_NO_MISS 4.4 GFLOPS



```
program matmul_cache
```

```
    !Just calling the following functions here  
end program matmul_cache
```

```
real function cache_miss(a, b, n)
```

```
    integer :: i, j
```

```
    real :: a(1024,1024), b(1024,1024)
```

```
    do i = 1, n
```

```
        do j = 1, n
```

```
            a(i,j) = b(i,j)
```

```
        enddo
```

```
    enddo
```

```
end function cache_miss
```

```
real function cache_no_miss(a, b, n)
```

```
    integer :: i, j
```

```
    real :: a(1024,1024), b(1024,1024)
```

```
    do j = 1, n
```

```
        do i = 1, n
```

```
            a(i,j) = b(i,j)
```

```
        enddo
```

```
    enddo
```

```
end function cache_no_miss
```




Hardware Counters

- To measure more counters at the same time,

```
export TAU_METRICS=TIME:PAPI_FP_INS:PAPI_L1_DCM
```

- Each counter will generate one profile in separate subdirectories that look like,

```
MULTI__GET_TIME_OF_DAY, MULTI__PAPI_L1_DCM
```

- Intel and AMD nodes have different counters
- Up to 25 counters/events can then be recorded at a time



Tracing

- What happens in my code at a given time? When?
- Use **Jumpshot** to visualize results
- Significant overhead. Turn off profiling!

```
$ export TAU_MAKEFILE=/panfs/storage.local/\
> opt/hpc/gnu/openmpi/tau/x86_64/lib/\
> Makefile.tau-papi-mpi-pdt-trace
And compile your code, run
After job is finished, cd to TRACEDIR
$ tau_treemerge.pl
$ tau2slog2 tau.trc tau.edf -o tau.slog2
$ jumpshot tau.slog2
```



Tracing – Jumpshot view

