Shared-memory Parallel Programming with Cilk Plus (Parts 2-3)

John Mellor-Crummey

Department of Computer Science
Rice University

johnmc@rice.edu

Last Thursday

- **Threaded programming models**
- **Introduction to Cilk Plus**
  - —tasks
  - —algorithmic complexity measures
  - —scheduling
  - —performance and granularity
  - —task parallelism examples
    - – vector addition using divide and conquer
    - – nqueens: exploratory search

Outline for Today

- Cilk Plus
  - —explore speedup and granularity
  - —task parallelism example
    - – cilksort
  - —parallel loops
  - —reducers
- Data race detection with cilkscreen
- Assessing Cilk Plus performance with cilkview

Cilk Plus Adoption is Growing

- **Intel icpc supports the full Cilk Plus language as well as the cilkscreen and cilkview tools**
- **GCC 4.8-cilkplus branch supports cilkscreen and cilkview, but lacks support for parallel loops**
- **GCC 5 and later support the full Cilk Plus language, but do not support cilkscreen or cilkview**
- **Clang supports the full Cilk Plus language, but does not support cilkscreen and cilkview**

For your assignment you will use the Intel compiler, which supports spawn, parallel loops, reducers, the Cilk Plus vector notation, along with the Cilk tools.

Review: Cilk Plus Parallel Performance Model

$$c_1 = \frac{T_1}{T_s}$$ **work overhead**

$$T_p \le c_1 \frac{T_s}{P} + c_\infty T_\infty$$

"Minimize work overhead (c_1) at the expense of a larger critical path overhead (c_∞), because work overhead has a more direct impact on performance"

$$T_p \approx c_1 \frac{T_s}{P}$$ **assuming parallel slackness**

Speedup Demo

- Explore speedup of naive fibonacci program
  - —`cp -r /projects/comp422/cilkplus-examples/fib ~/fib`
  - —`cd ~/fib`
  - —fib.cpp: a program for computing n^{th} fibonacci #
  - —build the examples: `make`
  - —experiment with the fibonacci program
    - – `make runt W=n` *computes fib(41) with n workers*
    - – compute fib(41) for different values of W, 1 ≤ W ≤ 12
    - – what value of W yields the lowest execution time?
    - – what is the speedup vs. the execution time of "`./fib-serial 41`"?
    - – how does this speedup compare to the total number of HW threads?

Granularity Demo

- **Explore how changing increasing the granularity of parallel work in fib improves performance (by reducing c_1)**
  - **—fib-trunc.cpp: a program for computing n^th fibonacci #**
    - **this version differs in that one can execute subtrees of height H sequentially rather than spawning parallel tasks all the way down**
  - **—build the examples:** `make`
  - **—experiment with the fibonacci program with truncated parallelism**
    - `make runt H=h` *computes fib(41) with lowest H levels serial*
    - **compute fib(41) for different values of H, 2 ≤ H ≤ 41**
    - **what value of H yields the lowest execution time**
    - **what is the speedup vs. the execution time of "`./fib-serial 41`"?**
    - **how does this speedup compare to the total number of HW threads?**

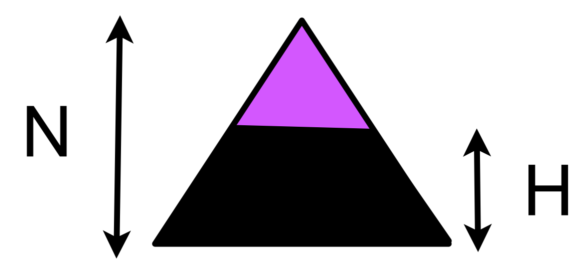

Cilksort Sketch

Variant of merge sort

```
cilksort(in[1..n], out[1..n]) =
   spawn cilksort(in[1..n/2], tmp[1..n/2])
   spawn cilksort(in[n/2..n], tmp[n/2..n])
   sync
   spawn cilkmerge(tmp[1..n/2], tmp[n/2..n], out[1..n])
```

Cilkmerge Sketch

```
cilkmerge(A[1..n], B[1..m], C[1..(n+m)]) =
    // Find the median of A union B using binary  search.
    // The binary search gives a pair (ma, mb) such that
    // ma + mb = (n + m)/2 and all elements in A[1..ma] are
    // smaller than B[mb..m], and all the B[1..mb] are smaller
    // than all elements in A[ma..n].
  spawn cilkmerge(A[1..ma], B[1..mb], C[1..(n+m)/2])
  spawn cilkmerge(A[ma..m], B[mb..n], C[(n+m)/2 .. (n+m)])
```

Median finding idea: S. G. Akl and N. Santoro, "Optimal Parallel Merging and Sorting Without Memory Conflicts", IEEE Trans. Comp., Vol. C-36 No. 11, Nov. 1987

Optimizing Performance of cilksort

- Recursively subdividing all the way to singletons is <u>expensive</u>
- When size(remaining sequence) to sort or merge is small (2K)
  - use sequential quicksort
  - use sequential merge

Cilksort in Practice

```
void cilksort(ELM *low, ELM *tmp, long size) {
    long quarter = size / 4;
    ELM *A, *B, *C, *D, *tmpA, *tmpB, *tmpC, *tmpD;
    if (size < QUICKSIZE) { seqquick(low, low + size - 1) return; }

    A = low; tmpA = tmp;
    B = A + quarter; tmpB = tmpA + quarter;
    C = B + quarter; tmpC = tmpB + quarter;
    D = C + quarter; tmpD = tmpC + quarter;

    cilk_spawn cilksort(A, tmpA, quarter);
    cilk_spawn cilksort(B, tmpB, quarter);
    cilk_spawn cilksort(C, tmpC, quarter);
    cilk_spawn cilksort(D, tmpD, size - 3 * quarter);
    cilk_sync;

    cilk_spawn cilkmerge(A, A + quarter - 1, B, B + quarter - 1, tmpA);
    cilk_spawn cilkmerge(C, C + quarter - 1, D, low + size - 1, tmpC);
    cilk_sync;

    cilk_spawn cilkmerge(tmpA, tmpC - 1, tmpC, tmpA + size - 1, A);
    cilk_sync;
}
```

base case

4-way split with 2-levels of merge delivers result in place

Cilksort

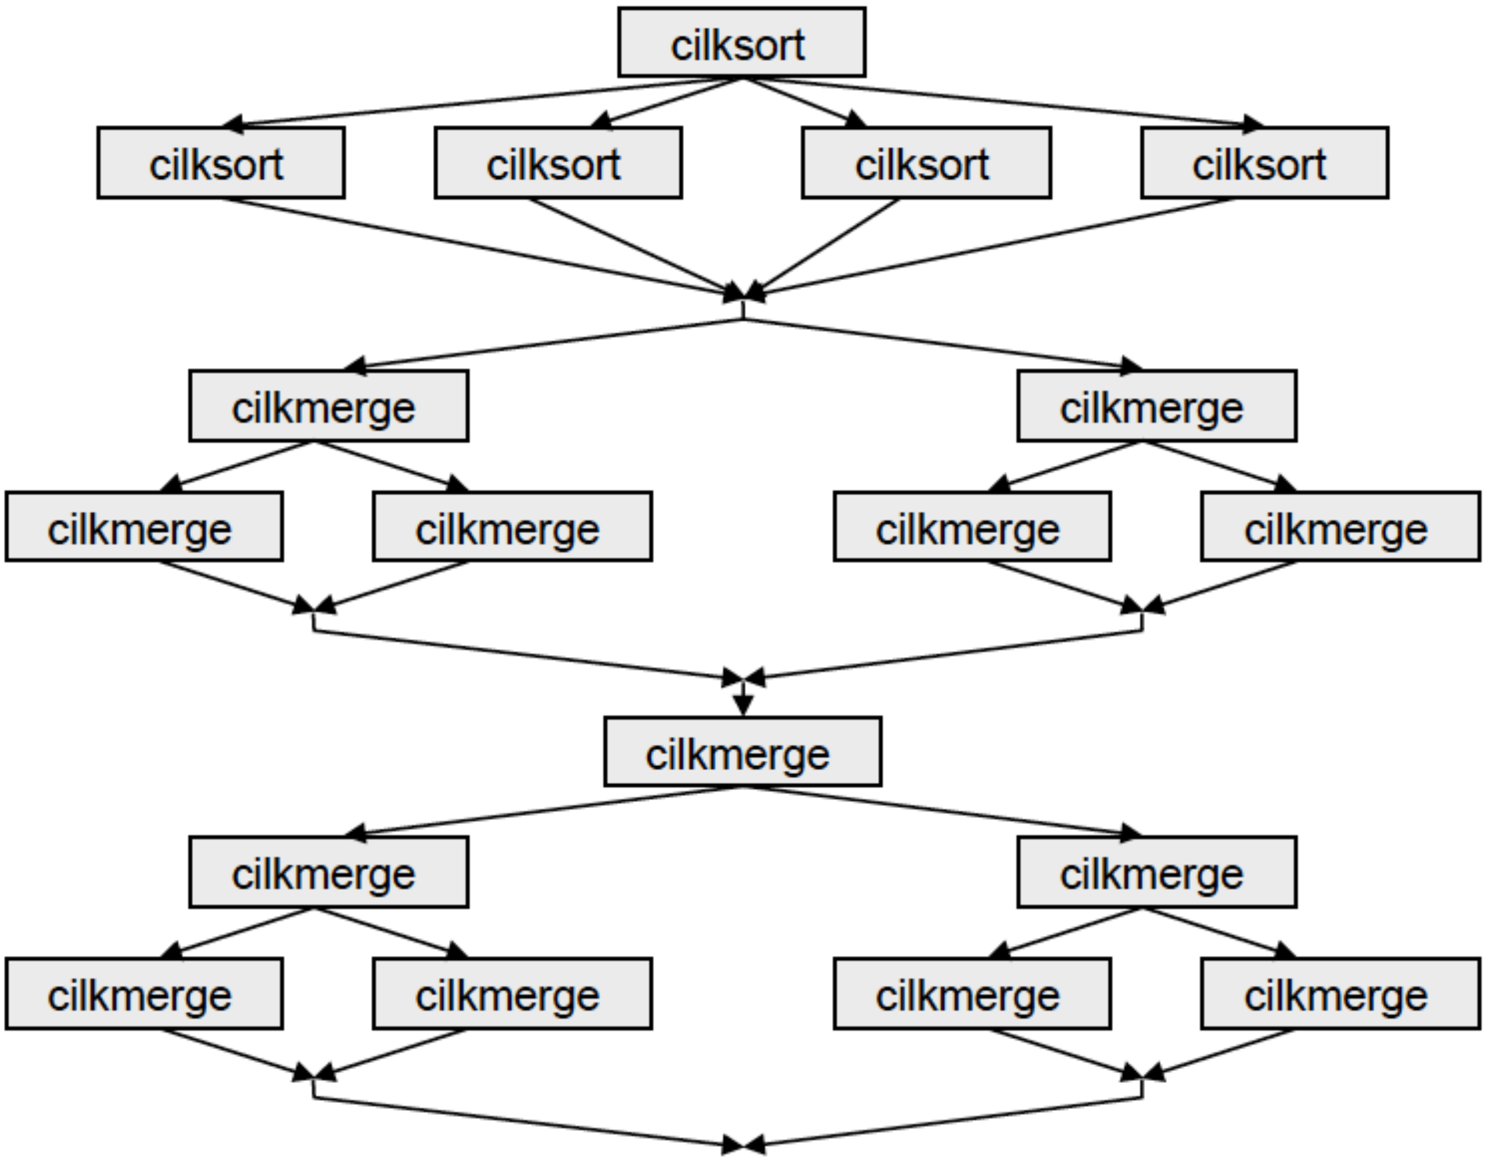

Figure credit: Stephen Olivier et al. “Scheduling Task Parallelism on Multi-Socket Multicore Systems.”
https://htor.inf.ethz.ch/ross2011/slides/ross2011-olivier.pdf

Cilk Plus Parallel Loop: cilk_for

```
cilk_for (T v = begin; v < end; v++) {
      statement_1;
      statement_2;
      ...
}
```

- Loop index v
  - type T can be an integer, ptr, or a *C++ random access iterator*
- Main restrictions
  - runtime must be able to compute total # of iterations on entry to cilk_for
    - must compare v with end value using <, <=, !=, >=, or >
    - loop increment must use ++, --, +=, v = v + incr, or v = v - incr

      if v is not a signed integer, loop must count up
- Implicit cilk_sync at the end of a cilk_for

Loop with a cilk_spawn vs. cilk_for

- for (int i = 0; i < 8; i++) { cilk_spawn work(i); } cilk_sync;

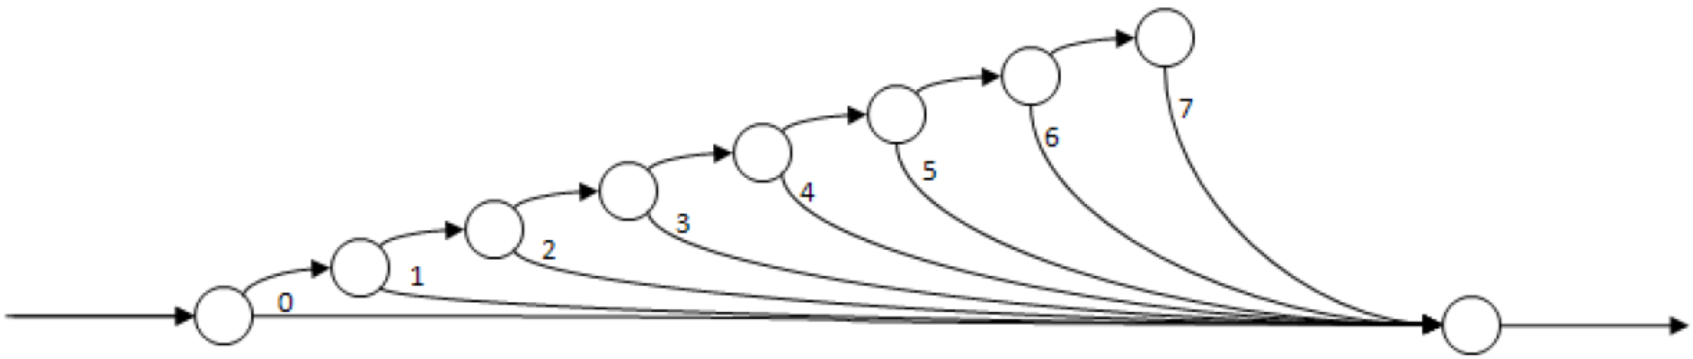

- cilk_for (int i = 0; i < 8; i++) { work(i);}

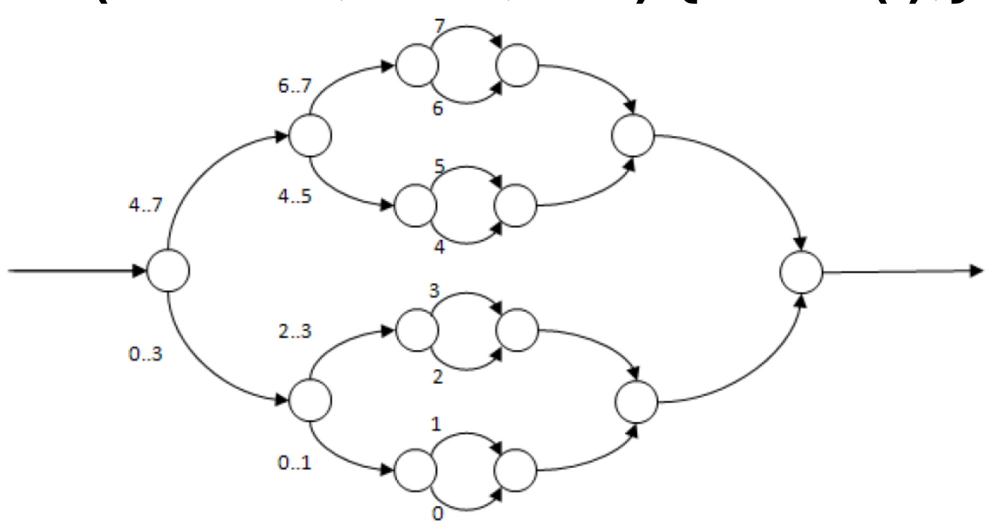

Note: computation on edges

cilk_for uses divide-and-conquer

Figure credits: Intel Cilk++ Programmer's Guide. Document # 322581-001US.

Some Restrictions for cilk_for

- Value added to or subtracted subtracted from loop control variable must not change from one iteration to the next
- No early exit
  - —no break or return statement within loop
  - —no goto in loop unless target is within loop body
- Loop induction variable restrictions
  - —cilk_for (unsigned int i, j = 42; j < 1; i++, j++) { ... }
    - – only one loop variable allowed
  - —cilk_for (unsigned int i = 1; i < 16; ++i) i = f();
    - – can't modify loop variable within loop
  - —cilk_for (unsigned int i = 1; i < x; ++i) x = f();
    - – can't modify end within loop

other restrictions in the manual

- Note: in Cilk Plus, the loop induction variable need not be declared in the loop header; in Cilk++, this was a requirement

cilk_for Implementation Sketch

- **Recursive bisection used to subdivide iteration space down to chunk size**

```
void run_loop(first, last)
{

    if (last - first) < grainsize)
    {
        for (int i=first; i<last ++i) LOOP_BODY;
    }
    else
    {
        int mid = (last-first)/2;
        cilk_spawn run_loop(first, mid);
        run_loop(mid, last);
    }
}
```

cilk_for Grain Size

- Iterations divided into *chunks* to be executed serially
  - chunk is sequential collection of one or more iterations
- Maximum size of chunk is called *grain size*
  - grain size too small: spawn overhead reduces performance
  - grain size too large: reduces parallelism and load balance
- Default grain size
  - #pragma cilk grainsize = min(2048, N / (8*p))
- Can override default grain size
  - #pragma cilk grainsize = expr
    - expr is any C++ expression that yields an integral type (e.g. int, long)

      e.g. #pragma cilk grainsize = n/(4*__cilkrts_get_nworkers())
  - pragma must immediately precede cilk_for to which it applies

Parallelizing Vector Addition

C

```
void vadd (real *A, real *B, int n){
  int i; for (i=0; i<n; i++) A[i]+=B[i];
}
```

Cilk Plus

```
void vadd (real *A, real *B, int n){
  if (n<=BASE) {
    int i; for (i=0; i<n; i++) A[i]+=B[i];
  } else {
    cilk_spawn vadd (A, B, n/2);
    vadd (A+n/2, B+n/2, n-n/2);
    cilk_sync;
  }
}
```

```
void vadd (real *A, real *B, int n){
  int i; cilk_for (i=0; i<n; i++) A[i]+=B[i];
}
```

The Problem with Non-local Variables

- **Nonlocal variables are a common programming construct**
  - global variables = nonlocal variables in outermost scope
  - nonlocal = declared in a scope outside that where it is used
- **Example**

```
int sum = 0;
for(int i=1; i<n; i++) {
  sum += i;
}
```

- **Rewriting parallel applications to avoid them is painful**

Collision Detection

Automaker: hierarchical 3D CAD representation of assemblies

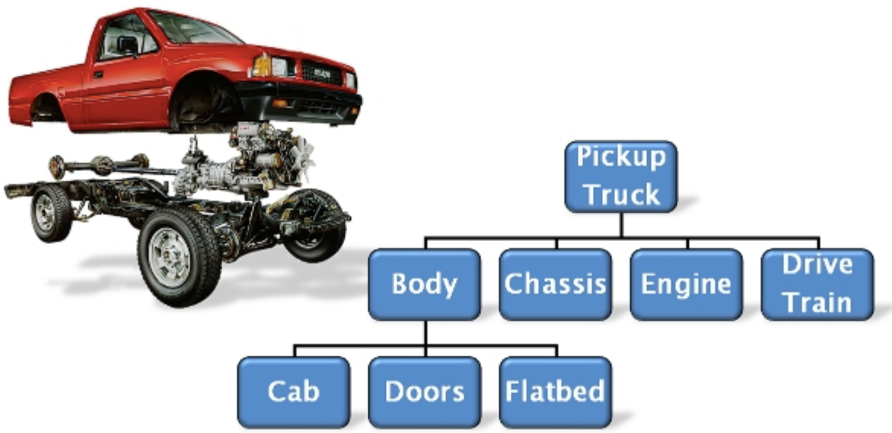

Computing a cutaway view

```
Node *target;
std::list<Node *> output_list;
...
void walk(Node *x)  {
  switch (x->kind) {
  case Node::LEAF:
    if (target->collides_with(x))
        output_list.push_back(x);
    break;
  case Node::INTERNAL:
    for (Node::const_iterator
                child = x->begin();
                child != x->end();
                ++child)
        walk(child);
    break;
  }
}
```

Adding Cilk Plus Parallelism

Computing a cutaway view in parallel

```
Node *target;
std::list<Node *> output_list;
...
void walk(Node *x)  {
  switch (x->kind) {
  case Node::LEAF:
    if (target->collides_with(x))
        output_list.push_back(x);
    break;
  case Node::INTERNAL:
    cilk_for (Node::const_iterator
              child = x->begin();
              child != x->end();
              ++child)
        walk(child);
    break;
  }
}
```

Global variable causes data races!

Solution 1: Locking

Computing a cutaway view in parallel

```
Node *target;
std::list<Node *> output_list;
mutex m;
...
void walk(Node *x)  {
  switch (x->kind) {
  case Node::LEAF:
    if (target->collides_with(x))
    { m.lock(); output_list.push_back(x); m.unlock(); }
    break;
  case Node::INTERNAL:
    cilk_for (Node::const_iterator
              child = x->begin();
              child != x->end();
              ++child)
      walk(child);
    break;
  }
}
```

- **Add a mutex to coordinate accesses to output_list**
- **Drawback: lock contention can hurt parallelism**

Solution 2: Refactor the Code

```
Node *target;
        std::list<Node *> output_list;
        ...
        void walk(Node *x, std::list<Node *> &o_list) {
          switch (x->kind) {
          case Node::LEAF:
            if (target->collides_with(x))
                o_list.push_back(x);
            break;
          case Node::INTERNAL:
            std::vector<std::list<Node *>>
                  child_list(x.num_children);
            cilk_for (Node::const_iterator
                        child = x->begin();
                        child != x->end();
                        ++child)
                walk(child, child_list[child]);
            for (int i=0; i < x.num_children; ++i)
                o_list.splice(o_list.end(), child_list[i]);
            break;
          }
```

- **Have each child accumulate results in a separate list**
- **Splice them all together**
- **Drawback: development time, debugging**

Solution 3: Cilk Plus Reducers

```
Node *target;
cilk::reducer_list_append<Node *> output_list;
        ...
        void walk(Node *x) {
          switch (x->kind) {
          case Node::LEAF:
            if (target->collides_with(x))
              output_list.push_back(x);
            break;
          case Node::INTERNAL:
            cilk_for (Node::const_iterator
                       child = x->begin();
                       child != x->end();
                       ++child)q
               walk(child);
            break;
          }
        }
```

- **Resolve data races without locking or refactoring**
- **Parallel strands may see different views of reducer, but these views are combined into a single consistent view**

Cilk Plus Reducers

- Support update of nonlocal variables without races
  - —deterministic update using associative operations
    - – e.g., global sum, list and output stream append, ...
    - – result using is same as serial version
      - independent of # processors or scheduling
- Can be used without significant code restructuring
- Can be used independently of the program's control structure
  - — unlike constructs defined only over loops
- Implemented efficiently with minimal overhead
  - —they don't use locks in their implementation
    - – avoids loss of parallelism from enforcing mutual exclusion when updating shared variables

Reducers

Serial execution (depth first):

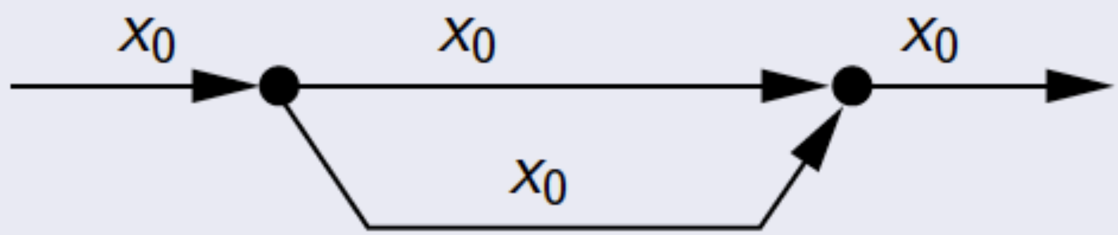

Parallel execution:

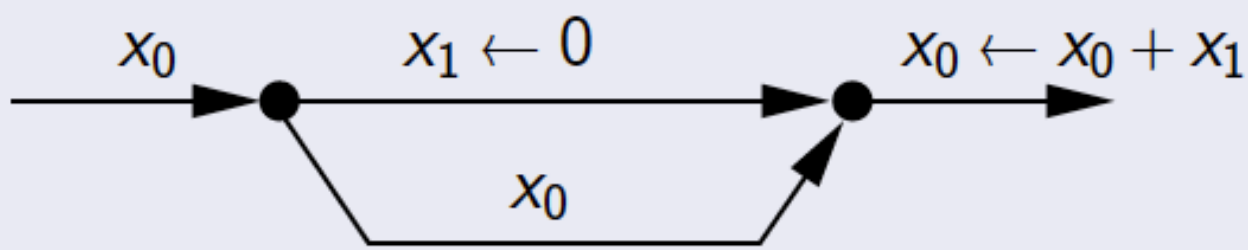

Matteo Frigo, Pablo Halpern, Charles E. Leiserson, Stephen Lewin-Berlin, Reducers and other Cilk++ hyperobjects. Slides for *SPAA'09,* August 11–13, 2009, Calgary, Alberta, Canada.

Reducing Over List Concatenation

Program:

```
x.append(0);
cilk_spawn x.append(1);
x.append(2);
x.append(3);
cilk_sync;
```

Serial execution:

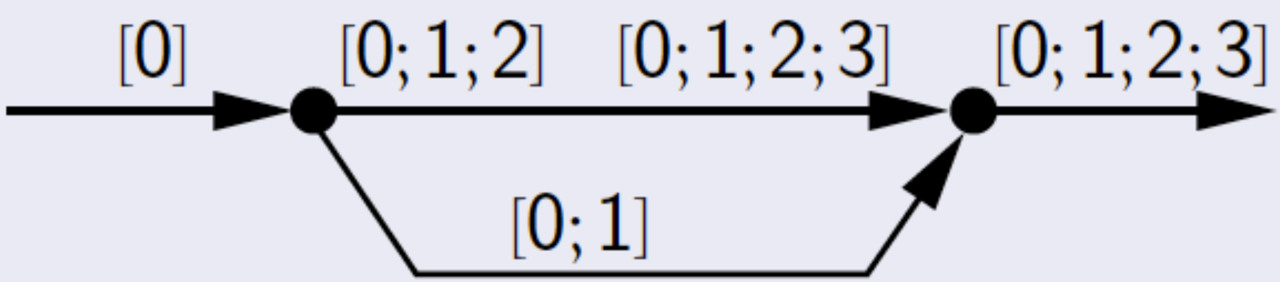

Matteo Frigo, Pablo Halpern, Charles E. Leiserson, Stephen Lewin-Berlin, Reducers and other Cilk++ hyperobjects. Slides for *SPAA'09,* August 11–13, 2009, Calgary, Alberta, Canada.

Reducing Over List Concatenation

Program:

```
x.append(0);
cilk_spawn x.append(1);
x.append(2);
x.append(3);
cilk_sync;
```

Parallel execution:

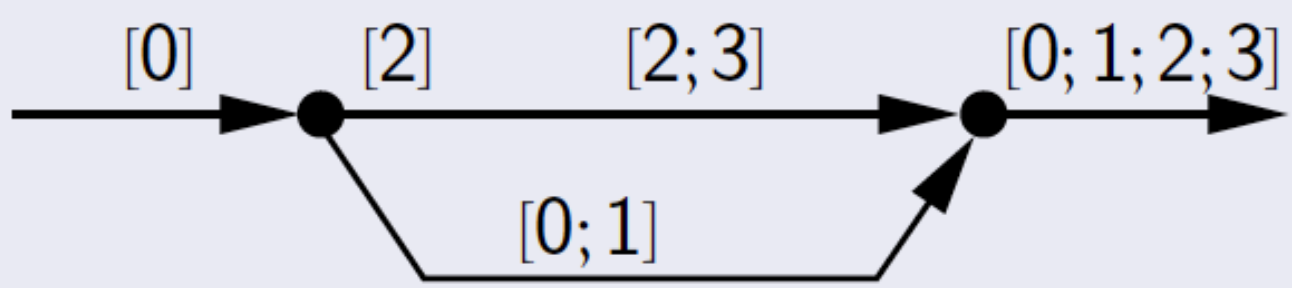

Matteo Frigo, Pablo Halpern, Charles E. Leiserson, Stephen Lewin-Berlin, Reducers and other Cilk++ hyperobjects. Slides for *SPAA'09,* August 11–13, 2009, Calgary, Alberta, Canada.

Cilk++ Reducers Under the Hood

- **If no steal occurs, a reducer behaves like a normal variable**
- **If a steal occurs**
  - **the continuation receives a view with an identity value**
  - **the child receives the reducer as it was prior to the spawn**
  - **at the corresponding cilk_sync**
    - **the value in the continuation is merged into the reducer held by the child using the reducer's `reduce` operation**
    - **the new view is destroyed**
    - **the original (updated) object survives**

Using Cilk Plus Reducers

- **Include the appropriate Cilk Plus reducer header file**

  `reducer_list.h, reducer_max.h reducer_min.h, reducer_opadd.h reducer_opand.h, reducer_opor.h, reducer_opxor.h reducer_ostream.h, reducer_string.h, reducer_vector.h`

- **Declare a variable as a reducer rather than a standard type**
  - global sum
    - `cilk::reducer_opadd<unsigned long> sum`
  - list reducer
    - instead of "`std::list<int> sequence`", use `cilk::reducer_list_append<int> sequence`

- **Use reducers in the midst of work that includes parallelism created with `cilk_spawn` or `cilk_for`**

- **Retrieve the reducer's terminal value with `var.get_value()` after the parallel updates to the reducer are complete**

Reducer Demo - I

- See /projects/comp422/cilkplus-examples/sum
- Compare a program with a racing reduction, a mutex protecting the race, and a reducer
- Versions:
  - race.cpp: code with a racing sum reduction
  - lock.cpp: code with a mutex to avoid the race
  - reducer.cpp: code with a reducer to avoid the race
- Compare performance of the various versions
  - ./race 100000000
  - ./lock 100000000
  - ./reducer 100000000
  - how does the performance of the parallel summation using reducers compare to
    - parallel summation with races?
    - parallel summation with locks?
    - serial summation?

Reducer Demo - II

- See /projects/comp422/cilkplus-examples/order/order.cpp
- order.cpp is a program that contains
  - a parallel loop where iterations race to write output
  - a parallel loop where iterations write output using an ostream reducer
  - a divide and conquer computation where recursive tasks write using an ostream reducer
- See how the output differs as parallel work is mapped to cores using work stealing

Concurrency Cautions

- **Only limited guarantees between descendants or ancestors**
  - —**DAG precedence order maintained and nothing more**
  - —**don’t assume atomicity between different procedures!**

Race Conditions

- **Data race**
  - —two parallel strands access the same data
  - —at least one access is a write
  - —no locks held in common
- **General determinacy race**
  - —two parallel strands access the same data
  - —at least one access is a write
  - —a common lock protects both accesses

A Data Race Example

- Example

```
int sum = 0;
cilk_for(int i=1; i<n; i++) {
  sum += i;
}
```

- What can go wrong?
  - concurrent reads and writes can interleave in unpredictable ways

time ↓

read sum
read sum
write sum + i_j
write sum + i_k

legend
thread n
thread m

  - the update by thread m is lost!

Cilkscreen

- **Detects and reports <u>data races</u> at runtime**
  - **—finds all data races even those by third-party or system libraries**
- **Does not report determinacy races**
  - **—e.g. two concurrent strands use a lock to access a queue**
    - **enqueue & dequeue operations could occur in different order**

      potentially leads to different result

Race Detection Strategies in Cilkscreen

- **Lock covers**
  - —two conflicting accesses to a variable don't race if some lock L is held while each of the accesses is performed by a strand
- **Access precedence**
  - —two conflicting accesses do not race if one must precede the other
    - access A is by a strand X, which precedes the cilk_spawn of strand Y which performs access B
    - access A is performed by strand X, which precedes a cilk_sync that is an ancestor of strand Y

Cilkscreen Race Example

```
#include <stdio.h>
#include "mutex.h"

long sum = 0;
mutex m;

#ifdef SYNCH
#define LOCK m.lock()
#define UNLOCK m.unlock()
#else
#define LOCK
#define UNLOCK
#endif
```

```
void do_accum(int l, int u)
{
        if (u == l) { LOCK; sum += l; UNLOCK; }
        else {
          int mid = (u+l)/2;
          cilk_spawn do_accum(l, mid);
          do_accum(mid+1, u);
        }
}
int main()
{
        do_accum(0, 1000);
        printf("sum = %d\n", sum);

        long ssum = 0;
        for (int i = 0; i <= 1000; i++) ssum +=i;
        printf("serial sum = %d\n", ssum);
}
```

note: mutex class coded using pthread_mutex lock primitives

Cilkscreen Limitations

- Only detects races between Cilk Plus strands
  —depends upon their strict fork/join paradigm
- Only detects races that occur given the input provided
  —does not prove the absence of races for other inputs
  —choose your testing inputs carefully!
- Runs serially, 15-30x slower
- Increases the memory footprint of an application
  —could cause an error if memory demand is too large
- If you build your program with debug information (compile with -g), cilkscreen will associate races with source line numbers

Cilkscreen Output

```
cilkscreen ./race 20000
Cilkscreen Race Detector V2.0.0, Build 4501
summing integers from 0 to 20000

Race condition on location 0x601870
  write access at 0x400c86: (/projects/comp422/cilkplus-examples/races/race.c:22, do_accum+0x1f0)
  read access at 0x400c80: (/projects/comp422/cilkplus-examples/races/race.c:22, do_accum+0x1ea)
    called by 0x400e0a: (/projects/comp422/cilkplus-examples/races/race.c:26, do_accum+0x374)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400fb6: (/projects/comp422/cilkplus-examples/races/race.c:35, main+0x85)

Race condition on location 0x601870
  write access at 0x400c86: (/projects/comp422/cilkplus-examples/races/race.c:22, do_accum+0x1f0)
  write access at 0x400c86: (/projects/comp422/cilkplus-examples/races/race.c:22, do_accum+0x1f0)
    called by 0x400e0a: (/projects/comp422/cilkplus-examples/races/race.c:26, do_accum+0x374)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400d22: (/projects/comp422/cilkplus-examples/races/race.c:25, do_accum+0x28c)
    called by 0x400fb6: (/projects/comp422/cilkplus-examples/races/race.c:35, main+0x85)
sum = 200010000
serial sum = 200010000
2 errors found by Cilkscreen
```

cilkscreen Demo

- **Explore cilkscreen race detection**
  - `cp -r /projects/comp422/cilkplus-examples/races ~/races`
  - `cd ~/races`
  - `programs:`
    - `race.c -`

      `a task parallel summation with a race`

      `race can be suppressed with -DSYNCH using a mutex)`
    - `race2.c -`

      `a cilk_for version with overlapping memset and memcpy operations`
    - `race3.c -`

      `a cilk_for loop that calls a function that mutates a static variable`

Performance Measures

- T_s = serial execution time
- T_1 = execution time on 1 processor (total work), $T_1 \geq T_s$
- T_p = execution time on P processors
- T_∞ = execution time on infinite number of processors
  - longest path in DAG
    - length reflects the cost of computation at nodes along the path
  - known as "critical path length"

Work and Critical Path Example

Procedure oriented view

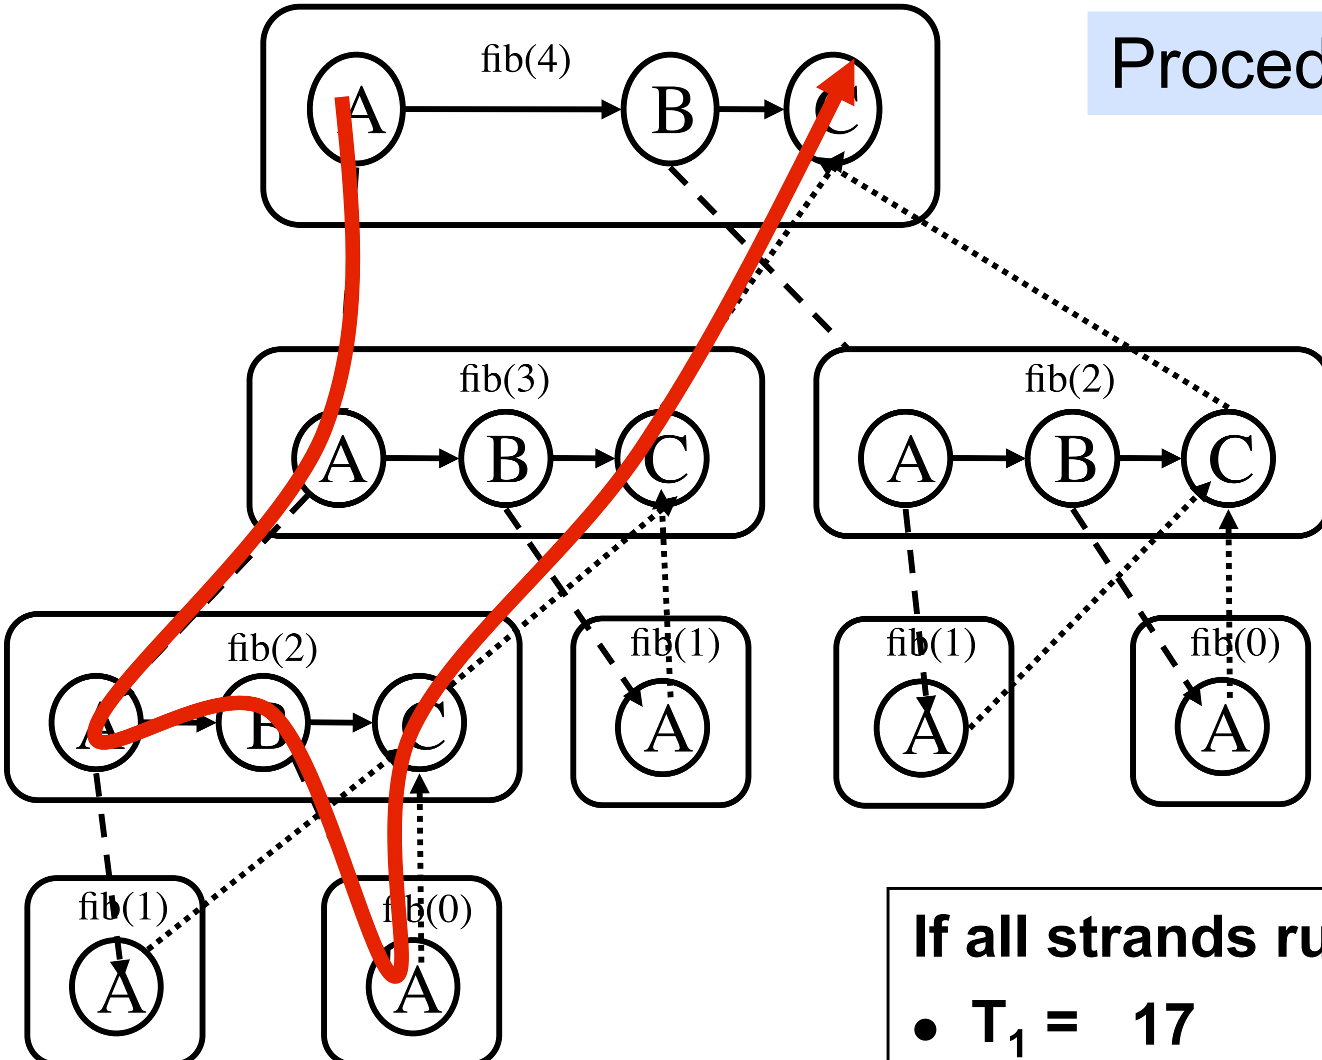

If all strands run in unit time

- **T_1 = 17**
- **T_∞ = 8 (critical path length)**

Execution DAG View

- Cilk Plus uses the word “strand” for a serial section of the program
- A “knot” is a point where three or more strands meet
- Two kinds of knots
  - spawn knots: one input strand, two output strands
  - sync knots: two or more input strands, one output strand

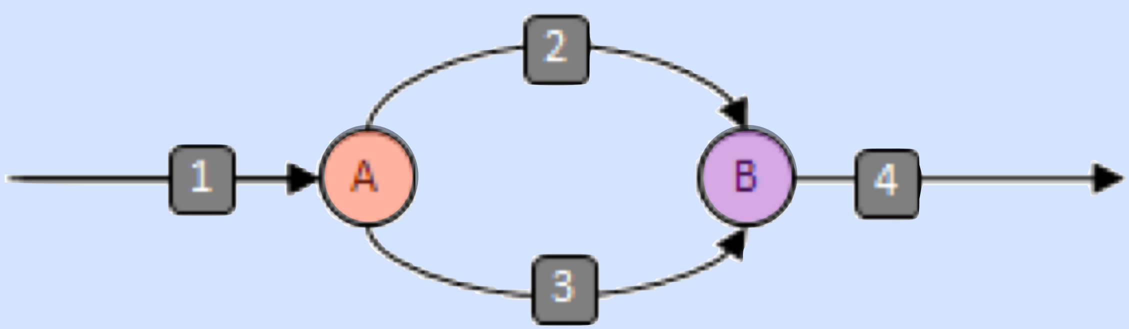


```
...
do_stuff1();
cilk_spawn func3();
do_stuff2();
cilk_sync;
do_stuff4();
...
```

Another Execution DAG

- DAG represents the series-parallel structure of the execution of a Cilk Plus program
- Example:
  - two spawns (A) & (B)
  - one sync (C)

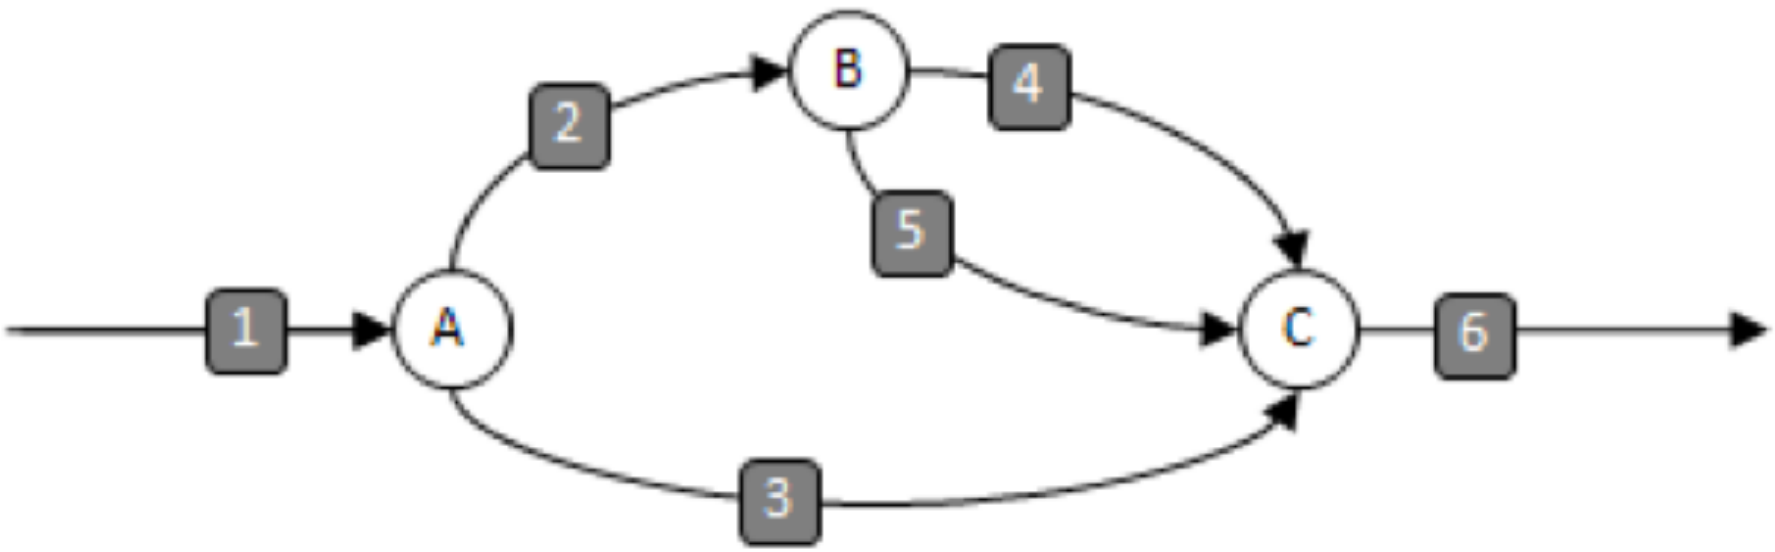

Note: computation on edges

Work and Span

- **Edges represent serial computation (work)**

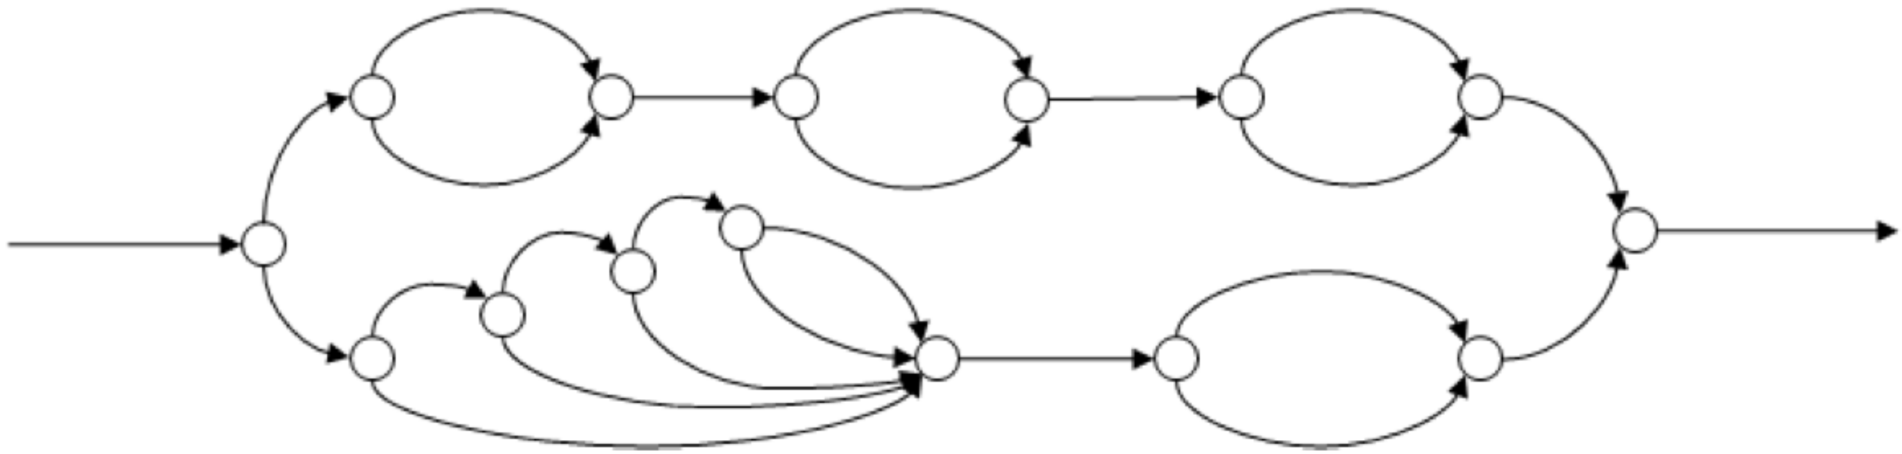

- **Span: most expensive path from beginning to end**
  - — also known as critical path length

work = 181ms

span = 68ms

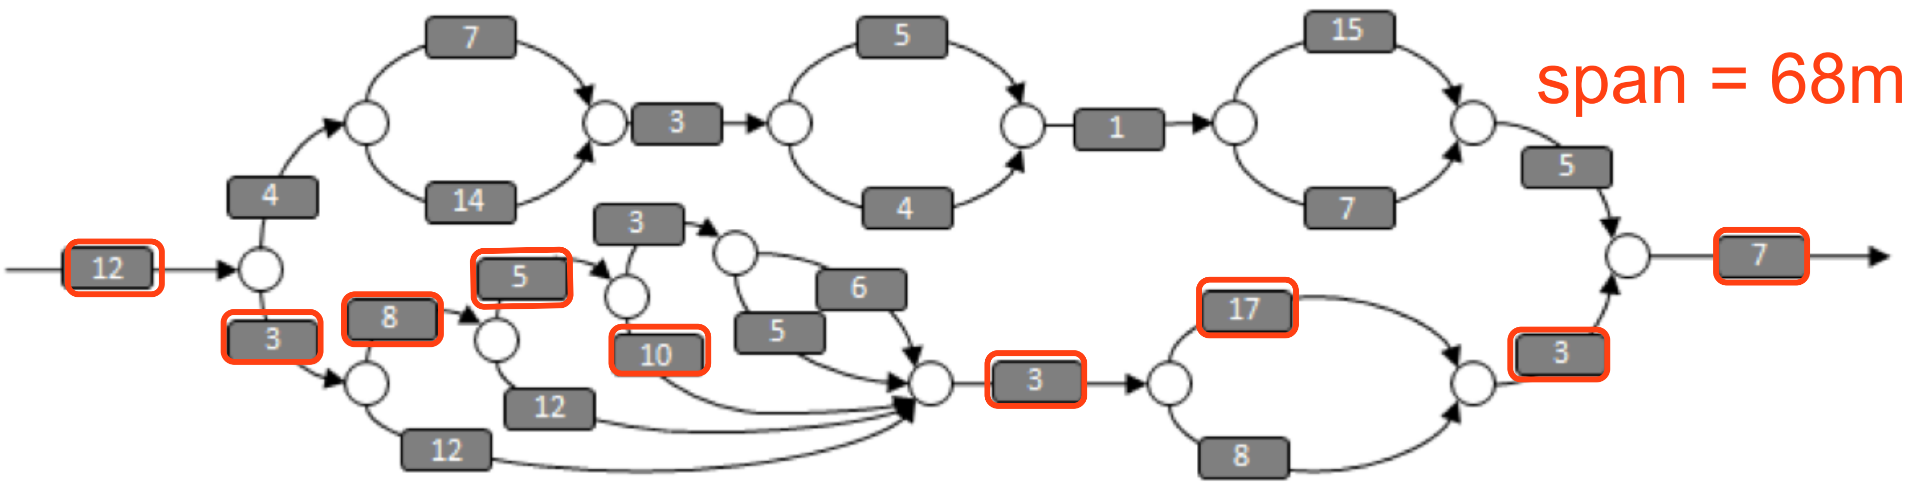

Note: computation on edges

cilkview

- **Rewrites executable to measure execution in terms of work and span**
  - measures
    - work - total # instructions executed, w/o parallel ovhd
    - span - # instructions executed on the critical path (w/o ovhd)
    - burdened span - # instructions executed on critical path (incl ovhd)
    - parallelism - work/span (max speedup on infinite cores, w/o ovhd)
    - burdened parallelism - work/(burdened span)
    - number of spawns/syncs
    - average instructions per strand - work/strands
    - strands along span - # strands in the critical path
    - average instructions / strand on span = work/(strands along span)
    - total number of atomic instructions - e.g., used for locks
    - frame count
- **Predicts speedup on various numbers of processors based on work and span**

cilkview Demo

- **Explore cilkview for performance analysis using fib example** **`/projects/comp422/cilkplus-examples/fib`**
  - `cilkview ./fib 20`
  - `cilkview ./fib 30`
  - `cilkview ./fib 35`
  - `cilkview ./fib-trunc 35 10`

Cilk Plus Array Notation

- Elementwise arithmetic

```
c[:] = a[:] + 5;
```

- Set even rows in a 2D array

```
b[0:5:2][:] = 12;
```

- Vector conditionals

```
// Check and report each element containing 5 w/ Array Notation
if (5 == a[:]) an_results[:] = "Matched";
else an_results[:] = "Not Matched";
```

- Vector conditionals

```
// Call a fn on each element of a vector using Array Notation
fn(a[:]);
```

– On STIC, see /projects/comp422/cilkplus/examples/features-tutorial

More Cilk Plus Features

- **See /projects/comp422/cilkplus-features-tutorial**
  - **—array_notations: vector notation in Cilk Plus**
  - **—reducers: more reducer examples**
- **Each directory contains a Makefile that can build and run all examples**

Recall: Task Scheduling in Cilk

Strategies

- **Work-stealing: processor looks for work when it becomes idle**
- **Lazy parallelism: don't realize parallelism until necessary**
  - **benefits:**
    - **executes with precisely as much parallelism as needed**
    - **minimizes the number of threads that must be set up**
    - **runs with same efficiency as serial program on uniprocessor**

Compilation Strategy

MIT Cilk generates two copies of each procedure

- **Fast clone: for optimized execution on a single processor**
  - —spawned threads are fast
- **Slow clone: triggered by work stealing, full parallel support**
  - —used to handle execution of "stolen procedure frames"
  - —supports Cilk's work-stealing scheduler
  - —few steals when enough parallel slackness exists
    - speed of slow copy is not critical for performance
- **"Work-first" principle: minimize cost in fast clone**

Two Schedulers

- **Nanoscheduler: compiled into cilk program**
  - **—execute cilk function and spawns in exactly the same order as C**
  - **—on one PE: when no microscheduling needed, same order as C**
  - **—efficient coordination with microscheduler**
- **Microscheduler**
  - **—schedule procedures across a fixed set of processors**
  - **—implementation: randomized work-stealing scheduler**
    - **when a processor runs out of work, it becomes a thief**
    - **steals from victim processor chosen uniformly at random**

Nanscheduler Sketch

- Upon entering a cilk function
  - allocate a frame in the heap
  - initialize frame to hold function's state
  - push the frame on the bottom of a deque
    - frame on stack ↔ frame in deque
- At a spawn
  - save function state into the frame
    - only live, dirty variables
  - save the entry number into the frame
  - call spawned procedure as a function
- After each spawn
  - check to see if if parent has been stolen
    - if frame is still in the deque, it has not
  - if so, clean up C stack
- Each sync becomes a no-op
- When the procedure returns

Fast clone

```
int fib (int n)
{
    fib_frame *f;                       frame pointer
    f = alloc(sizeof(*f));              allocate frame
    f->sig = fib_sig;                   initialize frame
    if (n<2) {
        free(f, sizeof(*f));            free frame
        return n;
    }
    else {
        int x, y;
        f->entry = 1;                   save PC
        f->n = n;                       save live vars
        *T = f;                         store frame pointer
        push();                         push frame
        x = fib (n-1);                  do C call
        if (pop(x) == FAILURE)          pop frame
            return 0;                   frame stolen
        ...                             second spawn
        ;                               sync is free!
        free(f, sizeof(*f));            free frame
        return (x+y);
    }
}
```

Fast Clone and Nanoscheduler

- **Fast clone is never stolen**
  - —converted to slow when steal occurs
  - —enables optimizations
- **No sync needed in fast clone**
  - —no children have been spawned
- **Frame saves state:**
  - —PC (entry number)
  - —live, dirty variables
- **Push and pop must be fast**

Nanoscheduler Overheads

Basis for comparison: serial C

- **Allocation and initialization of frame, push onto 'stack'**
  - — a few assembly instructions
- **Procedure's state needs to be saved before each spawn**
  - — entry number, live variables
- **Check whether frame is stolen after each spawn**
  - — two reads, compare, branch
- **On return, free frame - a few instructions**
- **One extra variable to hold frame pointer**

Runtime Support for Scheduling

Each processor has a ready deque (doubly ended queue)

- Tail: worker adds or removes procedures (like C call stack)
- Head: thief steals from head of a victim's deque

Deque for a Process

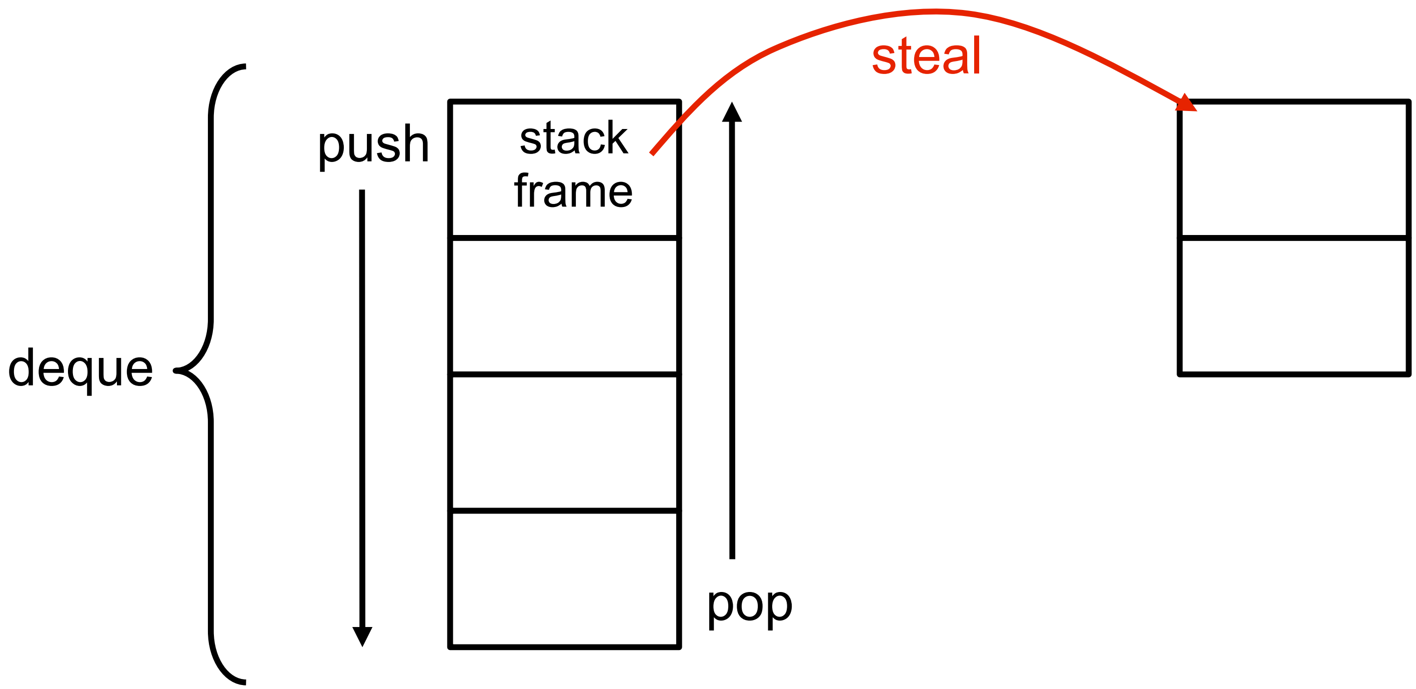

- *Deque* grows downward
- *Stack frame* contains local variables for a procedure invocation
  - Procedure call → new frame is pushed onto the bottom of the deque
  - Procedure return → bottom frame is popped from the deque

Cilk's Cactus Stacks

A cactus stack enables sharing of a C function's local variables

```
void A() { B(); C(); }
void B() { D(); E(); }
void C() { F(); }
void D() {}
void E() {}
void F() {}
```

call tree

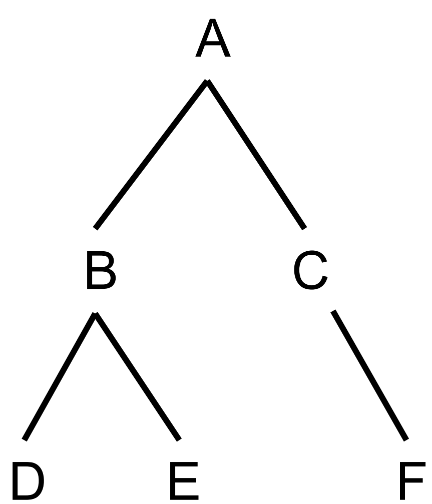

each procedure's view of stack

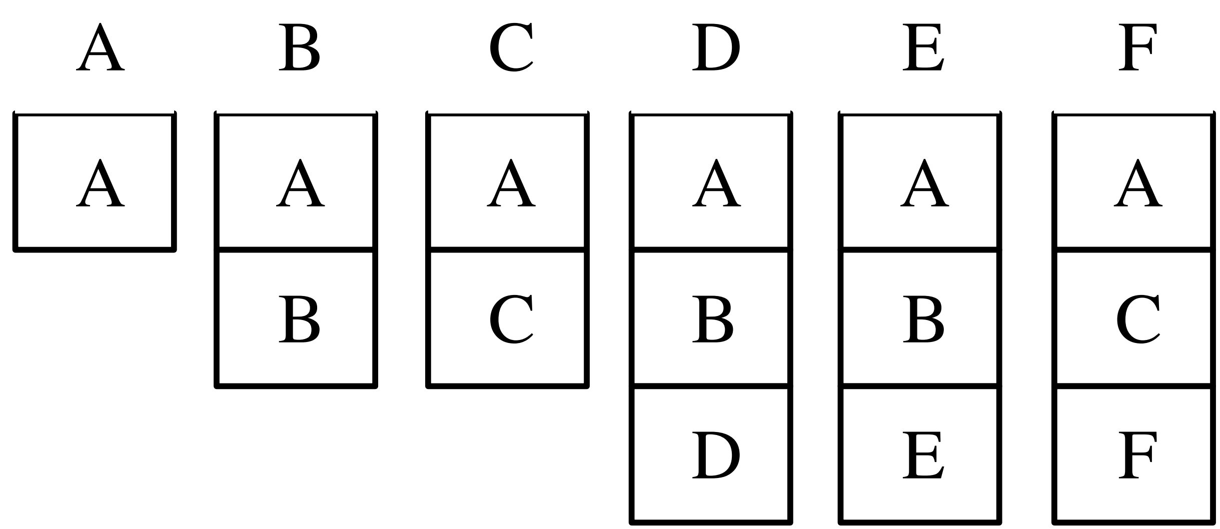

Rules

- —pointers can be passed down call chain
- —only pass pointers up if they point to heap
  - – functions <u>cannot</u> return ptrs to local variables

Microscheduler

Schedule procedures across a fixed set of processors

- **When a processor runs out of work, it becomes a thief**
  — **steals from victim processor chosen uniformly at random**
- **When it finds victim with frames in its deque**
  — **takes the topmost frame (least recently pushed)**
  — **places frame into its own deque**
  — **gives the corresponding procedure to its own nanoscheduler**
- **Microscheduler executes slow clone**
  — **receives only pointer to frame as argument**
    – **real args and local state in frame**
  — **restores pgm counter to proper place using switch stmt (Duff's device)**
  — **at a sync, must wait for children**
  — **before the procedure returns, place return value into frame**

Coordinating Thief and Worker

Options

- Always use a lock to manipulate each worker's deque
- Use protocol that only relies on atomicity of read and write
  - based on ideas from a locking protocol by Dijkstra

Simplified THE Protocol (Without the 'E')

- **Shared memory deque**
  - T: first unused
  - H: head
  - E: exception
- **Work-first**
  - move costs from worker to thief
- **One worker per deque**
- **One thief at a time**
  - enforced by lock

Worker/Victim

```
push() {
  T++;
}

pop() {
  T--;
  if (H > T) {
    T++;
    lock(L);
    T--;
    if (H > T) {
      T++;
      unlock(L);
      return FAILURE;
    }
    unlock(L);
  }
  return SUCCESS;
}
```

Thief

```
steal() {
  lock(L);
  H++;
  if (H > T) {
    H--;
    unlock(L);
    return FAILURE;
  }
  unlock(L);
  return SUCCESS;
}
```

- actions on tail contribute to work overhead
- actions on head contribute only to critical path overhead

Deque Pop

Three cases

(a) **no conflict**

(b) At least one (thief or victim) finds (H > T) and backs up; other succeeds

(c) Deque is empty, both threads return

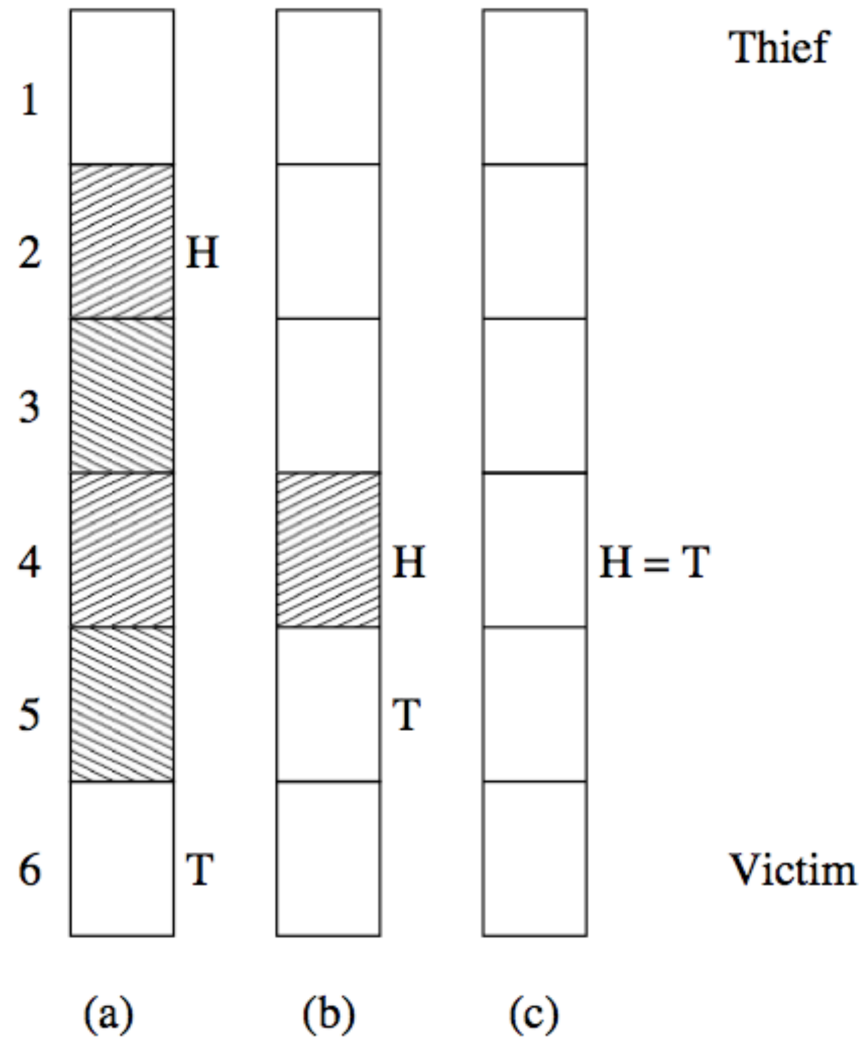

Work Overhead for `fib`

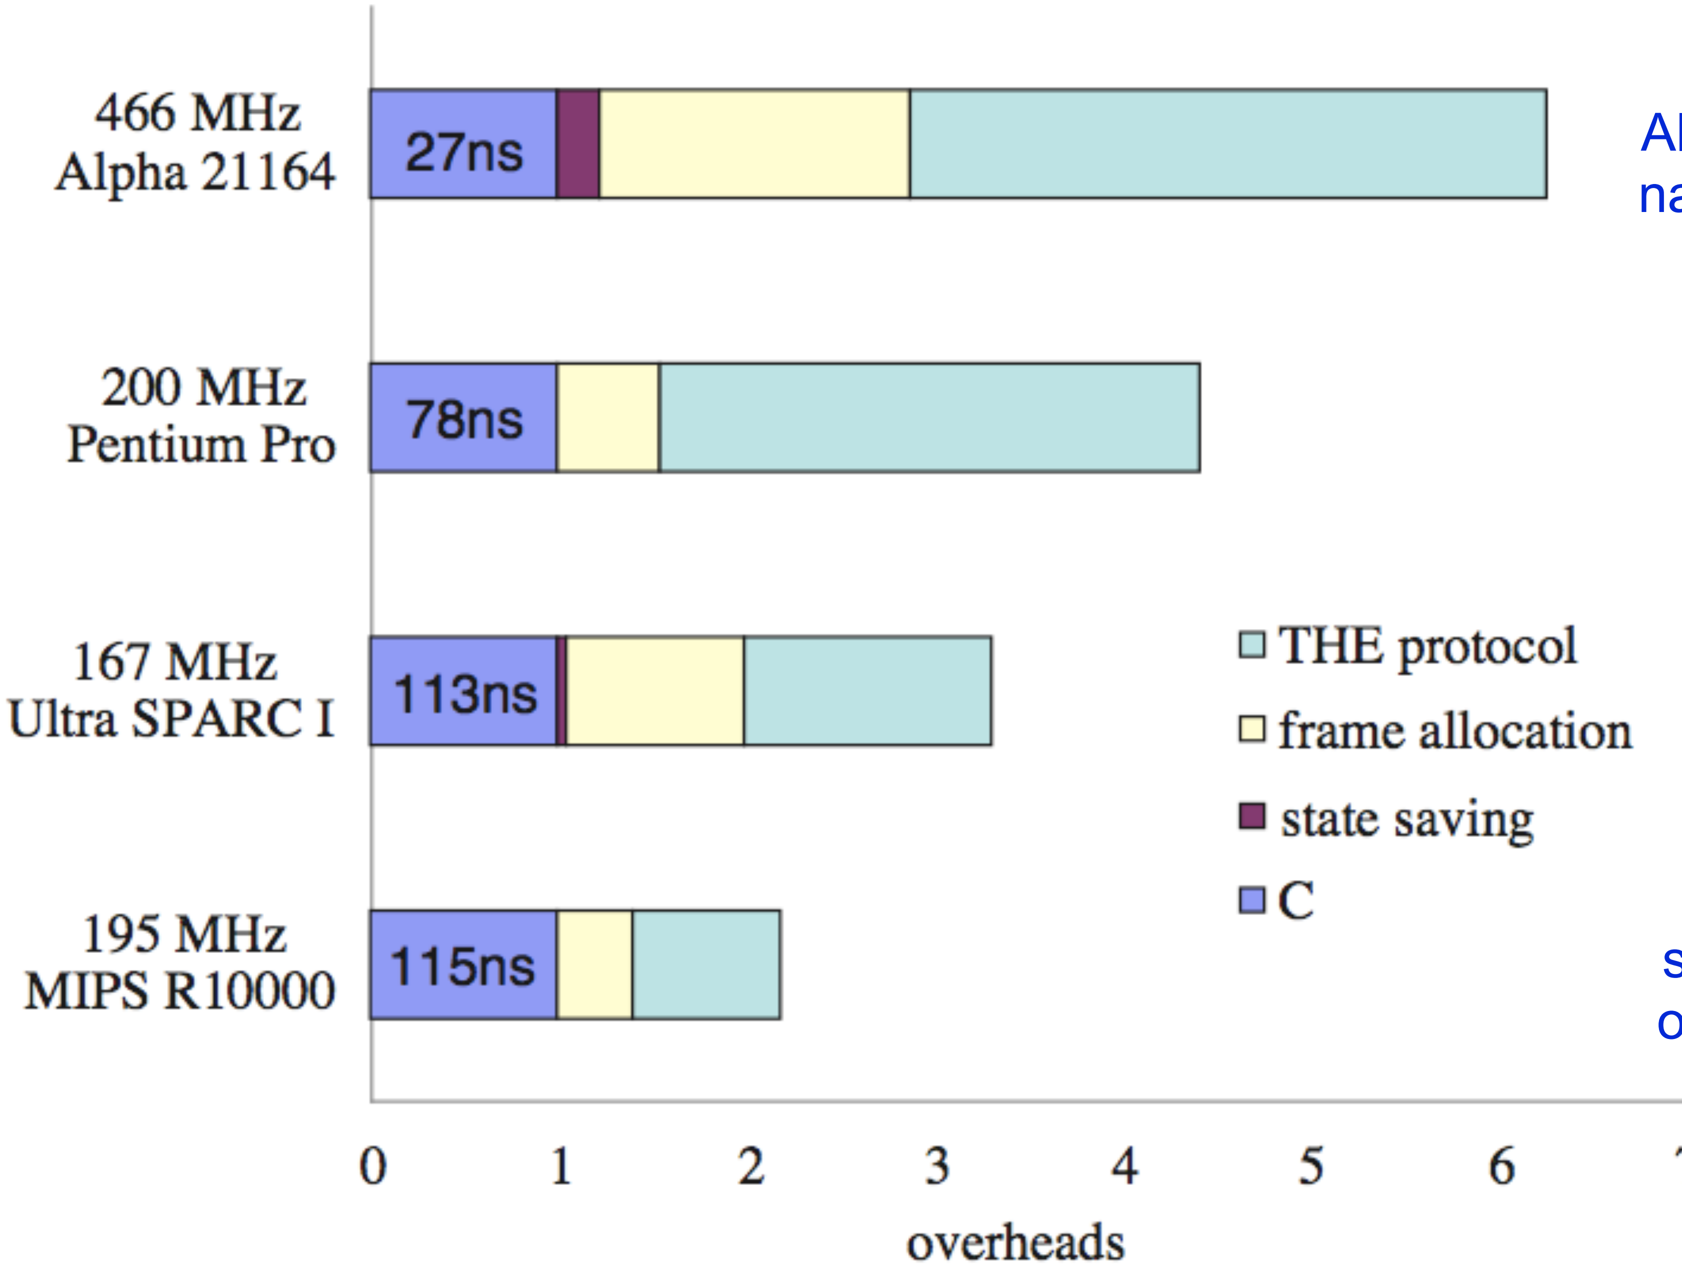

Alpha has fast native function calls

state saving overhead small because of write buffers

References - I

- **Matteo Frigo, Charles Leiserson, and Keith Randall. The implementation of the Cilk-5 multithreaded language. In PLDI (Montreal, Quebec, Canada, June 17 - 19, 1998), 212-223.**
- **Mingdong Feng and Charles E. Leiserson. 1997. Efficient detection of determinacy races in Cilk programs. In *Proceedings of the ninth annual ACM symposium on Parallel algorithms and architectures* (SPAA '97). ACM, New York, NY, USA, 1-11.**
- **Guang-Ien Cheng, Mingdong Feng, Charles E. Leiserson, Keith H. Randall, and Andrew F. Stark. 1998. Detecting data races in Cilk programs that use locks. In *Proceedings of the tenth annual ACM symposium on Parallel algorithms and architectures* (SPAA '98). ACM, New York, NY, USA, 298-309.**

References - II

- Yuxiong He, Charles E. Leiserson, and William M. Leiserson. 2010. The Cilkview scalability analyzer. In Proc. of the 22nd annual ACM symposium on Parallelism in algorithms and architectures (SPAA '10). ACM, New York, NY.
- Charles E. Leiserson. Cilk LECTURE 1. Supercomputing Technologies Research Group. Computer Science and Artificial Intelligence Laboratory. http://bit.ly/mit-cilk-lec1
- Matteo Frigo, Pablo Halpern, Charles E. Leiserson, and Stephen Lewin-Berlin. Reducers and other Cilk++ hyperobjects. SPAA '09, 79-90. Talk Slides. April 11, 2009. http://bit.ly/reducers
- Charles Leiserson, Bradley Kuzmaul, Michael Bender, and Hua-wen Jing. MIT 6.895 lecture notes - Theory of Parallel Systems. http://bit.ly/mit-6895-fall03
- Intel Cilk++ Programmer’s Guide. Document # 322581-001US.