
Programming Shared-memory Platforms with Pthreads

John Mellor-Crummey

**Department of Computer Science
Rice University**

`johnmc@rice.edu`

Threaded Programming Models

- **Library-based models**
 - all data is shared, unless otherwise specified
 - examples: Pthreads, Intel Threading Building Blocks, Java Concurrency, Boost, Microsoft .Net Task Parallel Library
- **Directive-based models, e.g., OpenMP**
 - shared and private data
 - pragma syntax simplifies thread creation and synchronization
- **Programming languages**
 - Cilk Plus (Intel, GCC)
 - CUDA (NVIDIA)
 - Habanero-Java (Rice)

Topics for Today

- The POSIX thread API (Pthreads)
- Synchronization primitives in Pthreads
 - mutexes
 - condition variables
 - reader/writer locks
- Thread-specific data

POSIX Thread API (Pthreads)

- **Standard threads API supported on almost all platforms**
- **Concepts behind Pthreads interface are broadly applicable**
 - largely independent of the API
 - useful for programming with other thread APIs as well
 - Windows threads
 - Java threads
 - ...
- **Threads are peers, unlike Linux/Unix processes**
 - no parent/child relationship

Why Should I Care About Pthreads?

**Pthreads is the foundation for
multithreaded programming models**

- **Used to implement higher-level threading libraries such as Boost and Intel's Threading Building Blocks**
- **Used to implement runtime systems for directive-and language-based programming models such as OpenMP and Cilk Plus**

PThread Creation

Asynchronously invoke **thread_function** in a new thread

```
#include <pthread.h>
int pthread_create(
    pthread_t *thread_handle, /* returns handle here */
    const pthread_attr_t *attribute,
    void * (*thread_function)(void *),
    void *arg); /* single argument; perhaps a structure */
```

attribute created by **pthread_attr_init**:
specifies the size for the thread's stack and how
the thread should be managed by the OS

Thread Attributes

Special functions exist for getting/setting each attribute property
e.g., `int pthread_attr_setstacksize(pthread_attr_t *attr, size_t stacksize)`

- Stack size
- Detach state
 - `PTHREAD_CREATE_DETACHED`, `PTHREAD_CREATE_JOINABLE`
 - reclaim storage at termination (detached) or retain (joinable)
- Scheduling policy
 - `SCHED_OTHER`: standard round robin (priority must be 0)
 - `SCHED_FIFO`, `SCHED_RR`: real time policies
 - `FIFO`: re-enter priority list at head; `RR`: re-enter priority list at tail
- Scheduling parameters
 - only priority
- Inherit scheduling policy
 - `PTHREAD_INHERIT_SCHED`, `PTHREAD_EXPLICIT_SCHED`
- Thread scheduling scope
 - `PTHREAD_SCOPE_SYSTEM`, `PTHREAD_SCOPE_PROCESS`

Wait for Pthread Termination


Suspend execution of calling thread until **thread** terminates

```
#include <pthread.h>
int pthread_join (
    pthread_t thread, /* thread id */
    void **ptr); /* ptr to location for return code a terminating
                  thread passes to pthread_exit */
```


Running Example: Monte Carlo Estimation of Pi

Approximate Pi

- generate random points with $x, y \in [-0.5, 0.5]$
- test if point inside the circle, i.e.,
 $x^2 + y^2 < (0.5)^2$
- ratio of circle to square =
 $\pi r^2 / 4r^2 = \pi / 4$
- $\pi \approx 4 * (\text{number of points inside the circle}) / (\text{number of points total})$



Example: Creation and Termination (main)

```
#include <pthread.h>
#include <stdlib.h>
#define NUM_THREADS 32
void *compute_pi (void *);
...
int main(...) {
    ...
    pthread_t p_threads[NUM_THREADS];
    pthread_attr_t attr;
    pthread_attr_init(&attr);
    for (i=0; i< NUM_THREADS; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void*) &hits[i]);
    }
    for (i=0; i< NUM_THREADS; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
    ...
}
```

The diagram consists of three rectangular boxes with arrows pointing to specific parts of the code. The box labeled 'default attributes' has an arrow pointing to the `&attr` argument in the `pthread_create` call. The box labeled 'thread function' has an arrow pointing to the `compute_pi` argument. The box labeled 'thread argument' has an arrow pointing to the `(void*) &hits[i]` argument.

default attributes

thread function

thread argument

Example: Thread Function (compute_pi)

```
void *compute_pi (void *s) {
    int seed, i, *hit_pointer;
    double x_coord, y_coord;
    int local_hits;
    hit_pointer = (int *) s;
    seed = *hit_pointer;
    local_hits = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        x_coord = (double)(rand_r(&seed))/(RAND_MAX) - 0.5;
        y_coord = (double)(rand_r(&seed))/(RAND_MAX) - 0.5;
        if ((x_coord * x_coord + y_coord * y_coord) < 0.25)
            local_hits++;
    }
    *hit_pointer = local_hits;
    pthread_exit(0);
}
```

tally how many random
points fall in a unit circle
centered at the origin

`rand_r`: reentrant
random number
generation in
[0,RAND_MAX]

Example: Thread Function (compute_pi)

```
void *compute_pi (void *s) {
    int seed, i, *hit_pointer;
    double x_coord, y_coord;
    int local_hits;
    hit_pointer = (int *) s;
    seed = *hit_pointer;
    local_hits = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        x_coord = (double)(rand_r(&seed))/(RAND_MAX) - 0.5;
        y_coord = (double)(rand_r(&seed))/(RAND_MAX) - 0.5;
        if ((x_coord * x_coord + y_coord * y_coord) < 0.25)
            local_hits++;
    }
    *hit_pointer = local_hits;
    pthread_exit(0);
}
```

avoid false sharing by using a local accumulator

Critical Sections and Mutual Exclusion

- Critical section = code executed by only one thread at a time

/ threads compete to update global variable best_cost */*

```
if (my_cost < best_cost)
    best_cost = my_cost;
```

- Mutex locks enforce mutual exclusion in Pthreads
 - mutex lock states: locked and unlocked
 - only one thread can lock a mutex lock at any particular time
- Using mutex locks
 - request lock before executing critical section
 - enter critical section when lock granted
 - release lock when leaving critical section

created by
`pthread_mutex_attr_init`
specifies mutex type

- Operations

```
int pthread_mutex_init (pthread_mutex_t *mutex ,
                        const pthread_mutexattr_t *lock_attr)
int pthread_mutex_lock(pthread_mutex_t *mutex_lock)
int pthread_mutex_unlock(pthread_mutex_t *mutex_lock)
```

atomic operation

Mutex Types

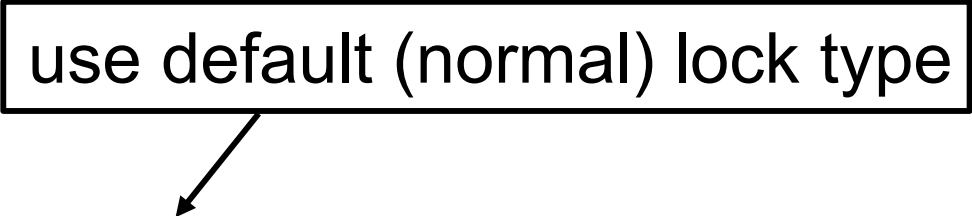
- **Normal**
 - thread deadlocks if tries to lock a mutex it already has locked
- **Recursive**
 - single thread may lock a mutex as many times as it wants
 - increments a count on the number of locks
 - thread relinquishes lock when mutex count becomes zero
- **Errorcheck**
 - report error when a thread tries to lock a mutex it already locked
 - report error if a thread unlocks a mutex locked by another

Example: Reduction Using Mutex Locks

```
pthread_mutex_t cost_lock;
...
int main() {
    ...
    pthread_mutex_init(&cost_lock, NULL);
    ...
}

void *find_best(void *list_ptr) {
    ...
    pthread_mutex_lock(&cost_lock);    /* lock the mutex */
    if (my_cost < best_cost)            critical section
        best_cost = my_cost;
    pthread_mutex_unlock(&cost_lock); /* unlock the mutex */
}
```

use default (normal) lock type



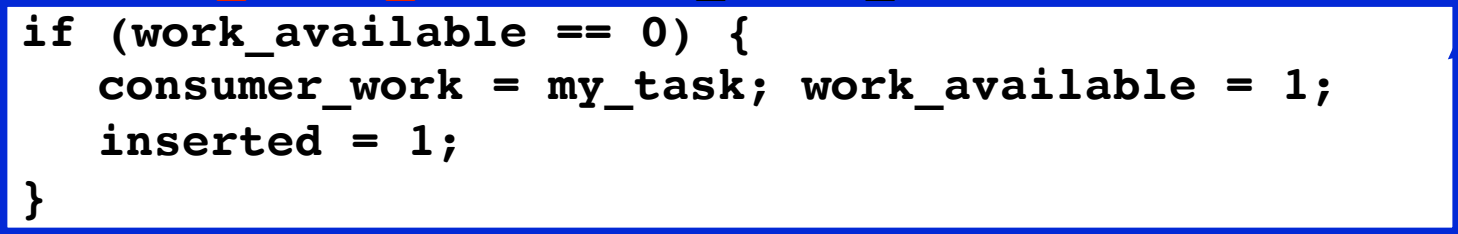
Producer-Consumer Using Mutex Locks

Constraints

- **Producer thread**
 - must not overwrite the shared buffer until previous task has picked up by a consumer
- **Consumer thread**
 - must not pick up a task until one is available in the queue
 - must pick up tasks one at a time

Producer-Consumer Using Mutex Locks

```
pthread_mutex_t task_queue_lock;
int task_available;
...
main() {
    ...
    task_available = 0;
    pthread_mutex_init(&task_queue_lock, NULL);
    ...
}
void *producer(void *producer_thread_data) {
    ...
    while (!done()) {
        inserted = 0;
        create_task(&my_task);
        while (inserted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (work_available == 0) {
                consumer_work = my_task; work_available = 1;
                inserted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
    }
}
```



A blue rectangular box highlights the code block from `pthread_mutex_lock(&task_queue_lock);` to `pthread_mutex_unlock(&task_queue_lock);`. A blue arrow points from the text "critical section" to this box.

critical section

Producer-Consumer Using Locks

```
void *consumer(void *consumer_thread_data) {
    int extracted;
    struct task my_task;
    /* local data structure declarations */
    while (!done()) {
        extracted = 0;
        while (extracted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (work_available == 1) {
                my_task = consumer_work;
                work_available = 0;
                extracted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
        process_task(my_task);
    }
}
```

critical section



Overheads of Locking

- Locks enforce serialization
 - threads must execute critical sections one at a time
- Large critical sections can seriously degrade performance
- Reduce overhead by overlapping computation with waiting

```
int pthread_mutex_trylock(pthread_mutex_t *mutex_lock)
```

- acquires lock if available
- returns EBUSY if not available
- enables a thread to do something else if a lock is unavailable

Condition Variables for Synchronization

Condition variable: associated with a **predicate** and a **mutex**

- Using a condition variable
 - thread can block itself until a condition becomes true
 - thread locks a mutex
 - tests a predicate defined on a shared variable
 - if predicate is false, then wait on the condition variable
 - waiting on condition variable unlocks associated mutex
 - when some thread makes a predicate true
 - that thread can signal the condition variable to either
 - wake one waiting thread
 - wake all waiting threads
 - when thread releases the mutex, it is passed to first waiter

Pthread Condition Variable API

/* initialize or destroy a condition variable */

```
int pthread_cond_init(pthread_cond_t *cond,  
    const pthread_condattr_t *attr);  
int pthread_cond_destroy(pthread_cond_t *cond);
```

/* block until a condition is true */

```
int pthread_cond_wait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex);  
int pthread_cond_timedwait(pthread_cond_t *cond,  
    pthread_mutex_t *mutex,  
    const struct timespec *wtime);
```

abort wait if time exceeded

/* signal one or all waiting threads that condition is true */


```
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_cond_broadcast(pthread_cond_t *cond);
```

wake one

wake all

Condition Variable Producer-Consumer

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;
/* other data structures here */
main() {
    /* declarations and initializations */
    task_available = 0;
    pthread_init();
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
}
```



A diagram consisting of a rectangular box labeled "default initializations". Three arrows originate from the bottom of this box and point to the arguments "NULL" in the following three lines of code: `pthread_init()`, `pthread_cond_init(&cond_queue_empty, NULL)`, and `pthread_mutex_init(&task_queue_cond_lock, NULL)`.

Producer Using Condition Variables

```
void *producer(void *producer_thread_data) {
    int inserted; task_t *t;
    while (!done()) {
        t = create_task();
        pthread_mutex_lock(&task_queue_cond_lock);
        while (work_available == 1)
            pthread_cond_wait(&cond_queue_empty,
                             &task_queue_cond_lock);
        consumer_work = t;
        work_available = 1;
        pthread_cond_signal(&cond_queue_full);
        pthread_mutex_unlock(&task_queue_cond_lock);
    }
}
```

releases mutex on wait



reacquires mutex when woken



Why Loop When Awaiting A Condition?

```
...
pthread_mutex_lock(&task_queue_cond_lock);
note { while (work_available == 1)
loop {   pthread_cond_wait(&cond_queue_empty,
                        &task_queue_cond_lock);
        ...
        pthread_cond_signal(&cond_queue_full);
        pthread_mutex_unlock(&task_queue_cond_lock);
    }
```

When using condition variables there is always a **boolean predicate** that indicates if the thread should proceed or wait

Spurious wakeups may occur when waiting on condition variables.

Thus, waking up from a wait on a condition variable doesn't imply anything about the value of the boolean predicate; the predicate must be re-evaluated when a conditional wait completes

Why Allow Spurious Wakeups?

- **Defining condition variable waits to permit spurious forces correct/robust code by requiring predicate loops.**
 - "Religiously" using a loop protects the application against its own imperfect coding practices.
- **Making condition wakeup completely predictable might substantially slow all condition variable operations.**
 - It isn't difficult to imagine machines and implementation code that could exploit this semantics to improve the performance of average condition wait operations.

-- David R. Butenhof - author of "Programming with POSIX Threads"

Consumer Using Condition Variables

```
void *consumer(void *consumer_thread_data) {  
    while (!done()) {  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (work_available == 0)  
            pthread_cond_wait(&cond_queue_full,  
                             &task_queue_cond_lock);  
        my_task = consumer_work;  
        work_available = 0;  
        pthread_cond_signal(&cond_queue_empty);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
        process_task(my_task);  
    }  
}
```

note loop {

releases mutex on wait

reacquires mutex when woken

Reader-Writer Locks

- **Purpose: access to data structure when**
 - frequent reads
 - infrequent writes
- **Acquire read lock**
 - OK to grant when other threads already have acquired read locks
 - if write lock on the data or queued write locks
 - reader thread performs a condition wait
- **Acquire write lock**
 - if multiple threads request a write lock
 - must perform a condition wait

Read-Write Lock Sketch

- While pthreads provides a pthread_rwlock, you could build your own using basic primitives
- Use a data type with the following components
 - a count of the number of active readers
 - 0/1 integer specifying whether a writer is active
 - a condition variable readers_proceed
 - signaled when readers can proceed
 - a condition variable writer_proceed
 - signaled when one of the writers can proceed
 - a count pending_writers of pending writers
 - a mutex read_write_lock
 - controls access to the reader/writer data structure

Thread-Specific Data

Goal: associate some state with a thread

- **Choices**

- pass data as argument to each call thread makes
 - not always an option, e.g. when using predefined libraries
- store data in a shared variable indexed by thread id
- using thread-specific keys

- **Why thread-specific keys?**

- libraries want to maintain internal state
- don't want to require clients to know about it and pass it back
- substitute for static data in a threaded environment

- **Operations**

associate NULL with key in each active thread

```
int pthread_key_create(pthread_key_t *key, void (*destroy)(void *))
```

```
int pthread_setspecific(pthread_key_t key, const void *value)
```

```
void *pthread_getspecific(pthread_key_t key)
```

retrieve value for current
thread from key

associate (key,value)
with current thread

Thread-Specific Data Example: Key Creation

Example: remember performance information for a thread

```
#include <pthread.h>
```

```
static pthread_key_t profiler_state;
```

```
initialize_profiler_state() {
```

```
...
```

```
pthread_key_create(&profiler_state,  
                  (void *) free_profile);
```

```
...
```

```
}
```

```
void free_profile(profile *my_profile) {
```

```
    free(my_profile);
```

```
}
```

opaque handle
used to locate
thread-specific data

The diagram consists of two rectangular boxes. The top box contains the text 'opaque handle used to locate thread-specific data'. An arrow points from this box to the 'profiler_state' variable in the code line 'static pthread_key_t profiler_state;'. Another arrow points from the same box to the '&profiler_state' argument in the code line 'pthread_key_create(&profiler_state, (void *) free_profile);'.

destructor for key value

The diagram consists of a rectangular box containing the text 'destructor for key value'. An arrow points from this box to the 'free_profile' argument in the code line 'pthread_key_create(&profiler_state, (void *) free_profile);'. Another arrow points from the same box to the 'free_profile' function definition below.

Thread-Specific Data Example: Specific Data

Example: remember profiler state for a thread

```
void init_thread_profile(...) {  
    profile *my_profile = (profile *) malloc(...);  
    pthread_setspecific(profiler_state, (void *) my_profile);  
    ...  
}  
  
void update_thread_profile(...) {  
    profile *my_profile = (profile *)  
        pthread_getspecific(profiler_state);  
    // update profile  
}
```

References

- **Adapted from slides “Programming Shared Address Space Platforms” by Ananth Grama.**
- **Bradford Nichols, Dick Buttlar, Jacqueline Proulx Farrell. “Pthreads Programming: A POSIX Standard for Better Multiprocessing.” O'Reilly Media, 1996.**
- **Chapter 7. “Introduction to Parallel Computing” by Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar. Addison Wesley, 2003**