

---

# Collective Communication

**John Mellor-Crummey**

**Department of Computer Science  
Rice University**

**`johnmc@rice.edu`**

# Group Communication

---

- **Motivation: accelerate interaction patterns within a group**
- **Approach: collective communication**
  - group works together *collectively* to realize a communication
  - constructed from pairwise point-to-point communications
- **Implementation strategy**
  - standard library of common collective operations
  - leverage target architecture for efficient implementation
- **Benefits of standard library implementations**
  - reduce development effort and cost for parallel codes
  - improve performance through efficient implementations
  - improve quality of scientific applications

# Topics for Today

---

- **One-to-all broadcast and all-to-one reduction**
- **All-to-all broadcast and reduction**
- **All-reduce and prefix-sum operations**
- **Scatter and gather**
- **All-to-all personalized communication**
- **Optimizing collective patterns**

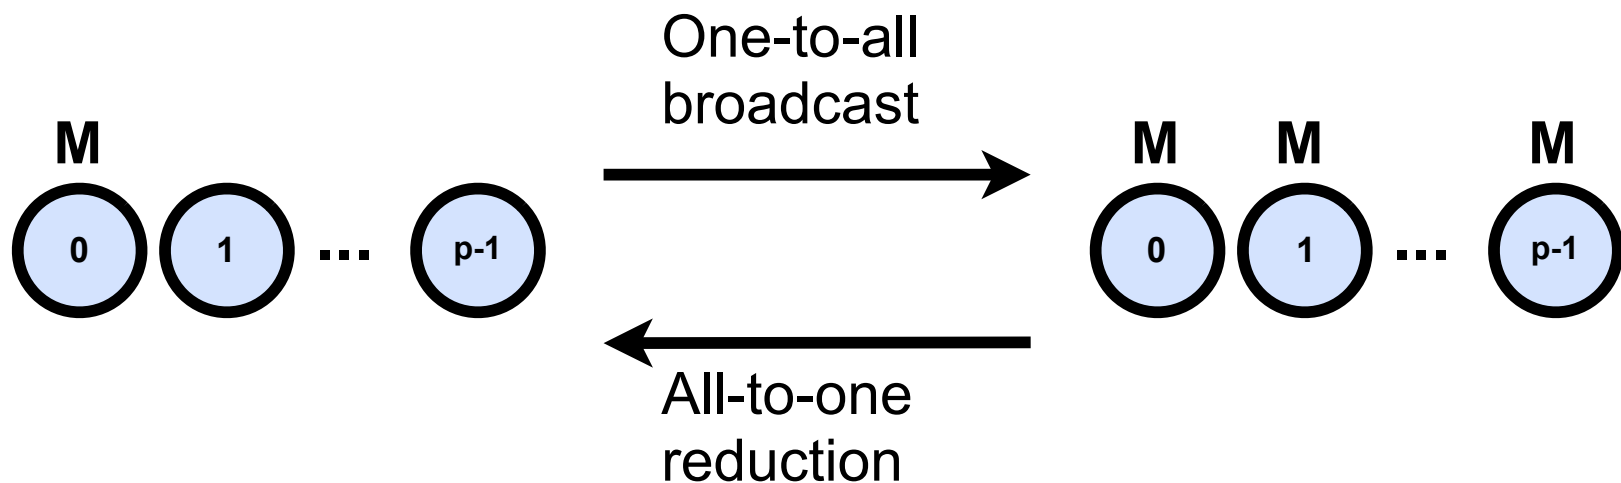
# Assumptions

---

- Network is bidirectional
- Communication is single-ported
  - node can receive only one message per step
- Communication cost model
  - message of size  $m$ , no congestion, time =  $t_s + t_w m$
  - congestion: model by scaling  $t_w$

# One-to-All and All-to-One

- **One-to-all broadcast**
  - a processor has  $M$  units of data that it must send to everyone
- **All-to-one reduction**
  - each processor has  $M$  units of data
  - data items must be combined using some associative operator
    - e.g. addition, min, max
  - result must be available at a target processor



# One-to-All and All-to-One on a Ring

- **Broadcast**


- naïve solution

- source sends  $p - 1$  messages to the other  $p - 1$  processors

- use recursive doubling


- source sends a message to a selected processor

yields two independent problems over halves of the machine




- **Reduction**

- invert the process




# Broadcast on a Balanced Binary Tree

- Consider processors arranged in a dynamic binary tree
  - processors are at the leaves
  - interior nodes are switches
- Assume leftmost processor is the root of the broadcast
- Use recursive doubling strategy:  $\log p$  stages



# Broadcast and Reduction on a 2D Mesh

- Consider a square mesh of  $p$  nodes
  - treat each row as a linear array of  $p^{1/2}$  nodes
  - treat each column as a linear array of  $p^{1/2}$  nodes
- Two step broadcast and reduction operations
  1. perform the operation along a row
  2. perform the operation along each column concurrently




broadcast on  
4 x 4 mesh

- Generalizes to higher dimensional meshes



# Broadcast and Reduction on a Hypercube

- Consider hypercube with  $2^d$  nodes
  - view as  $d$ -dimensional mesh with two nodes in each dimension
- Apply mesh algorithm to a hypercube
  - $d$  ( $= \log p$ ) steps



# Broadcast and Reduction Algorithms

---

- Each of aforementioned broadcast/reduction algorithms
  - adaptation of the same algorithmic template
- Next slide: a broadcast algorithm for a hypercube of  $2^d$  nodes
  - can be adapted to other architectures
  - in the following algorithm
    - *my\_id* is the label for a node
    - *X* is the message to be broadcast

# One-to-All Broadcast Algorithm


```

1.  procedure GENERAL_ONE_TO_ALL_BC( $d, my\_id, source, X$ )
2.  begin
3.       $my\_virtual\_id := my\_id \text{ XOR } source;$ 
4.       $mask := 2^d - 1;$ 
5.      for  $i := d - 1$  downto 0 do /* Outer loop */
6.           $mask := mask \text{ XOR } 2^i;$  /* Set bit  $i$  of  $mask$  to 0 */
7.          if  $(my\_virtual\_id \text{ AND } mask) = 0$  then
8.              if  $(my\_virtual\_id \text{ AND } 2^i) = 0$  then // even
9.                   $virtual\_dest := my\_virtual\_id \text{ XOR } 2^i;$ 
10.                 send  $X$  to  $(virtual\_dest \text{ XOR } source);$ 
11.                 /* Convert  $virtual\_dest$  to the label of the physical destination */
12.             else // odd
13.                  $virtual\_source := my\_virtual\_id \text{ XOR } 2^i;$ 
14.                 receive  $X$  from  $(virtual\_source \text{ XOR } source);$ 
15.                 /* Convert  $virtual\_source$  to the label of the physical source */
16.             endelse;
17.         endfor;
18.     end GENERAL_ONE_TO_ALL_BC

```

I am communicating on behalf of a  $2^i$  subcube

position relative to source



One-to-all broadcast of a message  $X$  from  $source$  on a hypercube

# All-to-One Reduction Algorithm

```
1.  procedure ALL_TO_ONE_REDUCE( $d, my\_id, m, X, sum$ )
2.  begin
3.      for  $j := 0$  to  $m - 1$  do  $sum[j] := X[j]$ ;
4.       $mask := 0$ ;
5.      for  $i := 0$  to  $d - 1$  do
6.          /* Select nodes whose lower  $i$  bits are 0 */
7.          if  $(my\_id \text{ AND } mask) = 0$  then
8.              if  $(my\_id \text{ AND } 2^i) \neq 0$  then // odd
9.                   $msg\_destination := my\_id \text{ XOR } 2^i$ ;
10.                 send  $sum$  to  $msg\_destination$ ;
11.             else // even
12.                  $msg\_source := my\_id \text{ XOR } 2^i$ ;
13.                 receive  $X$  from  $msg\_source$ ;
14.                 for  $j := 0$  to  $m - 1$  do
15.                      $sum[j] := sum[j] + X[j]$ ;
16.                 endwhile;
17.              $mask := mask \text{ XOR } 2^i$ ; /* Set bit  $i$  of  $mask$  to 1 */
18.         endfor;
19.     end ALL_TO_ONE_REDUCE
```

I am communicating on  
behalf of a  $2^i$  subcube



All-to-One sum reduction on a  $d$ -dimensional hypercube

Each node contributes msg  $X$  containing  $m$  words, and node 0 is the destination 12

# Broadcast/Reduction Cost Analysis

---

## Hypercube

- Log p point-to-point simple message transfers  
—each message transfer time:  $t_s + t_w m$

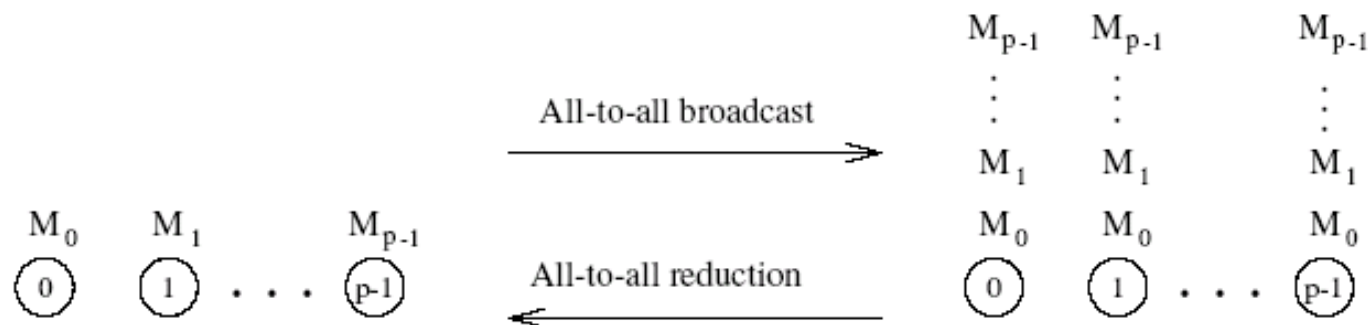
- Total time

$$T = (t_s + t_w m) \log p.$$

# All-to-All Broadcast and Reduction

Each processor is the source as well as destination


- **Broadcast**
  - each process broadcasts its own  $m$ -word message all others
- **Reduction**
  - each process gets a copy of the result



# All-to-All Broadcast/Reduction on a Ring

---


```
1.  procedure ALL_TO_ALL_BC_RING(my_id, my_msg, p, result)
2.  begin
3.      left := (my_id - 1) mod p;
4.      right := (my_id + 1) mod p;
5.      result := my_msg;
6.      msg := result;
7.      for i := 1 to p - 1 do
8.          send msg to right;
9.          receive msg from left;
10.         result := result ∪ msg;
11.     endfor;
12. end ALL_TO_ALL_BC_RING
```



message size  
stays constant

Also works for a linear array with bidirectional communication channels

# All-to-All Broadcast on a Ring



For an all-to-all reduction


- combine (rather than append) each incoming message into your local result
- at each step, forward your incoming msg to your successor




# All-to-all Broadcast on a Mesh

## Two phases

- Perform row-wise all-to-all broadcast as for linear array/ring
  - each node collects  $p^{1/2}$  messages for nodes in its own row
  - consolidates into a single message of size  $mp^{1/2}$
- Perform column-wise all-to-all broadcast of merged messages



(a) Initial data distribution




(b) Data distribution after rowwise broadcast

# All-to-all Broadcast on a Hypercube


- Generalization of the mesh algorithm to  $\log p$  dimensions
- Message size doubles in each of  $\log p$  steps

1 value @ each




(a) Initial distribution of messages

2 values @ each




(b) Distribution before the second step

4 values @ each



(c) Distribution before the third step

8 values @ each



(d) Final distribution of messages

# All-to-all Broadcast on a Hypercube

---

```
1.  procedure ALL_TO_ALL_BC_HCUBE(my_id, my_msg, d, result)
2.  begin
3.      result := my_msg;
4.      for i := 0 to d - 1 do
5.          partner := my_id XOR  $2^i$ ;
6.          send result to partner;
7.          receive msg from partner;
8.          result := result  $\cup$  msg;
9.      endfor;
10. end ALL_TO_ALL_BC_HCUBE
```

# All-to-all Reduction

---

- Similar to all-to-all broadcast, except for the merge
- Algorithm sketch

`my_result = local_value`

`for each round`

`send my_result to partner`

`receive msg`

`my_result = my_result  $\oplus$  msg`

`post condition: each my_result now contains global result`

# Cost Analysis for All-to-All Broadcast

---

- Ring

— $(t_s + t_w m)(p-1)$

- Mesh

—phase 1:  $(t_s + t_w m)(p^{1/2} - 1)$

—phase 2:  $(t_s + t_w m p^{1/2})(p^{1/2} - 1)$

—total:  $2t_s(p^{1/2} - 1) + t_w m(p - 1)$

- Hypercube

$$T = \sum_{i=1}^{\log p} (t_s + 2^{i-1} t_w m)$$

$$= t_s \log p + t_w m(p - 1).$$

**Above algorithms are asymptotically optimal in msg size**

# Prefix Sum

---

- **Pre-condition**
  - given  $p$  numbers  $n_0, n_1, \dots, n_{p-1}$  (one on each node)
    - node labeled  $k$  contains  $n_k$
- **Problem statement**
  - compute the sums  $s_k = \sum_{i=0}^k n_i$  for all  $k$  between 0 and  $p-1$
- **Post-condition**
  - node labeled  $k$  contains  $s_k$

# Prefix Sum

---

- Can use all-to-all reduction kernel to implement prefix sum
- Constraint
  - prefix sums on node  $k$ : values from  $k$ -node subset with labels  $\leq k$
- Strategy
  - implemented using an additional result buffer
  - add incoming value to result buffer on node  $k$ 
    - only if the msg from a node  $\leq k$

# Prefix Sum on a Hypercube

---

```
1.  procedure PREFIX_SUMS_HCUBE(my_id, my_number, d, result)
2.  begin
3.      result := my_number;
4.      msg := result;
5.      for i := 0 to d - 1 do
6.          partner := my_id XOR  $2^i$ ;
7.          send msg to partner;
8.          receive number from partner;
9.          msg := msg + number;
10.         if (partner < my_id) then result := result + number;
11.     endfor;
12. end PREFIX_SUMS_HCUBE
```




# Scatter and Gather

- **Scatter**


- a node sends a unique message of size  $m$  to every other node
  - AKA one-to-all personalized communication
- algorithmic structure is similar to broadcast
  - scatter: message size get smaller at each step
  - broadcast: message size stay constant

- **Gather**


- single node collects a unique message from each node
- inverse of the scatter operation; can be executed as such




# Scatter on a Hypercube




(a) Initial distribution of messages



(b) Distribution before the second step



(c) Distribution before the third step



(d) Final distribution of messages

# Cost of Scatter and Gather

---

- **Log  $p$  steps**

- in each step

- machine size halves
    - message size halves

- **Time**


$$T = t_s \log p + t_w m(p - 1).$$

- **Note: time is asymptotically optimal in message size**


# All-to-All Personalized Communication

## Total exchange

- Each node: distinct message of size  $m$  for every other node



# All-to-All Personalized Communication



# All-to-All Personalized Communication

---

- Every node has  $p$  pieces of data, each of size  $m$
- Algorithm sketch for a ring

for  $k = 1$  to  $p - 1$

send message of size  $m(p - k)$  to neighbor

select piece of size  $m$  out of message for self

- Cost analysis

$$\begin{aligned} T &= \sum_{i=1}^{p-1} (t_s + t_w m(p - i)) \\ &= t_s(p - 1) + t_w m \sum_{i=1}^{p-1} i \\ &= (t_s + t_w mp/2)(p - 1) \end{aligned}$$

# Optimizing Collective Patterns

## Example: one-to-all broadcast of large messages on a hypercube

- Consider broadcast of message  $M$  of size  $m$ , where  $m$  is large
- Cost of straightforward strategy  $T = (t_s + t_w m) \log p$
- Optimized strategy
  - split  $M$  into  $p$  parts  $M_0, M_1, \dots, M_p$  of size  $m/p$  each
    - want to place  $M_0 \cup M_1 \cup \dots \cup M_p$  on all nodes
  - scatter  $M_i$  to node  $i$
  - have nodes collectively perform all-to-all broadcast
    - each node  $k$  broadcasts its  $M_k$
- Cost analysis
  - scatter time =  $t_s \log p + t_w(m/p)(p-1)$  (slide 27)
  - all-to-all broadcast time =  $t_s \log p + t_w(m/p)(p-1)$  (slide 21)
  - total time =  $2(t_s \log p + t_w(m/p)(p-1)) \approx 2(t_s \log p + t_w m)$   
(faster than slide 13)

# References

---

- **Adapted from slides “Principles of Parallel Algorithm Design” by Ananth Grama**
- **Based on Chapter 4 of “Introduction to Parallel Computing” by Ananth Grama, Anshul Gupta, George Karypis, and Vipin Kumar. Addison Wesley, 2003**