
Mutual Exclusion: Classical Algorithms for Locks

John Mellor-Crummey

**Department of Computer Science
Rice University**

`johnmc@rice.edu`

Motivation

- **Problem: ensure that a data structure is maintained consistently**
 - avoid conflicting accesses to shared data (data races)
 - read/write conflicts
 - write/write conflicts
- **Locks guarantee consistency by providing exclusion**
 - acquire lock before manipulating the shared data
 - release lock when finished manipulating the shared data

Problems with Locks

- **Conceptual**
 - coarse-grained: poor scalability
 - fine-grained: hard to write
- **Semantic**
 - deadlock
 - priority inversion
- **Performance**
 - intolerance of page faults and preemption

Alternatives to Locks

- **Transactional memory (TM)**

- support arbitrary atomic actions on multi-word shared data

```
atomic (entries > 0) {  
    node *first = head; head = head->next;  
    entries--; return first;  
}
```

- transactions that don't conflict run uninterrupted in parallel
- transactions that conflict abort and retry
 - benefit: no need for programmer to worry about deadlock!
 - cost: repeated aborts can waste resources and hurt performance

+ easy to use, well-understood metaphor

– high overhead in software; HTM on Blue Gene/Q, Intel Haswell, IBM Power8

± subject of much active research

- **Ad hoc non-blocking synchronization (NBS)**

+ thread failure/delay cannot prevent progress

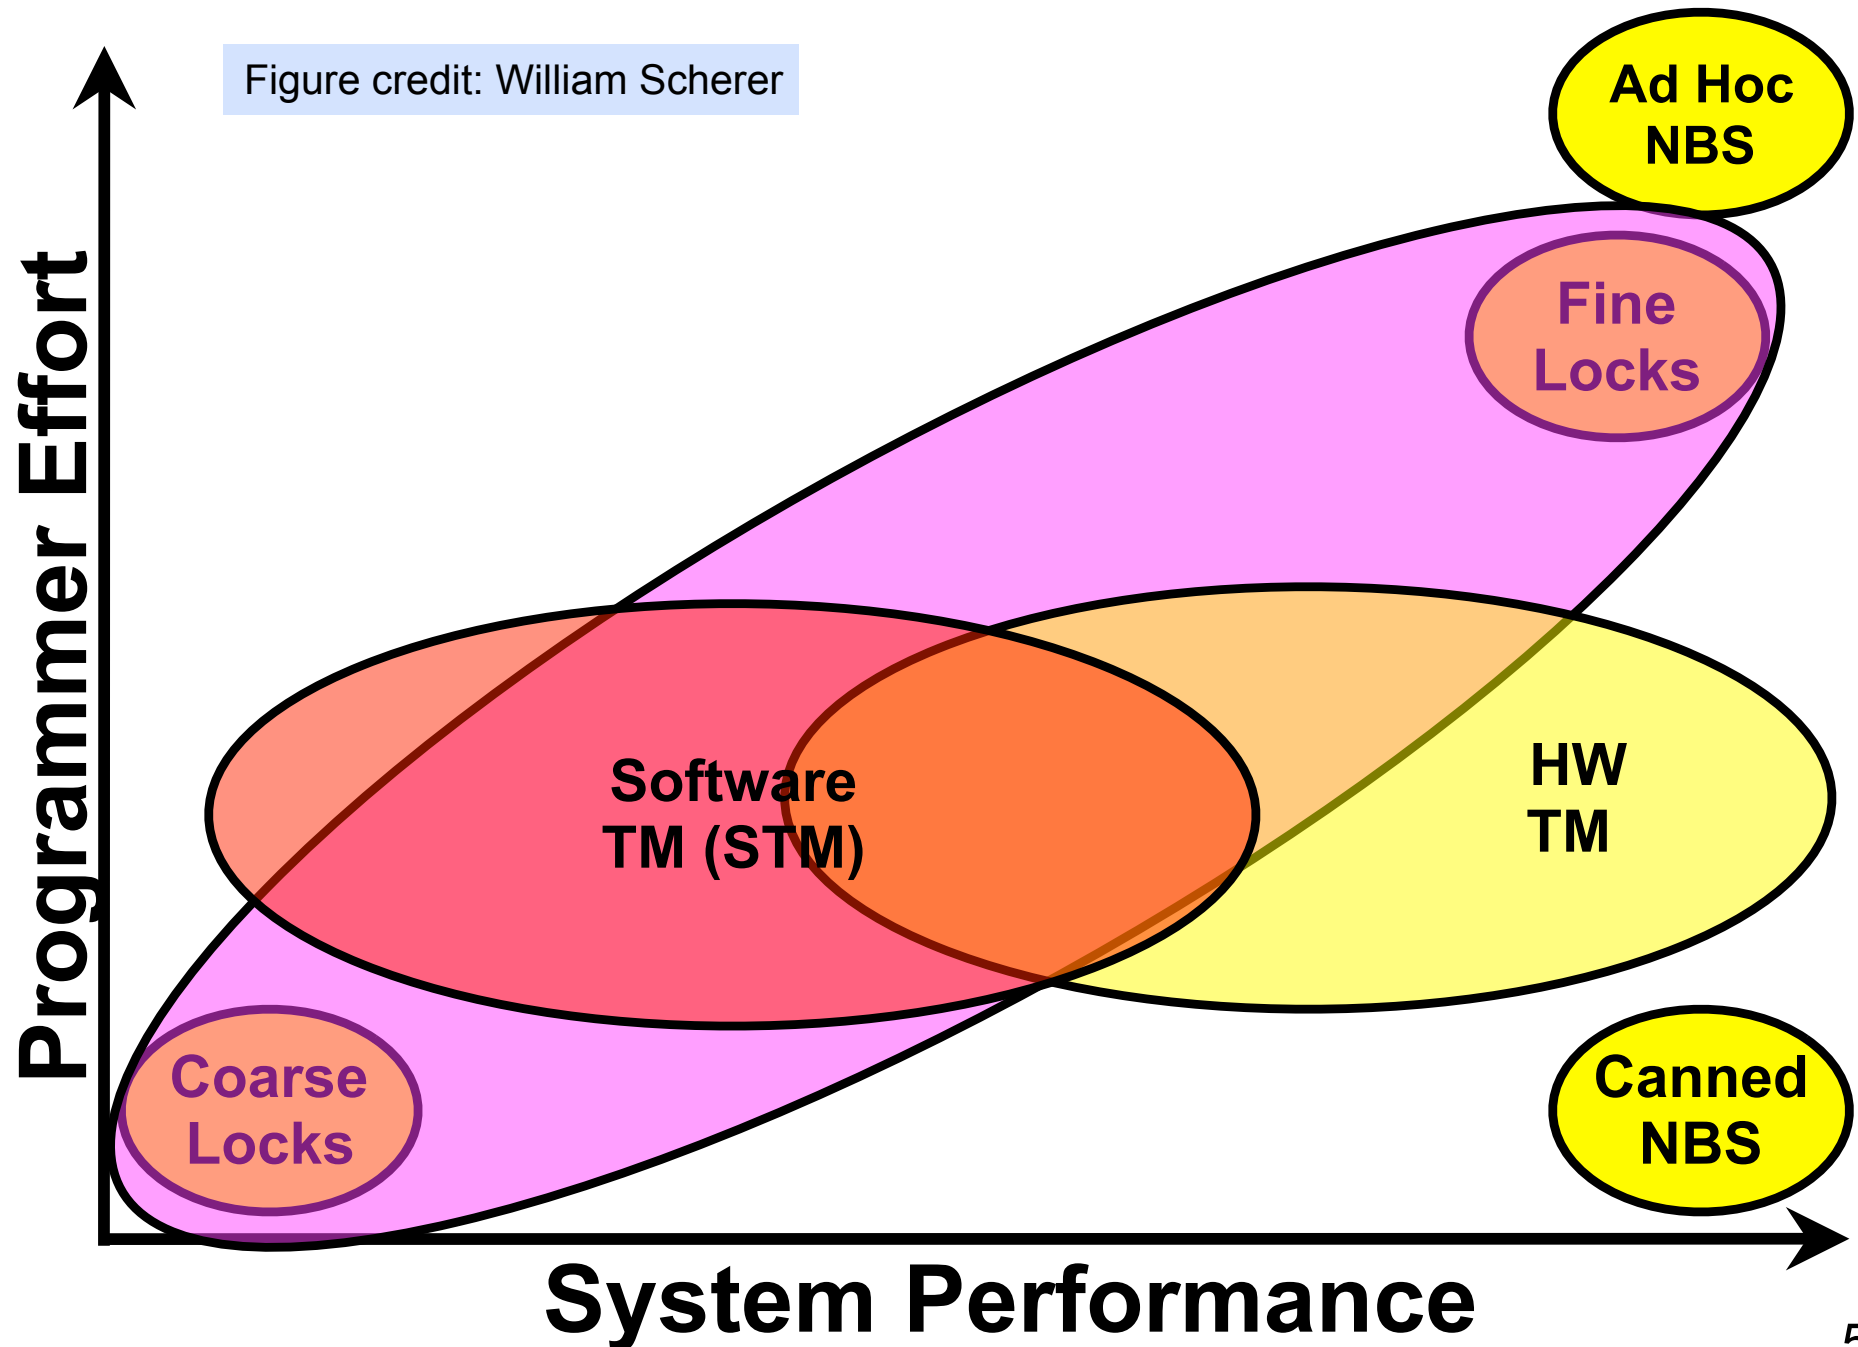
+ can be faster than locks (stacks, queues)

– difficult to write: every new algorithm is a publishable result

+ can be “canned” in libraries (e.g. java.util.concurrent's ConcurrentLinkedQueue)

Synchronization Landscape

Figure credit: William Scherer



Properties of Good Lock Algorithms

- **Mutual exclusion (*safety* property)**
 - critical sections of different threads do not overlap
 - cannot guarantee integrity of computation without this property
- **No deadlock**
 - if some thread *attempts* to acquire the lock, then some thread *will* acquire the lock
- **No starvation**
 - every thread that attempts to acquire the lock eventually succeeds
 - implies no deadlock

Notes

- **Deadlock-free locks do not imply a deadlock-free program**
 - e.g., can create circular wait involving a pair of “good” locks
- **Starvation freedom is desirable, but not essential**
 - practical locks: many permit starvation, although it is unlikely to occur
- **Without a real-time guarantee, starvation freedom is weak property**

Topics for Today

Classical locking algorithms using load and store

- Steps toward a two-thread solution
 - two partial solutions and their properties
- Peterson's algorithm: a two-thread solution
- Tree lock for n threads
- Lamport's bakery lock for n threads

Classical Lock Algorithms

- Use atomic load and store only, no stronger atomic primitives
- Not used in practice
 - locks based on stronger atomic primitives are more efficient
- Why study classical lock algorithms?
 - understand the principles underlying synchronization
 - ubiquitous in parallel programs
 - appreciate their subtlety
 - understand the motivation for atomic operations in hardware

Toward a Classical Lock for Two Threads

- First, consider two inadequate but interesting lock algorithms
 - use load and store only
- Assumptions
 - only two threads
 - each thread has a unique value of `self_threadid` $\in \{0,1\}$

Lock1

```
class Lock1: public Lock {
private:
    volatile bool flag[2];
public:
    void acquire() {
        int other_threadid = 1 - self_threadid;
        flag[self_threadid] = true;
        while (flag[other_threadid] == true);
    }
    void release() {
        flag[self_threadid] = false;
    }
}
```

set my flag



wait until other flag
is false



Using Lock1

assume that initially
both flags are false

thread 0

thread 1

flag[0] = true

while(flag[1] == true);

CS_0

flag[0] = false

flag[1] = true

while(flag[0] == true);

wait

CS_1

flag[1] = false

Lock1 Provides Mutual Exclusion

Proof

- Suppose not. Then $\exists j, k \in \text{integers}$

$$CS_0^j \not\rightarrow CS_1^k \quad \text{and} \quad CS_1^k \not\rightarrow CS_0^j$$

- Consider each thread's acquire before its j^{th} (k^{th}) critical section

$$\text{write}_0(\text{flag}[0] = \text{true}) \rightarrow \text{read}_0(\text{flag}[1] == \text{false}) \rightarrow CS_0 \quad (1)$$

$$\text{write}_1(\text{flag}[1] = \text{true}) \rightarrow \text{read}_1(\text{flag}[0] == \text{false}) \rightarrow CS_1 \quad (2)$$

- However, once $\text{flag}[1] == \text{true}$, it remains *true* while thread 1 in CS_1
- So (1) could not hold unless

$$\text{read}_0(\text{flag}[1] == \text{false}) \rightarrow \text{write}_1(\text{flag}[1] = \text{true}) \quad (3)$$

- From (1), (2), and (3)

$$\text{write}_0(\text{flag}[0] = \text{true}) \rightarrow \text{read}_0(\text{flag}[1] == \text{false}) \rightarrow \quad (4)$$

$$\text{write}_1(\text{flag}[1] = \text{true}) \rightarrow \text{read}_1(\text{flag}[0] == \text{false})$$

- By (4) $\text{write}_0(\text{flag}[0] = \text{true}) \rightarrow \text{read}_1(\text{flag}[0] == \text{false})$: a contradiction

Lock1

```
class Lock1: public Lock {
private:
    volatile bool flag[2];
public:
    void acquire() {
        int other_threadid = 1 - self_threadid;
        flag[self_threadid] = true;
        while (flag[other_threadid] == true);
    }
    void release() {
        flag[self_threadid] = false;
    }
}
```

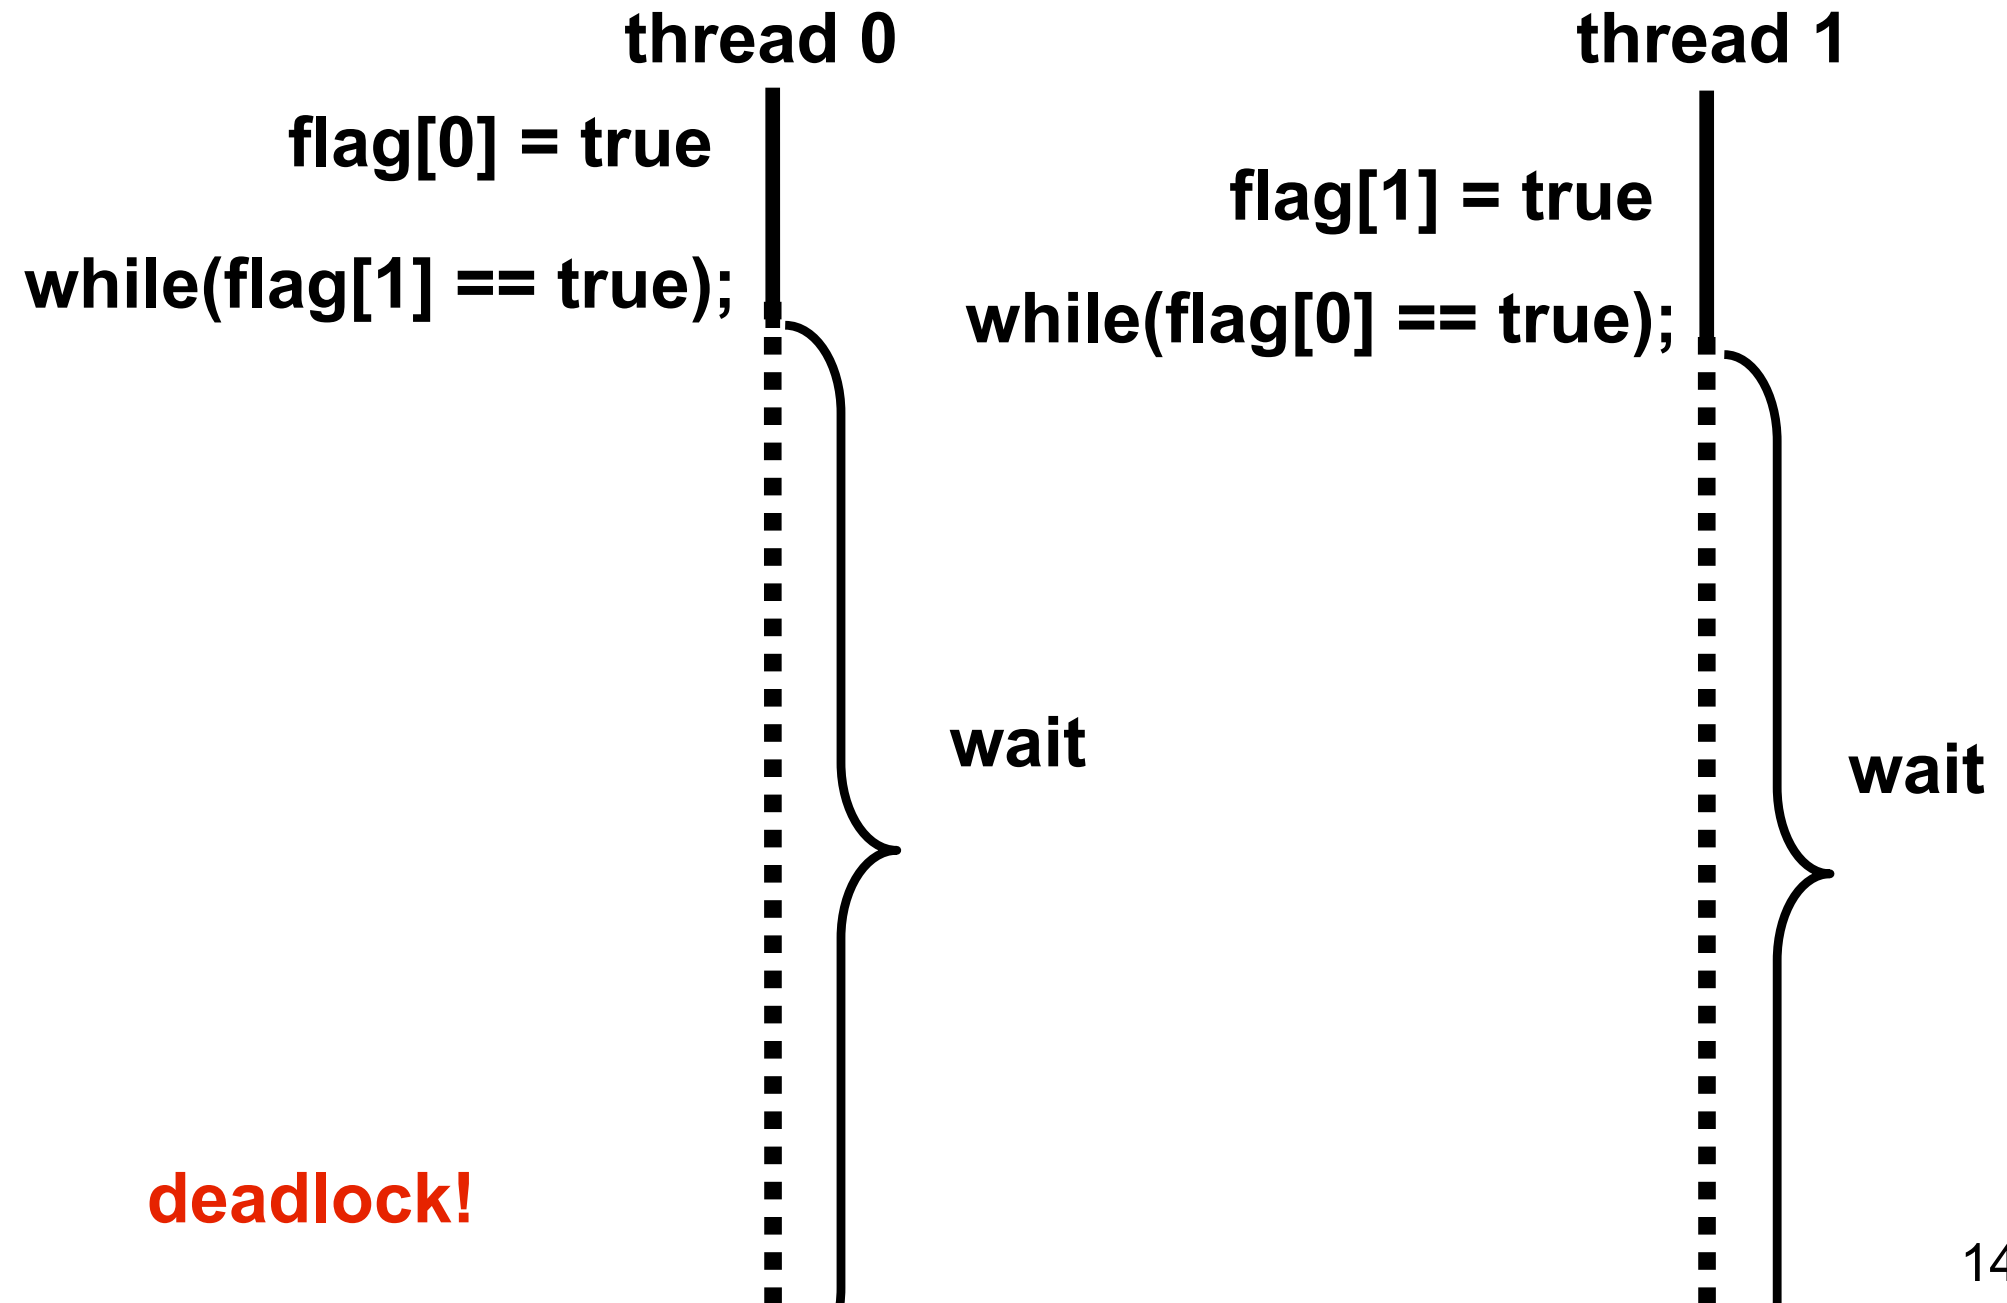
set my flag



wait until other flag
is false



Using Lock1



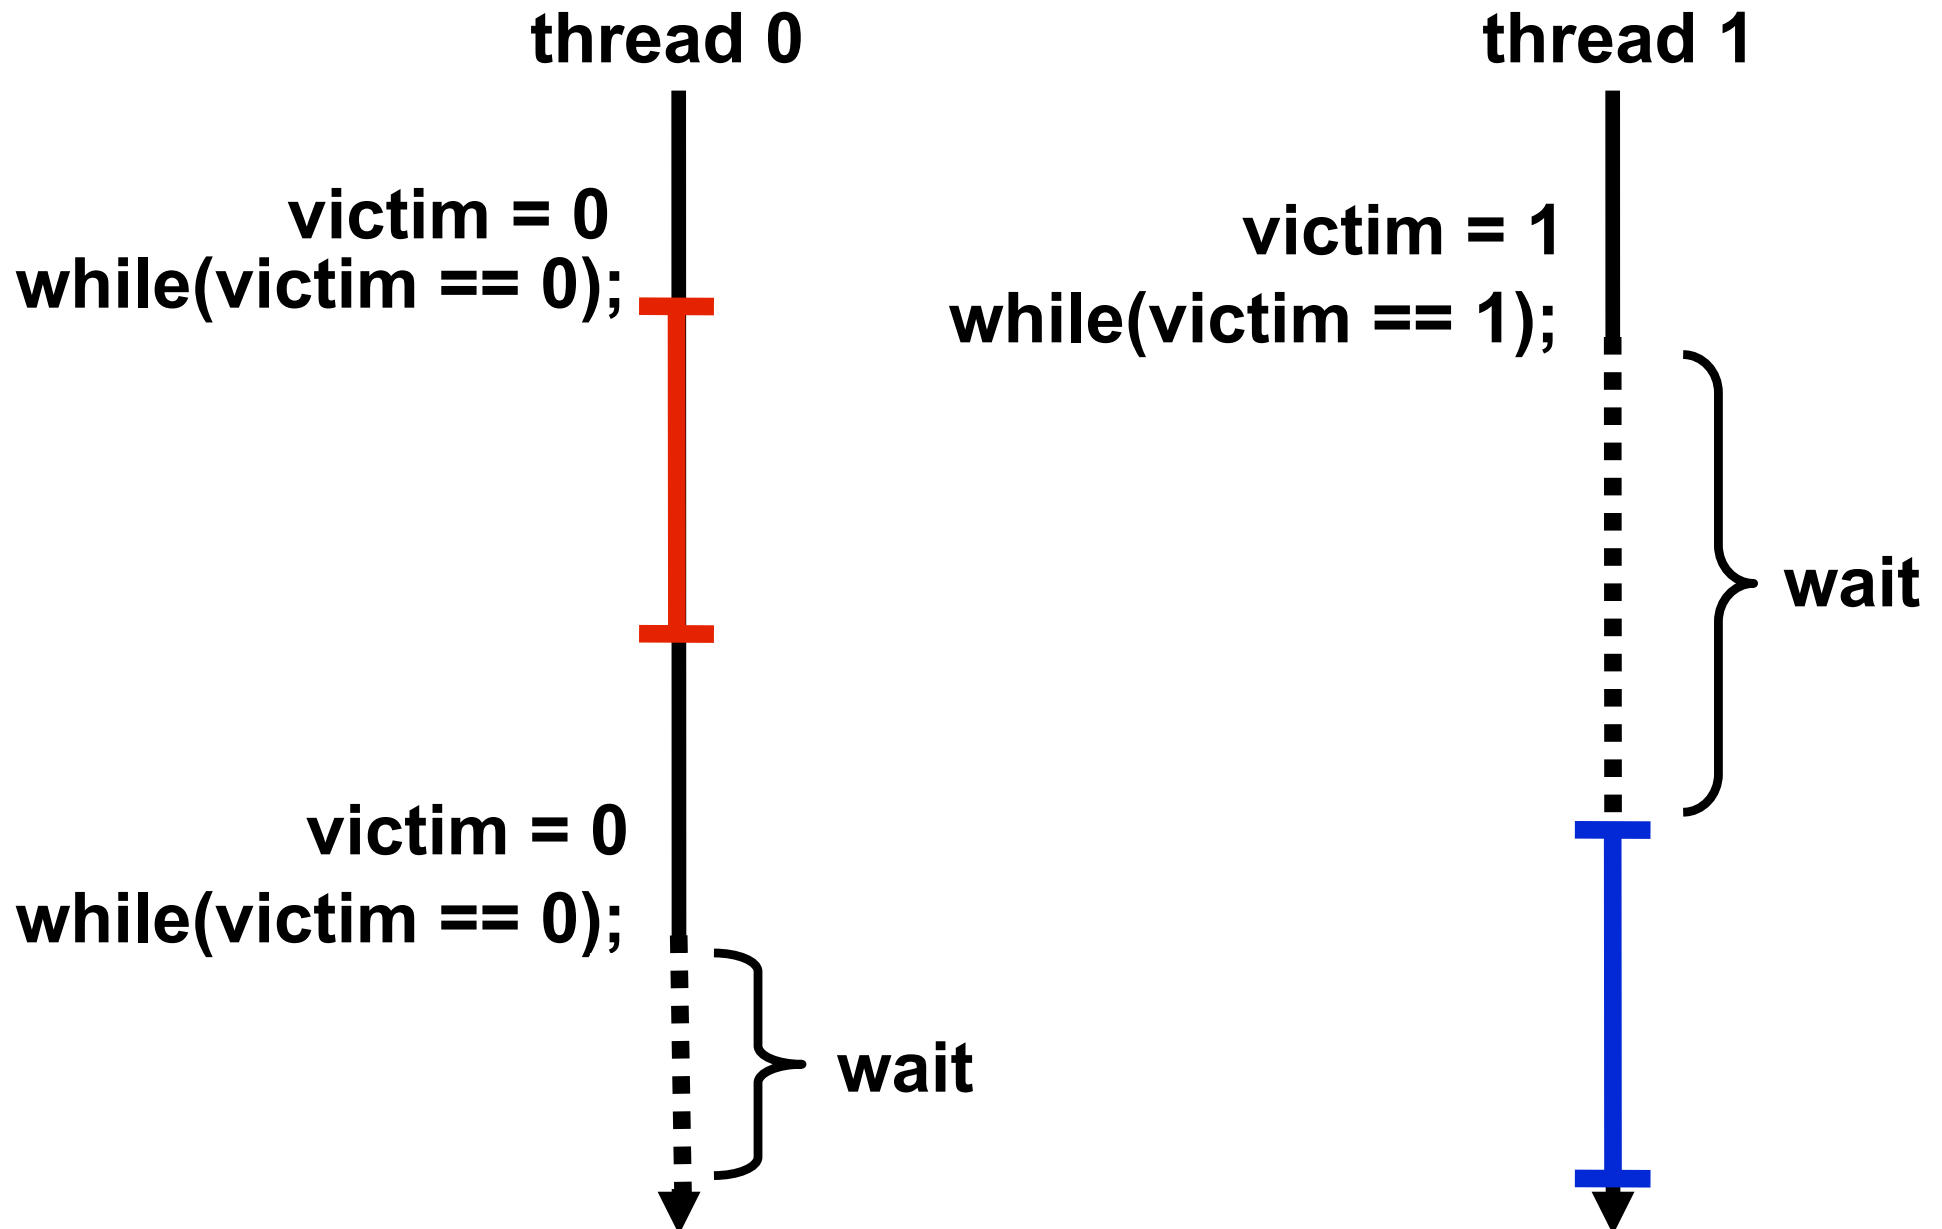
Summary of Lock1 Properties

- **Lock1 guarantees mutual exclusion**
- **Works if one thread completes its acquire before the other**
- **Deadlock if both threads write flags before either reads**
- **Since it admits deadlock, Lock1 is inadequate**

Lock2

```
class Lock2: public Lock {
private:
    volatile int victim;
public:
    void acquire() {
        victim = self_threadid;
        while (victim == self_threadid); // busy wait
    }
    void release() { }
}
```


Using Lock2



Lock2 Provides Mutual Exclusion

Proof

- Suppose not. Then $\exists j, k \in \text{integers}$

$$CS_0^j \not\Rightarrow CS_1^k \quad \text{and} \quad CS_1^k \not\Rightarrow CS_0^j$$

- Consider each thread's acquire before its j^{th} (k^{th}) critical section

$$\text{write}_0(\text{victim} = 0) \rightarrow \text{read}_0(\text{victim} \neq 0) \rightarrow CS_0 \quad (1)$$

$$\text{write}_1(\text{victim} = 1) \rightarrow \text{read}_1(\text{victim} \neq 1) \rightarrow CS_1 \quad (2)$$

- For thread 0 to enter the critical section, thread 1 must assign $\text{victim} = 1$

$$\text{write}_0(\text{victim} = 0) \rightarrow \text{write}_1(\text{victim} = 1) \rightarrow \text{read}_0(\text{victim} \neq 0) \rightarrow CS_0 \quad (3)$$

- Once $\text{write}_1(\text{victim} = 1)$ occurs, victim does not change

- Therefore, thread 1 cannot $\text{read}_1(\text{victim} \neq 1)$ and enter CS_1

- Contradiction!


```
void acquire() {  
    victim = self_threadid;  
    while (victim == self_threadid); // busy wait  
}
```

Lock2 protocol

Lock2

```
class Lock2: public Lock {
private:
    volatile int victim;
public:
    void acquire() {
        victim = self_threadid;
        while (victim == self_threadid); // busy wait
    }
    void release() { }
}
```

Using Lock2



Summary of Lock2 Properties

- **Guarantees mutual exclusion**
- **If two threads run concurrently: acquire succeeds for one**
- **Deadlock if one thread runs before the other**
- **Since it admits deadlock, Lock2 is inadequate**

Combining the Ideas

Lock1 and Lock2 complement each other

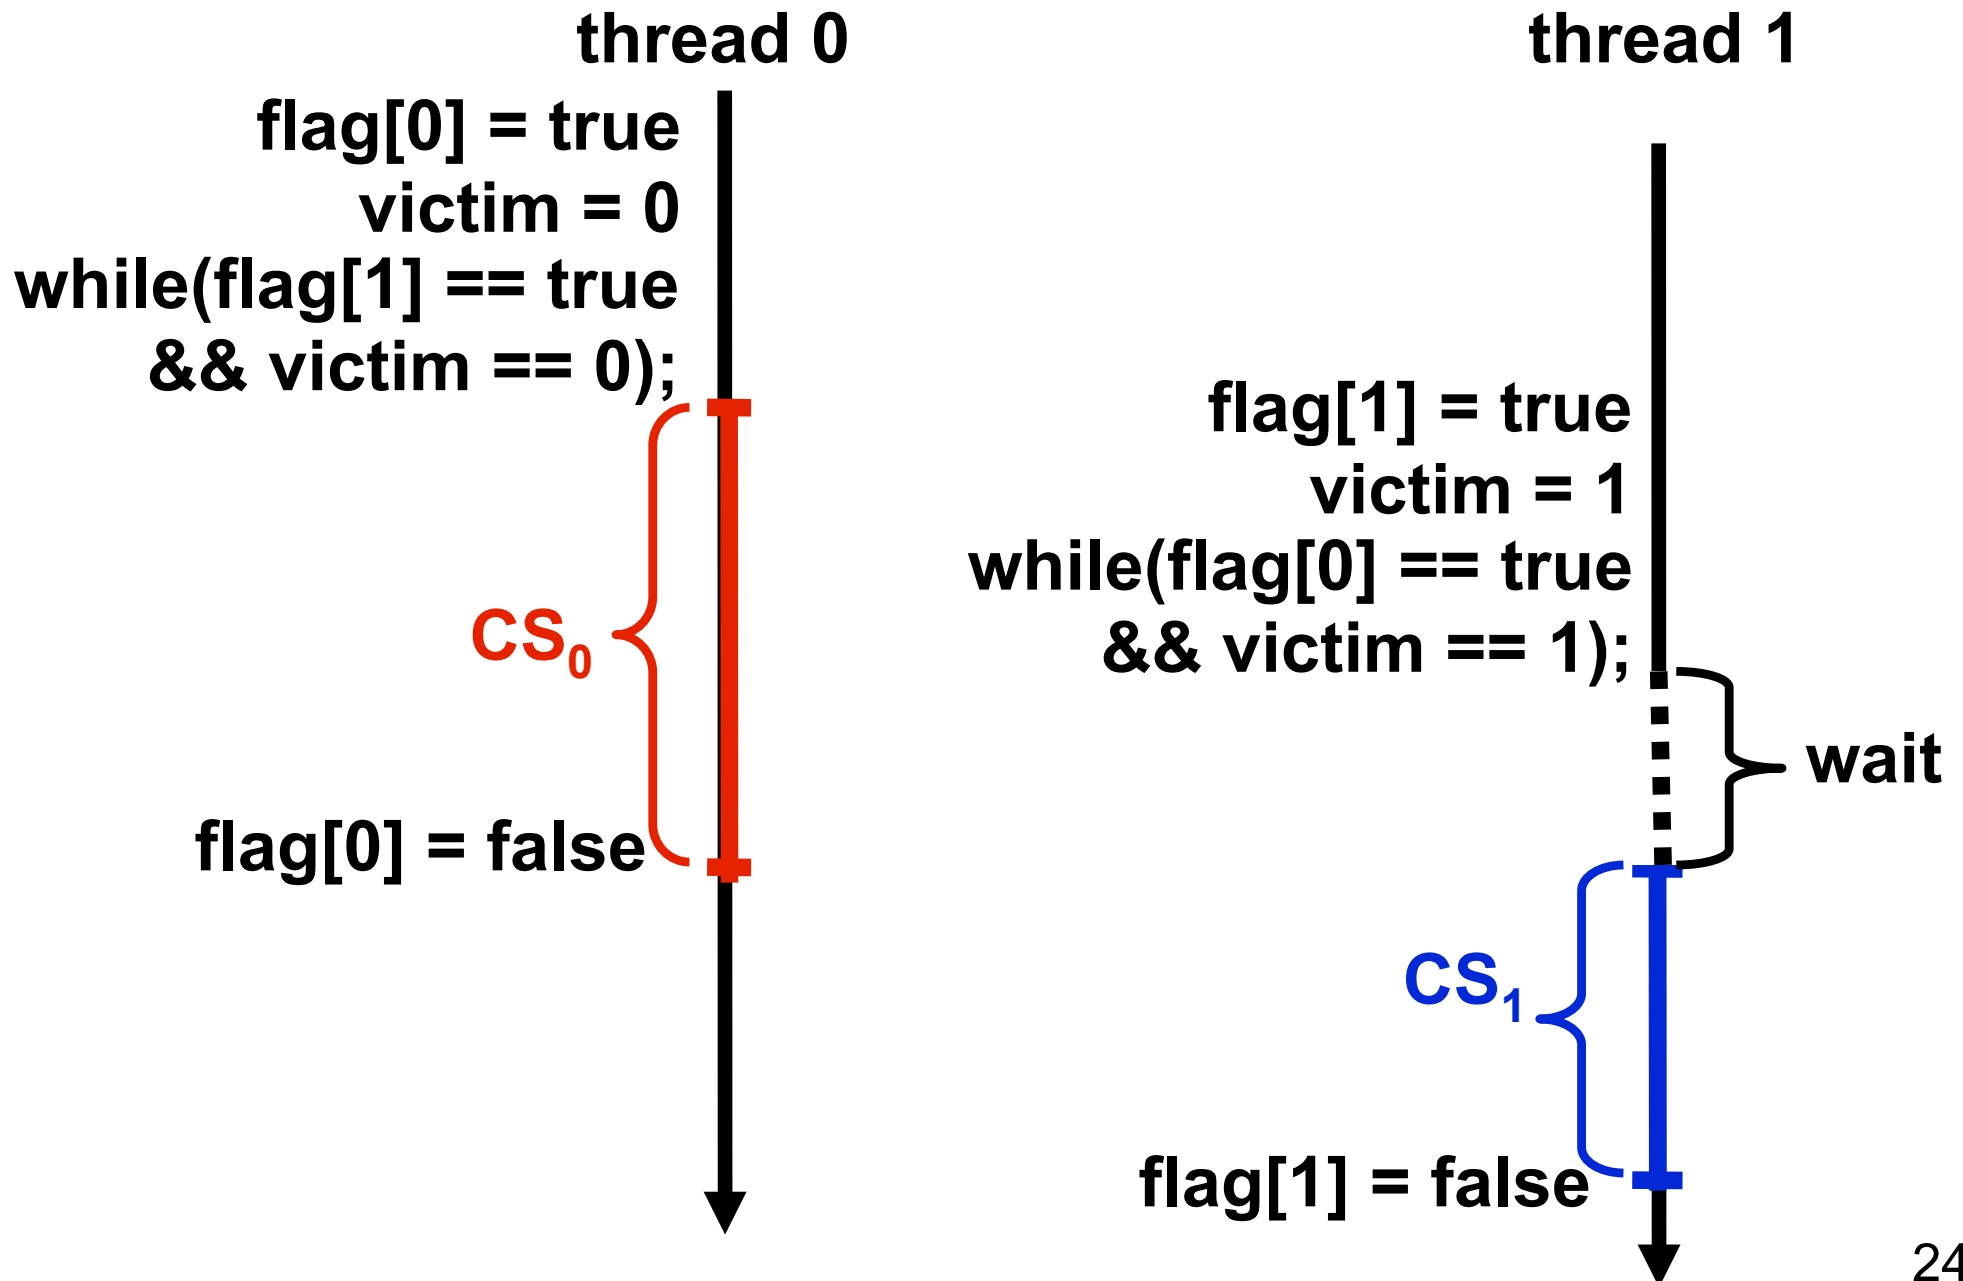
- Each succeeds under conditions that causes the other to fail
 - Lock1 succeeds when CS attempts **do not** overlap
 - Lock2 succeeds when CS attempts **do** overlap
- Design a lock protocol that leverages the strengths of both...

Peterson's Algorithm: 2-way Mutual Exclusion

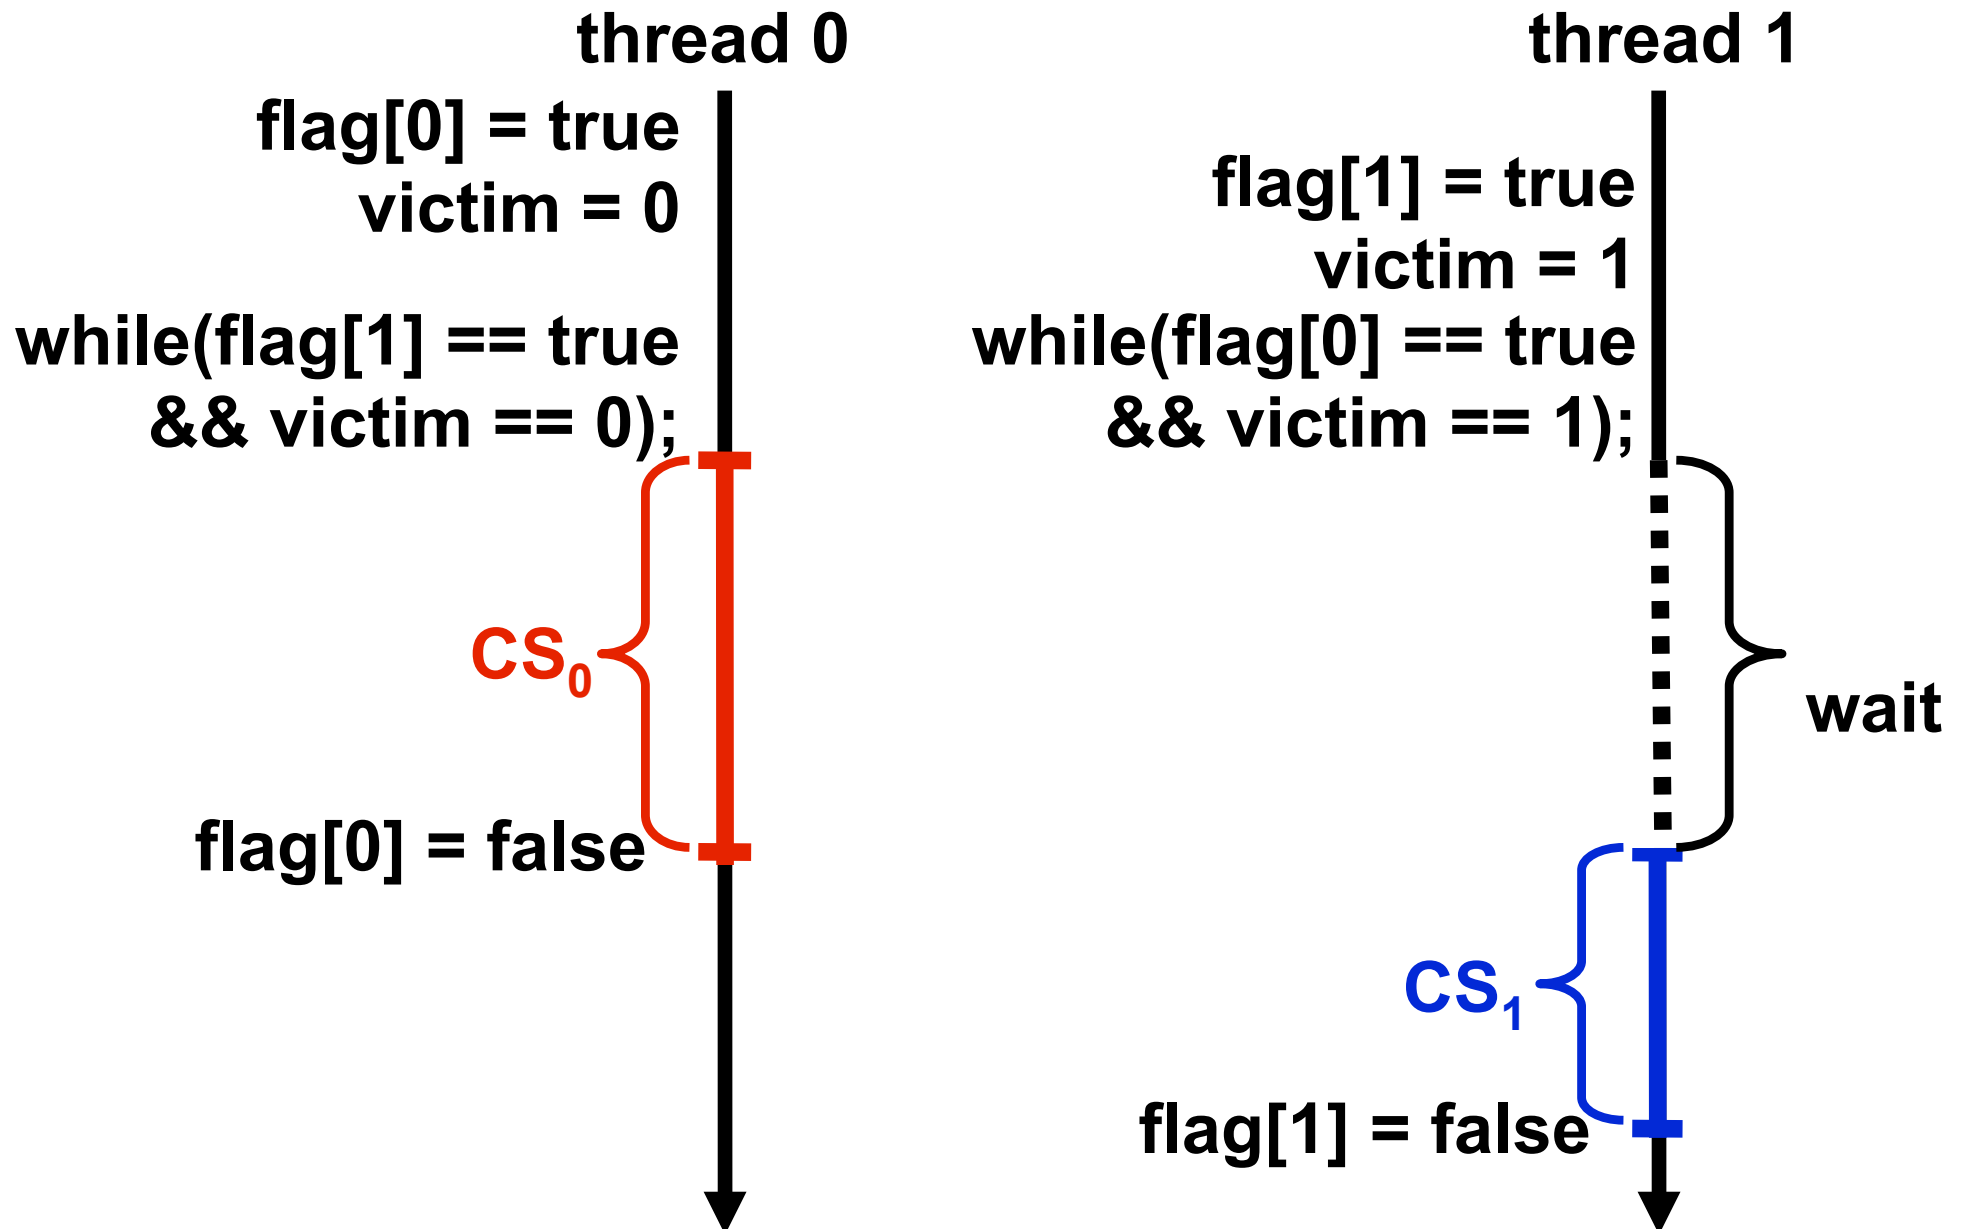
```
class Peterson: public Lock {
private:
    volatile bool flag[2];
    volatile int victim;
public:
    void acquire() {
        int other_threadid = 1 - self_threadid;
        flag[self_threadid] = true;      // I'm interested
        victim = self_threadid           // you go first
        while (flag[other_threadid] == true &&
               victim == self_threadid);
    }
    void release() {
        flag[self_threadid] = false;
    }
}
```

Gary Peterson. Myths about the Mutual Exclusion Problem.
Information Processing Letters, 12(3):115-116, 1981.

Peterson's Lock: Serialized Acquires



Peterson's Lock: Concurrent Acquires



Peterson's Algorithm Provides Mutual Exclusion

- Suppose not. Then $\exists j, k \in \text{integers}$

$$CS_0^j \not\Rightarrow CS_1^k \quad \text{and} \quad CS_1^k \not\Rightarrow CS_0^j$$

- Consider each thread's lock op before its j^{th} (k^{th}) critical section

$$\text{write}_0(\text{flag}[0] = \text{true}) \rightarrow \text{write}_0(\text{victim} = 0) \rightarrow \\ \text{read}_0(\text{flag}[1] == \text{false}) \text{ or } \text{read}_0(\text{victim} != 0) \rightarrow CS_0 \quad (1)$$

$$\text{write}_1(\text{flag}[1] = \text{true}) \rightarrow \text{write}_1(\text{victim} = 1) \rightarrow \\ \text{read}_1(\text{flag}[0] == \text{false}) \text{ or } \text{read}_1(\text{victim} != 1) \rightarrow CS_1 \quad (2)$$

- Without loss of generality, assume thread 0 was the last to write victim

$$\text{write}_1(\text{victim} = 1) \rightarrow \text{write}_0(\text{victim} = 0) \quad (3)$$

- From (1), (2), and (3), thread 0 must read **victim == 0** in (1)
- Since thread 0 nevertheless enters its CS, it must have read **flag[1]==false**
- From (1), it must be the case that **write₀(victim = 0) → read₀(flag[1] == false)**
- From (1), (2), and (3) and transitivity,

$$\text{write}_1(\text{flag}[1] = \text{true}) \rightarrow \text{write}_1(\text{victim} = 1) \rightarrow \\ \text{write}_0(\text{victim} = 0) \rightarrow \text{read}_0(\text{flag}[1] == \text{false}) \quad (4)$$

- From (4), it follows that **write₁(flag[1] = true) → read₀(flag[1] == false)**
- Contradiction!

Peterson's Algorithm is Starvation-Free


- Suppose not: WLG, suppose that thread 0 waits forever in acquire
 - it must be executing the while statement
 - waiting until $\text{flag}[1] == \text{false}$ or $\text{victim} != 0$
- What is thread 1 doing while thread 0 fails to make progress?
 - perhaps outside the critical section
 - $\text{flag}[1] == \text{true}$ only if thread 1 is awaiting or in the critical section
contradiction!
 - perhaps entering and leaving the critical section
 - if so, thread 1 will set victim to 1 when it tries to re-enter the CS
 - once it is set to 1, it will not change
 - thus, thread 0 must eventually return from acquire
contradiction!
 - waiting in acquire as well
 - waiting for $\text{flag}[0] == \text{false}$ or $\text{victim} == 0$
 - victim cannot be both 1 and 0, thus both threads cannot wait
contradiction!
- Corollary: Peterson's lock is deadlock-free as well

From 2-way to N-way Mutual Exclusion


- **Peterson's lock provides 2-way mutual exclusion**
- **How can we generalize to N-way mutual exclusion, $N > 2$?**
- **Several strategies that are generalizations of Peterson's lock**

An N-way Lock as a Tree of Peterson Locks

- For a lock involving N threads, construct a balanced binary tree with $N/2$ leaves. Assume $N = 2^k$
- Each thread uses Peterson's lock to compete against another thread in a leaf node of the tree
- When a thread acquires a lock, it moves up the tree to compete for the parent lock
- When a thread acquires the root lock, it may enter the critical section
- When a thread exits the critical section, it releases locks along the path from the root to its leaf



Properties of Tree of Peterson Locks



- **$O(N)$ space**
 - if $N = 2^k$, there are 2^{k-1} leaves and $N-1$ nodes in total
- **$\lg N$ steps to acquire or release the lock**

Lamport's N-way Bakery Algorithm

```
class LamportBakery: public Lock {
private:
    volatile bool flag[N]; volatile Label label[N];
public:
    void acquire() {
        int i = self_threadid;
        flag[i] = true;
        label[i] = max(label[0], ..., label[N-1]) + 1;
        while (exists k != i such that
            flag[k] && <label[k], k> <_ <label[i], i> );
    }
    void release() {
        flag[self_threadid] = 0;
    }
}
```

lexicographic ordering of
<label, thread_id> tuples;
thread id is used in tuple
to break labeling ties

Bakery Algorithm Intuition

- **Data structure components**
 - flag[A] = Boolean that indicate whether A wants to enter the CS
 - label[A] = integer that indicates the thread's turn to enter the bakery
- **Protocol operation**
 - when a thread tries to acquire the lock, it generates a new label
 - reads all other thread labels in some arbitrary order
 - generates a label greater than the largest it read
 - notes:
 - if 2 threads select labels concurrently, they may get the same
 - algorithm uses lexicographical order on pairs of (label, thread_id)
 - $(\text{label}[j], j) <_L (\text{label}[k], k)$
 - iff $(\text{label}[j] < \text{label}[k]) \vee ((\text{label}[j] == \text{label}[k]) \wedge j < k)$
 - in the waiting phase
 - a thread repeatedly rereads the labels
 - waits until
 - no thread with its flag set has a smaller (label, thread_id) pair
- **Proofs: See Herlihy and Shavit manuscript (deadlock-free, FIFO, ME)**

Spin Lock Performance: Maximal Contention

- Peterson-Buhr is a tree of Peterson's 2-party locks using load/store
- Spinlock uses test-and-set
- MCS lock uses SWAP and CAS

Figure credit: Peter A. Buhr, David Dice and Wim H. Hesselink. High-performance N-thread software solutions for mutual exclusion. Concurrency and Computation: Practice and Experience, 2014.




Figure 31. Critical section entry-counts, maximal contention: $N = 1.32$, SPARC, 20 s, measure of algorithm performance, where higher value is better.

Observations

- Bakery algorithm is concise, elegant and fair
- Why is it not practical?
 - must read N distinct locations; N could be very large
 - threads must be assigned unique ids between 0 and $n-1$
 - awkward for dynamic threads
 - value of a label is monotonically increasing & unbounded
- Are locking algorithms based on load/store commonly used?
 - no.
 - minimum space $O(N)$
 - uncontended acquisition latency is $O(\lg N)$
- Atomic primitives enable locks with
 - constant space
 - constant time acquisition in the uncontended case
 - maximum number of threads need not be known in advance

Spin Lock Performance: Minimal Contention

- Peterson-Buhr is a tree of Peterson's 2-party locks using load/store
- Spinlock uses test-and-set
- MCS lock uses SWAP and CAS

Figure credit: Peter A. Buhr, David Dice and Wim H. Hesselink. High-performance N-thread software solutions for mutual exclusion. Concurrency and Computation: Practice and Experience, 2014.




Figure 35. Critical section entry-counts, minimal contention: $N = 1.32$: SPARC, 20 s, measure of algorithm performance for zero contention, where higher value is better.

References

- Maurice Herlihy and Nir Shavit. “Art of Multiprocessor Programming” Chapter 2 “Mutual Exclusion,” Morgan Kaufmann, 2008.
- Gary Peterson. Myths about the Mutual Exclusion Problem. *Information Processing Letters*, 12(3), 115-116, 1981. <http://cs.nyu.edu/~lerner/spring12/Read03-MutualExclusion.pdf>