
Parallel Computing Wrapup

John Mellor-Crummey

**Department of Computer Science
Rice University**

`johnmc@rice.edu`

Course Objectives

- **Learn fundamentals of parallel computing**
 - principles of parallel algorithm design
 - programming models and methods
 - parallel computer architectures
 - modeling and analysis of parallel programs and systems
 - parallel algorithms
- **Develop skill writing parallel programs**
 - programming assignments employing a variety of models
- **Develop skill analyzing parallel computing problems**
 - develop parallelizations for different styles of computations

Review: Parallel Algorithm Design

Recipe to solve a problem using multiple processors

Typical steps for constructing a parallel algorithm

- identify what pieces of work can be performed concurrently
- partition and map work onto independent processors
- distribute a program's input, output, and intermediate data
- coordinate accesses to shared data: avoid conflicts
- ensure proper order of work using synchronization

Why “typical”? Some of the steps may be omitted.

- if data is in shared memory, distributing it may be unnecessary
- if using message passing, there may not be shared data
- the mapping of work to processors can be done statically by the programmer or dynamically by the runtime

Principles of Parallel Algorithm Design

- **Algorithm models**
 - data-parallel task graph work pool
 - master slave, pipeline, hybrid
 - **Decomposition techniques**
 - recursive
 - data driven: input data, output data, intermediate data
 - hybrid decomposition
 - exploratory decomposition
 - speculative decomposition
 - **Task generation**
 - static vs. dynamic
- assignment 4
- assignment 1
- assignment 2
- assignment 3
- assignment 1
- 2,3,4

Implementation Techniques

- **Concurrency and mapping**
 - static mapping strategies for regular problems
 - dynamic mapping
 - centralized task queue
 - work stealing
- **Communication model**
 - one-sided vs. two sided
- **Collective communication**
 - flavors
 - one-to-all: broadcast
 - all-to-one: reduce
 - all-to-all
 - parallel prefix computations: scan
 - gather/scatter
 - implementation techniques
 - broadcast of large messages as scatter + all-to-all

Programming Models

- **Shared-memory parallel programming**
 - Cilk/Cilk++
 - OpenMP
 - Pthreads
- **Global address space programming models**
 - Unified Parallel C (UPC)
- **Message passing and MPI**
- **GPU programming with CUDA**
- **MapReduce**

Parallel Architectures

- **Control structure and communication models**
 - control structure: SIMD, MIMD
 - communication models
 - shared address space
 - message passing platforms
- **Network topologies**
 - static/direct vs. dynamic/indirect networks
 - bus, crossbar, omega, hypercube, fat tree, mesh, Kautz graph
 - hybrid interconnects
 - evaluation metrics
 - degree, diameter, bisection width, channel width & rate, cost
- **Coherence, routing, and network embeddings**
 - blocking vs. non-blocking networks
 - routing techniques: store & forward, packet, wormhole
 - cache coherence: protocols, snoopy caches, directories, SCI
 - embeddings: dilation, congestion

Synchronization

- **Insufficient synchronization causes data races**
 - unordered, conflicting operations
- **Mutual exclusion: classical algorithms for locks**
 - explore formal reasoning about concurrent operations
- **Lock synchronization with atomic primitives**
 - practical algorithms for pairwise coordination
- **Barrier synchronization**
 - separate phases to prevent overlap of conflicting operations
 - strategies for fast, primitive collective synchronization

Parallel Algorithms

- **Parallel sorting**
- **Dense matrix algorithms**
 - **Cannon's algorithm**
 - **2.5D matrix multiply**

Top Ten Tips for Parallel Computing

It's all about the performance

- **Use an efficient algorithm**
 - clever implementation will yield to asymptotic inefficiency at scale
- **Partition your data and computation carefully**
 - the wrong data partitioning can yield high communication volume
 - the wrong computation partitioning can lead to load imbalance
 - work stealing can help
- **Choose your programming model judiciously**
 - shared-memory models make irregular problems easier
- **Avoid serialization**
 - efficiency requires all processors and cores to be computing
 - may require changes to algorithm and partitioning of data & computation
- **Choose the proper grain size for computation**
 - wrong grain size can lead to excessive communication frequency

Top Ten Tips for Parallel Computing

- **Design carefully to avoid race conditions**
 - an ounce of design is worth a pound of debugging
- **Avoid contention**
 - shared variable “hot spots”
 - msg passing: contention for interconnect links or destinations
- **Use the cache**
 - on microprocessor-based systems, memory hierarchy is **IMPORTANT**
- **Don't forget the microarchitecture**
 - an efficient algorithm kernel can boost performance by integer factors
- **Exploit parallelism at all levels**
 - SIMD instructions
 - instruction-level parallelism on pipelined processors
 - multiple cores; multiple threads per core (SMT, SIMT)
 - multi-socket nodes (SMP)
 - hardware accelerators (GPU, manycore) in nodes
 - clusters and supercomputers: nodes + interconnect