

Further introductory topics of C

Xiaoqiang Wang

What is the memory

- Memory is like a big table of numbered slots where bytes can be stored.
- The number of a slot is its **Address**. One byte **Value** can be stored in each slot.
- Array data values span more than one slot, like the character string "Hello\n"

Addr	Value
0	
1	
2	
3	
4	'H' (72)
5	'e' (101)
6	'l' (108)
7	'l' (108)
8	'o' (111)
9	'\n' (10)
10	'\0' (0)
11	
12	

72?

Data type

- A **Type** names a logical meaning to a span of memory. Some simple types are:

```
char  
char [10]  
int  
float  
double
```

```
a single character (1 slot)  
an array of 10 characters  
signed 4 byte integer  
4 byte floating point  
signed 8 byte floating point
```

- Use **sizeof()** to calculate the size of the Type:

```
double x[10];
```

```
printf("x spans %d bytes memory\n", sizeof(x));
```

What is a variable?

A **Variable** names a place in memory where you store a **Value** of a certain **Type**.

You first **Define** a variable by giving it a name and specifying the type, and optionally an initial value

declare vs define?

Initial value of x is undefined

```
char x;  
char y='e';
```

Initial value

Name

What names are legal?

Type is single character (char)

extern? static? const?

symbol table?

Symbol	Addr	Value
	0	
	1	
	2	
	3	
x	4	?
y	5	'e' (101)
	6	
	7	
	8	
	9	
	10	
	11	
	12	

The compiler puts them somewhere in memory.

Multi-byte variables

- Different types consume different amounts of memory. No necessary to be adjacent immediately.

```
char x;  
char y='e';  
int z = 0x01020304;
```

0x means the constant is written in hex

padding

An int consumes 4 bytes

Symbol	Addr	Value
	0	
	1	
	2	
	3	
x	4	?
y	5	'e' (101)
	6	
	7	
z	8	4
	9	3
	10	2
	11	1
	12	

What is the Stack?

- It's a special region of your computer's memory that stores temporary variables created by each function (including the `main()` function).
- The stack is a "FILO" (first in, last out) data structure
- A key to understanding the stack is the notion that **when a function exits**, all of its variables are popped off of the stack (and hence lost forever). Thus stack variables are **local** in nature. This is related to a concept we saw earlier known as **variable scope**, or local vs global variables.

Stack

- The stack grows and shrinks as functions push and pop local variables
- There is no need to manage the memory yourself, variables are allocated and freed automatically
- The stack has size limits
- Stack variables only exist while the function that created them, is running

Stack

Recall lexical scoping. If a variable is valid “within the scope of a function”, what happens when you call that function recursively? Is there more than one “exp”?

Yes. Each function call allocates a “stack frame” where Variables within that function’s scope will reside.

float x	5.0	
int exp	0	Return 1.0
float x	5.0	
int exp	1	Return 5.0
int argc	1	
char **argv	0x2342	
float p	5.0	

↑
Grows

```
#include <stdio.h>
#include <inttypes.h>

float pow(float x, int exp)
{
    /* base case */
    if (exp == 0) {
        return 1.0;
    }

    /* “recursive” case */
    return x*pow(x, exp - 1);
}

int main(int argc, char **argv)
{
    float p;
    p = pow(5.0, 1);
    printf("p = %f\n", p);
    return 0;
}
```


Challenge Problem

- Write `pow()` so it requires $\log(n)$ iterations.

Scoping summary

(Returns nothing)

- The scope of Function Arguments is the complete body of the function.
- The scope of Variables defined inside a function starts at the definition and ends at the closing brace of the containing block
- The scope of Variables defined outside a function starts at the definition and ends at the end of the file. Called “**Global**” Vars.

```
void p(char x)
{
    char y;           /* p, x */
    char z;           /* p, x, y */
                    /* p, x, y, z */
}
                    /* p */
char z;              /* p, z */

void q(char a)
{
    char b;           /* p, z, q, a, b */
    {
        char c;       /* p, z, q, a, b, c */
    }
    char d;           /* p, z, q, a, b, d (not c) */
}

/* p, z, q */
```

char b?

legal?

Referencing Data from Other Scopes

- So far, all of our examples all of the data values we have used have been defined in our lexical scope

```
float pow(float x, uint exp)
{
    float result=1.0;
    int i;
    for (i=0; (i < exp); i++) {
        result = result * x;
    }
    return result;
}

int main(int argc, char **argv)
{
    float p;
    p = pow(10.0, 5);
    printf("p = %f\n", p);
    return 0;
}
```


Nothing in this scope

Uses any of these variables

Can a function modify its arguments?

What if we wanted to implement a function `pow_assign()` that *modified* its argument, so that these are equivalent:

```
float p = 2.0;  
/* p is 2.0 here */  
p = pow(p, 5);  
/* p is 32.0 here */
```



```
float p = 2.0;  
/* p is 2.0 here */  
pow_assign(p, 5);  
/* p is 32.0 here */
```

Would this work?

```
void pow_assign(float x, uint exp)  
{  
    float result=1.0;  
    int i;  
    for (i=0; (i < exp); i++) {  
        result = result * x;  
    }  
    x = result;  
}
```

NO!

Remember the stack!

```
void pow_assign(float x, uint exp)
{
    float result=1.0;
    int i;
    for (i=0; (i < exp); i++) {
        result = result * x;
    }
    x = result;
}

{
    float p=2.0;
    pow_assign(p, 5);
}
```

In C, all arguments are passed as values

But, what if the argument is the *address* of a variable?

float x	32.0
uint exp	5
float result	32.0
float p	2.0

↑
Grows

Passing Addresses

- Recall our model for variables stored in memory
- What if we had a way to find out the address of a symbol, and a way to reference that memory location by address?

```
address_of(y) == 5  
memory_at[5] == 101
```

```
void f(address_of_char p)  
{  
    memory_at[p] = memory_at[p] - 32;  
}
```

```
char y = 101;    /* y is 101 */  
f(address_of(y)); /* i.e. f(5) */  
/* y is now 101-32 = 69 */
```

Symbol	Addr	Value
	0	
	1	
	2	
	3	
char x	4	'H' (72)
char y	5	'e' (101)
	6	
	7	
	8	
	9	
	10	
	11	
	12	


“Pointers”

This is exactly how “pointers” work.

“address of” or reference operator: `&`
“memory_at” or dereference operator: `*`

```
void f(address_of_char p)
{
    memory_at[p] = memory_at[p] - 32;
}
```

```
char y = 101;      /* y is 101 */
f(address_of(y)); /* i.e. f(5) */
/* y is now 101-32 = 69 */
```



A “pointer type”: pointer to char

```
void f(char * p)
{
    *p = *p - 32;
}
```

```
char y = 101;      /* y is 101 */
f(&y);             /* i.e. f(5) */
/* y is now 101-32 = 69 */
```

Pointers are used in C for many other purposes:

- Passing large objects without copying them
- Accessing dynamically allocated memory
- Referring to functions

Pointer Validity

- A **Valid** pointer is one that points to memory that your program controls.
- Using invalid pointers will cause non-deterministic behavior, and will often cause Linux to kill your process (SEGV or Segmentation Fault).
- There are two general causes for these errors:
 - Program errors that set the pointer value to a strange number
 - Use of a pointer that was at one time valid, but later became invalid
- Will ptr be valid or invalid?

```
char * get_pointer()
{
    char x=0;
    return &x;
}


{
    char * ptr = get_pointer();
    *ptr = 12; /* valid? */
}
```


Answer: Invalid!

- A pointer to a variable allocated on the stack becomes invalid when that variable goes out of scope and the stack frame is “popped”. The pointer will point to an area of the memory that may later get reused and rewritten.

```
char * get_pointer()
{
    char x=0;
    return &x;
}

{
    char * ptr = get_pointer();
    *ptr = 12; /* valid? */
    other_function();
}
```



But now, `ptr` points to a location that's no longer in use, and will be reused the next time a function is called!

Arrays

- Arrays in C are composed of a particular type, laid out in memory in a repeating pattern. Array elements are accessed by stepping forward in memory from the base of the array by a multiple of the element size.

```
/* define an array of 10 chars */  
char x[5] = {'t', 'e', 's', 't', '\0'};
```

Brackets specify the count of elements. Initial values optionally set in braces.

```
/* accessing element 0 */  
x[0] = 'T';
```

Arrays in C are 0-indexed (here, 0..9)

```
/* pointer arithmetic to get elt 3 */  
char elt3 = *(x+3); /* x[3] */
```

$x[3] == *(x+3) == 't'$ (NOT 's'!)

```
/* x[0] evaluates to the first element;  
 * x evaluates to the address of the  
 * first element, or &(x[0]) */
```

What's the difference
between `char x[]` and
`char *x`?

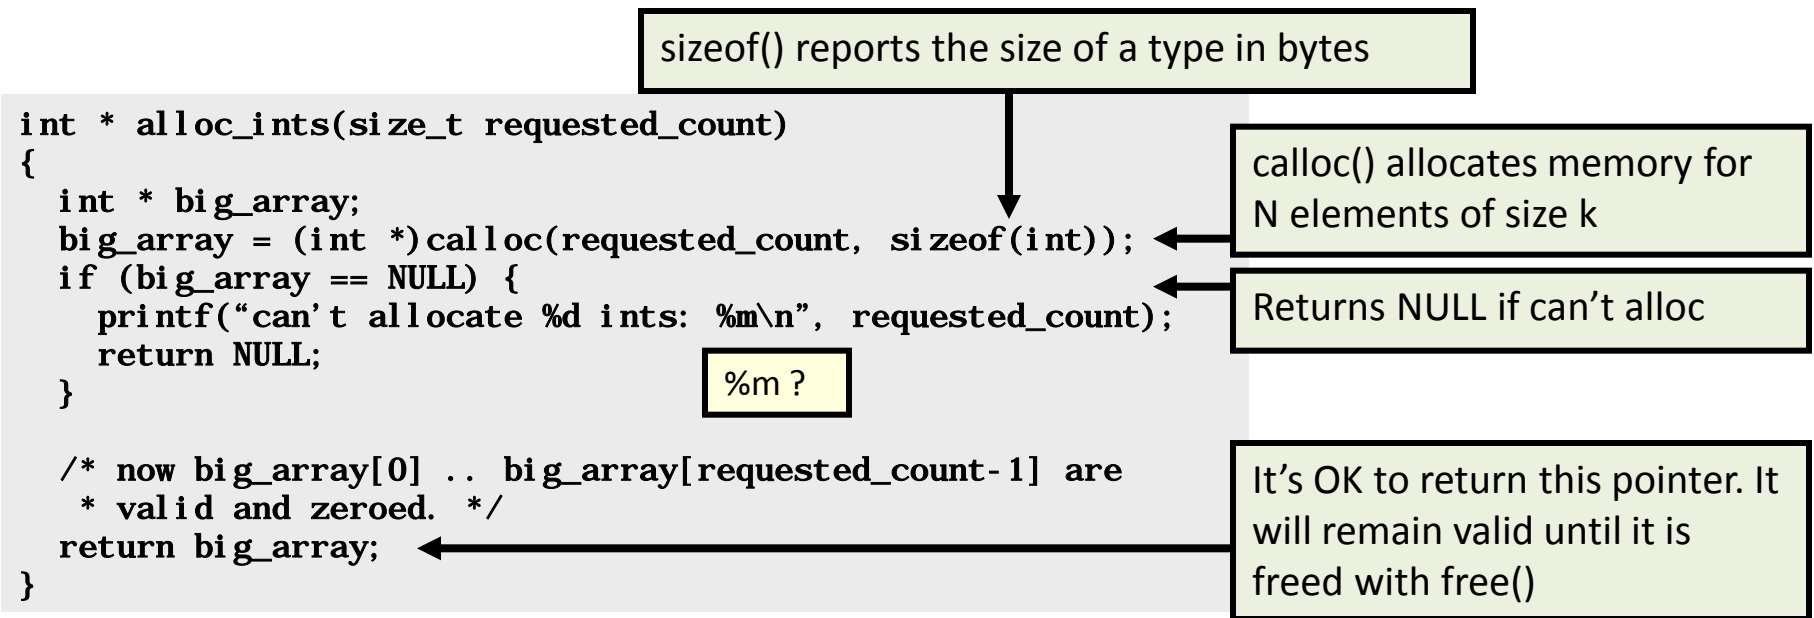
```
/* 0-indexed for loop idiom */  
#define COUNT 10  
char y[COUNT];  
int i;  
for (i=0; i<COUNT; i++) {  
    /* process y[i] */  
    printf("%c\n", y[i]);  
}
```

For loop that iterates from
0 to COUNT-1.
Memorize it!

Symbol	Addr	Value
char x [0]	100	't'
char x [1]	101	'e'
char x [2]	102	's'
char x [3]	103	't'
char x [4]	104	'\0'

Dynamic Memory Allocation

- So far all of our examples have allocated variables **statically** by defining them in our program. This allocates them in the stack.
- But, what if we want to allocate variables based on user input or other dynamic inputs, at run-time? This requires **dynamic** allocation.



Dynamic Memory Functions

Function	Description
<code>malloc</code>	allocates the specified number of bytes
<code>realloc</code>	increases or decreases the size of the specified block of memory. Reallocates it if needed
<code>calloc</code>	allocates the specified number of bytes and initializes them to zero
<code>free</code>	releases the specified block of memory back to the system

`malloc` returns a [void pointer](#) (`void *`), `malloc` allocates based on byte count but not on type. One may "cast" this pointer to a specific type:

```
int *ptr;  
ptr = malloc(10 * sizeof (*ptr)); /* without a cast */  
ptr = (int *)malloc(10 * sizeof (*ptr)); /* with a cast */
```

Do not forget to free your dynamic allocated memory!

Caveats with Dynamic Memory

- Dynamic memory is useful. But it has several caveats:
 - Whereas the stack is automatically reclaimed, dynamic allocations must be tracked and `free()`'d when they are no longer needed. With every allocation, be sure to plan how that memory will get freed. Losing track of memory is called a “memory leak”.
 - Whereas the compiler enforces that reclaimed stack space can no longer be reached, it is easy to accidentally keep a pointer to dynamic memory that has been freed. Whenever you free memory you must be certain that you will not try to use it again. It is safest to erase any pointers to it.
 - Because dynamic memory always uses pointers, there is generally no way for the compiler to statically verify usage of dynamic memory. This means that errors that are detectable with static allocation are not with dynamic

Memory Functions

- Include <mem.h>
- `void *memchr(const void *ptr, int ch, size_t len)`
 - `memchr` finds the first occurrence of `ch` in `ptr` and returns a pointer to it (or a null pointer if `ch` was not found in the first `len` bytes)
- `int memcmp(const void *ptr1, const void *ptr2, size_t len)`
 - `memcmp` compares two memory byte by byte. Return 0 if equal.
- `void *memcpy(void *dst, const void *src, size_t len)`
 - `memcpy` copies `len` characters from `src` to `dst` and returns the original value of `dst`
 - The result of `memcpy` is undefined if `src` and `dst` point to overlapping areas of memory
- `void *memmove(void *dst, const void *src, size_t len)`
 - `memmove` is just like `memcpy` except that `memmove` is guaranteed to work even if the memory areas overlap
- `void *memset(void *ptr, int byteval, size_t len)`
 - `memset` sets the first `len` bytes of the memory area pointed to by `ptr` to the value specified by `byteval`

String functions

- Include <string.h>
- `size_t strlen(const char *str)`
 - `strlen` returns the number of characters in `str` that precede the terminating null (`'\0'`) character
- `char *strcpy(char *dst, const char *src)`
 - Copy characters from `src` to `dst` (up to and including the terminating null character (`'\0'`) of `src`)
- `char *strncpy(char *dst, const char *src, size_t len)`
 - Copy up to `len` characters from `src` to `dst`
 - If `src` is shorter than `len`, `dst` is filled to `len` characters with null characters
 - If `src` is longer than `len` characters, the string in `dst` is **not** terminated with a null character
- `strcpy` and `strncpy` return a pointer to `dst`
- If `src` and `dst` are overlapping string, the results of `strcpy` and `strncpy` are undefined

String functions

- `char *strcat(char *dst, const char *src)`
 - Append (catenate) the string `src` onto the end of `dst`, beginning by overwriting the terminating null in `dst` and continuing until the terminating null character of `src` is copied to `dst`
- `int strcmp(const char *str1, const char *str2)`
 - Compare each character in `str1` to the corresponding character in `str2` until either a dissimilar character or a null terminator (`'\0'` character) is found
 - If both strings are identical up to and including the terminating `'\0'`, 0 is returned
 - If `str1[n] < str2[n]`, -1 is returned
 - If `str1[n] > str2[n]`, 1 is returned
 - If `str1` is shorter than `str2` (for example, `str1` points to "abc" and `str2` points to "abcdefg"), -1 is returned
 - Likewise, if `str1` is longer than `str2`, 1 is returned
- `char *strstr(const char *str, const char *substr);`
 - `strstr` finds the first occurrence of `substr` in `str` and returns a pointer to the first character of the substring in `str`
 - If `substr` is not found in `str`, a null pointer is returned.
 - If `substr` is the empty string (that is, if the first character in `substr` is a null character) a pointer to the first character in `str` is returned.
- Conversion Functions: `atoi`, `atol` and `atof`
 - `int atoi(const char *str)`, convert string to integer

File I/O

- `#include <stdio.h>`
- For files you want to read and write, you need a file pointer, e.g.:
`FILE *fp;`
- Opening a file:
`fp = fopen(filename, mode);`
 - mode: “r” - Read only
 - “w” – Create a file to write
 - “a” – Append
 - “b” – binary mode
- Always close a file when you no need to use it:
`fclose(fp);`
- Writing to a file:
 - Text mode: `fprintf(...), fscanf(...), fgetc(fp), fgets(buf, n, fp);`
 - Binary mode: `fwrite(ptr, size_of_elements, number_of_elements, fp);`
`fread(ptr, size_of_elements, number_of_elements, fp);`
- you can set the position where to do the write and read:
`fseek(fp, offset, whence);`
whence can be `SEEK_SET` – beginning of file,
`SEEK_CUR` - Current position, `SEEK_END` - End of file.
`ftell(fp)`: get the current position of the file.

File I/O Example

- You can treat a text file as a binary file, and read the entire file into memory:

```
/* example: read an entire file */
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```
int main ()
```

```
{
```

```
    FILE * pFile;
```

```
    long ISize;
```

```
    char * buffer;
```

```
    size_t result;
```

```
    pFile = fopen ( "myfile.bin" , "rb" );
```

```
    if (pFile==NULL) {fputs ("File error",stderr); exit (1);}
```

```
    // obtain file size:
```

```
    fseek (pFile , 0 , SEEK_END);
```

```
    ISize = ftell (pFile);
```

```
    fseek (pFile , 0 , SEEK_SET);
```

```
    // allocate memory to contain the whole file:
```

```
    buffer = (char*) malloc (sizeof(char)*ISize);
```

```
    if (buffer == NULL) {fputs ("Memory error",stderr); exit (2);}
```

```
    // copy the file into the buffer:
```

```
    result = fread (buffer,1,ISize,pFile);
```

```
    if (result != ISize) {fputs ("Reading error",stderr); exit (3);}
```

```
    /* the whole file is now loaded in the memory buffer. */
```

```
    // terminate
```

```
    fclose (pFile);
```

```
    free (buffer);
```

```
    return 0;
```

```
}
```