

HOMEWORK: Fortran

In this homework, write a Fortran program that demonstrates the use of modules, interfaces, types, and operator overloading.

Consider the problem of storing a sparse matrix, and the creation of operators to add two sparse matrices.

A 2D matrix of size $n \times m$ has mn elements a_{ij} where $i = 1, \dots, m$ and $j = 1, \dots, n$. This matrix is sparse if most of its elements are zero. Note that the matrix is not necessarily a square matrix.

We need two integer variables: `nr=m` and `nc=n` (not necessarily equal). We will not store the full matrix, since with $m = n = 10^6$, would require memory of 10^{12} floats, which is more than you have available on your desktop computers. Instead, we will simply track each non-zero element by its row, column and array value at that location.

This matrix will be defined by the Fortran 90 type:

```
type sparse
  integer :: nr, nc    ! number of rows and columns
  integer :: nel       ! number of nonzero elements

  ! row (ia) and column (ja) of each nonzero element
  integer, dimension(:), pointer :: ia, ja

  ! double precision: nel elements
  real(8), dimension(:), pointer :: a
end type sparse
```

To better understand the storage format, consider the matrix of `real(8)`

$$A = \begin{pmatrix} 0. & 2. & 5. & 0. & 0. & 3. \\ 0. & 0. & 1. & 2. & 0. & 4. \\ 1. & 5. & 0. & 0. & 2. & 1. \end{pmatrix}$$

Instead of storing $18 \text{ real}(8) = 144$ bytes, we will represent the matrix with three arrays. First, `a` is a 1D array whose size is `nel`, equal to the number of nonzero elements of A , written out in the same order of appearance as found in A (scanning along rows).

$$a = (2., 5., 3., 1., 2., 4., 1., 5., 2., 1.)$$

Second, the columns of each nonzero element of A are stored in `ja`, also of size `nel`:

$$ja = (2, 3, 6, 3, 4, 6, 1, 2, 5, 6)$$

Finally for each of the three rows of A , we identify the first nonzero element: row 1 gives 2, row two gives 1, and row 3 gives 1. Each of these elements is found in the array \mathbf{a} , and its position is stored in array \mathbf{ia} :

$$\mathbf{ia} = (1, 4, 7, 11)$$

The size of \mathbf{ia} is $\mathbf{nel}+1$, and the last value of \mathbf{ia} is $\mathbf{nel}+1=11$. In this example, $\mathbf{nel}=10$. So we have stored 10 real(8) in \mathbf{a} , 10 integers in \mathbf{ja} and 4 integers in \mathbf{ia} , for a total of $(10*8+10*4+4*4=136)$ bytes. We have already saved 8 bytes. Not much. But in this case, there are 10 nonzero elements in a matrix of size 18. In the homework, the matrix will be of size 10^7 , and there will be many less nonzero elements. Thus the savings are much more significant.

Consider row i (where $i = 1, 2, 3$). The nonzero elements of A can be found in $\mathbf{a}(\mathbf{ia}(i))$ to $\mathbf{a}(\mathbf{ia}(i+1)-1)$. Thus for row 2, $\mathbf{ia}(2)=4$ and $\mathbf{ia}(3)=7$. Therefore, the nonzero elements of A are $\mathbf{a}(4)=3$, $\mathbf{a}(5)=4$, and $\mathbf{a}(6)=6$. The columns of these three non-zero elements are $\mathbf{ja}(4)=3$, $\mathbf{ja}(5)=4$, and $\mathbf{ja}(6)=6$.

In summary, to scan this array, scan each row: $i = 1, nr$. For each row i , calculate $\mathbf{ia}(i)$ and $\mathbf{ia}(i+1)-1$. The nonzero columns are $\mathbf{ja}(\mathbf{ia}(i))$ through $\mathbf{ja}(\mathbf{ia}(i+1)-1)$, and the nonzero elements of A are $\mathbf{A}(i, \mathbf{ja}(\mathbf{ia}(i)))$ through $\mathbf{A}(i, \mathbf{ja}(\mathbf{ia}(i+1)-1))$.

When adding two arrays, for each row, identify the nonzero columns and add the corresponding elements of the two arrays. Be careful: while the two matrices to be added are both sparse, the nonzero elements are not necessarily in the same locations. You should check that the addition works properly on small 3×3 or 4×4 matrices with nonzero elements in different locations. .

Under non circumstances are you to store the full matrix. The matrix should on be stored in compressed format. However, you are allowed to allocate temporary storage for a single row or column if you feel that is necessary.

- Create a module called `sparse_mod`, that defines the `sparse_matrix` type which contains all necessary information on the matrix: number of rows and columns (nr, nc), number of nonzero elements, and pointers to arrays \mathbf{a} , \mathbf{ia} , \mathbf{ja} .
- Allocate \mathbf{nel} elements for \mathbf{a} and \mathbf{ja} and $\mathbf{nr}+1$ elements for \mathbf{ia} .
- Create a matrix of size 1,000 by 10,000, and fill it with 10,000 elements. Use a constructor for this purpose. Thus on average, each row contains 10 elements. You should use a random number generator to assign the columns for each row, and you should use a random number generator to fill the nonzero elements of the matrix with values between 0 and 10.
- Create an operator(+) in the module to add two sparse matrices. For example, consider the two matrices A and B :

$$A = \begin{pmatrix} 0 & 3 & 0 \\ 5 & -7 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

$$B = \begin{pmatrix} 0 & 3 & 0 \\ 2 & -6 & 0 \\ 4 & 0 & -3 \end{pmatrix}$$

Notice that A and B matrices do not have zeros in the same elements.
The sum is:

$$C = A + B = \begin{pmatrix} 0 & 6 & 0 \\ 7 & -13 & 0 \\ 4 & 0 & -2 \end{pmatrix}$$

- When adding two sparse matrices of equal dimension, it is therefore important to consider all the non-zero elements of both matrices.
- Create a main program. This main program should have a structure similar to:

```

program main
use sparse_mod
implicit none
type(sparse) :: mat1, mat2, matsum
....
! Initialize mat1 and mat2, which should have "n" non-zero elements, randomly
! distributed across columns and rows. Write your own routine to create random
! numbers between 1 and N, where N is the size of the array you wish to construct.
! The nonzero elements should be at the same locations.

...
    call initialize (mat1)
    call initialize (mat2)

    matsum = mat1 + mat2

print results      ! (e.g. the nonzero elements of rows 5 for mat1, mat2
                  !                               and matsum
end program main

```

- Repeat the above summation with the same size matrix but with a variable number of nonzero elements. Three cases with: 100 nonzero elements, 10000 nonzero elements, and 1,000,000 nonzero elements. Do this once you are convinced your code produces the correct results. Time (benchmark) the addition for each of these three cases. Present the results in a table, in your pdf file. To benchmark accurately, consider doing

```

! start_timer
  do nb = 1,100
    ! perform the addition
  end do
! end_timer

or

  do nb = 1,100
    ! start_timer
    !   perform the addition
    ! end_timer

```

Make sure you provide me the time for a single addition. Tell me what computer you are running on, and what compiler options you use. It is a good habit to get into.

- Make sure you have NOTES/README/INSTALL files to describe ideas you may have had while you were working (NOTES), what the program does (README), and what to do to get the program to run (INSTALL).
- Make sure the code is documented (you are graded on this)